

Are All Slugfester Assessments on the Same Scale?

A narrower level difference, persistent process warnings, and a better test for the next review.

2.35

points lower on average
in the later process

56 / 71

returning people with
lower later means

7,032

verified comparable
move scores

The process-level difference has narrowed, but comparability still needs testing.

PHIL STILWELL · SLUGFESTER

October 9, 2026 | Frozen archive: 308 assessments

Recalculated scores; direct slogan evidence retains its September scope.

Executive summary

Correct arithmetic does not establish a common measuring standard. The 291 comparable assessments reconstruct successfully from 7,032 scored moves. Yet the earlier and later processes differ in score levels, annotation rates, and the range of marks assigned to individual scoring areas.

The earlier group contains 179 debates with an average side score of 81.32. The later group contains 112 with an average of 78.98. The later average is 2.35 points lower, compared with 2.82 points in September. Of 71 people present in both groups, 56 have lower later averages, 13 higher, and 2 unchanged.

These are warning signs for pooling, not proof that the later process causes lower scores. Different topics, opponents, sources, and performances enter the two groups. The appropriate response is a shared-source comparability test, not subtracting a convenient offset from public scores or assuming a newer model must be more accurate.

Reader's guide

Executive summary	2
1. What the reconstruction verifies	4
2. Score levels differ, though the difference has narrowed	4
3. Returning speakers help, but do not solve the problem	5
4. Equal scales can conceal different score vocabularies	6
5. Annotation practice is another comparability warning	7
6. What a useful reassessment should test	7
7. A qualified conclusion	8
Methods and sources	8

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. Score results refer to the October 9 snapshot. Direct slogan counts retain their September source coverage; they have not been extended to new transcripts.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. What the reconstruction verifies

The analysis checks 112 final-ledger fingerprints, 7,032 move-score calculations, 3,206 section-side scores, 582 overall side scores, and all 7,032 public move-score matches. It preserves move importance, section weighting, the published rounding sequence, and any burden-completion adjustment.

This establishes that the analytic records reproduce the published calculations. It does not establish that a source was perfectly transcribed, that every important contribution was selected, or that each judgment is correct. A calculator can accurately add marks that were assigned under different standards.

The two groups are defined by the structure of their scoring records: earlier detailed findings and later standalone final dimensions. They are not randomized treatments or a simple partition by debate recording year. All comparable records in this analysis are labeled 5.6 Sol; this research refresh is not an Astra reassessment of the debates. The historical slogan reading is a separate study with its own provenance.

2. Score levels differ, though the difference has narrowed



Figure 1. The two processes have different observed mean levels. Each debate contributes one midpoint.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population.

For each debate, average the two side scores, then average those midpoints within the process. With two sides per debate, this gives each debate equal influence. The later-minus-earlier difference is -2.35 points, with a resampling range of -3.08 to -1.62.

The later group's expanded average is closer to the earlier group's than it was in September. That change comes from added assessed performances, not a recalibration of old marks. It would be misleading to retain the old 2.82-point difference merely because the larger claim about comparability remains relevant.

A process-level mean cannot isolate the judging procedure. If the new records contain more difficult motions, different speakers, or less complete replies, lower scores might be appropriate. If a later instruction uses the scale more strictly, some change could come from judgment practice. This archive does not randomly separate those explanations.

3. Returning speakers help, but do not solve the problem

The 71 people appearing in both groups provide a useful bridge. Their mean later-minus-earlier change is -2.14 points when each person receives equal influence. 56 have lower later averages, 13 higher, and 2 the same.




Figure 2. Most returning people have lower later averages, but these are different debates, not duplicate judgments of one performance.

How to read this graph. The horizontal coordinate is the earlier average and the vertical coordinate the later average. Below the diagonal means lower later. Both axes enlarge the observed score range. These are different performances, not the same debate judged twice.

This reduces concern that the level difference is solely caused by entirely different people entering the catalogue. It does not control each person's opponents, questions, preparation, or which performances were

selected. A returning speaker could face a stronger opponent later, or appear in a different type of event. Their change cannot automatically be attributed to the reviewer.

The bridge also weights people differently from the overall comparison. The process mean gives each debate equal weight; the bridge summary gives each person equal weight. They are related descriptive checks, not interchangeable estimates or independent experiments.

4. Equal scales can conceal different score vocabularies

Earlier clarity marks are reconstructed from a small set of categorical findings. Later final dimensions can use a more varied set of numbers. The chart counts the actual distinct marks and shows their frequency rather than assuming both processes use the nominal 0-100 scale in the same way.




Figure 3. The panels compare exact Clarity and Task values. Tall spikes reveal frequently reused marks.

How to read this graph. Teal is earlier; rust is later. Each bar shows the share of moves assigned that exact area score. Both panels use identical axes. More distinct marks mean finer numerical resolution, not necessarily better judgments.

Finer numerical resolution is not automatically better resolution of argument quality. Forty or more distinct marks can imply unjustified precision just as four marks can hide meaningful differences. The issue is whether similar argument features receive similar scores and whether meaningful differences receive defensibly different scores.

The six dimensions also move together more closely in the later records. The leading shared statistical pattern accounts for about 73% of later dimension variation versus about 57% earlier. This is a descriptive property of the ratings, not an accuracy score. Correlation can arise from real argumentative quality, overlapping definitions, a general impression affecting several judgments, or all three.

5. Annotation practice is another comparability warning

Any fallacy or cognitive-bias tag appears on 2.54% of earlier moves and 15.21% of later moves. These are pooled move rates. Equal-debate rates are 2.48% and 15.53%, so longer assessments do not explain the whole recorded difference.




Figure 4. The much larger later tag rate persists when the denominator is all assessed moves in each process.

How to read this graph. A move counts once if it has any fallacy or bias tag. Bar length is the percentage of all moves in that process. Different annotation rates can reflect procedure as well as different arguments.

The analysis does not assume that more tags mean more mistakes, more accurate detection, or more biased reviewers. Those interpretations require reading shared passages under the same acceptance rules. The contrast is large enough that readers should avoid interpreting a low personal tag rate as an inherently superior reasoning record without considering which process assessed it.

6. What a useful reassessment should test

The next site-wide reassessment is tentatively planned for spring 2027, when stronger AI agents may be available. A newer model name should motivate testing, not substitute for it. A practical comparability exercise would use a fixed set of complete sources covering different topics, speaker positions, formats, score bands, and disputed cases.

The shared sources should be judged independently without access to the previous scores. The comparison should separate move selection, interpretation of the argument, dimension ratings, tag acceptance, and final arithmetic. Otherwise a changed total could be mistaken for a change in strictness when the models actually selected different parts of the debate.

Review disagreements against the source, not by automatically averaging them. Publish examples that explain what earns a high or low mark in each area. Include strong religious cases, weak skeptical cases,

reversals of PRO and CON, and creator-only versus fuller-theism burdens. Independent checking should be capable of finding that the favored procedure is worse on some tasks.

If a systematic offset remains after genuine like-for-like comparisons, a calibration method may be warranted. The present mixed archive does not justify one. Subtracting each process's observed mean is only a sensitivity exercise because it removes real case-mix differences along with possible measurement differences.

7. A qualified conclusion

Reason 1: score calculations are reproducible, but that is a different property from fair judgment. **Reason 2:** several observed features vary by assessment process. **Reason 3:** returning-speaker comparisons reinforce the warning without eliminating topic and opponent differences.

Therefore: the scores are suitable for careful, source-linked exploration but should not be pooled as perfectly interchangeable measurements. Small cross-process differences warrant less confidence than their numerical precision might suggest.

Therefore: the next review should measure comparability on shared material, preserve prior records, and explain substantive revisions. This edition changes the statistical reports, not anyone's underlying score or biography.

Methods and sources

This is a fixed October 9, 2026 research edition, based on source revision e65bc3c0b6. The score studies reconstruct published records rather than commission fresh assessments. The direct slogan counts retain the September 5 review of 187 transcripts; this edition recalculates those saved records and adds a separately labeled current score check. No new slogan reviews or debate assessments were performed.

The source manifest records file fingerprints. The analysis uses 20,000 deterministic resampling draws with baseline seed 20260904, retaining it for edition comparability. Whole debates are drawn for paired debate estimates; shared speakers for cross-role estimates. Ranking models and their dependence checks use the units explained in paper seven. These calculations do not capture every selection effect, repeated-speaker dependency, caption error or systematic judging bias.

The [data and methods page](#) links the source records, classification decisions, calculations, chart inputs and reading guides. Download the [current research package](#) and [historical direct slogan study](#). Earlier editions remain preserved in the repository's history. Debate links identify published assessments, not fresh independent evidence for this research.