

# Are All Slugfester Assessments on the Same Scale?

What shifts between assessment processes - and how the next major-model review should test comparability.

---

**2.82**

point lower mean score  
in the later process

**45 / 51**

repeated speakers whose  
mean score is lower later

**237**

comparable assessments with  
5,282 verified scored moves

The score levels, dimension patterns, and tag rates differ enough to make unqualified pooling unsafe. These are warning signs of non-comparability, not proof that a particular process caused the changes.

**PHIL STILWELL · SLUGFESTER**

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

## Executive summary

**The scores should not be treated as though every assessment came from one proven uniform measuring instrument.** The September 4 archive contains 237 comparable locked assessments, all labeled 5.6 Sol, but their records divide into two assessment-process families. The earlier 179 debates average **81.32 points per side**; the later 58 average **78.51**. The later mean is **2.82 points lower**, with a debate-resampling interval for the difference of approximately  $-3.67$  to  $-1.93$  points.

The change is not removed by looking at the same people. Among **51 speakers appearing in both process families**, 45 have lower mean scores later, five have higher means, and one is unchanged. Their equal-speaker mean change is  **$-3.23$  points**. This is strong evidence of a comparability problem worth investigating, but it is not a controlled measurement of an effect caused by the production process: these speakers face different opponents and propositions in different recordings.

Other features change too. Responsiveness is 4.43 dimension points lower later, precision 3.68 lower, and relevance/burden 3.37 lower, while confidence/fairness is almost unchanged. The six dimensions are more tightly related in later moves. Any fallacy-or-bias tag appears on **2.54% of earlier moves and 17.27% of later moves**, about a 6.8-fold relative rate. More frequent tagging is not, by itself, evidence of worse arguments or better judging.

The practical conclusion is not to discard the archive. **Within-debate comparisons are generally better protected than comparisons of absolute scores across processes**, because the two sides share the same assessment context. Even that protection is conditional: a common shift can cancel in a margin, but a change in scale or side-specific treatment need not. The religious-versus-skeptical mean gaps are reassuringly similar across the two families - 6.32 and 6.41 points - without proving universal comparability.

This Astra-era edition establishes a frozen baseline of **253 published assessments, 237 comparable locked records, and 5,282 verified scored moves**. It is new analysis, not an Astra re-judging of the debates. The intended next substantive review is the next major GPT-model transition. That review should test improvements in accuracy and fairness through shared source material and independent checks, rather than assuming that a newer model produces a more objective scale.

# Reader's guide

|                                                               |    |
|---------------------------------------------------------------|----|
| Executive summary                                             | 2  |
| 1. The same number can mean different things                  | 4  |
| 2. The data actually available                                | 4  |
| 3. The score-level shift                                      | 5  |
| 4. The same-speaker bridge preserves the warning              | 6  |
| 5. The dimension profile changes, not just the overall level  | 7  |
| 6. The dimensions move together more strongly later           | 8  |
| 7. Annotation drift is especially large                       | 9  |
| 8. Why within-debate margins are better protected             | 10 |
| 9. What the arithmetic audit establishes                      | 11 |
| 10. What to do with the existing archive                      | 11 |
| 11. The next major-model review should be an anchor study     | 12 |
| 12. A practical release standard                              | 12 |
| 13. Conclusion: trust requires a scale that can be questioned | 13 |
| Methods and sources                                           | 13 |

## Reading the evidence

Start with the summary and figures for the main argument. The worked examples explain the statistical ideas; the final methods and sources show how the numbers can be checked. Counts and results refer to the frozen September 4 archive, not to later additions to the site.

# 1. The same number can mean different things

Imagine two teachers both grading on a 100-point scale. One routinely awards 85 for work the other would grade 80. The numbers have the same range and may use the same category names, yet comparing a student from one class directly with a student from the other can be misleading. The issue is not necessarily dishonesty or incompetence. It is whether the instruments have been shown to behave comparably.

A scoring rubric helps by making the criteria explicit. It does not automatically ensure that different implementations, prompts, source preparations, or judgment procedures apply those criteria identically. Even a single model name can conceal different task framing or different ways of converting qualitative judgments into numbers.

This matters for SLUGFESTER because readers use the scores for several purposes: comparing the two sides of one debate, comparing different debates, reading trends, interpreting tag frequencies, and ranking speakers by their average performance. Those uses require different degrees of comparability. A useful audit should identify where the evidence is strongest rather than either granting blanket trust or declaring every score meaningless.

The present paper examines the recorded scale empirically. It asks whether score levels, dimension profiles, and annotation rates behave similarly across the two process families represented in the locked data. It does not presume that every observed difference is a defect in the process, and it does not use the differences as grounds for quietly changing the original assessments.

# 2. The data actually available

The full archive contains **253 assessments and 5,492 public moves**. Of these, **237 assessments have the comparable locked records used for the main numerical analysis**. They account for 5,282 moves and 474 overall side scores. Sixteen other-format records are retained in the full archive but are not treated as equivalent inputs to the detailed scale analysis.

The two process families are identified from their record structures, not guessed from whether their scores are high or low. Earlier records preserve detailed closed findings from which some dimensions are derived. Later standalone records store final judgments for all six dimensions. Both use the same nominal rubric label and the same published six-dimension weights.


| Record family                         | Debates | Scored moves | Published model label                                        |
|---------------------------------------|---------|--------------|--------------------------------------------------------------|
| Earlier detailed-findings process     | 179     | 3,423        | 5.6 Sol                                                      |
| Later standalone process              | 58      | 1,859        | 5.6 Sol                                                      |
| Other formats outside this comparison | 16      | 210          | Five labeled 5.6 Terra Extra High; eleven lack this metadata |

The last row should not be interpreted as zero-quality or evidence-free assessment. It indicates a lack of the same comparable record format in this analysis. Missing model metadata is missing information, not evidence that no model was used. Likewise, a locked record is an auditable fixed record, not proof that its judgments are correct.

Debate numbers describe the archive's sequence. They are not the years in which the debates occurred. A later-numbered assessment may concern an older recording. Consequently, a change across these groups should not be described as a historical decline in the quality of public debate without additional evidence.

### 3. The score-level shift

To compare levels while keeping the two sides together, we average PRO and CON within each debate. That produces one midpoint per debate. Averaging those midpoints is equivalent to averaging all side scores within a family, but resampling the debate midpoint correctly keeps the paired scores together in the uncertainty calculation.



**Figure 1.** Each dot is the midpoint of one debate's two overall scores. Diamonds mark family means. This is a focused score axis, not a bar chart from zero. The later distribution is centered lower, although the distributions overlap considerably. Scope: 237 debate midpoints; focused score scale.

The earlier mean is **81.32**, with a resampling interval of about 80.92 to 81.73. The later mean is **78.51**, with an interval of about 77.74 to 79.28. The later-minus-earlier difference is **-2.82 points**, and the resampling interval for that difference is -3.67 to -1.93.

The shift is large enough to matter for cross-debate comparisons. Paper seven finds that the median gap between adjacent displayed speaker ranks is only 0.17 points. A process-associated shift measured in whole points is therefore much larger than many of the differences that determine exact leaderboard places. That


does not mean every rank is wrong. It means a raw mean may mix performance with unequal exposure to assessment processes.

The source populations could also differ in true argumentative quality. Perhaps later debates contain harder motions or different speakers. The level shift alone cannot distinguish that possibility from a change in measurement. This is why the same-speaker bridge and the other diagnostic features are needed.

## 4. The same-speaker bridge preserves the warning

A *bridge* here means speakers observed under both process families. For each of the 51 such speakers, we calculate their mean later score minus their mean earlier score. This asks whether the overall change is merely the result of replacing a high-scoring set of people with a lower-scoring set.

### 51 speakers bridge the two process families



**Figure 2.** The diagonal marks equal earlier and later means. Forty-five of the 51 shared speakers fall below it. Each dot is a speaker average, not a reassessment of the same transcript under both processes.

Scope: One dot per speaker with appearances in both process families.

The equal-speaker average change is **−3.23 points**, with a resampling interval of approximately −4.09 to −2.37. Forty-five speakers are lower later, five higher, and one unchanged. The direction is not concentrated in just a few prolific participants.

This makes a simple explanation based only on different identities less convincing. It does not establish that the process itself caused the difference. A speaker's later assessed appearances may involve more difficult


claims, stronger opponents, less familiar topics, or different stages of their career. A change in source selection can affect the same people differently across groups.

The strongest bridge would use **the same debate transcript under both processes**, with comparable source material and independent judgments. The current bridge uses the same people in different debates. That is a meaningful improvement over comparing unrelated populations, but it is not the anchor needed to estimate a clean conversion between scales.

For this reason, subtracting or adding 3.23 to every later score would be unjustified. The average bridge difference is evidence of risk, not a calibrated correction formula. A blanket adjustment could remove genuine performance differences or introduce new errors into particular cases.

## 5. The dimension profile changes, not just the overall level

An equal downward shift in every dimension would be relatively simple to understand. The observed pattern is more complicated. The later process has lower mean scores in several dimensions, but the changes are not uniform.



**Figure 3.** Each bar is the later mean minus the earlier mean, with one equal-weight observation per debate after preserving the official within-debate weights. The near-zero confidence/fairness change contrasts with the larger shifts in responsiveness, clarity, and burden. Scope: Equal debate weighting; 179 earlier, 58 later.

Responsiveness falls by **4.43 dimension points**; precision and clarity by **3.68**; relevance and burden by **3.37**; evidence and warrant by **2.98**; and coherence by **1.46**. Confidence and fairness changes by only **-0.04**. This is not well described as every judge simply becoming three points stricter about everything.

Some of the difference may come from how the dimensions are recorded. Earlier precision and confidence scores are derived in part from categorical findings. Later scores are stored as final dimension values. A


dimension with a limited set of derived values can have a different distribution from one assigned directly, even when the descriptive label is unchanged.

The implication is important for a reader who compares a speaker's profile across debates. A lower later responsiveness score may reflect a weaker reply, a difference in how targets are identified, a different scoring practice, or a combination. The dimension name alone does not establish that a ten-point difference means exactly the same thing in both families.

This does not make the dimensions useless. They can still help explain the local assessment. It does mean that cross-family profile comparisons should retain the process context and should not be interpreted as measured changes in a speaker's ability without further checks.

## 6. The dimensions move together more strongly later

Correlation measures how two quantities tend to vary together. A value near +1 means that higher values of one tend to accompany higher values of the other; near zero means little linear relationship in the observed data; near -1 means an opposing relationship. Correlation does not show that one quantity causes the other.



**Figure 4.** Both panels use the same -1 to +1 color scale. Later dimensions are more strongly related across several pairs, particularly those involving relevance/burden. These are descriptive move-level correlations, not independent statistical tests on thousands of unrelated cases.

Scope: Unweighted move-level descriptive correlations; repeated observations.

In the earlier moves, coherence and evidence are strongly related, with a correlation of about 0.87. But coherence and relevance/burden are essentially unrelated, near -0.01. In later moves, coherence and evidence remain strongly related, around 0.82, while coherence and relevance/burden rise to about 0.65. Responsiveness and relevance/burden rise from roughly 0.16 to 0.76.

There are at least two broad interpretations. Later moves may genuinely have strengths and weaknesses that cluster more tightly together. Alternatively, the judging process may be allowing an overall impression to


influence several dimensions, or the assessment structure may make the dimensions less distinct. Both processes could operate at once.

A summary calculation finds that the strongest common statistical pattern accounts for **56.9% of standardized dimension variation earlier and 70.5% later**. This is the first-component share from a principal-component analysis - a way of finding the single weighted combination that captures the most variation among related measurements. It is not a percentage of scores that are correct, and it does not identify a psychological “halo effect” by itself.

The important practical lesson is simpler than the technical calculation: six numerical dimensions should not be mistaken for six independent witnesses. If they move together, their agreement provides less independent corroboration than a reader may assume. The scale can still be informative, but its apparent multidimensional precision needs to be interpreted cautiously.

## 7. Annotation drift is especially large

Any accepted fallacy-or-bias tag appears on **87 of 3,423 earlier moves**, or **2.54%**, and **321 of 1,859 later moves**, or **17.27%**. The later rate is about **6.8 times** the earlier rate. Named-fallacy tags alone show a similar contrast: approximately 1.90% versus 12.80%.



**Figure 5.** Pooled move rates differ sharply between the process families. A tag indicates an accepted annotation, not an independently verified census of all errors. More tags do not automatically mean worse debate performance or better assessment.

Scope: 5,282 locked moves; pooled rates.

The later records also contain more moves per debate on average: about 32.1 versus 19.1. That is why rates rather than raw counts are necessary. But the gap is not simply an artifact of the number of recorded moves. Averaging each debate's tagged share gives **2.48% earlier and 18.12% later**, preserving the contrast when long records lose their extra influence.

This is a direct warning against mixing tag frequencies as though all records used a stable detection system. A participant who appears mostly in later debates has more exposure to a process in which tags are more

common. A raw count of their tags would mix performance, number of appearances, number of recorded moves, and assessment-family differences.

Paper five shows the consequence for a headline: 80.8% of earlier-process losses lack a named-fallacy tag, compared with only 14.5% later. The combined archive still has a majority of untagged losses because the earlier family is larger. That pooled majority is a fact about the mixture, not a stable finding about all future assessments.

The data do not tell us whether earlier tags were too sparse, later tags too liberal, both appropriate to different material, or some mixture of those possibilities. Answering that question requires independent review of representative source passages under a shared tagging standard.

## 8. Why within-debate margins are better protected

If a scoring process adds or subtracts the same amount from both sides, the margin stays unchanged. For example, scores of 84 and 78 differ by six points. If both are reduced by three, 81 and 75 still differ by six. This is why comparing two sides within one assessment is less exposed to a common shift than comparing an 84 from one process with a 78 from another.

But not every change is additive. If a process compresses scores toward the middle, the gap can shrink. If it treats critical and constructive moves differently, the two sides may shift by different amounts. If it identifies different reply targets or records different moves, the local assessment itself changes. Sharing an assessment context reduces some risks; it does not remove all of them.

The observed margins provide a useful check. Mean absolute margins are **6.76 points earlier and 6.09 later**. That is a smaller contrast than the score-level shift. Within the broad religious-versus-skeptical set, the mean skeptical advantage is **6.32 earlier and 6.41 later**. The headline position gap therefore survives this major process split particularly well.

This stability is reassuring for that specific comparison. It is not a guarantee that every subgroup, speaker, or dimension is calibrated. The process families contain different debate samples. Similar group means can coexist with important changes in individual assessments, just as different group means can arise without a process effect.

**A useful hierarchy of caution:** reading the two sides of one debate is generally safer than comparing raw scores across debates; comparing many debates within one process is generally safer than pooling processes without qualification. “Safer” means less exposed to particular measurement differences - not known to be error-free.

## 9. What the arithmetic audit establishes

The analysis reconstructs all 5,282 comparable move scores from their six dimension values, using the same weights and rounding behavior as the site's calculator. It verifies the 2,558 side-section scores and all 474 overall side scores. It also checks that every comparable public move matches its underlying record and verifies the stored file hashes for all 58 later final records.

These checks matter because an incorrect join, missing move, or wrong weighting could create a spurious difference before interpretation even begins. They show that the statistics analyze the published scoring records rather than an accidentally altered version of them.

They do **not** establish that the source transcript was complete, that a statement was interpreted fairly, that the selected reply target was correct, or that the dimension judgment was well calibrated. A perfectly reproducible computation can preserve a mistaken judgment. Numerical consistency is necessary for trustworthy reporting, but insufficient for trustworthy measurement.

This distinction should shape the site's public language. A locked score is transparent and revisable evidence about an assessed performance. It is not an objective fact about the strength of an argument in the same sense as a verified arithmetic total. Readers should be able to challenge the source interpretation without being told that reproducibility settles the matter.

## 10. What to do with the existing archive

The present evidence supports **stratification before correction**. In everyday terms, keep the process families visible and compare like with like where possible. Report their separate means, tag rates, and sample sizes. Show both the raw result and the process-specific result when a conclusion depends on pooling.

Do not retroactively subtract a universal offset from every earlier score or add one to every later score. The dimension shifts are not uniform, and the same-transcript bridge needed to estimate a reliable conversion does not yet exist. A cosmetic correction could create a false sense of precision while obscuring the original record.

For speaker rankings, display appearance counts and uncertainty, and make process exposure inspectable. Paper seven finds that subtracting process means changes the broad rank ordering only modestly, but individual places can move. That check is a sensitivity analysis, not a validated replacement leaderboard.

For tag comparisons, report rates rather than totals and keep the process denominator near the claim. Even a within-process rate remains a record of accepted annotations, not a definitive error prevalence. A speaker's count should not become a personal quality rating detached from the actual arguments.

For source corrections, prioritize errors that could alter a central inference: an omitted qualification, a misunderstood burden, a misassigned reply target, or an unfaithful transcript segment. These are more important than preserving a previously published average for the sake of continuity.

## 11. The next major-model review should be an anchor study

The earlier research plan discussed reassessing perhaps twice a year as AI improved. This edition marks a different practical milestone: a frozen Astra-era series intended to remain the baseline until the next major GPT-model transition. That is a publication policy, not a claim that the current scores have become final or that the next model will necessarily be more objective.

The next review should begin with an **anchor set**: the same source debates assessed under the old and proposed new procedures. It should cover different topics, formal orientations, narrow and ambitious burdens, close and large margins, and both well-known and less frequent speakers. Selection should be specified before new scores are inspected.

The old and new runs should use the same source material and avoid seeing each other's totals, winner labels, or rationales. Where practical, speaker identity should be hidden without damaging the meaning of the transcript. Human review should focus on source fidelity, burden assignment, quality of support, and disagreement cases rather than treating either model as the answer key.

A sound comparison would report three things separately. First, does the new process better capture the source arguments according to independent review? Second, is it more repeatable when the same material is judged again? Third, do its scores remain comparable to the old scale, or do levels, spreads, and dimensions shift? A process can improve on one of these questions without improving on all three.

The study should also include planned negative controls. Harmless relabeling of equivalent positions should not systematically change the winner. Removing a genuine supporting premise should not improve the argument's evidence rating. Correcting a misquotation should have an interpretable effect. Such checks make “improved judging” a testable claim rather than a slogan attached to a new model launch.

## 12. A practical release standard

Before a new generation replaces or supplements the old one, SLUGFESTER should publish the anchor results, disagreements, and known limitations. If the scale shifts, the site should retain generation labels and explain whether old scores were recalibrated, freshly reassessed, or simply left as historical records. Those operations are not interchangeable.

A newly written paper should not imply a newly judged corpus. A refreshed graph should not imply independent confirmation. An added debate should not be allowed to silently change the denominator of a previously frozen report. Versioned source data, dates, and a correction history let readers tell which kind of change occurred.

The existing research files support that practice by retaining a source revision, file hashes, normalized data, calculation scripts, and chart source contracts. The seven PDFs all refer to the same 253-assessment snapshot. Later additions to the public site can coexist with these reports as long as the snapshot date and coverage remain visible.

## 13. Conclusion: trust requires a scale that can be questioned

The audit finds a coherent warning pattern: lower later score levels, a similar change among repeated speakers, different dimension profiles, stronger coupling between dimensions, and substantially more frequent tags. It is not reasonable to ignore those differences merely because the rubric and model labels match.

It is equally unreasonable to treat them as proof that the archive has no value. The arithmetic is traceable, individual judgments can be inspected, and some central within-debate findings remain stable across the process families. The right response is to distinguish uses that are relatively well supported from comparisons that require calibration or stronger qualification.

**The archive is most useful when its measuring instrument is part of the evidence, not an invisible authority behind it.** Readers should know what a score means, which procedure produced it, how it compares with another score, and what could require revision. The next major model should face those questions rather than inherit an assumption of superiority.

The best outcome of a model transition is not a new set of numbers presented with greater confidence. It is a better-tested relationship between the numbers, the source arguments, and the judgments readers can independently defend.

## Methods and sources

This audit uses source revision **76d006b37**, frozen September 4, 2026. The earlier/later split follows record structure: detailed closed findings versus later standalone final dimensions. Both families are labeled 5.6 Sol and use the same nominal rubric weights. Sixteen other-format assessments are excluded from detailed scale comparisons; their public counts and metadata limitations are reported separately.

Score-level uncertainty resamples debate midpoints, preserving paired sides. The level-difference interval uses independent resamples within the two process families. The repeated-speaker bridge gives equal weight to each of 51 speakers and resamples speakers. All reported resampling uses 20,000 draws and fixed documented seeds. These intervals do not remove source-selection bias, repeated-opponent dependence, unmeasured task difficulty, or systematic judging error.

Dimension levels first preserve the official move-importance and section weights, then average equally across debates. Correlation matrices use unweighted moves within each family and are descriptive. The first-component share is the largest eigenvalue of the six-dimension correlation matrix divided by six. Pooled tag rates use tagged moves divided by all moves; equal-debate tag rates are supplied as a check against record-length differences.

The [analysis directory](#) includes the full 51-speaker bridge, score and dimension summaries, source hashes, validation counts, and executable calculations. The [Backend library](#) links all seven papers. No automatic reassessment schedule is created by this report, and no source assessment is changed by its calculations.