

Are All Slugfester Assessments on the Same Scale?

Why earlier and later marks need careful comparison, and a practical plan for the next full assessment.

2.82

points lower on average
in later assessments

45 / 51

people whose average score
is lower in later records

4 vs 41

different clarity marks used
in earlier vs later moves

Later scores are lower on average, and the six scoring areas move together more closely. Shared test debates and clearer scoring examples can help the next review separate genuine differences from changes in judging.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

A score of 80 may not mean quite the same thing in every Slugfester assessment. The earlier 179 comparable assessments average 81.32 points per side. The later 58 average 78.51. The later average is **2.82 points lower**. This is not just a change in who appears: of 51 speakers found in both groups, 45 have lower averages later, five have higher averages, and one is unchanged.

Other differences matter too. Scores for answering the other side fall by 4.43 points, while scores for care and fairness barely change. Several scoring areas rise and fall together more closely in the later records. Fallacy-or-bias labels appear on 2.54% of earlier moves and 17.27% of later moves. We should not mix the two groups without explaining these differences.

The cause is not yet known. The debates, opponents, claims, and assessment procedures differ. This paper adds one concrete clue: earlier clarity scores use only four different values, while later scores use 41. A more flexible score scale, a changed reading of the rules, one overall impression affecting several marks, or real differences in the arguments could each play a part. These are possible explanations, not findings about what the model was thinking.

The next complete reassessment should score old and new debates under one tested procedure. It should use common benchmark debates, mix earlier and later records throughout the work, repeat some judgments, and check each scoring area separately. It should **not** start by adding 2.82 points to later scores or by trying to make the two group averages equal. Equal averages would not prove fair judging.

The data remain the September 4 snapshot: 253 published assessments, including 237 comparable one-on-one records and 5,282 scored moves. These are existing 5.6 Sol judgments. Astra rewrote and analyzed the papers; it did not judge all the debates again.

Reader's guide

Executive summary	2
1. Two teachers can use the same scale differently	4
2. What we compared	4
3. Later scores are lower on average	5
4. The change also appears among the same people	6
5. The six scoring areas do not change equally	7
6. What it means for scores to move together	8
7. Why might the later marks move together more closely?	9
8. Fallacy and bias labels change even more	11
9. Comparing two sides in one debate is better protected	12
10. Correct arithmetic does not prove correct judgment	13
11. What to do with the current scores	13
12. How to improve the next complete reassessment	13
13. What should be published with the new results?	15
14. Conclusion: three arguments for a fairer next assessment	16
Methods and sources	17

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. All results refer to the September 4 snapshot, not later additions to the site.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. Two teachers can use the same scale differently

Imagine two teachers marking work out of 100. One gives 85 to work the other would give 80. Both can use the same category names and both can add the marks correctly. Yet a student from one class cannot safely be compared with a student from the other by looking only at the final number.

Debate scores can have the same problem. A written scoring rule helps, but instructions, transcript preparation, and the way an argument is divided into parts can change how the rule is used. Even the same AI model can receive a different task.

This matters in three common situations. A reader may compare the two sides of one debate, compare scores from different debates, or rank speakers by their averages. The first comparison keeps more conditions the same. The other two need stronger evidence that the scores mean the same thing across records.

We therefore ask a practical question: **what must be checked before these scores can be compared fairly?** We are not trying to prove that the archive is worthless, nor to protect every score already published.

2. What we compared

The full archive contains 253 assessments and 5,492 public moves. A move is a claim, reason, objection, or reply selected for assessment. It is not every sentence in the recording. The detailed comparison uses 237 one-on-one assessments whose saved records can be checked in the same way.

Assessment group	Debates	Scored moves	Published model name
Earlier records	179	3,423	5.6 Sol
Later records	58	1,859	5.6 Sol
Other formats, kept outside this comparison	16	210	Five list 5.6 Terra Extra High; eleven do not list a model

The earlier and later groups are identified by how their records are built, not by whether the scores are high or low. In the earlier group, some marks are calculated from fixed answer categories. In the later group, final marks for all six areas are stored directly. Both groups use the same published weights for the six areas.

The other 16 assessments are not being called poor or unsupported. They simply lack the same record format used in this comparison. Likewise, a record with a missing model name does not show that no model was used.

“Earlier” and “later” refer to the assessment process, not the filming date. A newly assessed debate may have been recorded years ago. These findings do not show that public debate has become worse over time.

3. Later scores are lower on average

The 2.82-point difference is calculated before rounding the group averages. Subtracting the printed values, 81.32 and 78.51, gives 2.81 instead. That one-hundredth difference is a rounding effect, not a different finding.

For each debate, we first average the two sides' overall scores. An 84 and a 78 become one debate value of 81. We then average those debate values within each group. This keeps a debate's two sides together when checking how the result changes with the mix of records.




Figure 1. The groups overlap, but the later scores sit lower on average. A dot is one debate, not one speaker or one argument. The plot enlarges the part of the score scale where the observations lie; it does not start at zero.

How to read this graph. Each small dot is one debate's average of its two side scores. Teal is earlier; rust is later. Dark diamonds mark the group averages. Small vertical offsets only separate overlapping dots; they carry no score information. The horizontal axis enlarges the observed part of the 0-100 score scale.

The earlier average is 81.32; the later average is 78.51. The difference is 2.82 points after rounding from the full numbers. Repeatedly drawing debates from each group gives a middle-95% range of **1.93 to 3.67 points lower later**. Drawing allows repeats, and the calculation is run 20,000 times.

That range asks how the result changes when the existing debates receive different amounts of influence. It does not include every possible transcript mistake, unfair judgment, or gap in the selection of debates. The average of the fixed archive is already known; the range is a check on how much it depends on the mix of records.

A difference of a few points matters when speaker averages are close. Paper seven finds a typical gap of only 0.17 points between neighboring ranks. This does not prove those ranks are wrong. It means a reader needs to know whether a speaker's average mostly comes from earlier or later assessments.

Perhaps the later debates really contain weaker arguments. Perhaps their claims are harder to defend. Perhaps the scoring procedure changed. The two group averages alone cannot tell us which explanation is right.

4. The change also appears among the same people

Fifty-one speakers appear in both assessment groups. For each, we compare their average later score with their average earlier score. We then give each speaker equal influence in the result.

51 speakers bridge the two assessment groups




Figure 2. Forty-five dots fall below the equal-score line: those speakers have lower averages in the later records. This compares the same people in different debates, not the same debate judged twice.

How to read this graph. One teal dot represents one of the 51 people assessed in both groups. The horizontal position is their earlier average; the vertical position is their later average. The dashed line means equal averages. Below it means lower later. Both axes enlarge scores 62-94. These are different debates by the same people, not repeated judgments of one debate.

The average change is **3.23 points lower later**. Repeated draws of these speakers give a range of roughly 2.37 to 4.09 points lower. The direction is widespread: 45 lower, five higher, and one unchanged.

This weakens the explanation that the archive simply replaced high-scoring people with low-scoring people. But it does not settle the cause. The same person may face a stronger opponent, argue a harder claim, or discuss a less familiar subject in the later group.

The stronger test would judge **the same transcript under both procedures**. That would hold the speaker, opponent, claim, and recorded performance still while changing the assessment method. Our present comparison does not do that.

We therefore should not add 3.23 points to every later score. A difference between group averages is not a proven correction for any particular debate.

5. The six scoring areas do not change equally

The records score six things. This paper uses short labels in the graphs: **Logic** means consistent reasoning; **Support** means evidence and reasons for accepting the claim; **Replies** means answering the other side; **Task** means staying on the question and doing what one's claim requires; **Clarity** means clear, exact wording; **Care** means matching confidence to support and describing the other side fairly.

The change is not an equal shift in every scoring area




Figure 3. Every bar compares the later average with the earlier average in one scoring area. The leftward bars mean lower later marks. Care and fairness barely change, unlike replies, clarity, and the task the speaker needs to complete.

How to read this graph. Each rust bar is the later average minus the earlier average, giving every debate equal influence. Left of zero means lower later scores. The six areas use their own 0-100 scales here; the bars are not weighted parts of the overall 2.82-point change. There are 179 earlier and 58 later debates.

Area key: Logic = consistent reasoning; Support = evidence and reasons; Replies = answering the other side; Task = staying on the question and doing what the claim requires; Clarity = clear, exact wording; Care = fitting confidence to support and treating the other side fairly.

Later reply scores are 4.43 points lower. Clarity is 3.68 lower; staying on task is 3.37 lower; support is 2.98 lower; and logic is 1.46 lower. Care and fairness differ by only 0.04 points. Each area is scored out of 100 before its weight is applied to the overall result.

This is not simply one flat three-point reduction. It suggests that we need to examine the areas separately. A rule that changes how replies are judged need not change how wording is judged.

The numbers describe the marks, not a confirmed change in speakers' abilities. A lower reply score might mean a weaker reply. It might also mean the assessment identified a different objection as the one the speaker needed to answer. Those possibilities require inspection of the source passage.

6. What it means for scores to move together

Suppose arguments that score highly for logic also tend to score highly for support. Arguments low on one tend to be low on the other. We say the two marks **move together**. A correlation number summarizes how closely this happens.

A value near **+1** means the marks line up closely as they rise and fall. A value near **0** means there is little straight-line pattern. A value near **-1** means one tends to rise as the other falls. Zero does not rule out every possible relationship. And none of these values, by itself, tells us what caused the pattern.




Figure 4. Read one row label and one column label to identify the two areas being compared. Dark blue means they rise and fall closely together. The same color scale is used for both assessment groups. The diagonal compares each area with itself, so its 1.0 values are automatic, not evidence of good judging.

How to read this graph. Match a row to a column to compare two scoring areas. Blue near +1 means they rise and fall together; red near -1 means opposite movement; pale near zero means little straight-line relationship. Both panels use one color scale. Diagonal 1.0 values compare an area with itself and are automatic. These are score relationships, not accuracy percentages.

Area key: Logic = consistent reasoning; Support = evidence and reasons; Replies = answering the other side; Task = staying on the question and doing what the claim requires; Clarity = clear, exact wording; Care = fitting confidence to support and treating the other side fairly.

Logic and support are closely related in both groups: about 0.87 earlier and 0.82 later. The bigger changes involve the task score. Its link with logic rises from about -0.01 to 0.65. Its link with replies rises from about 0.16 to 0.76.

In plain language, later arguments given a high task score are much more likely also to receive high logic and reply scores. Earlier task scores reveal far less about those other marks. This is a description of the recorded scores, not proof that one mark influenced another.

An additional calculation looks for the single combined pattern that captures the largest share of differences across all six areas, after putting their spreads on the same scale. Its share rises from **56.9% earlier to 70.5% later**. The calculation is called principal-component analysis. Readers need not learn its name to grasp the warning: later marks behave more like several views of one overall rating. The percentages are not accuracy rates.

Six marks are not necessarily six separate confirmations. One unsupported claim can rightly receive lower marks for logic, support, and confidence. Or a judge's poor overall impression can spill into all three marks. The scores alone do not tell us which happened.

7. Why might the later marks move together more closely?

The following are possible explanations. More than one could be true. Each needs a test that uses the same source arguments; none should be treated as a finding about the model's private reasoning.

Some earlier scores were placed into broad categories

Earlier clarity marks take only **four different values** across 3,423 moves. Later clarity marks take **41 different values** across 1,859 moves. Also, **2,179 earlier moves, or 63.7%, receive exactly 87** for staying on task. These counts come directly from the saved marks.




Figure 5. Earlier clarity and task scores bunch at a few values. Later scores are spread across more possible marks. A flexible scale can show relationships that a coarse scale hides, but it can also create an appearance of precision that the judgments do not deserve.

How to read this graph. Teal bars show earlier moves; rust bars show later moves. Each panel counts the share of moves receiving each exact score, on the full 0-100 scale. The percentage scale is identical in both panels. A tall isolated bar means many moves received that same mark; more distinct marks do not automatically mean more accurate judging.

If many quite different arguments receive the same task score, that mark cannot track all their differences in logic or support. A later procedure with more possible marks may reveal links the earlier scale hid. This is one reason to test score definitions before blaming every increase in correlation on careless judging.

But finer marks are not automatically better. Giving 83 rather than 84 needs a sound reason. More possible numbers can add detail, or merely add unsupported precision. The common-transcript test should compare both the numerical marks and the reasons given for them.

One overall impression may influence several marks

A judge might first decide that an argument is strong and then give it high marks in most areas. A weak overall impression might pull them all down. This would make the marks move together even when the source contains a mixture of strengths and weaknesses.

For example, an unsupported claim can be expressed very clearly. A well-supported claim can be worded badly. The scoring process should be able to keep those cases apart. A useful test would judge each area without showing the other area's marks, then compare the results with a run that scores everything together.

The rules may be read as asking the same question

“Did the speaker back up the claim?”, “Did the speaker meet the task?”, and “Was the speaker too confident?” partly overlap. If the instructions do not explain the differences, a reviewer may give the same weakness nearly the same treatment in each area.

Some overlap is proper. The problem is not correlation itself. The problem would be several marks that add no distinct information. Test passages should include clear wording with weak support, strong support with a missed reply, and a modest claim that succeeds without proving a larger worldview.

The selected moves may be different kinds of material

The later records contain about 32.1 assessed moves per debate, compared with 19.1 earlier. This does not establish that later transcripts are longer or better covered. It shows that the recorded unit of analysis may differ.

Splitting a broad passage into shorter claims can change what counts as a complete answer or an unsupported claim. Topic and speaker selection can also make strengths cluster. To test this explanation, reviewers should first use the same passages and move boundaries, then separately test the effect of changing those boundaries.

Review and averaging may make marks more alike

When reviewers resolve disagreements, they may rightly correct isolated mistakes. They might also, without intending to, make each argument's marks look more internally consistent. Averaging two judgments can change the pattern too.

The next assessment should therefore keep first judgments and final judgments. Compare how the six marks relate before and after disagreement review. If the links tighten, inspect the reasons: were inconsistent judgments repaired, or were genuine mixed strengths smoothed away? This proposal is not a claim that either happened in the current archive.

8. Fallacy and bias labels change even more

At least one fallacy-or-bias label appears on **87 of 3,423 earlier moves: 2.54%**. The later count is **321 of 1,859: 17.27%**, about 6.8 times the earlier rate. Named fallacies alone show a similar contrast: about 1.90% earlier and 12.80% later.




Figure 6. The bars show percentages, not raw totals, so the groups' different numbers of recorded moves do not determine the comparison. More labels do not, by themselves, show either worse arguments or better reviewing.

How to read this graph. Teal shows earlier moves and rust later moves. Each fraction is tagged moves divided by all moves in that group. A move counts once if it has any fallacy or bias tag, even if it has several. Bar length shows percentage, allowing the groups' different totals to be compared.

The difference remains when each debate gets equal influence: the average labeled share is 2.48% earlier and 18.12% later. It is therefore not just a result of longer assessment records getting counted more often.

A speaker assessed mainly in the later group faces a setting where labels are much more common. Counting their labels without showing the number of appearances, moves, and assessment group would be misleading.

We do not yet know whether earlier reviewers missed errors, later reviewers applied labels too freely, the arguments genuinely differ, or several of these explanations apply. A common set of passages needs independent review under one clear rule for each label.

9. Comparing two sides in one debate is better protected

Scores of 84 and 78 differ by six points. Subtract three from both and 81 and 75 still differ by six. Thus a shared upward or downward shift can change score levels without changing the gap between the sides.

Not every change works that way. If a procedure squeezes scores toward the middle, the gap can shrink. If it treats a positive case and a reply differently, the sides can move by different amounts. Sharing a debate helps, but it does not remove every risk.

The average size of the gap, ignoring which side leads, is 6.76 points earlier and 6.09 later. In the religious-versus-skeptical comparisons, the average non-theist advantage is **6.32 earlier and 6.41 later**. That particular finding survives the process split well.

We should draw the limited conclusion those numbers support: the overall position gap is quite stable across these two groups. We should not conclude that every topic, speaker, or scoring area has been shown to be directly comparable.

10. Correct arithmetic does not prove correct judgment

The audit rebuilt all 5,282 comparable move scores from their six marks and checked the weights and rounding. It also verified 2,558 side-section scores and all 474 overall side scores. Every comparable public move matched its saved record. All 58 later final files matched their recorded digital fingerprints.

These checks rule out several simple explanations for the reported differences: missing rows, wrongly matched records, or a different calculation from the one used on the site. They show that we analyzed the published numbers correctly.

They do not prove that a transcript was complete, a claim was understood fairly, or a reply was matched to the right objection. A calculator can add mistaken judgments without making an arithmetic error. Readers must still be able to challenge the assessment by pointing to the source.

11. What to do with the current scores

Keep the earlier and later groups visible. Show their separate averages and label rates when a comparison uses both. Let readers inspect which group supplied a speaker's scores.

Do not apply one blanket increase or decrease. The six areas did not move equally, and we lack the same-transcript test needed to estimate a reliable conversion. Adding 2.82 points would make the averages match by design, not establish that the arguments were judged fairly.

Keep the original scores and record any real correction with a reason. A missing qualification, wrongly assigned objection, or incorrect transcript passage deserves repair. An inconvenient group average is not, by itself, an error to repair.

12. How to improve the next complete reassessment

The aim should be **one tested procedure for the full eligible archive**, not a number adjustment that makes old and new averages look alike. The following is a proposed plan, not a reassessment already carried out or a paid job now running.

First, choose common benchmark debates

Before seeing new scores, choose a manageable set of the same debates to judge under the old and proposed procedures. For example, 30-40 benchmarks could cover several topics, both role assignments, close and large gaps, and frequent and less frequent speakers. Include records from both assessment groups.

Use part of the set to improve the instructions. Keep the rest untouched until the procedure is fixed. Otherwise, repeated adjustments can make the method look good on examples used to design it while failing on fresh material.

Second, hold the source and the task still

Give both procedures the same checked transcript and enough surrounding text to interpret each passage. Keep speaker turns, move boundaries, reply targets, and the claim each side must defend the same for the first comparison.

Then test source preparation separately. If a new extraction method identifies more claims, compare those additions with the recording. Do not change the transcript, move selection, judging instructions, and scoring calculation all at once and then attribute the result to the new model alone.

Third, make each scoring area earn its own mark

Write short examples that separate the six areas. An argument can be clear but unsupported; supported but off-topic; relevant but fail to answer the objection; or cautious yet mistaken. Each area's rule should say which features matter and which do not.

Require a brief source-based reason for each mark. Test scoring one area at a time without showing the other marks. Compare that with scoring all six together. Prefer the approach that independent review finds more faithful to the arguments, not simply the one with lower correlations.

Use the same rule for translating judgments into numbers throughout the new run. In particular, decide whether clarity uses a few defined bands or a finer scale. Test whether reviewers can explain and repeat the distinctions before treating every single point as meaningful.

Fourth, keep earlier and later debates mixed during the run

Do not assess all older records first and all newer records last. Mix topics, positions, speakers, and original assessment groups across batches. Hide old scores and winners from fresh judges. Hide identities where that does not remove information needed to understand the argument.

This reduces the risk that a changed instruction, a different session, or an emerging habit affects one group more than the other. Keep a record of the instructions and model settings for every batch. If they change materially, stop and test the change rather than silently continuing.

Fifth, repeat a small set and check for drift

Judge a small unseen set at least three times in fresh runs. Put some unchanged benchmark passages into different batches. Check whether marks, reply targets, and labels remain reasonably stable when the source is unchanged.

Set acceptable differences after the pilot but before the full run. A practical warning threshold should reflect how much variation the tested procedure normally produces. Do not choose a threshold afterward merely because the finished run passes it.

Sixth, test whether the system notices meaningful changes

Create clearly marked test versions of selected passages, separate from the real archive. Remove a crucial reason; correct a misquotation; add a missing limit; or change a reply so it answers the actual objection. Check whether the relevant mark changes for an understandable reason.

Also make harmless changes, such as swapping PRO and CON labels while keeping the actual claims clear. These should not systematically move the result. The purpose is to test both kinds of error: missing a real improvement and reacting to an irrelevant change.

Seventh, review disagreements without treating agreement as truth

Use independent first judgments and preserve them. Review disagreements against the source before calculating final totals. Ask a separate reviewer to check a spread of cases, including some where the judges agree. Two matching mistakes are still mistakes.

Compare marks and their relationships before and after review. Check whether the final process preserves cases with mixed strengths. For fallacy labels, inspect both labeled passages and passages left unlabeled; otherwise the audit can find false labels while overlooking missed ones.

Finally, reassess the full archive before replacing its scale

If the pilot and untouched tests succeed, apply the same procedure to all eligible old and new debates. Publish the new scores as a dated assessment generation, with links to the old records and explanations of material changes.

If some formats cannot yet be assessed in the same way, keep them clearly separate. If only part of the archive is rerun, do not present the mixed old and new scores as one freshly tested scale. Report incomplete coverage openly.

The target is not equal averages for theist and non-theist sides, PRO and CON, or older and newer recordings. Real performance differences may remain. The target is that the **same quality of reasoning receives comparable treatment**, regardless of when the record entered the site.

13. What should be published with the new results?

Publish four kinds of checks: whether the source was captured correctly, whether reviewers judged it fairly, whether repeated runs agree, and whether old and new scores can be compared. Better performance on one does not guarantee better performance on the others.

Show average changes and the spread of changes, broken down by scoring area and original assessment group. A near-zero overall change can hide large changes in opposite directions. A stable winner count can hide changes in the reasons for the scores.

Keep dates and coverage clear. Slugfester may rerun assessments roughly twice a year as models and quality controls permit. These seven papers have a separate plan: preserve the September 4 baseline until the next major GPT review, apart from genuine corrections. Rewriting a paper does not mean its debates were freshly judged.

14. Conclusion: three arguments for a fairer next assessment

Argument one: the present scores need a warning about comparison

1. Scores can be compared directly only when there is good reason to think the same marks mean the same things.
2. Later averages are 2.82 points lower. Forty-five of 51 shared speakers score lower later. The six areas change by different amounts, and label rates differ sharply.
3. A common model name and correct arithmetic do not explain those differences or rule out a changed standard.

Therefore: readers should keep the two assessment groups visible and avoid treating every cross-debate point difference as a proved difference in performance.

Argument two: one fixed score adjustment is not justified

1. A fair correction must separate a change in judging from a real change in the arguments being judged.
2. The current groups contain different debates, and no common-transcript experiment isolates the effect of the procedure.
3. The changes are not a flat shift: reply scores fall by 4.43 points while care and fairness barely change.

Therefore: adding 2.82 or 3.23 points to later scores would be an untested adjustment. The next review should judge the same sources under controlled conditions and then reassess the full eligible archive consistently.

Argument three: closer agreement between marks needs explanation

1. Different scoring areas are useful only if each contributes a defensible judgment about the argument.
2. Later marks move together more closely. Earlier marks also show strong bunching, including only four clarity values and a task score of 87 for 63.7% of moves.
3. That pattern could reflect a better scale, overlapping rules, a shared overall impression, different arguments, or several causes together.

Therefore: the next assessment should test those explanations with common passages, separate-area judgments, repeat runs, and review of the reasons. It should not force the marks apart or call their agreement proof of accuracy.

The practical endpoint is clear: test the procedure, apply it consistently, and keep the evidence open to challenge. A new model should earn trust by handling the same arguments better, not by producing newer numbers.

Methods and sources

The fixed source revision is **76d006b37**, September 4, 2026. The main comparison contains 179 earlier and 58 later assessments. The 16 other-format records remain outside it. Both comparable groups list 5.6 Sol and the same overall scoring weights.

Score averages keep the two sides of each debate together. Ranges use 20,000 repeated draws of debates; the 51-speaker comparison instead draws speakers. These checks do not remove selective debate collection, repeated opponents, source errors, or unfair judging. Area averages keep official within-debate weights. Correlations use individual moves, without those weights. The combined-pattern share is the largest value from the correlation-matrix decomposition divided by six.

The added score-use check counts distinct marks and exact frequencies in the saved 5,282 moves. It does not infer how fine a scale the model could have used. The [analysis directory](#) contains the source records, full calculations, new score-use counts, chart data, and validation notes. The [Backend library](#) links all seven papers. The proposed remedies are not tested interventions in the current dataset, and no public assessment is changed by this paper.