

Are Theist Arguments More Often Slogan-Like?

Testing a measurable pattern of weak warrant and overclaim, with close readings of Lennox and contrasting cases.

28.6%

theist moves meeting the original three-part risk rule

7.9%

non-theist moves meeting that same rule

146

debates with the required fields; 187 in a separate score-based check

The measured risk pattern is markedly more common on the theist side. It is not a literal census of unfalsifiable slogans, and it does not establish that emotional reinforcement caused the difference.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

The archive supports a narrower and more defensible version of the slogan hypothesis: **the religious side more often combines weak public support, overstated confidence, and unstable or missing qualification**. In the 146 debates that retain the necessary detailed findings, 398 of 1,393 theist moves meet all three conditions, compared with 111 of 1,407 non-theist moves. That is **28.6% versus 7.9%**, a 3.62-fold relative rate.

When each debate receives equal weight, the mean difference is **20.65 percentage points**, with a 95% debate-resampling interval of 16.28 to 25.01. The pattern is not produced by John Lennox alone: removing his eligible debates leaves a 21.05-point difference. His contributions nevertheless provide unusually clear examples for close reading, including affirmations of creator-given worth, postmortem compensation, and miracle claims protected by a diagnosis of the critic's prior assumptions.

The crucial limitation belongs beside the headline. **The measured three-part pattern is not a direct count of literally unfalsifiable slogans**. Some flagged moves are poorly supported arguments that could be improved or tested. Some unflagged statements may still be insulated from correction. The findings measure an argumentative risk pattern, not its true-positive rate as a detector of non-falsifiability.

The archive now contains 253 assessments, but the original test cannot honestly be expanded to all of them: newer scoring records do not retain the same finding fields. A separate, explicitly different score-based check covers all 187 comparable religious-versus-skeptical debates. It finds a 23.62-percentage-point difference in the same direction. The later-process subset alone shows a 21.39-point difference. That strengthens the broader concern about weak warrant and overclaim without pretending the two tests are interchangeable.

Phil Stilwell's hypothesis also proposes a cause: that religion is often sustained through emotional affirmation rather than rigorous public argument. The observed pattern is consistent with that explanation. **It does not establish how religious communities enforce belief, why individuals believe, or what emotional effect a particular statement had on its audience**. Those are additional questions requiring different evidence.

The paper's conclusion is therefore both strong and bounded: the archive contains a substantial asymmetry in a measurable pattern associated with slogan-like reasoning, and concrete cases show how claims can be protected from correction. The appropriate response is to demand clear support and clear vulnerability to revision from both religious and secular affirmations.

Reader's guide

Executive summary	2
1. Why slogans are an argumentative problem	4
2. Four concepts that should not be collapsed	4
3. The original three-part test	5
4. The measured asymmetry is substantial	5
5. Does the pattern survive narrower checks?	6
6. What the new assessments can - and cannot - add	7
7. Lennox example one: infinite worth through creation	8
8. Lennox example two: compensation after death	9
9. Lennox example three: miracles and the critic's assumptions	9
10. Other examples: private confirmation and protected motives	10
11. The same standard catches secular overclaim	11
12. Not every theist argument fits the pattern	11
13. The causal hypothesis remains open	12
14. Conclusion: confidence must be able to lose	12
Methods and sources	13

Reading the evidence

Start with the summary and figures for the main argument. The worked examples explain the statistical ideas; the final methods and sources show how the numbers can be checked. Counts and results refer to the frozen September 4 archive, not to later additions to the site.

1. Why slogans are an argumentative problem

A slogan is not weak merely because it is brief. “Correlation is not causation” is concise and useful, although applying it correctly requires explanation. Nor is a claim weak merely because it moves an audience. A compelling account of suffering can be emotionally powerful and evidentially serious. The problem begins when the persuasive force of a statement substitutes for the work needed to establish it.

An affirmation can make a conclusion feel settled without showing why an unconvinced listener should accept it. It may unite a group, reassure a person, express a commitment, or frame an experience. Those functions can be legitimate. They become problematic when the statement is offered as public evidence for the very worldview that gives the affirmation its authority.

The hypothesis examined here predicts that this substitution is more common on the theist side. One proposed explanation is that standards acceptable for faith maintenance - trust, personal revelation, inherited authority, existential consolation - can enter debate as if they already supplied shared reasons. The question is not whether religious speech contains emotion. It is whether the move from conviction to justification has been completed.

The same issue arises outside religion. “Science will eventually explain everything,” “markets always know best,” or “everyone who disagrees is protecting privilege” can be used as correction-resistant affirmations. These illustrative formulations are not quotations from the corpus. They show why the test must target an argumentative pattern rather than a favored set of beliefs.

2. Four concepts that should not be collapsed

The earlier paper's theme becomes clearer when four judgments are separated. A claim can be rhetorical without being false, unsubstantiated without being unfalsifiable, or non-empirical without being irrational. Conflating those categories would turn a critique of sloppy reasoning into an instance of it.

Concept	What it means here
Rhetorically forceful	The presentation encourages confidence, identification, or emotional commitment. This can accompany good or poor reasoning.
Weakly substantiated	The contribution does not supply enough support for the strength of its conclusion. Additional evidence or a better argument might repair it.
Slogan-like	A compact or repeatable affirmation does work that its supporting reasoning has not earned. The description concerns its use, not simply its length.
Insulated from correction	No clear result is allowed to count against the claim, or adverse results are routinely redescribed as confirmation.

Literal empirical falsifiability asks whether an observation could conflict with a claim. It is important in scientific testing, but it is not a universal test of meaningful thought. Mathematical arguments are assessed through proof; moral claims through reasons, implications, and consistency; metaphysical proposals

through coherence, explanatory comparison, and their relationship to experience. They need not all operate like laboratory hypotheses.

The broader standard relevant to this paper is **answerability to reasons**. Does the speaker identify a burden that could remain unmet? Would a failed premise, a stronger alternative, a counterexample, or an unexpected observation lower confidence? A claim need not be refutable by a single experiment to be exposed to rational correction.

The right question is not simply “Can you falsify this?” Ask: “What would count against it, and why?” The answer may involve an observation, a logical objection, a historical source, or a competing explanation. What matters is whether the proposed support can genuinely succeed or fail.

3. The original three-part test

The earlier assessment process records detailed judgments about warrant, the strength asserted, and the stability of qualifications. The test uses those existing fields. A move is flagged only when it satisfies all three conditions below.

- Its evidence-and-warrant score is below 70.
- Its asserted force is judged materially or radically overstated relative to the support supplied.
- Its terms or scope are unstable, or an important qualification is missing or materially misleading.

The conjunction matters. Weak support alone does not make a slogan. A compressed statement alone does not make one either. The test aims to identify the more concerning pattern in which a contribution carries limited support but presents a stronger or less qualified conclusion than that support earns.

No new reader has independently coded all of these passages for this edition. The calculation applies a transparent rule to recorded assessment judgments. That makes it reproducible, but it also means the result inherits the assessments' interpretations. A model that systematically overreads confidence on one side could inflate the measured difference.

The eligible set remains **146 debates and 2,800 moves**. The full archive's growth to 253 does not create missing fields in later records. The unchanged denominator is an important result of the data audit, not a failure to update the statistics. The original quantities were recalculated; their unchanged values reflect an unchanged eligible measurement set.

4. The measured asymmetry is substantial

Weak warrant appears in 57.2% of theist moves and 19.7% of non-theist moves. Material overclaim appears in 36.3% and 10.0%, respectively. Compression or qualification problems appear in 53.4% and 29.9%. The combined pattern is narrower: **398 theist moves and 111 non-theist moves**, or 28.6% and 7.9% of their respective totals.




Figure 1. The bars show pooled move rates in the 146-debate eligible set. The final row requires all three conditions together. The separate components overlap and must not be added to produce a total.

Scope: 146 earlier-process debates; original finding fields available.

There are two useful ways to describe the combined difference. Subtracting the pooled percentages gives about **20.7 percentage points**. Dividing them gives a **3.62-fold relative rate**. The first tells us the absolute difference in prevalence; the second tells us the multiplier relative to the smaller baseline. Reporting only the multiplier can make a result difficult to evaluate without knowing the baseline.

The primary paired analysis avoids letting longer assessments dominate. Within each debate, we calculate each side's flagged share and subtract the non-theist share from the theist share. Averaging these 146 differences gives **20.65 percentage points**. The direction favors a higher theist rate in 103 debates, ties in 26, and reverses in 17. Thus the pooled result is not just the product of one or two unusually large records.

The 95% resampling interval, 16.28 to 25.01 points, describes sensitivity to reselecting observed debates. It does not estimate how many flags are genuine slogans, and it cannot eliminate systematic judging error. The distinction between a precisely reproducible measurement and a valid interpretation of that measurement remains essential.

5. Does the pattern survive narrower checks?

The narrower checks ask whether the pattern is confined to opening claims, peripheral statements, a particular role orientation, or Lennox's contributions. They preserve the same three-part rule while changing which recorded moves enter the comparison. The changing denominators are part of the result, not a technical detail to conceal.

The paired three-part pattern under narrower checks




Figure 2. The same three-part rule is applied to narrower selections. The theist-CON row has only 15 debates and an interval crossing zero; it is not conclusive evidence that role orientation is irrelevant.
Scope: Earlier process only; original three-part rule.

The paired difference is 27.20 percentage points for constructive moves across 78 eligible debates, 17.47 for replies across 145, and 22.09 for high-importance moves across 146. Each comparison requires both sides to have the relevant kind of move. These are related but not identical populations, so their differences should not be treated as a causal effect of making a reply rather than an opening claim.

Removing Lennox leaves 140 debates and a 21.05-point gap. His examples are therefore illustrative, not the numerical foundation of the finding. In his six eligible debates, **15 of 56 theist-side moves** meet the rule. Those counts concern particular assessed contributions. They are not a judgment about every statement in his work or his personal intentions.

The theist-CON check is weaker. In 15 eligible debates where the religious side occupies CON, the mean difference is 11.00 points, but the interval runs from about -4.65 to 25.47. The data are compatible with a positive pattern but are too limited and variable to establish it securely in that subgroup. This is a meaningful restraint on the conclusion: the argument should not claim to have ruled out burden or role effects completely.

6. What the new assessments can - and cannot - add

Later records contain final dimension scores but not the earlier rule's separate judgments of asserted force and qualification. Reconstructing those missing judgments from the final scores would falsely suggest that the same measurement had been performed. Instead, this edition introduces a separate descriptive check.

The new check flags a move when evidence/warrant is below 70, precision/clarity below 80, and confidence/fairness below 80. This is a **score-based conjunction**, not a translation of the original rule. The last dimension combines confidence with fair representation, so a low score can reflect uncharitable treatment even when the confidence component is not the main problem.



Figure 3. The score-based conjunction shows a higher theist rate in both assessment-process families. It supports continuity of the broader low-support pattern, but it does not convert the newer records into a literal slogan census.

Scope: Evidence <70, precision <80, confidence/fairness <80 together; NOT literal non-falsifiability.

Across the 187 relevant debates, the average within-debate flagged rates are 36.7% for theist moves and 13.1% for non-theist moves: a **23.62-percentage-point paired gap**. The earlier-process gap is 24.25 points, and the later-process gap is 21.39. The later interval is approximately 12.00 to 30.66 points across 41 debates.

The two tests overlap conceptually and rely on the same underlying scoring framework. They are not independent experiments whose evidence can simply be multiplied. Their value is more modest: the broader pattern of low support coupled with other presentation weaknesses is not confined to the oldest records. The precise prevalence of correction-resistant slogans remains unmeasured.

7. Lennox example one: infinite worth through creation

In debate 170, John Lennox grounds human meaning and worth in divine intention and image-bearing. A short source phrase captures the move: **“infinite value because it bears the image of its creator.”** The assessment gives the move an evidence-and-warrant score of 61. See [Lennox-Singer, 93:20](#).

The affirmation performs a clear emotional and moral function. It tells a person that their existence is wanted and their worth is not reducible to productivity, ability, or social approval. Recognizing that appeal is not a refutation of the proposition. Nor should a reader infer that the statement is insincere or manipulative merely because it reassures.

The argumentative difficulty is that the explanation assumes a disputed creator and leaves “infinite value” insufficiently specified. What establishes divine intention? How does creation confer the relevant moral status? Why would a secular account of equal standing or finite purpose fail? Without those steps, the statement offers a compelling description of life *within Christian belief* more readily than a public reason to adopt that belief.

This is an example of **unfalsifiable slogan proffering in its protective use**: every person's existence can be described as compatible with being wanted by God, while no contrary condition is specified in the move. That does not show that human worth is false or that moral value must be empirically measurable. The issue is whether creator-given worth is being offered as an established explanation without identifying what could undermine it.

A stronger version would separate the moral commitment from its proposed grounding. It could defend equal worth through explicit moral reasoning, then explain why divine intention adds a necessary or superior foundation. Alternatively, it could present the statement honestly as a faith-based affirmation rather than as evidence that should settle a public debate.

8. Lennox example two: compensation after death

In the same debate, Lennox appeals to divine solidarity and postmortem restoration. The source passage concludes that **“God can compensate.”** The assessed contribution receives an evidence/warrant score of 64. See [Lennox-Singer, 83:00](#).

The humane aspiration is clear: a victim's suffering need not be the final fact about that person's life. Christianity's account of Christ entering suffering and of an eventual life without pain offers a response to abandonment and despair. A careful critic should acknowledge that this is more than a random change of subject.

Yet three claims need separation. Compensation may improve a victim's eventual condition. It may outweigh an earlier harm. It may or may not justify permitting the harm in the first place. Showing one does not automatically establish the others. In ordinary moral reasoning, a later benefit does not invariably authorize an earlier avoidable injury.

The correction problem arises if any amount of suffering can be absorbed by an unspecified future compensation. More suffering can always be met with more unseen restoration. Without independent support for that restoration and an account of permissible divine action, the proposal risks becoming an answer that no observed pattern can weaken.

This is not an empirical disproof of an afterlife. It is a diagnosis of why the particular reply can become insulated from the evidence it is supposed to address. A stronger defense would explain what the afterlife claim rests on, distinguish compensation from justification, and identify whether some pattern of suffering would still count against the defended conception of God.

9. Lennox example three: miracles and the critic's assumptions

In debate 169, Lennox describes miracles as exceptional divine acts within an ordinarily regular universe. He says **“science has nothing to say about singularities”** and connects categorical rejection of the supernatural to prior exclusion. The recorded evidence/warrant score is 54. See [Lennox-Krauss, 24:34](#).

There is a legitimate point inside this move. A critic who rules out supernatural action by definition has not independently shown that a particular miracle report is false. Likewise, an exception to an otherwise regular pattern is not simply the same claim as a world without any regularities. Those distinctions deserve engagement.

The problem is the stronger protective use. A miracle claim needs a way to distinguish divine action from error, legend, coincidence, or an unknown natural cause. Identifying a possible prior bias in the critic does not supply that discrimination. And the uniqueness of an event does not prevent reasoning from traces and evidence about it; much historical inquiry concerns singular events.

If nonacceptance is routinely explained as prior naturalism, the critic's objection can be converted from a challenge to the evidence into evidence of the critic's closed worldview. That is the slogan-like escape route. It is not established merely by pointing out that philosophical assumptions exist. The fair test is whether the speaker also answers the actual evidential reasons for nonacceptance.

10. Other examples: private confirmation and protected motives

Ray Comfort provides two distinct cases in debate 142. At 95:13, private manifestation of the risen Christ, promised after repentance and trust, is offered as support for resurrection. The evidence/warrant rating is 39. A participatory test is not automatically worthless: some experiences require a person to undertake a practice. But the test needs intelligible success and failure conditions.

What counts as manifestation? How is it distinguished from expectation or a conflicting religious experience? What follows if a sincere seeker undertakes the practice and does not receive confirmation? If failure can always be assigned to insufficient sincerity, the procedure protects the claim from precisely the result that ought to test it.

At 25:40, Comfort explains rejection of God through avoidance of moral accountability, comparing it with criminals' attitude toward law enforcement. The evidence/warrant score is 30. The assessment's concern is not that avoidance of responsibility never occurs. It is the general attribution of motive without an independent way to distinguish it from sincere disagreement.

The danger is a self-sealing interpretation: agreement confirms the religious diagnosis, and disagreement confirms it too by revealing resistance. A psychological claim could in principle be investigated. In the argumentative use at issue, however, the speaker has not identified what evidence would show that the diagnosis is wrong about a particular critic.

These examples show why “unfalsifiable” should describe the **way the claim is protected in the exchange**, not function as a casual insult meaning “I do not believe it.” The relevant question is whether the explanation permits the challenge to make any rational difference.

11. The same standard catches secular overclaim

Peter Atkins's categorical treatment of scientific practice and religious belief in debate 85 supplies a non-theist counterexample. The assessment records an evidence/warrant score of 40 for a move treating religion as superstition or wishful thinking and asserting intellectual incompatibility. The broader [Atkins-Fox scorecard](#) gives the theist side a 15-point advantage.

A claim of incompatibility must specify what is incompatible with what. Particular doctrines may conflict with particular findings; that does not establish that every religious interpretation conflicts with scientific reasoning in the same way. If religious scientists are dismissed as insufficiently thoughtful whenever they appear as counterexamples, the generalization is being protected rather than properly tested.

Another recorded non-theist move invokes an unrestricted space of possible worlds to make a godless multiverse overwhelmingly probable. The problem is not that multiverse proposals are automatically irrational. It is that a probability claim requires a justified way of assigning weight to the possibilities. Merely saying there are enormously many possibilities does not supply that weighting.

An easy mathematical illustration helps. There are infinitely many positive integers and infinitely many even positive integers. Counting each set as “infinite” does not establish that a fair lottery over all integers exists, or tell us how to draw from it. A claim about the probability of whole universes requires more care, not less. This is a methodological criticism that applies equally to unsupported fine-tuning probabilities on the religious side.

12. Not every theist argument fits the pattern

The archive also supplies strong theist-side counterexamples. In debate 137, William Lane Craig argues for the authenticity of a passage in First Corinthians using manuscripts, early citation, vocabulary, context, and the effect of removing the passage. Its evidence/warrant score is 88. Those considerations can be examined and challenged. Whether the larger resurrection case succeeds is separate from the quality of that bounded historical move. See [Craig-Price](#).

In debate 139, Matthew Adelstein explains that hypotheses should be compared using both their initial plausibility and how strongly they lead us to expect the evidence. The move receives an evidence/warrant score of 87. In debate 144, Plantinga carefully distinguishes establishing compatibility from establishing that divine guidance actually occurred; that move receives an overall score of 92.

These cases matter for the hypothesis. If the proposed mechanism were simply “religion makes rational argument impossible,” well-supported religious-side moves would be inconvenient exceptions. A more precise hypothesis predicts variation: religious arguments should improve when they supply independent support, expose assumptions, compare alternatives, and limit their conclusions. The counterexamples fit that more useful formulation.

The newer Sechler-Clifton debate provides another kind of counterweight. The assessment credits Sechler for revising or narrowing inherited positions about salvation and hell rather than defending every traditional claim unchanged. His side still scores lower overall. Concession and flexibility can be strengths without

guaranteeing victory; a lower score should not erase evidence that a speaker is responsive to reasons. See [debate 253](#).

13. The causal hypothesis remains open

The proposed explanation is that emotional affirmation helps sustain religious belief and that this habit can carry into argument. The close readings show how an affirmation can offer reassurance, belonging, or protection from uncertainty. The statistical analysis shows that a related low-support pattern is more common on the religious side. These are compatible observations, but there is a gap between them.

The archive has no systematic measures of audience emotion, institutional pressure, childhood instruction, identity threat, private belief formation, or sincerity. It therefore cannot establish that theism is *largely enforced* through emotional devices. The debate transcripts are not a representative sample of religious life. Even an unmistakably reassuring statement might have been offered sincerely as a reason rather than deliberately as an emotional substitute.

Other explanations remain plausible: more ambitious religious claims, asymmetrical burdens, conventions of public apologetics, unequal speaker skill, or a judging framework more comfortable with skeptical caution. The measured conjunction itself partly overlaps with the dimensions used to assign scores, so its association with low scores is not independent evidence for a psychological cause.

A stronger follow-up would randomly select passages, hide speaker and side identities where feasible, and have separate teams code emotional appeals, correction conditions, and adequacy of support. The teams should not see each other's judgments or the public scores. Their disagreements would be published. Only then could the analysis ask whether emotional framing predicts missing support beyond topic, role, and conclusion strength.

14. Conclusion: confidence must be able to lose

The most defensible claim is not that all theist arguments are slogans. It is that **a substantial, explicitly measured pattern of weak support joined to overclaim and poor qualification is more common on the theist side of this archive**. The result persists when Lennox is removed and appears in a separate score-based check of the enlarged comparison set. Concrete passages show why the concern is more than a label attached to a low score.

The strongest examples share a protective structure. Creation guarantees worth without independent support for the grounding. Future compensation absorbs present suffering without a specified limit. Rejection of a miracle becomes evidence of the critic's prejudice. Failure to receive religious confirmation can be assigned to the seeker; disagreement can be assigned to moral evasion. In each case, the question is what would make the proposed support fail.

That question also disciplines the present paper. Its rule is a risk indicator, not a validated census of unfalsifiability. Its examples are selected illustrations, not a representative sample of every speaker. Its causal explanation remains testable and incomplete. A paper criticizing insulated beliefs should not insulate its own interpretation from contrary evidence.

An argument earns confidence when the same procedure that can support it can also expose its limits. That principle neither privileges atheism nor exempts theism. It asks every worldview to distinguish the statement that sustains its adherents from the reason that should persuade someone who does not already agree.

Methods and sources

The source revision is **76d006b37**, frozen September 4, 2026. The archive has 253 assessments; the broad position comparison has 187 debates and 4,086 moves. The original detailed-findings rule applies to 146 earlier-process debates, comprising 1,393 theist and 1,407 non-theist moves. The separate final-score conjunction applies to all 187 comparisons, including 41 later-process debates. The two eligible sets must not be combined as though the measurements were identical.

Pooled percentages divide flagged moves by all eligible moves on the same substantive side. Paired percentage-point differences first subtract within each debate and then average across debates. Intervals use 20,000 complete-debate replacement resamples, seed 20260904. They reflect variation in observed debate composition, not measurement validity, transcript accuracy, all repeated-speaker dependence, or the uncertainty of the causal hypothesis.

The Lennox, Comfort, Atkins, Craig, Adelstein, and Plantinga examples are traceable to locked source excerpts and public scorecards. Short quotations above are quotations; the surrounding case descriptions are paraphrases and interpretations. Timestamp links point to the relevant recorded contribution, although caption timing and the boundaries of an exchange may be approximate. The examples were selected to explain mechanisms and counterexamples, not to estimate their prevalence.

The [reproducible analysis directory](#) contains the rule, classifications, source hashes, calculations, and chart data. The [Backend library](#) links the companion papers on score gaps, topic differences, role effects, fallacy labels, measurement comparability, and rankings. This Astra-era edition reanalyzes existing assessments; it does not claim that Astra independently rescored the corpus.