

Are Theist Arguments More Often Slogan-Like?

A dated transcript study, with a separate updated score-based comparison.

77 / 19

protected uses: religious
versus skeptical side

187

historically reviewed transcripts
September 5 direct study

234

current score comparisons
not new slogan reviews

The direct finding concerns criticism-blocking slogans - and remains a dated study.

PHIL STILWELL · SLUGFESTER

October 9, 2026 | Frozen archive: 308 assessments

Recalculated scores; direct slogan evidence retains its September scope.

Executive summary

The strongest direct finding remains about slogans that both substitute for reasons and block criticism. The September 5 review of 187 retained transcripts found 77 protected uses on the religious side and 19 on the skeptical side. The rates, allowing for speech length and giving each debate equal weight, were 0.56 and 0.16 uses per 10,000 words. Recalculating the published debate-level records reproduces those findings.

Those counts have not expanded to the current catalogue. This October edition rewrites their interpretation and adds a separately labeled current score-based comparison. It does not claim that the 47 new religious-versus-skeptical comparisons have received direct slogan coding. Their scores cannot supply missing transcript judgments.

The broader unsupported-slogan difference remains less secure: 0.97 versus 0.65 uses per 10,000 words, with a repeated-draw range for the difference that crosses zero. Most reviewed debates - 144 of 187 - contain no detected protected slogan on either side. A group difference should not be converted into a description of every believer or skeptic.

The new score-based warning rule finds simultaneous weak support, imprecision, and insufficiently calibrated confidence more often on the religious side. That is consistent with an argumentative weakness worth investigating, but it is not a new count of slogans and does not establish the emotional or psychological causes of belief.

Reader's guide

Executive summary	2
1. Two bodies of evidence, two different questions	4
2. A memorable sentence is not automatically a defective argument	4
3. Allow for how much each side speaks	5
4. Concentration and uncertainty matter	6
5. What the expanded score archive adds	7
6. What can responsibly be concluded?	8
Methods and sources	8

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. Score results refer to the October 9 snapshot. Direct slogan counts retain their September source coverage; they have not been extended to new transcripts.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. Two bodies of evidence, two different questions

The direct study returns to the retained transcripts and examines uses in context. Each of 187 debates received one fresh GPT-6 Astra reading at low reasoning effort, without the old scores or expected results. The archived annotations and source-linked incidents remain unchanged. The five earlier two-reader pilot pairs were excluded from the final study and are not silently added back here.

The current score study instead reads existing public assessments. It contains 234 religious-versus-skeptical comparisons and 5,652 assessed moves. These moves are selected argumentative contributions; they are not complete transcripts. Numerical thresholds applied to their scores can detect a pattern in those judgments, but cannot establish the wording or function of a slogan in the surrounding conversation.

Evidence	Coverage	What the refresh does
Direct transcript annotations	187 September-source transcripts	Recalculates saved counts and word-normalized rates; no new readings
Current score records	234 comparisons	Recalculates a separate three-score warning rule
New direct annotations	None in this edition	Not inferred, extrapolated, or represented as completed

More scorecards are not more slogan reviews. A complete new transcript study would require applying the same contextual reading rules to the new eligible sources and validating the resulting annotations.

2. A memorable sentence is not automatically a defective argument

An unsupported slogan must be a compact or repeatable affirmation that presents a conclusion as settled **and substitutes for a needed reason**. A protected slogan must also contain positive evidence of shielding the claim from criticism. The latter is a subset of the former; their counts must not be added.

Brevity, confidence, repetition, humor, religious language, or emotional force alone is not enough. A short sentence can summarize a carefully supported argument. A story can communicate real evidence. A metaphysical claim can be open to rational criticism even if no single laboratory experiment could settle it.

Protection means something stronger: for example, treating every possible outcome as confirmation, or declaring disagreement evidence that the critic is morally defective regardless of the objection. These are illustrative descriptions of the rule, not new quotations from the archive. The problem is the function of the claim in context, not its tone alone.

A statement such as “my opponent has not disproved this” can be appropriate against an overconfident denial. It becomes inadequate if used as the sole support for a positive conclusion. Equally, a confident skeptical dismissal can substitute a label for engagement. The same coding rule must apply to both.

3. Allow for how much each side speaks

The historical files contain 3,496,979 caption words. Of these, 1,433,125 were attributed to religious-side speakers and 1,439,245 to skeptical-side speakers. Moderator speech, introductions, other excluded material, and uncertain attribution account for the remainder. These are caption-based denominators, not verified word-for-word transcripts of every sound in the recordings.

The rate for each side is its counted uses divided by attributed substantive words, multiplied by 10,000. Compare those rates within each debate, then average with equal weight per debate. A longer recording does not automatically count more heavily in the principal result.

For illustration only, four uses in 20,000 words give a rate of two per 10,000; three in 10,000 give a rate of three. The side with the larger count can therefore have the smaller rate. Immediate repetition within one uninterrupted argument counts once; renewed use after an actual response can count again.




Figure 1. These historical rates remain reproducible from the saved direct-review records. They do not describe newly added transcripts.

How to read this graph. Rust is theist; teal is non-theist. Each debate receives equal weight after allowing for speech length. Both panels start at zero on the same scale. Protected slogans are a subset of unsupported slogans; do not add them. These dated counts have not been extended to new transcripts.

The protected counts are 77 and 19, nested inside unsupported counts of 128 and 82. The absolute rates are low even where their ratio looks large. A ratio near 3.5 for protected uses should not hide the fact that most debates contribute zero.

4. Concentration and uncertainty matter

Of the 187 reviewed debates, 32 contain protected uses only on the religious side, nine only on the skeptical side, two on both, and 144 on neither. The middle within-debate difference is zero. An average can be positive because a minority of cases have larger values while the middle case remains zero.

The historical supplementary calculations also show concentration. Ten Bruggencate, Comfort, and Turek contribute 31 of the 77 religious-side protected uses; Atkins, Hitchens, and Ehrman contribute 14 of the 19 skeptical-side uses. These are contributions to totals, not fair personal rankings. Appearances, speaking time, and opponents differ.

Removing the two debates with the largest religious-side protected counts leaves a historical difference of about 0.23 uses per 10,000 words, with a repeated-draw range of about 0.02-0.43. The overall result is not wholly produced by those debates, but its magnitude depends on the mixture. This was a supplementary check chosen after inspecting the counts, not a separately predicted experiment.




Figure 2. The broader unsupported-slogan difference crosses zero under resampling; the stricter protected-slogan difference stays positive.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population. The unsupported-slogan range crosses zero; the protected-slogan range does not. Both use the September direct annotations, not new scorecard judgments.

The protected difference is 0.40, with a middle-95% range of 0.11 to 0.74. The unsupported difference is 0.32, with a range of -0.13 to 0.78. All units are uses per 10,000 words, religious minus skeptical.

Whole-debate resampling preserves the actual opposing sides but cannot correct a reading error, turn one AI reader into independent double coding, or remove repeated-speaker dependence. File fingerprints and located quotations establish traceability, not infallible interpretation. “None detected” is not a certificate of perfect reasoning.

5. What the expanded score archive adds

The separate warning rule requires all three scores on one assessed move: Support below 70, Clarity below 80, and Care below 80. Within each debate, calculate the share of each side's moves satisfying all three; then average the side differences. This procedure gives debates equal influence despite different numbers of assessed moves.

The current religious-side share is 33.87%; the skeptical-side share is 13.40%. Their difference is 20.47 percentage points across 234 comparisons. Earlier-process comparisons show a 24.25-point difference; later-process comparisons show 14.21.




Figure 3. This current warning rule is based on scores, not on direct identification of slogans. Its percentages must not be combined with the word-normalized transcript rates.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population. A flag requires all three score thresholds on the same move. A low score does not demonstrate a criticism-blocking slogan. Never combine these percentages with the transcript word rates.

These thresholds are conventional screening choices, not natural boundaries. A weak, unclear, overconfident argument may contain no slogan at all. A protected slogan may receive scores above one threshold. The scores also share the assessment model's judgments, so their agreement is not independent validation of the direct study.

The older detailed-findings rule remains available in the archived analysis, but this edition does not revive it as a substitute for direct reading. The purpose of the current check is to describe a related assessment pattern while being explicit about the evidential gap.

6. What can responsibly be concluded?

Reason 1: the historical direct study identifies a stronger directional pattern for criticism-blocking slogans than for unsupported slogans in general. **Reason 2:** the behavior is unevenly distributed, and most reviewed debates have no protected use detected. **Reason 3:** the current scores reveal a related weakness pattern but do not extend the direct count.

Therefore: the strongest claim remains a qualified one about recorded uses under a particular contextual rule. It is not that religious people merely feel while skeptics reason, nor that every religious assertion is insulated from criticism. Emotional wording does not establish private motivation, belief formation, or effects on listeners.

Therefore: a future extension should preserve the same definitions, read new transcripts without being shown the expected direction, validate counts and speaker attribution, and publish the uncertainty and exclusions. Until then, the direct finding should carry its September date prominently. Rewriting a report must not manufacture evidence that was never collected.

Methods and sources

This is a fixed October 9, 2026 research edition, based on source revision e65bc3c0b6. The score studies reconstruct published records rather than commission fresh assessments. The direct slogan counts retain the September 5 review of 187 transcripts; this edition recalculates those saved records and adds a separately labeled current score check. No new slogan reviews or debate assessments were performed.

The source manifest records file fingerprints. The analysis uses 20,000 deterministic resampling draws with baseline seed 20260904, retaining it for edition comparability. Whole debates are drawn for paired debate estimates; shared speakers for cross-role estimates. Ranking models and their dependence checks use the units explained in paper seven. These calculations do not capture every selection effect, repeated-speaker dependency, caption error or systematic judging bias.

The [data and methods page](#) links the source records, classification decisions, calculations, chart inputs and reading guides. Download the [current research package](#) and [historical direct slogan study](#). Earlier editions remain preserved in the repository's history. Debate links identify published assessments, not fresh independent evidence for this research.