

Are Theist Arguments More Often Slogan-Like?

A fresh reading of 187 debates: unsupported slogans, closed-off claims, and what the evidence says about emotional persuasion.

187

complete retained transcripts
reviewed under the same rule

77 / 19

protected slogan uses detected:
theist / non-theist

144

debates with no protected slogan
detected on either side

The clearest difference concerns slogans that shut out criticism, not slogans in general. This supports a narrower part of the hypothesis, while leaving its proposed emotional cause unproved.

PHIL STILWELL · SLUGFESTER

September 5, 2026 | 187 relevant debates

Fresh slogan review. Original debate scores have not been changed.

Executive Summary

The strongest result concerns a particular kind of slogan: one that replaces a needed reason and then shuts out criticism. A fresh review of all 187 relevant debates found 77 such uses on the theist side and 19 on the non-theist side. After allowing for how much each side spoke and giving each debate equal weight, the rates were 0.56 and 0.16 uses per 10,000 words. The average difference was +0.40 in the theist direction.

The broader claim about unsupported slogans is less settled. The review found 128 theist uses and 82 non-theist uses, with average rates of 0.97 and 0.65. However, the direction of that broader difference is not stable across repeated draws from these debates. The paper therefore supports the narrower criticism-blocking claim more clearly than a general claim that theists use more unsupported slogans.

Most debates contained no detected protected slogan on either side. That was true of 144 out of 187 debates. The pattern is uneven, with several speakers contributing many of the counted uses. A difference across the archive is not a description of every believer, every skeptic, or every exchange.

The proposed emotional cause remains unproved. Emotional wording appeared in many of the detected unsupported slogans on both sides. This study measures words used in debates. It does not measure why people become religious, what maintains their beliefs, or how listeners respond.

What is new in this edition?

The previous paper counted combinations of warning signs in existing argument assessments. This edition returns to the complete retained transcripts and asks the slogan question directly. All 187 relevant debates received the same fresh review, including the 41 later assessments that lacked the older detailed fields. These are new research annotations, not new debate scores.

The central finding: the evidence points most clearly to a difference in how some claims are protected from criticism. It does not justify the larger statement that one side reasons while the other merely feels.

The archive was frozen on September 4, 2026, at 253 assessments. The 187 relevant transcript reviews were completed for this September 5 report. All numerical comparisons below refer to that fixed set, not to later additions to the site.

1. A memorable line is not automatically a bad argument

The problem is not that a claim is short, forceful, religious, or emotional. The problem is what the line does in the argument. A brief sentence can summarize a careful case. A moving story can help explain real evidence. Neither should be counted as a defect merely because it is memorable.

The review used three separate questions. Each must be answered from the surrounding exchange, not from a quotation pulled out of context.

Question	What must be present?	What is not enough?
Is it slogan-like?	A compact or repeatable affirmation presents a general conclusion as settled.	Brevity, confidence, humor, or repetition alone.
Does it replace support?	The affirmation does the work of a reason the speaker has not supplied.	A summary of reasons already given, or an ordinary disputed premise.
Does it shut out criticism?	A stated move rules out objections or contrary outcomes regardless of their merits.	Faith, weak evidence, or failure to name a laboratory test.

An **unsupported slogan** meets the first two conditions. A **protected slogan**, the shorter name used here, meets all three. The second category is inside the first: every counted protected slogan is also a counted unsupported slogan. The two totals must not be added together.

What does “unfalsifiable” mean here?

In everyday use, the word can mean that nothing is allowed to count against a claim. This paper asks for a visible reason to apply that label. For example, a speaker might treat every possible result as confirmation, or say that disagreement itself proves the critic is dishonest. The claim gains a way to survive without answering the objection.

That is different from a claim that cannot be settled by one experiment. Questions about morality, meaning, and explanation can still be challenged by reasons, consistency, competing accounts, or consequences. The absence of a simple scientific test is not a finding of defective reasoning.

A useful distinction: “You have not disproved my claim” does not mean “I have supported my claim.” And “I can explain away your objection” does not mean “I have answered it with an independently supported reason.”

The labels describe particular uses. They do not settle the truth of a whole worldview. Uncertain cases remain uncertain rather than being pushed into the category that best fits the hypothesis.

2. The comparison allows for how much each side spoke

Raw counts can reward the side that speaks less. A fairer comparison asks how often a counted use appears within a similar amount of speech, then compares the two sides of the same debate.

The source files contain 3,496,979 caption words across about 335 hours of recordings. The review assigned 1,433,125 words to theist speakers and 1,439,245 to non-theist speakers. These nearly equal totals are reassuring, but individual debates can still have unequal contributions.

For each side of each debate, the calculation is:

Rate per 10,000 words = counted uses ÷ attributed words × 10,000. These are uses of a claim, not the number of individual words that are slogans.

Consider an invented example, used only to explain the arithmetic. One side makes four counted uses in 20,000 words. Its rate is 2 per 10,000. The other makes three uses in 10,000 words. Its rate is 3. Although the first side has the larger raw count, the second uses the device more often for the amount it speaks.

The main result averages these rates with **equal weight for each debate**. A three-hour debate does not get three times the influence of a one-hour debate. The difference is calculated as theist rate minus non-theist rate: a positive number means more theist uses.

What was counted once, and what was left out?

A distinct use within one uninterrupted argument counts once. Immediate repetition does not create a new use. A renewed assertion after an actual reply can count again. A promotional clip that replays later speech is not fresh argument and is excluded.

Moderator and audience speech, advertisements, introductions, and other excluded material account for 517,106 words. Another 107,503 words, about 3.1% of the source total, could not be assigned safely because speakers were unclear or mixed together. Substantive answers to audience questions remain included when the speaker can be identified.

The study also reports a **pooled rate**: add all uses on a side, then divide by all that side's words. That answers a different question about the combined body of speech and gives longer contributions more influence. Both calculations point in the same direction here, but their values are not interchangeable.

3. The clearer gap is in slogans that shut out criticism

Theist speakers have the higher recorded average on both tests, but the stricter test gives the clearer difference. For unsupported slogans, the average is 0.97 versus 0.65 uses per 10,000 words. For protected slogans, it is 0.56 versus 0.16.




Figure 1. Average rates in all 187 debates, using the September 4 archive and the fresh September 5 review.

How to read this graph. Rust = theist; teal = non-theist. Bar length shows detected uses per 10,000 caption words attributed to that side, averaged with equal weight for each of the 187 debates. Both panels use the same scale and start at zero. The right panel is a stricter subset of the left; do not add them together.

The protected-slogan difference is about **0.40 additional theist uses per 10,000 words**, averaged across debates. The theist rate is about 3.5 times the non-theist rate. That ratio sounds large, but both rates are low in absolute terms. Many debates contribute zero, while a few contribute much more.

The raw counts are 128 versus 82 unsupported slogans, of which 77 versus 19 also meet the protection rule. Pooled across all attributed words, the corresponding rates are 0.89 versus 0.57, and 0.54 versus 0.13. Neither way of allowing for speech length reverses the recorded direction.

Do not read 0.56 as 56%, or as a debate score. It is an average number of detected uses per 10,000 words. A decimal average does not mean that a fraction of a slogan was counted.

A count difference alone cannot show how reliable a broader conclusion would be. Sections 5 and 6 examine how much these averages depend on the particular debates included.

4. A group difference is not an everywhere pattern

Most of these debates did not contain a detected protected slogan on either side. There were 144 such debates, or 77.0% of the 187. This matters just as much as the difference between 77 and 19 uses.




Figure 2. Debates with at least one detected use, separated by which side used it. Each panel accounts for all 187 debates.

How to read this graph. Each bar counts debates, not slogans or speakers. Navy = at least one detected use on both sides; rust = only theist; teal = only non-theist; pale grey = neither. Each panel adds to 187. A debate with many uses still counts once. Neither means none detected under this rule, not proof of perfect reasoning.

For unsupported slogans, 61 debates contain at least one theist use and 49 contain at least one non-theist use. These totals overlap because 21 debates contain uses on both sides. The remaining 98 have none detected on either side.

For protected slogans, the corresponding side totals are 34 and 11 debates, including two with uses on both sides. That leaves 32 theist-only cases and nine non-theist-only cases. The difference therefore extends beyond one prolific speaker, but it remains concentrated in a minority of the archive.

The **middle within-debate difference is zero** on both tests. To find a middle value, arrange the 187 differences from lowest to highest and look at the central one. A few large positive or negative values can move the average while leaving this middle value unchanged.

“None detected” is not a certificate of sound reasoning. A debate can contain poor evidence, evasive replies, or weak inferences without meeting this paper's narrow slogan rule. A single AI reader may also miss a qualifying use.

5. How firmly does the average hold up?

The archive gives exact counts under this review, but not a guaranteed result for every other set of debates. To see how dependent an average is on this particular mix, the analysis repeatedly draws from the existing debates.

Imagine putting the 187 debate numbers in a bag. Draw one, record both sides together, and put the number back. Continue until there are 187 draws. Some debates appear more than once and some are left out. Calculate the average difference for that new mixture. Repeat the whole process 20,000 times.

This procedure is called **resampling**: drawing again from records already collected. Keeping both sides together preserves each actual debate comparison. The line shown for a result covers the middle 95% of those 20,000 averages. The lowest 2.5% and highest 2.5% are outside it.

Main test	Recorded average difference	Middle 95% of repeat-draw averages
Unsupported slogans	+0.32	-0.13 to +0.78
Protected slogans	+0.40	+0.11 to +0.74

All values are theist minus non-theist uses per 10,000 words. Negative means more non-theist uses; positive means more theist uses.

The first range crosses zero. The broader unsupported-slogan result can point either way within this exercise. That does not erase the observed count difference. It means that the particular mixture of debates matters enough to make a general directional claim less secure.

The protected-slogan range stays above zero. That gives stronger support to the narrower finding within the recorded material. It is not a 95% guarantee that the hypothesis is true.

What this exercise cannot fix

Drawing records again cannot correct a reading error copied into every draw. It cannot recover missing speech or turn selected YouTube debates into a random sample of humanity. The same speakers also appear repeatedly, so debates are not fully independent examples of personal behavior. These lines are useful checks on the mix of debates, not complete measures of every source of uncertainty.

6. The narrow result survives the planned checks

The protected-slogan difference remains in the theist direction in each planned selection. The broader unsupported-slogan result is less consistent in its repeat-draw ranges.




Figure 3. Average differences across the full set and four overlapping selections.

How to read this graph. Each dot is the average within-debate difference. Right of the dashed zero line means more theist uses; left means more non-theist uses. Each line covers the middle 95% of averages from 20,000 repeat draws of whole debates. Crossing zero means the direction changes in this exercise. Both panels use the same scale; n means debates. These are overlapping checks, not five independent studies, and the lines do not cover AI reading errors.

“Earlier” and “later” mean the original **assessment groups**, not necessarily when the debates were recorded. All received the same fresh slogan review. The earlier group contains 146 debates; the later group contains 41. The later group has eight protected theist uses and none detected on the non-theist side, spread across just five debates. A zero here does not establish that the behavior is absent outside these records.

The narrower selection removes broad religion-and-culture topics and specified mixed-position cases; it contains 146 debates. It is not the same 146 as the earlier group. Removing all eight Lennox debates leaves 179, with a protected-slogan difference of +0.44. The result is therefore not produced by Lennox's inclusion.

These selections overlap heavily. They are checks on one study, not five independent confirmations. Differences in speakers and subject matter remain possible explanations for differences between selections.

7. A few people matter, and uncertain cases matter

A total can exaggerate how widely a practice is shared. Sye Ten Bruggencate, Ray Comfort, and Frank Turek account for 31 of the 77 protected theist uses, or 40.3%. Peter Atkins, Christopher Hitchens, and Bart Ehrman account for 14 of the 19 protected non-theist uses, or 73.7%.

These are contributions to totals, not a fair speaker ranking. People appear different numbers of times, speak for different lengths, and face different opponents. A higher count is not automatically a higher personal tendency.

The two largest theist debate counts contribute 23 of the 77 protected uses: 12 in the Ten Bruggencate-Dillahunty debate and 11 in the Comfort-Dillahunty debate. As an extra check chosen after seeing these counts, removing those two entire debates leaves a smaller difference of +0.23, with a repeat-draw range of +0.02 to +0.43. They matter substantially, but they do not produce the entire recorded pattern.

The same removal reduces the broader unsupported-slogan difference from +0.32 to +0.11, with a range from -0.27 to +0.44. This is another reason not to turn the broader result into the paper's strongest claim.

What if every plausible borderline case counted?

The main totals require clear “yes” findings. A separate check adds every uncertain candidate for which none of the required conditions was clearly absent. It does not change the original labels.

Rule used	Unsupported uses: theist / non-theist	Protected uses: theist / non-theist
Clear cases only	128 / 82	77 / 19
Clear plus plausible uncertain cases	287 / 169	221 / 85

Under the inclusive rule, the average differences become +0.85 and +0.94. Their repeat-draw ranges are +0.25 to +1.45 and +0.47 to +1.44. The theist direction remains, but the size changes sharply. This is a sensitivity check, not permission to relabel uncertainty as fact.

There are also 13 candidates excluded because speech was mixed or the speaker uncertain. Seven meet the unsupported conditions and six the protected conditions before attribution. They are not silently assigned to either side. Nor do these detected candidates place an upper limit on everything the reader might have missed.

8. A claim can make doubt itself impossible

The clearest examples show the protective move explicitly. The following cases illustrate what the rule detects. They were selected to explain contrasting forms, not because they represent a random sample of speakers.

Sye Ten Bruggencate: doubt becomes confirmation

In his debate with Matt Dillahunty, Ten Bruggencate is asked whether any experience could change his position. He answers no, says that “doubt presupposes the existence of God reason presupposes God,” and says that changing his position is impossible. [Debate 5, about 1:29:49](#).

The defect is not confidence by itself. The reply makes the very act of questioning the claim depend on the truth of the claim. Whatever a critic says must supposedly borrow the religious foundation under dispute.

That connection needs support. Why must reasoning require this foundation? Why could an alternative account not explain our ability to reason? Simply affirming the dependence does not answer either question. In this use, the maxim both supplies a missing reason and prevents doubt from counting against the conclusion. It therefore meets the protected-slogan rule.

Christopher Hitchens: claimed knowledge becomes disqualification

Against John Lennox, Hitchens provisionally grants that a creator might exist, then asks how anyone could know that creator's wishes. The question is legitimate: establishing a creator would not, by itself, establish a moral instruction from that creator.

He goes further, dismissing people who claim this knowledge and saying that claimed certainty should be distrusted “a priori,” meaning in advance. [Debate 15, about 18:16](#).

The missing step is a reason that every proposed route to this knowledge must fail. Asking for evidence is fair. Disqualifying a claimant before assessing that evidence closes off the exchange. The review applies the same protected-slogan rule to this non-theist use.

The scope matters: this passage concerns claimed knowledge of God's wishes. It is not evidence that Hitchens rejected every kind of certainty in every subject. The criticism is of the advance dismissal, not of the legitimate demand that Lennox support his claim.

9. Lennox supplies an example, not the main explanation

The direct review does not make John Lennox the central source of the aggregate difference.

Across his eight included debates, it finds four unsupported slogan uses, only one of which also meets the protection rule. His opponents collectively have eight and five. Eight debates are too few for a sweeping comparative judgment, but these results clearly restrain the initial expectation that he would dominate this category.

The protected case: excluding scientific criticism

In his exchange with Lawrence Krauss, Lennox says, “But science has nothing to say about singularities.” Here “singularities” means exceptional or one-off events in an argument about miracles, not a technical claim about black holes. [Debate 169, about 24:59](#).

The surrounding exchange is essential. Krauss challenges the practice of suspending regularities whenever they conflict with a religious claim. Lennox answers that God created regularities and can intervene to draw attention to himself. He then excludes science from the exceptional case and attributes Krauss's rejection to a prior commitment to atheism.

The possibility of intervention is a proposed explanation; it does not, on its own, show that evidence about regularities has nothing relevant to contribute. Nor does identifying an opponent as an atheist answer that opponent's request for evidence. The short boundary claim supplies a missing reason for excluding criticism. That is why this particular use meets the stricter rule.

A fair reply might distinguish what scientific evidence can establish from what it cannot establish, then explain why independent historical or other evidence should outweigh the objection. The study does not require a miracle to be repeatable before it can even be discussed. It requires a reason for dismissing otherwise relevant evidence.

The contrast: a larger plan is not automatically a closed claim

In the Hitchens debate, Lennox answers the prospect of the universe's destruction by placing it within a larger divine plan. This may appear to rescue design from any disappointing outcome. Yet the surrounding response invokes a particular religious expectation, and he soon makes revelation stand or fall on evidence concerning Christ. [Debate 15, about 55:20](#).

The review leaves possible protection uncertain and does not count this passage as a confirmed slogan. A stated evidence condition does not prove that the evidence succeeds. It does mean that the claim should not be called immune to criticism solely because it invokes a future beyond present observation.

10. Three Lennox slogans leave a missing step

A slogan can overstate a conclusion without making correction impossible. These three uses meet the unsupported-slogan rule but not the stricter protection rule. Keeping those findings separate makes the criticism more exact.

Historical influence is not exclusive necessity

Lennox approvingly quotes Christianity as the sole foundation of Western liberties and adds, “to this day we have no other option.” He follows this with claims about universities, liberty, and law. [Debate 15, about 35:22](#).

Even accepting those historical claims would not show that no other present basis for liberty is possible. Something can help a practice begin without remaining its only possible support. The missing step runs from **historical contribution** to **exclusive continuing necessity**. The categorical affirmation does not establish that step.

Naming a cause is not establishing a cause

Against Michael Shermer, Lennox uses a quotation that explains mass deaths by saying, “simply say we have forgotten God.” The surrounding argument concerns morality's foundation. [Debate 222, about 1:08:49](#).

An account of where moral obligations come from does not by itself explain why a particular historical catastrophe happened. That requires evidence connecting the proposed cause to the events and considering other explanations. The quotation compresses a large causal claim into an affirmation. Shermer offers a counterexample involving Christian warfare; the exchange does not make such counterexamples inadmissible.

A group's good contribution is not the whole history

In the same debate, Lennox credits Christianity with freeing slaves and secularism with restoring slavery. [Debate 222, about 1:14:33](#).

He acknowledges atheist caregivers and offers a religious reason for valuing people. Those points do not establish the broad allocation of historical credit and blame. Shermer contests the allocation. The paper does not settle that history from their competing assertions; it identifies the missing causal account in this use.

The shared weakness is an unearned step: influence becomes necessity; a moral position becomes a historical cause; selected contributions become a whole group's record. None of these steps is secured merely by repeating the conclusion.

These passages help readers identify weak support. They should not be advertised as three further confirmed examples of an unfalsifiable slogan. The direct review found that additional feature in the Krauss passage, not in these three.

11. Emotional wording does not establish an emotional cause

Both sides use emotional language within the slogans detected here. Among the 128 unsupported theist uses, 102 include at least one recorded emotional device. Among the 82 non-theist uses, 73 do. That is 79.7% and 89.0%, respectively.




Figure 4. Emotional wording among the detected unsupported slogans, not among all spoken words.

How to read this graph. Rust = theist; teal = non-theist. Bars show the percentage of detected unsupported slogans also marked for comfort, fear, shame, identity, or ridicule. The counts beside each bar show the denominator. This is not the percentage of all speech that is emotional. It measures wording, not a speaker's motive, a listener's reaction, or what causes religious belief.

The review looked for comfort, fear, shame, identity, and ridicule. A single use can contain several of them. The strongest descriptive contrast within these labels is not “emotion versus reason”: comfort appears in 27 theist uses and one non-theist use, while ridicule appears in 34 and 41. These are overlapping counts, not exclusive parts of a total.

The non-theist percentage is higher in this selected category, even though the theist side has the larger absolute number. Neither fact establishes which side is more emotional in all its speech. The denominator is detected unsupported slogans, not everything said in the debates.

Most importantly, wording is not a measurement of causes. A comforting statement might reflect a belief, reinforce it, or persuade no one. A derisive statement might rally an audience or alienate it. This study observes neither those effects nor the speaker's private motive.

The original hypothesis contains two claims: a difference in debate behavior, and an explanation involving emotional enforcement of belief. Evidence for the first does not automatically become evidence for the second.

12. What this stronger study still cannot tell us

Direct reading improves the match between the question and the evidence. It does not remove judgment from the measurement. One fresh GPT-6 Astra reading at Light effort was used for every debate. The reader saw the retained transcript and source-based speaker clues, not the old scores, old slogan findings, or an expected direction of difference.

Speaker labels were varied independently of the expected result, but names and recognizable positions remained in the source. The reading was therefore not blind to identity. The same model can apply a preference or misunderstanding consistently across many debates. There was no full second-reader check, and no measure of agreement between independent readers.

Source fingerprints, complete event coverage, quotation locations, and arithmetic were checked. Printed examples received additional contextual inspection. These checks establish traceability to the saved text; they do not establish that every caption is correct or that every interpretation is beyond dispute. This was not a new review of all the recordings by ear.

Three limits that change what the conclusion can mean

Selected debates are not the general population. The 187 comparisons come from the site's fixed religious-versus-skeptical selection. Some contain wider political, cultural, or moral exchanges. The labels identify the sides in the study, not a claim that every counted sentence is specifically theological. The narrower-topic check reduces this concern but cannot remove selection bias.

Missing or mixed speech can matter. About 3.1% of source words were not safely assigned to a side. They were not treated as evidence that no slogans occurred. Missed uses in either assigned or unassigned text remain possible; the repeat-draw ranges do not cover that problem.

Counts are not scores, and frequency is not importance. One strategically important slogan may matter more to a debate than several minor ones. The study counts uses without assigning their impact. It does not show that slogan use caused the score gap discussed in another paper. All original debate marks and rankings remain unchanged.

A protected move can also occur inside an extended argument that does not qualify as a slogan. Such cases were kept separately, not added to the 77 and 19. That prevents a study of slogans from quietly becoming a count of every weak or evasive argument.

13. What readers can do with the finding

Use the result to ask sharper questions, not to prejudge the next speaker. A group average does not tell you whether the particular sentence in front of you is supported.

A short test for either side

- **State the actual claim.** Is the speaker saying that something is possible, probable, necessary, or the only option? Those are different commitments.
- **Find the missing reason.** What must be true to get from the reasons offered to that exact conclusion? A historical influence, for example, does not establish exclusive necessity.
- **Check the whole exchange.** Was the reason supplied earlier? Did the speaker later qualify the statement? Do not turn a short summary into a slogan by removing its support.
- **Ask what could count against it.** This need not be one laboratory experiment. It could be a contradiction, a failed prediction, a rival explanation, or contrary historical evidence.
- **Inspect the response to criticism.** Does the speaker address the objection, or explain why no objection from this critic could ever count?

A good challenge: “What reason supports that extra step, and what would make you reconsider it?” The question tests the argument without pretending to know the speaker's motives.

What the next comprehensive review should preserve

Keep direct slogan fields separate from debate scores: the short quotation, surrounding passage, speaker, missing reason, any specific protection from criticism, and emotional wording. Record uncertain and non-slogan cases as well as positive findings. Apply the same rule to both sides and retain a clear word denominator.

When a major new model warrants a new edition, rerun a common set under one fixed set of instructions. Use a small, balanced second-reader sample to find disagreements in interpretation before expanding the study. That is a future improvement, not a claim that it happened here.

The emotional-cause hypothesis needs different evidence. A future study could test whether an affirmation changes acceptance when the supporting reasons are held constant, and compare it with a reasoned version of the same claim. Research on how beliefs are maintained would also need observations beyond public debates. More slogan counting alone cannot answer those questions.

14. Conclusion: the argument the evidence supports

The hypothesis gains support in a narrower, more concrete form. The conclusion follows from the following reasons and findings.

1. If an affirmation replaces a needed reason and rules out criticism in advance, it obstructs rational correction. That is the defect tested here. Religious content, emotional force, and lack of a simple experiment are not sufficient conditions.

2. The same direct review found more such uses on the theist side. Across 187 debates, there were 77 protected theist slogans and 19 protected non-theist slogans. With speech length allowed for and debates weighted equally, the average difference was +0.40 uses per 10,000 words. The middle 95% of repeat-draw averages stayed above zero, from +0.11 to +0.74.

3. The difference is not just a Lennox effect or the product of the two largest theist debate counts. It remains in the planned selections and becomes smaller, but stays positive in the repeat-draw check, when those two debates are removed. At the same time, 144 debates have no protected slogan detected on either side. The finding describes an uneven archive-level pattern, not a universal theist habit.

4. The broader slogan claim is less secure. Unsupported slogans are more numerous on the theist side, 128 versus 82, but their average-difference range crosses zero. The paper cannot give that broader claim the same strength as the narrower protection finding.

5. The proposed emotional explanation was not tested as a cause. Emotional wording is common in detected slogans on both sides. The study does not observe what produces religious belief or keeps it in place. It therefore cannot establish that emotional enforcement caused the recorded difference.

Therefore: in these 187 reviewed debates, the evidence supports the claim that theist speakers more often use unsupported slogans that shut out criticism. It provides less secure support for a general excess of unsupported slogans and no direct test of the proposed emotional cause. This is a finding from one AI reader of selected transcripts, not a verdict on all believers and skeptics.

The practical standard is the same for everyone: an affirmation earns its place by the reasons that support it and the objections it can answer. A claim does not become better supported by making disagreement impossible.

That conclusion is narrower than the starting hypothesis, but it is also more useful. It names the observed behavior, shows its extent, identifies the missing argumentative steps, and tells readers how to challenge them without relying on a label for the person speaking.

Evidence and source guide

Report date: September 5, 2026. **Archive date:** September 4, 2026. **Scope:** 187 relevant debates from 253 assessments, using all complete retained transcripts in that selection. “Complete” describes the retained source reviewed, not a guarantee about the original upload or caption accuracy.

Exact speech accounting

Source-word category	Caption words
Attributed to theist speakers	1,433,125
Attributed to non-theist speakers	1,439,245
Unclear or mixed speaker	107,503
Excluded material	517,106
Total retained source	3,496,979

The two attributed totals are the denominators for pooled rates. The main averages instead use each side's words within each debate, then give the 187 debates equal weight. Repeat-draw ranges use 20,000 whole-debate draws and a fixed random seed, 20260904, so the published calculation can be reproduced.

How to inspect the evidence

Every printed example links to its YouTube passage. Short quotations preserve the retained captions; this is not a fresh independent transcription. The study does not independently certify historical claims or quotations attributed by speakers to other authors.

The [study folder](#) and [method](#) contain the rules, frozen source inventory, original review outputs, separate editorial corrections, calculations, and this manuscript. The [debate-level data](#) list counts and word totals for every comparison. The [source-linked candidate records](#) retain clear, uncertain, and excluded findings, not just examples favorable to the hypothesis.

Why the old percentages are not repeated

The earlier edition reported 28.6% and 7.9% of eligible assessed moves meeting a combined warning rule in 146 debates. Those were indirect risk measures, not direct slogan frequencies. This edition changes both what is observed and what counts as the denominator. Comparing the old percentages with the new rates would not show improvement or deterioration. The new direct study supersedes the old slogan evidence; it does not revise the other six papers or the debate scores.