ASTRA-ERA EDITION / PLAIN-LANGUAGE REVISION

Beyond the Fallacy Count

What named mistakes can tell us, what a missing label cannot tell us, and why the rest of the argument matters.

61.7%
lower-scoring sides with no named-fallacy tag overall

14.5%
the same rate in the later assessment group

74.8%
comparable untagged losses behind in five or six areas

Most lower-scoring sides have no fallacy label in the combined archive, but the opposite is true in later assessments. The lasting lesson is that a label alone does not settle how well a side has argued.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

A missing fallacy label is not a clean bill of reasoning health. In the 253-assessment archive, 150 of 243 lower-scoring sides in non-tied debates **- 61.7% -** have no named-fallacy tag. But the archive-wide majority conceals a dramatic split: the rate is **80.8% in the earlier assessment process and only 14.5% in the later one**. The older paper's headline, “Debates Are Usually Lost Without a Named Fallacy,” remains numerically true for the combined archive, but is misleading if read as a stable rule of debate or of the current assessment process.

The deeper finding survives. Among 147 comparable lower-scoring sides without a named-fallacy tag, **110 trail on five or six of the six scoring areas**. That is **74.8%**. Their losses are commonly distributed across evidence, consistent reasoning, replies, precision, and proportionate confidence. A side can be relevant, understandable, and free of an accepted named label while still leaving much of its case unsupported.

Nor does a label determine defeat. A higher-scoring side carries at least one named-fallacy tag in **51 of the 243 decisive assessments**. In 11 cases, only the higher-scoring side carries such a tag. That does not make the flagged error harmless; it shows that one local diagnosis does not exhaust the comparison of the full performances.

This edition therefore moves beyond a contest over whether most losses contain a named error. It asks what tags record, what they omit, how their use changes between assessment processes, and what the ordinary scoring areas reveal. **The most useful habit is to inspect the unmet burden, not merely search for a familiar fallacy name.**

The paper draws on all **253 published assessments and 5,492 public reasoning moves** for tag counts, with detailed scoring area analysis restricted to the 237 comparable locked assessments. It is an analysis of recorded labels and scores, not an independent determination that every untagged passage is fallacy-free or every accepted tag is correct.

Reader’s guide

Executive summary 2

1. Why the title has changed 4

2. What exactly is being counted? 4

3. The archive-wide count - and its reversal 5

4. How a majority can disappear without any old result changing 6

5. The four possible label patterns 6

6. What the untagged lower sides are missing 7

7. Ordinary weaknesses can be decisive in combination 8

8. Case study: stating an “ought” is not justifying it 9

9. Case study: calling something an axiom does not complete the proof 9

10. The pattern is not confined to older religious-versus-skeptical contests 10

11. What the six labels can and cannot tell us 10

12. A more useful way to inspect a debate 12

13. Conclusion: check the reasoning, not just the labels 12

Methods and sources 13

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. All results refer to the September 4 snapshot, not later additions to the site.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. Why the title has changed

The earlier paper made a useful corrective point: debates are often lost through cumulative weaknesses rather than a single spectacular mistake. Its headline emphasized the majority of lower-scoring sides with no named-fallacy tag. The enlarged archive now makes the headline's dependence on the assessment process impossible to ignore.

Across the whole snapshot, that majority remains. Yet within the later process, most lower-scoring sides do carry a fallacy label. A paper that continued to present the combined majority without putting the reversal near the front would invite readers to mistake the makeup of the archive for a universal feature of reasoning.

The new title, **Beyond the Fallacy Count**, preserves the original subject while strengthening its conclusion. Named errors are valuable local diagnoses. Their presence is not a complete measure of a side's performance, and their absence is not evidence that the side's reasoning is sound. Those claims remain defensible even when the archive-wide percentage changes.

This is also a useful test of the research series itself. Recalculation should be allowed to change the framing, not merely refresh the numerals beneath an inherited conclusion. The September 4 Astra-era edition treats the expanded evidence as an opportunity to improve the question as well as the answer.

2. What exactly is being counted?

A *named-fallacy tag* is a public move label whose type is explicitly “fallacy.” The analysis does not count a bias tag as a fallacy, infer a tag from critical prose, or search for any negative word in the assessment. Those choices keep the measurement reproducible.

“Lower-scoring side” means the side with the lower final performance score. It does not mean the side whose worldview is false, the side that lost an audience vote, or the side whose speaker admitted defeat. The ten tied assessments have no lower-scoring side and are excluded from that total being counted.

“Without a named fallacy” means **without an accepted named-fallacy tag in the public record**. It does not mean that a fresh analyst could find no fallacy in the transcript. The distinction is essential because label policy, evidence-based thresholds, reviewer attention, and the set of available labels can all affect whether a weakness receives a tag.

The archive contains six named-fallacy labels in this snapshot: equivocation, red herring, begging the question, argument from ignorance, appeal to authority, and special pleading. These are mainly informal patterns of reasoning. They should not all be described as formal logical fallacies, which concern invalid structures of reasoning step. A named criticism can be serious without belonging to that narrower technical category.

3. The archive-wide count - and its reversal

There are **243 decisive assessments**: 253 total minus ten ties. Of their lower-scoring sides, 150 carry no named-fallacy tag and 93 carry at least one. In 139 decisive debates, neither side carries one. The basic archive-level observation is therefore correct: a majority of recorded losses occur without an accepted named-fallacy label on the lower-scoring side.

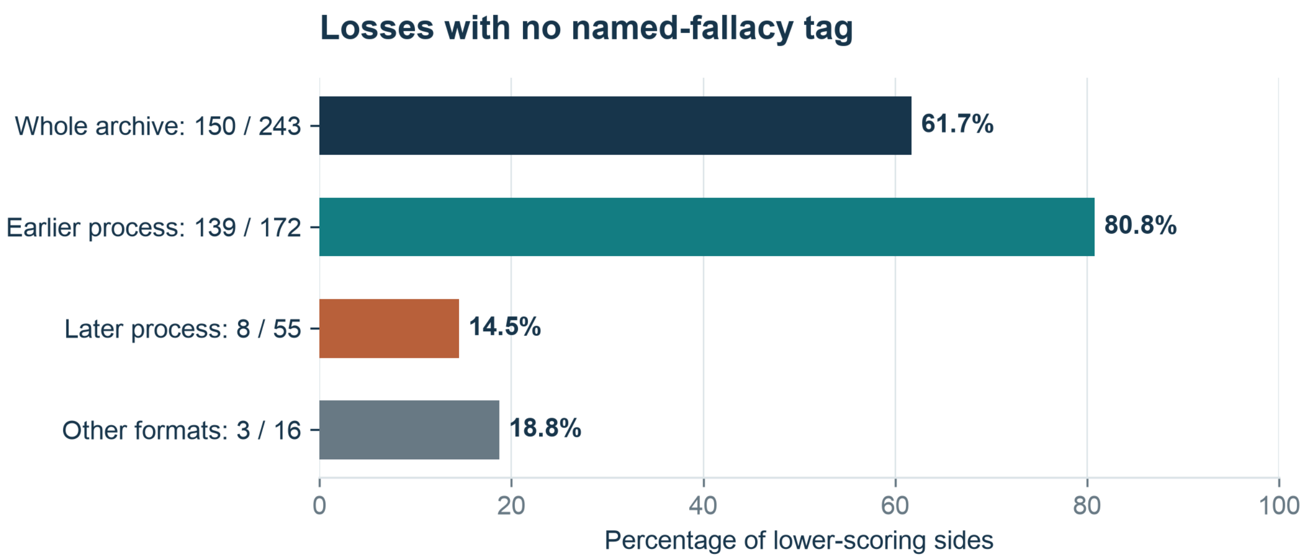

Figure 1. The combined archive and its assessment-process groups tell different stories. The later-process majority is the reverse of the earlier-process majority. Ties are excluded because they have no lower-scoring side.

How to read this graph. The dark bar combines the archive; teal is earlier, rust later, and gray other formats. Each printed fraction counts untagged lower-scoring sides over all lower-scoring sides in that group. Bar length is the percentage. Ten ties are excluded. The overall bar contains the other groups and is not an additional group.

The earlier process contributes **139 untagged losses out of 172**, or 80.8%. The later process contributes **eight out of 55**, or 14.5%. Other-format assessments contribute three out of 16, or 18.8%. Combining these groups produces the overall 61.7%.

That combined percentage is heavily influenced by the earlier records because they supply most of the decisive assessments. It should not be presented as though 243 examples were generated by one uniform label practice. The two assessment groups are described in detail in paper six; the differences in tags are among the strongest reasons for keeping them visible.

The contrast does not by itself prove that later debates contain more bad reasoning or that the newer reviewers are more accurate. More frequent labels could reflect different selection, more detailed records, a broader use of the available labels, changed standards for accepting tags, or real differences in the arguments. The current analysis observes the contrast but does not isolate its cause.

4. How a majority can disappear without any old result changing

The archive-wide rate is a weighted mixture. Each group's rate contributes in proportion to its number of decisive debates. That is why the earlier process can dominate the combined result even when later assessments show the opposite pattern.

An explicit calculation makes the issue concrete. Multiply 80.8% by the earlier group's share of the 243 decisive debates, 14.5% by the later share, and 18.8% by the other-format share. Adding those products reproduces the overall 61.7%, apart from rounding. The combined statistic is not wrong. It simply answers a different question from “What usually happens in later assessments?”

> **A hypothetical update, not a forecast:** suppose 110 further decisive assessments were added, with 16 untagged losses - the same 8-in-55 rate observed in the later group. The combined total would become 166 untagged losses out of 353, or 47.0%. The archive-wide majority would reverse even if no existing score or label changed.

This example shows why a headline based on a combined majority can be fragile. It depends not only on how arguments are assessed but on how many records each process contributes. The important question for a durable report is which conclusion would remain informative after that mixture changes.

Here the durable conclusion is that label counts are incomplete and process-sensitive indicators of quality. The exact proportion of untagged losses can move substantially while the underlying need to inspect support, reasoning step, and answering the other side remains.

5. The four possible label patterns

The relationship between tags and outcomes can be summarized without treating a tag as a verdict. Among the 243 decisive assessments, the four possible combinations are as follows.

Public named-fallacy pattern	Debates	What the count shows
Neither side tagged	139	Scores can differ without a named label on either side.
Only the lower-scoring side tagged	53	A label can accompany a lower score without explaining the whole gap.
Only the higher-scoring side tagged	11	A local tagged error does not necessarily decide the overall result.
Both sides tagged	40	Finding an error on one side does not settle the comparison.

A higher-scoring side has at least one tag in 11 + 40 = **51 debates**, or 21.0% of decisive assessments. A lower-scoring side has one in 53 + 40 = **93 debates**, or 38.3%. Tags are more common on lower-scoring sides, as one might expect, but they are neither necessary nor sufficient for a lower overall score.

The word "sufficient" needs care here. It means that the presence of a tag does not guarantee the lower score. It does not mean the error is acceptable or that one can ignore it because the side happened to win. A good overall performance can contain a defective move, just as a weak overall performance can contain several excellent ones.

Nor is this table an experiment in the causal effect of committing a fallacy. The labels and scores are judgments about the same underlying material, sometimes produced within related assessment processes. Their association may reflect real reasoning defects, correlated judging decisions, or both.

6. What the untagged lower sides are missing

Detailed scoring area comparisons are available for **147** lower-scoring sides with no named-fallacy tag. The other three untagged losses come from formats not included in the comparable locked analysis. For each eligible debate, we ask how many of the six officially weighted side-scoring area averages are lower on the losing side.

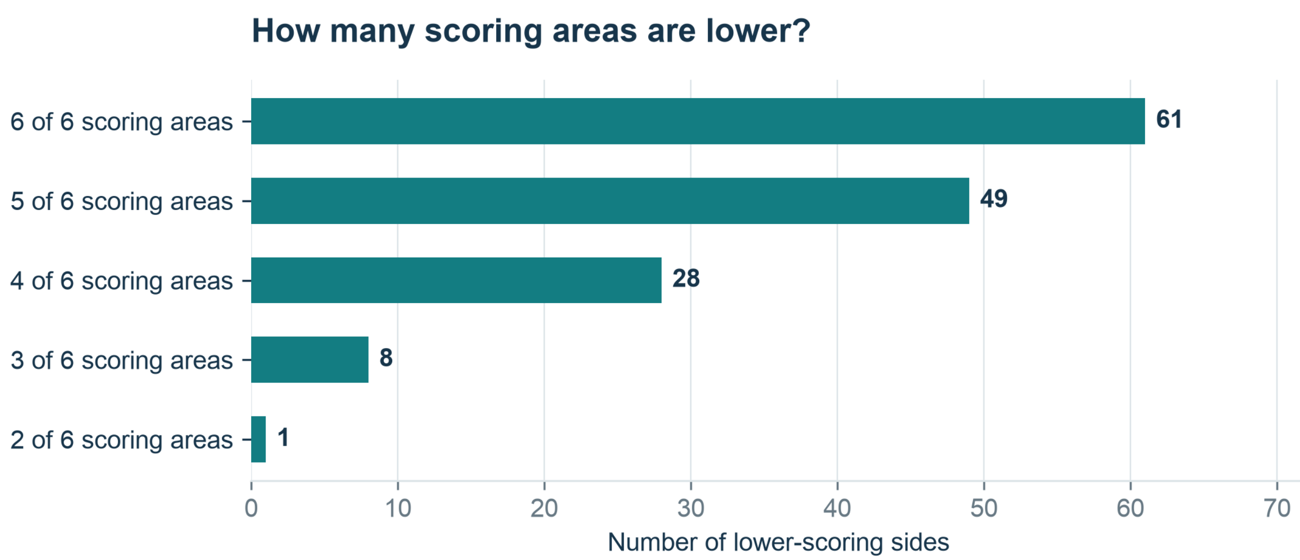

Figure 2. Most comparable untagged losses trail across several scoring areas, rather than only one. The six scoring areas are related aspects of the same assessed performance, not six independent experiments.

How to read this graph. Each bar counts lower-scoring sides with no named-fallacy tag that trail in exactly the stated number of scoring areas. For example, '5 of 6' means five area averages are lower than the opponent's. The bars sum to 147. No eligible case trails in only zero or one area.

Sixty-one trail on all six scoring areas, 49 on five, 28 on four, eight on three, and one on two. Thus **110 of 147, or 74.8%, trail on at least five scoring areas**. The pattern is usually broad rather than confined to a single conspicuous defect.

This result should not be oversold as an independent validation of the losing verdict. The final scores are built from the scoring areas, so selecting the lower-scoring side already selects a performance with an unfavorable weighted combination. The informative feature is how widely the deficits are distributed, not the mere fact that some deficits exist.

The result is also weighted heavily toward earlier-process records: 139 of the 147 untagged comparable losses come from that process. Only eight are later. Five of those eight trail on five or six scoring areas, but that is far too small a group to establish a precise later-process rate. The overall breadth finding should retain its total being counted and process context.

7. Ordinary weaknesses can be decisive in combination

The largest mean gaps in untagged comparable losses are 8.28 scoring area points for care and fairness, 7.84 for support, 6.95 for consistent reasoning, and 6.79 for answering the other side. Precision differs by 5.77 points. Relevance and burden differ by only 0.98 points on average.

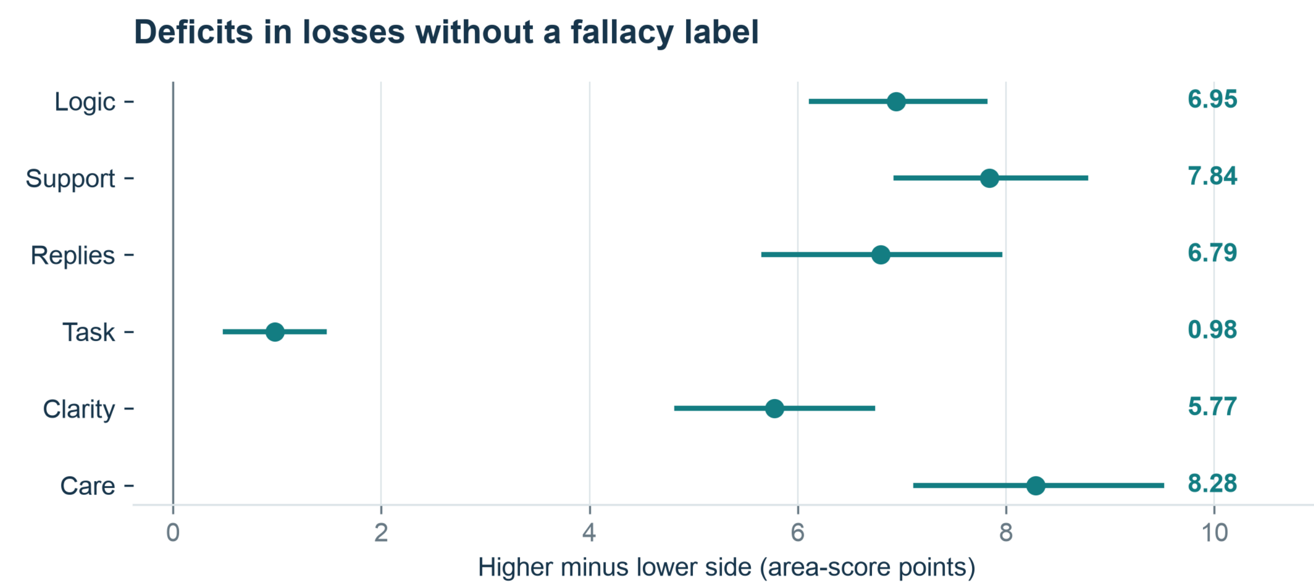

Figure 3. These are raw scoring area-point gaps between the higher and lower side, averaged across 147 untagged comparable losses. They are not each scoring area's weighted contribution to the final score. Relevance is often relatively close even when support and replies are not.

How to read this graph. Each dot is the average higher-side-minus-lower-side difference among 147 comparable untagged losses. Lines hold the middle 95% from repeated draws of those debates. Right of zero means a deficit for the lower side. These are area-score points before weighting, not parts that can be added to the overall gap.
Area key: Logic = consistent reasoning; Support = evidence and reasons; Replies = answering the other side; Task = staying on the question and doing what the claim requires; Clarity = clear, exact wording; Care = fitting confidence to support and treating the other side fairly.

That distribution helps explain the subjective experience of watching a debate in which no single obvious error seems to settle matters. A speaker may remain recognizably on topic and state an intelligible case. Yet the case may rely on underdeveloped starting claims, make stronger claims than its reasons support, omit important alternatives, or answer only part of an objection. The deficiencies accumulate.

Consider a purely illustrative scoring comparison. If a side trails by eight points in consistent reasoning, ten in evidence, and six in answering the other side, those gaps contribute 2.00, 2.00, and 1.20 points to the overall difference under weights of 25%, 20%, and 20%. That is already **5.20 overall points** before the remaining scoring areas and rounding are considered. No named fallacy is needed for the arithmetic to produce a substantial loss.

The example is not a measured individual debate. It demonstrates why an error-hunting approach can miss the structure of the scoring system. The final score evaluates how well a reasoning route is constructed and defended, not merely whether a contribution triggers a familiar label.

8. Case study: stating an “ought” is not justifying it

In Slick-Clifton, debate 127, Matt Slick's side scores 70 and Scott Clifton's 89. Neither side carries a named-fallacy tag. The losing side trails on five of the six scoring area averages, with particularly large differences in evidence, consistent reasoning, and answering the other side.

One assessed move argues that Christian discourse contains prescriptions while an atheistic system cannot derive binding obligations from descriptions alone. The assessment credits it for recognizing the relevant issue: an account of morality needs more than facts about what happens. The limitation is equally specific: **possessing a prescription is not the same thing as justifying it**. The bridge from divine character, authority, or human ends to obligation still needs to be defended.

The public move has no fallacy or bias tag, yet its critique identifies a substantial reasoning gap. A reader gains more by asking what would justify the prescription than by insisting on finding the correct label for the omission. The absence of a label does not make the gap disappear.

A second move criticizes a crude preference-based account of morality. The assessment notes that the opposing position included additional considerations - impartiality, welfare, and rules - that the reply did not fully address. Again, the issue is not merely the name of the mistake. It is whether the actual alternative has been represented and answered.

9. Case study: calling something an axiom does not complete the proof

In Fischer-Dillahunty, debate 141, the scores are 68 and 86. Neither side carries a named-fallacy tag, and the lower side trails on all six scoring area averages.

A recorded move presents a first cause as an axiomatic requirement across conceivable worlds. The assessment recognizes the intended universality but asks how imagining worlds with first causes establishes that all possible worlds require one. It also notes that conceivability, possibility, necessity, and axiomatic status are not interchangeable concepts.

Those distinctions matter even for readers who never use the word “modal,” the technical adjective for reasoning about possibility and necessity. A statement about what one can imagine is not automatically a statement about what must be true. If the argument depends on that upgrade, the speaker has to justify it.

Another contribution describes resurrection as historically substantiated without supplying the historical particulars needed for that assessment. Identifying a concrete candidate event is an improvement over speaking only about miracles in general. It does not complete the evidence-based task. These are recognizable strengths and weaknesses within an untagged move, not a binary choice between flawless reasoning and a named fallacy.

10. The pattern is not confined to older religious-versus-skeptical contests

The later-process Swinburne-Goff debate, number 240, ends 76 to 83. Both sides defend purposive explanations, so it is excluded from the papers' main religious-versus-skeptical comparison. It is included here because the question concerns argument quality across the archive.

Swinburne's lower-scoring side has no named-fallacy tag and trails on all six scoring area averages. One assessed reply grants that Goff's alternatives fit relevant facts, then asks whether their tailored structures need further explanation. The assessment recognizes the legitimate comparative concern but notes that the extra complexity is not specified and that parallel questions about divine attributes remain insufficiently addressed.

Another move argues that a life-producing multiverse would need a suitably structured generator. That is a relevant conditional challenge. The weakness is the claim that viable generators are themselves comparably fine-tuned without identifying or comparing the relevant models. A coherent question can remain under-supported as a positive conclusion.

The eight later untagged losses also include close one-point results in discussions of souls, purgatory, moral realism, and resurrection. Not every untagged loss is a dramatic collapse. The category ranges from near ties to a seven-point difference. That variation is another reason not to treat the absence of a tag as a substantive rating of the whole performance.

11. What the six labels can and cannot tell us

Across 5,492 public moves, **345 carry at least one named-fallacy tag**, or 6.3%. There are **354 label instances**, because some moves have more than one. Counts of moves and counts of labels should not be confused.

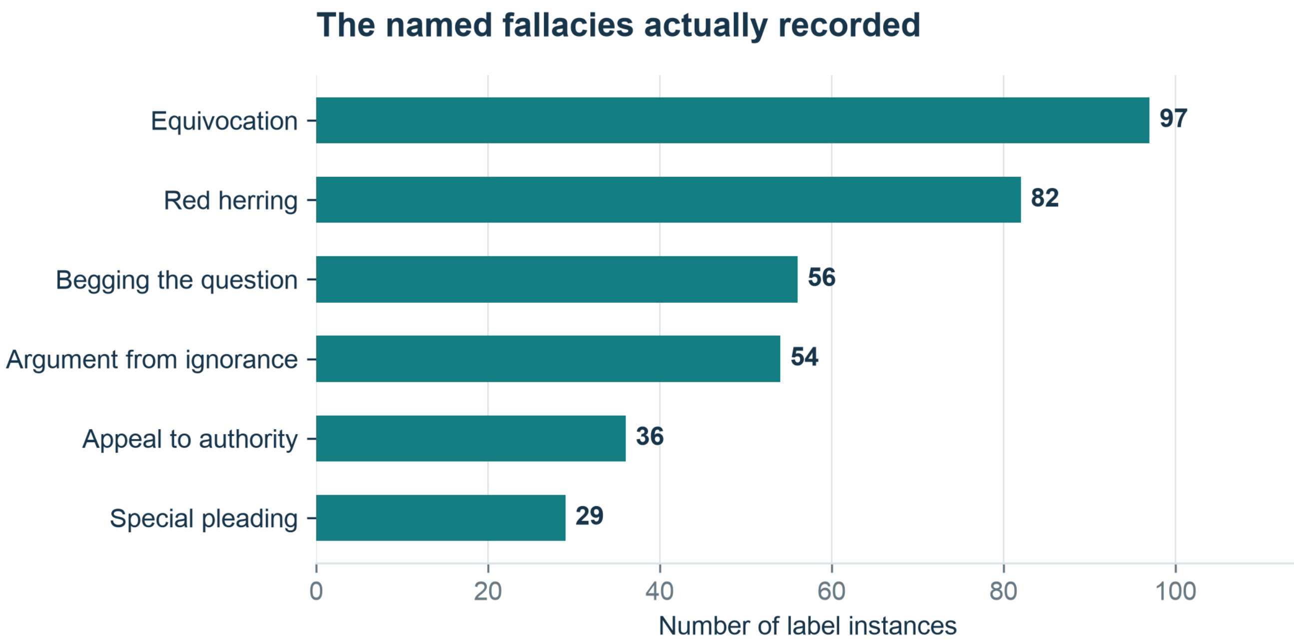

Figure 4. The archive records six named-fallacy labels. Equivocation and red herring are the most frequent. This is an inventory of accepted public labels, not an exhaustive topic grouping or a count of every possible error in the transcripts.

How to read this graph. Bar length counts accepted public uses of each named-fallacy label. One move can have more than one label: 354 label uses occur on 345 moves. Teal has no side or belief meaning in this chart. Missing labels do not show that no other errors occurred.

Equivocation draws attention to a word or concept shifting meaning within an argument. A red herring points to diversion from the issue that needs to be addressed. Begging the question concerns dependence on what is at issue. Argument from ignorance concerns an unsupported move from what is not known to a positive conclusion. The remaining labels concern problematic appeals to authority and inconsistent treatment of comparable cases.

Each can help a reader locate a defect. But naming a pattern is not the same thing as demonstrating it. A proper diagnosis must identify the relevant claim, the unsupported reasoning step or mismatch, and why the apparent weakness is not a reasonable move in the actual context. An appeal to a qualified expert, for example, is not automatically a fallacy.

Labels can also become rhetorical weapons. Announcing “special pleading” without showing that the cases are relevantly alike may be less informative than the untagged criticism it replaces. A reader who learns the names without learning the reasoning can become more confident while becoming no more accurate.

The best use of the site's reference links is therefore educational rather than prosecutorial: learn the pattern, return to the exact contribution, and ask whether the diagnosis fits. The public critique and source material should do the work needed to support a claim.

12. A more useful way to inspect a debate

Start with the conclusion the side needs to establish. Then identify its supporting route: which starting claims matter, which reasoning steps connect them, and which alternatives could weaken it? A named fallacy may reveal a broken link, but the route can also remain incomplete because a starting claim is merely asserted or an alternative is never compared.

Next, separate the quality of a reply from its relevance. A response can concern the same general subject while leaving the objection untouched. “Religion gives meaning” may be relevant to a discussion of religious life without answering whether the supernatural explanation is true. “Science has limits” may be relevant to an ultimate-explanation debate without establishing which proposed explanation should replace it.

Then check confidence. Does the speaker distinguish what has been established from what remains possible or speculative? An appropriately limited conclusion can be a strength. A categorical statement supplied with the same evidence may be weaker because it undertakes more than the evidence supports.

Finally, inspect local errors in proportion to their role. A peripheral lapse need not defeat an otherwise strong case; a central unsupported starting claim can undermine a whole route even without a named tag. Importance, coverage, and the relation between claims matter. Counting mistakes without considering those features can misrepresent the comparative performance.

13. Conclusion: check the reasoning, not just the labels

Argument one: the archive-wide majority is not a stable rule

1. A result should not be treated as a general rule when it changes sharply between major parts of the data.

2. Overall, **150 of 243 lower-scoring sides (61.7%)** have no named-fallacy tag. But the rate is **80.8% in earlier assessments** and **14.5% in later assessments**.

3. A missing tag records what the assessment labeled. It does not prove that the passage contains no fallacy. The different groups also contain different debates, so the figures alone do not identify why tagging changed.

Therefore: “debates are usually lost without a named fallacy” describes the combined snapshot, but not a dependable rule for later assessments or debate in general. The split belongs beside the headline, not in a footnote.

Argument two: a label is neither required for a loss nor enough to explain one

1. If a label fully settled argument quality, untagged sides would not regularly fall behind across several scoring areas, and a tagged side would not regularly score higher overall.

2. Among 147 comparable untagged losses, **110 (74.8%)** trail on five or six scoring areas. Weak support, unfinished reasoning, and incomplete replies can add up without a familiar fallacy name.

3. In **51 of 243 decisive assessments**, the higher-scoring side has at least one fallacy tag. A local mistake can matter without making that side's complete performance worse than its opponent's.

Therefore: inspect the named mistakes, but also inspect the rest of the argument. No tag means no recorded tag - not sound reasoning. One tag means a specific recorded criticism - not automatic defeat.

Ask what the speaker has actually established and what still lacks a reason. A fallacy name is useful when it helps answer that question. It is not a replacement for the answer.

Methods and sources

The source snapshot is revision **76d006b37**, dated September 4, 2026. Public tag counts include all 253 assessments and 5,492 reasoning cards. Cards stored as arrays in multi-speaker exchanges are flattened without duplicating them. Only tags explicitly typed as fallacies enter the named-fallacy counts; bias tags are kept separate. Ten tied assessments are excluded from lower-side rate.

The detailed scoring area analysis uses 147 lower-scoring sides without fallacy tags from the 237 comparable locked assessments. Side-scoring area means preserve official section and move-importance weights. “Trails on a dimension” means a strictly lower weighted mean, allowing a small numerical tolerance for floating-point arithmetic. Scoring area intervals use 20,000 whole-debate replacement resamples. The main archive counts are complete counts of this snapshot, not estimates requiring a sampling interval to establish their arithmetic.

The cohort comparison is descriptive. It does not prove that the process caused the label difference, and the sparse tags in earlier records prevent interpreting absence as a uniform statement about the underlying reasoning. Correlated scoring areas, selective source collection, recurring speakers, and systematic judging errors limit broader conclusions.

Calculations, label inventories, the lower-side table, source hashes, and chart data are in the September 4 analysis directory. The Backend library links the companion scale audit, which explains why label rates should remain separated by process. This paper reanalyzes existing judgments and does not independently relabel the complete transcript archive.