

Beyond the Fallacy Count

Cumulative weaknesses, uneven annotation, and the limits of treating debate as a hunt for named errors.

61.7%

lower-scoring sides with
no named-fallacy tag overall

14.5%

the same rate in the
later assessment process

74.8%

untagged comparable losses that
trail on five or six dimensions

The archive-wide majority is real, but it reverses in newer assessments. What survives that reversal is the deeper conclusion: named fallacies neither exhaust argument quality nor determine who scores lower.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

A missing fallacy label is not a clean bill of argumentative health. In the 253-assessment archive, 150 of 243 lower-scoring sides in non-tied debates - **61.7%** - have no named-fallacy tag. But the archive-wide majority conceals a dramatic split: the rate is **80.8% in the earlier assessment process and only 14.5% in the later one**. The older paper's headline, "Debates Are Usually Lost Without a Named Fallacy," remains numerically true for the combined archive, but is misleading if read as a stable rule of debate or of the current assessment process.

The deeper finding survives. Among 147 comparable lower-scoring sides without a named-fallacy tag, **110 trail on five or six of the six scoring dimensions**. That is **74.8%**. Their losses are commonly distributed across evidence, coherence, replies, precision, and proportionate confidence. A side can be relevant, understandable, and free of an accepted named label while still leaving much of its case unsupported.

Nor does a label determine defeat. A higher-scoring side carries at least one named-fallacy tag in **51 of the 243 decisive assessments**. In 11 cases, only the higher-scoring side carries such a tag. That does not make the flagged error harmless; it shows that one local diagnosis does not exhaust the comparison of the full performances.

This edition therefore moves beyond a contest over whether most losses contain a named error. It asks what tags record, what they omit, how their use changes between assessment processes, and what the ordinary scoring dimensions reveal. **The most useful habit is to inspect the unmet burden, not merely search for a familiar fallacy name.**

The paper draws on all **253 published assessments and 5,492 public argumentative moves** for tag counts, with detailed dimension analysis restricted to the 237 comparable locked assessments. It is an analysis of recorded labels and scores, not an independent determination that every untagged passage is fallacy-free or every accepted tag is correct.

Reader's guide

Executive summary	2
1. Why the title has changed	4
2. What exactly is being counted?	4
3. The archive-wide count - and its reversal	5
4. How a majority can disappear without any old result changing	6
5. The four possible label patterns	6
6. What the untagged lower sides are missing	7
7. Ordinary weaknesses can be decisive in combination	8
8. Case study: stating an “ought” is not justifying it	9
9. Case study: calling something an axiom does not complete the proof	9
10. The pattern is not confined to older religious-versus-skeptical contests	10
11. What the six labels can and cannot tell us	10
12. A more useful way to inspect a debate	11
13. Conclusion: reasoning is larger than its error labels	12
Methods and sources	13

Reading the evidence

Start with the summary and figures for the main argument. The worked examples explain the statistical ideas; the final methods and sources show how the numbers can be checked. Counts and results refer to the frozen September 4 archive, not to later additions to the site.

1. Why the title has changed

The earlier paper made a useful corrective point: debates are often lost through cumulative weaknesses rather than a single spectacular mistake. Its headline emphasized the majority of lower-scoring sides with no named-fallacy tag. The enlarged archive now makes the headline's dependence on the assessment process impossible to ignore.

Across the whole snapshot, that majority remains. Yet within the later process, most lower-scoring sides do carry a fallacy label. A paper that continued to present the pooled majority without putting the reversal near the front would invite readers to mistake the makeup of the archive for a universal feature of reasoning.

The new title, **Beyond the Fallacy Count**, preserves the original subject while strengthening its conclusion. Named errors are valuable local diagnoses. Their presence is not a complete measure of a side's performance, and their absence is not evidence that the side's reasoning is sound. Those claims remain defensible even when the archive-wide percentage changes.

This is also a useful test of the research series itself. Recalculation should be allowed to change the framing, not merely refresh the numerals beneath an inherited conclusion. The September 4 Astra-era edition treats the expanded evidence as an opportunity to improve the question as well as the answer.

2. What exactly is being counted?

A *named-fallacy tag* is a public move annotation whose type is explicitly “fallacy.” The analysis does not count a bias tag as a fallacy, infer a tag from critical prose, or search for any negative word in the assessment. Those choices keep the measurement reproducible.

“Lower-scoring side” means the side with the lower final performance score. It does not mean the side whose worldview is false, the side that lost an audience vote, or the side whose speaker admitted defeat. The ten tied assessments have no lower-scoring side and are excluded from that denominator.

“Without a named fallacy” means **without an accepted named-fallacy tag in the public record**. It does not mean that a fresh analyst could find no fallacy in the transcript. The distinction is essential because label policy, evidential thresholds, reviewer attention, and the set of available labels can all affect whether a weakness receives a tag.

The archive contains six named-fallacy labels in this snapshot: equivocation, red herring, begging the question, argument from ignorance, appeal to authority, and special pleading. These are mainly informal patterns of reasoning. They should not all be described as formal logical fallacies, which concern invalid structures of inference. A named criticism can be serious without belonging to that narrower technical category.

3. The archive-wide count - and its reversal

There are **243 decisive assessments**: 253 total minus ten ties. Of their lower-scoring sides, 150 carry no named-fallacy tag and 93 carry at least one. In 139 decisive debates, neither side carries one. The basic archive-level observation is therefore correct: a majority of recorded losses occur without an accepted named-fallacy label on the lower-scoring side.




Figure 1. The combined archive and its assessment-process groups tell different stories. The later-process majority is the reverse of the earlier-process majority. Ties are excluded because they have no lower-scoring side.

Scope: 243 decisive assessments; 10 ties excluded.

The earlier process contributes **139 untagged losses out of 172**, or 80.8%. The later process contributes **eight out of 55**, or 14.5%. Other-format assessments contribute three out of 16, or 18.8%. Combining these groups produces the overall 61.7%.

That pooled percentage is heavily influenced by the earlier records because they supply most of the decisive assessments. It should not be presented as though 243 examples were generated by one uniform annotation practice. The two process families are described in detail in paper six; the differences in tags are among the strongest reasons for keeping them visible.

The contrast does not by itself prove that later debates contain more bad reasoning or that the newer reviewers are more accurate. More frequent labels could reflect different selection, more detailed records, a broader use of the available labels, changed standards for accepting tags, or real differences in the arguments. The current analysis observes the contrast but does not isolate its cause.

4. How a majority can disappear without any old result changing

The archive-wide rate is a weighted mixture. Each group's rate contributes in proportion to its number of decisive debates. That is why the earlier process can dominate the combined result even when later assessments show the opposite pattern.

An explicit calculation makes the issue concrete. Multiply 80.8% by the earlier group's share of the 243 decisive debates, 14.5% by the later share, and 18.8% by the other-format share. Adding those products reproduces the overall 61.7%, apart from rounding. The combined statistic is not wrong. It simply answers a different question from “What usually happens in later assessments?”

A hypothetical update, not a forecast: suppose 110 further decisive assessments were added, with 16 untagged losses - the same 8-in-55 rate observed in the later group. The pooled total would become 166 untagged losses out of 353, or 47.0%. The archive-wide majority would reverse even if no existing score or label changed.

This example shows why a headline based on a pooled majority can be fragile. It depends not only on how arguments are assessed but on how many records each process contributes. The important question for a durable report is which conclusion would remain informative after that mixture changes.

Here the durable conclusion is that label counts are incomplete and process-sensitive indicators of quality. The exact proportion of untagged losses can move substantially while the underlying need to inspect support, inference, and responsiveness remains.

5. The four possible label patterns

The relationship between tags and outcomes can be summarized without treating a tag as a verdict. Among the 243 decisive assessments, the four possible combinations are as follows.

Public named-fallacy pattern	Debates	What the count shows
Neither side tagged	139	Scores can differ without a named label on either side.
Only the lower-scoring side tagged	53	A label can accompany a lower score without explaining the whole gap.
Only the higher-scoring side tagged	11	A local tagged error does not necessarily decide the overall result.
Both sides tagged	40	Finding an error on one side does not settle the comparison.

A higher-scoring side has at least one tag in $11 + 40 = \mathbf{51 \text{ debates}}$, or 21.0% of decisive assessments. A lower-scoring side has one in $53 + 40 = \mathbf{93 \text{ debates}}$, or 38.3%. Tags are more common on lower-scoring sides, as one might expect, but they are neither necessary nor sufficient for a lower overall score.

The word “sufficient” needs care here. It means that the presence of a tag does not guarantee the lower score. It does not mean the error is acceptable or that one can ignore it because the side happened to win. A good overall performance can contain a defective move, just as a weak overall performance can contain several excellent ones.

Nor is this table an experiment in the causal effect of committing a fallacy. The labels and scores are judgments about the same underlying material, sometimes produced within related assessment processes. Their association may reflect real argumentative defects, correlated judging decisions, or both.

6. What the untagged lower sides are missing

Detailed dimension comparisons are available for **147** lower-scoring sides with no named-fallacy tag. The other three untagged losses come from formats not included in the comparable locked analysis. For each eligible debate, we ask how many of the six officially weighted side-dimension averages are lower on the losing side.




Figure 2. Most comparable untagged losses trail across several dimensions, rather than only one. The six dimensions are related aspects of the same assessed performance, not six independent experiments.

Scope: 147 comparable losses with no named-fallacy tag.

Sixty-one trail on all six dimensions, 49 on five, 28 on four, eight on three, and one on two. Thus **110 of 147, or 74.8%, trail on at least five dimensions**. The pattern is usually broad rather than confined to a single conspicuous defect.

This result should not be oversold as an independent validation of the losing verdict. The final scores are built from the dimensions, so selecting the lower-scoring side already selects a performance with an

unfavorable weighted combination. The informative feature is how widely the deficits are distributed, not the mere fact that some deficits exist.

The result is also weighted heavily toward earlier-process records: 139 of the 147 untagged comparable losses come from that process. Only eight are later. Five of those eight trail on five or six dimensions, but that is far too small a group to establish a precise later-process rate. The overall breadth finding should retain its denominator and process context.

7. Ordinary weaknesses can be decisive in combination

The largest mean gaps in untagged comparable losses are 8.28 dimension points for confidence/fairness, 7.84 for evidence/warrant, 6.95 for coherence, and 6.79 for responsiveness. Precision differs by 5.77 points. Relevance and burden differ by only 0.98 points on average.




Figure 3. These are raw dimension-point gaps between the higher and lower side, averaged across 147 untagged comparable losses. They are not each dimension's weighted contribution to the final score. Relevance is often relatively close even when support and replies are not.

Scope: 147 locked lower-scoring sides without named-fallacy tags.

That distribution helps explain the subjective experience of watching a debate in which no single obvious error seems to settle matters. A speaker may remain recognizably on topic and state an intelligible case. Yet the case may rely on underdeveloped premises, make stronger claims than its reasons support, omit important alternatives, or answer only part of an objection. The deficiencies accumulate.

Consider a purely illustrative scoring comparison. If a side trails by eight points in coherence, ten in evidence, and six in responsiveness, those gaps contribute 2.00, 2.00, and 1.20 points to the overall difference under weights of 25%, 20%, and 20%. That is already **5.20 overall points** before the remaining dimensions and rounding are considered. No named fallacy is needed for the arithmetic to produce a substantial loss.

The example is not a measured individual debate. It demonstrates why an error-hunting approach can miss the structure of the scoring system. The final score evaluates how well an argumentative route is constructed and defended, not merely whether a contribution triggers a familiar label.

8. Case study: stating an “ought” is not justifying it

In [Slick-Clifton, debate 127](#), Matt Slick's side scores 70 and Scott Clifton's 89. Neither side carries a named-fallacy tag. The losing side trails on five of the six dimension averages, with particularly large differences in evidence, coherence, and responsiveness.

One assessed move argues that Christian discourse contains prescriptions while an atheistic system cannot derive binding obligations from descriptions alone. The assessment credits it for recognizing the relevant issue: an account of morality needs more than facts about what happens. The limitation is equally specific: **possessing a prescription is not the same thing as justifying it**. The bridge from divine character, authority, or human ends to obligation still needs to be defended.

The public move has no fallacy or bias tag, yet its critique identifies a substantial inferential gap. A reader gains more by asking what would justify the prescription than by insisting on finding the correct label for the omission. The absence of a label does not make the gap disappear.

A second move criticizes a crude preference-based account of morality. The assessment notes that the opposing position included additional considerations - impartiality, welfare, and rules - that the reply did not fully address. Again, the issue is not merely the name of the mistake. It is whether the actual alternative has been represented and answered.

9. Case study: calling something an axiom does not complete the proof

In [Fischer-Dillahunty, debate 141](#), the scores are 68 and 86. Neither side carries a named-fallacy tag, and the lower side trails on all six dimension averages.

A recorded move presents a first cause as an axiomatic requirement across conceivable worlds. The assessment recognizes the intended universality but asks how imagining worlds with first causes establishes that all possible worlds require one. It also notes that conceivability, possibility, necessity, and axiomatic status are not interchangeable concepts.

Those distinctions matter even for readers who never use the word “modal,” the technical adjective for reasoning about possibility and necessity. A statement about what one can imagine is not automatically a statement about what must be true. If the argument depends on that upgrade, the speaker has to justify it.

Another contribution describes resurrection as historically substantiated without supplying the historical particulars needed for that assessment. Identifying a concrete candidate event is an improvement over speaking only about miracles in general. It does not complete the evidential task. These are recognizable strengths and weaknesses within an untagged move, not a binary choice between flawless reasoning and a named fallacy.

10. The pattern is not confined to older religious-versus-skeptical contests

The later-process [Swinburne-Goff debate, number 240](#), ends 76 to 83. Both sides defend purposive explanations, so it is excluded from the papers' main religious-versus-skeptical comparison. It is included here because the question concerns argument quality across the archive.

Swinburne's lower-scoring side has no named-fallacy tag and trails on all six dimension averages. One assessed reply grants that Goff's alternatives fit relevant facts, then asks whether their tailored structures need further explanation. The assessment recognizes the legitimate comparative concern but notes that the extra complexity is not specified and that parallel questions about divine attributes remain insufficiently addressed.

Another move argues that a life-producing multiverse would need a suitably structured generator. That is a relevant conditional challenge. The weakness is the claim that viable generators are themselves comparably fine-tuned without identifying or comparing the relevant models. A coherent question can remain under-supported as a positive conclusion.

The eight later untagged losses also include close one-point results in discussions of souls, purgatory, moral realism, and resurrection. Not every untagged loss is a dramatic collapse. The category ranges from near ties to a seven-point difference. That variation is another reason not to treat the absence of a tag as a substantive rating of the whole performance.

11. What the six labels can and cannot tell us

Across 5,492 public moves, **345 carry at least one named-fallacy tag**, or 6.3%. There are **354 label instances**, because some moves have more than one. Counts of moves and counts of labels should not be confused.




Figure 4. The archive records six named-fallacy labels. Equivocation and red herring are the most frequent. This is an inventory of accepted public annotations, not an exhaustive taxonomy or a count of every possible error in the transcripts.

Scope: 354 label instances on 345 of 5,492 public moves; multiple tags possible.

Equivocation draws attention to a word or concept shifting meaning within an argument. A red herring points to diversion from the issue that needs to be addressed. Begging the question concerns dependence on what is at issue. Argument from ignorance concerns an unsupported move from what is not known to a positive conclusion. The remaining labels concern problematic appeals to authority and inconsistent treatment of comparable cases.

Each can help a reader locate a defect. But naming a pattern is not the same thing as demonstrating it. A proper diagnosis must identify the relevant proposition, the unsupported inference or mismatch, and why the apparent weakness is not a reasonable move in the actual context. An appeal to a qualified expert, for example, is not automatically a fallacy.

Labels can also become rhetorical weapons. Announcing “special pleading” without showing that the cases are relevantly alike may be less informative than the untagged criticism it replaces. A reader who learns the names without learning the reasoning can become more confident while becoming no more accurate.

The best use of the site's reference links is therefore educational rather than prosecutorial: learn the pattern, return to the exact contribution, and ask whether the diagnosis fits. The public critique and source material should do the argumentative work.

12. A more useful way to inspect a debate

Start with the conclusion the side needs to establish. Then identify its supporting route: which premises matter, which inferences connect them, and which alternatives could weaken it? A named fallacy may reveal a broken link, but the route can also remain incomplete because a premise is merely asserted or an alternative is never compared.

Next, separate the quality of a reply from its relevance. A response can concern the same general subject while leaving the objection untouched. “Religion gives meaning” may be relevant to a discussion of religious life without answering whether the supernatural explanation is true. “Science has limits” may be relevant to an ultimate-explanation debate without establishing which proposed explanation should replace it.

Then check confidence. Does the speaker distinguish what has been established from what remains possible or speculative? An appropriately limited conclusion can be a strength. A categorical statement supplied with the same evidence may be weaker because it undertakes more than the evidence supports.

Finally, inspect local errors in proportion to their role. A peripheral lapse need not defeat an otherwise strong case; a central unsupported premise can undermine a whole route even without a named tag. Importance, coverage, and the relation between claims matter. Counting mistakes without considering those features can misrepresent the comparative performance.

13. Conclusion: reasoning is larger than its error labels

The updated archive requires a more careful conclusion than the old headline suggested. **Most lower-scoring sides in the combined archive lack a named-fallacy tag, but most lower-scoring sides in the later assessment process do not.** The difference is too large to treat as a minor footnote. A durable account of argument quality must survive that change in annotation patterns.

What survives is the structural lesson. A side can lose through several ordinary shortcomings: insufficient support, an unfinished inference, a partly answered objection, unstable wording, or confidence that exceeds the evidence. Most comparable untagged losses trail across five or six dimensions. Conversely, an accepted fallacy on the higher-scoring side does not automatically overturn the total assessment.

The reader should neither ignore named errors nor worship them. A label is useful when it helps explain why an argument fails. It is not a substitute for that explanation, and its absence is not a certificate that the argument succeeds.

The decisive question is not “Can I name the mistake?” It is “Has the speaker supplied enough reason to cross the gap between the premises and the conclusion?” That question remains useful whether a debate contains no labels, one label, or many - and whether future assessment practices make tags more common or less.

Methods and sources

The source snapshot is revision **76d006b37**, dated September 4, 2026. Public tag counts include all 253 assessments and 5,492 argumentative cards. Cards stored as arrays in multi-speaker exchanges are flattened without duplicating them. Only tags explicitly typed as fallacies enter the named-fallacy counts; bias tags are kept separate. Ten tied assessments are excluded from lower-side prevalence.

The detailed dimension analysis uses 147 lower-scoring sides without fallacy tags from the 237 comparable locked assessments. Side-dimension means preserve official section and move-importance weights. “Trails on a dimension” means a strictly lower weighted mean, allowing a small numerical tolerance for floating-point arithmetic. Dimension intervals use 20,000 whole-debate replacement resamples. The main archive counts are complete counts of this snapshot, not estimates requiring a sampling interval to establish their arithmetic.

The cohort comparison is descriptive. It does not prove that the process caused the label difference, and the sparse tags in earlier records prevent interpreting absence as a uniform statement about the underlying reasoning. Correlated dimensions, selective source collection, recurring speakers, and systematic judging errors limit broader inference.

Calculations, label inventories, the lower-side table, source hashes, and chart data are in the [September 4 analysis directory](#). The [Backend library](#) links the companion scale audit, which explains why annotation rates should remain separated by process. This paper reanalyzes existing judgments and does not independently relabel the complete transcript corpus.