

Beyond the Fallacy Count

The overall majority shrinks, while the difference between assessment processes persists.

54.4%

lower-scoring sides with
no named-fallacy tag

17.8%

same rate in the
later process

75

higher-scoring sides
with a fallacy tag

Named fallacies neither exhaust argument quality nor determine the result.

PHIL STILWELL · SLUGFESTER

October 9, 2026 | Frozen archive: 308 assessments

Recalculated scores; direct slogan evidence retains its September scope.

Executive summary

A missing fallacy label is not an acquittal, and a present label does not decide the result. In the expanded archive, 161 of 296 lower-scoring sides have no named-fallacy tag: 54.4%. That is still a majority, but smaller than September's 61.7%.

The difference between assessment processes remains much larger than the change in the overall percentage. Earlier assessments have 139 untagged losses out of 172 (80.8%); later assessments have 19 out of 107 (17.8%). The overall majority therefore depends heavily on the mix of records.

Among 158 comparable losses with no named-fallacy tag, 118 trail in five or six scoring areas. 75 higher-scoring sides carry at least one fallacy tag. These observations support the paper's durable conclusion: argument quality is broader than a count of accepted labels. They do not establish that the label system is exhaustive, neutral, or independently validated.

Reader's guide

Executive summary	2
1. Count the right thing	4
2. The pooled majority is a mixture, not a law	4
3. No label can coexist with a broadly weak performance	5
4. A labeled error does not erase everything else	6
5. The inventory is narrow and inspectable	6
6. What should readers and reviewers do?	8
Methods and sources	8

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. Score results refer to the October 9 snapshot. Direct slogan counts retain their September source coverage; they have not been extended to new transcripts.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. Count the right thing

This study uses all 308 public assessments for annotation counts. It excludes 12 tied results when identifying a lower-scoring side, leaving 296 decisive assessments. “Loss” is shorthand for the lower published score; it is not a declaration that a worldview is false.

A named-fallacy tag is an accepted public annotation on an assessed move. It is not every weakness described in a critique, every bias tag, or every sentence in a transcript. A move with several fallacy labels counts once when asking whether a side has any label, but contributes several instances to a label-frequency inventory.

The catalogue has 307 unique video links rather than 308 independent recordings. Assessments 13 and 125 share a source under different formats. They remain separate observations in this assessment-level annotation inventory; the comparable one-on-one subset contains only one of them. The full-archive rate should not be presented as the percentage of independent videos with an error-free losing side.

2. The pooled majority is a mixture, not a law




Figure 1. The earlier and later groups give very different answers to the same label-count question.

How to read this graph. Fractions count lower-scoring sides without a fallacy label over all decisive results in each group. The dark total includes the other groups. No tag does not mean no error. Counts are assessments, including the disclosed shared-video pair.

The current combined rate is 54.4%, compared with 61.7% in the previous edition. The later-process rate is 17.8%, compared with 14.5% previously. The earlier group's 80.8% is unchanged. Thus the overall rate can fall even while the later group's own rate rises: the later group now occupies more of the archive.

For a simple illustration unrelated to the real counts, a catalogue with many records from a process that seldom tags fallacies will have a high untagged rate. Adding records from a process that tags them more

often can lower the overall rate without any change to the old judgments. This is why the process-specific rows are essential, not optional caveats below a dramatic headline.

The records alone cannot tell us whether later arguments contain more qualifying errors, later reviewers recognize more of the same errors, acceptance rules differ, or several factors contribute. Describing the contrast as annotation drift is shorthand for a change in recorded practice, not proof of its cause.

3. No label can coexist with a broadly weak performance

Of the 158 comparable untagged losses, 118 - 74.7% - trail in at least five of the six weighted scoring areas. Some arguments fall short through cumulative weaknesses: incomplete support, an unanswered objection, an unclear scope, or confidence stronger than the supplied reasons justify.




Figure 2. The number of areas behind the opponent is a broader description than a count of named fallacies.

How to read this graph. Each bar counts losses with no named-fallacy label that trail in exactly the stated number of weighted scoring areas. Bars add to the eligible total; these are not independent deficiencies.

These dimensions are not independent discoveries. One unsupported inference can depress several related ratings, and the dimensions themselves help define the overall score. The result nevertheless explains why an absence of a named label cannot be treated as a contradiction in a low score. The rubric rewards positive argumentative work as well as avoiding recognizable mistakes.

Imagine a speaker who stays on topic, never equivocates, and never explicitly appeals to an irrelevant authority, but supplies little evidence for the central claim and never answers a specific counterexample.

That illustrative performance can reasonably score below a well-supported answer without a dramatic named fallacy. The appropriate scrutiny is whether the critique identifies the missing work accurately, not whether it can attach a familiar Latin label.

4. A labeled error does not erase everything else

There are 75 decisive assessments in which the higher-scoring side has a fallacy tag. That does not excuse the tagged move. It means the final score incorporates the rest of the performance and the importance of the error in context.

Readers should ask what proposition the mistake affects, whether the speaker later repairs it, and whether the central conclusion depends on it. A local digression and a circular premise supporting the main case need not have equal consequences. The site's move importance and section weights attempt to capture differences of argumentative role, but those judgments remain open to challenge.

Conversely, a high total does not validate every argument uttered by the higher-scoring speaker. The most useful reading is local: identify the particular reason, inference, objection, and response. A scorecard is more informative when it directs attention back to those elements than when it becomes a substitute for them.

5. The inventory is narrow and inspectable

The snapshot records 550 named-fallacy label instances on 536 assessed moves. Six labels occur: Argument from ignorance, Appeal to authority, Equivocation, Begging the question, Special pleading, and Red herring. Their relative frequencies reflect what was assessed and accepted under the site's process, not the prevalence of every reasoning error in public discourse.




Figure 3. Label instances are not unique moves, and their frequency is not the same as their impact on a conclusion.

How to read this graph. A move may carry several labels, so label instances exceed tagged moves. The six names are the accepted inventory in this snapshot, not an exhaustive catalogue of reasoning errors.

Multiple tags can share a move. The denominator for a rate per assessed move is also different from a rate per spoken word or minute. The catalogue selects and scores argumentative moves rather than transcribing every sentence into a separate score unit. Comparisons with the direct slogan study therefore require attention to both the event definition and the denominator.

6. What should readers and reviewers do?

Reason 1: untagged lower scores remain common but are strongly dependent on assessment process.

Reason 2: many untagged losses trail on several broader dimensions. **Reason 3:** many higher-scoring sides have a tag, so presence alone does not determine the aggregate result.

Therefore: “Beyond the Fallacy Count” remains a better title than a universal claim that debates are usually lost without a fallacy. The stable lesson concerns the limits of the count, not a majority percentage that moves with the catalogue.

A useful correction report should identify a concrete passage, explain the proposed error or missing justification, and state how it affects the particular assessment. Reviewers should separately check the presence of the named error, its importance, and any downstream numerical effect. Counting more labels is not automatically stricter or fairer judging.

Therefore: read tags as pointers into an argument, not verdicts detached from it. The best future audit would apply consistent acceptance rules to a shared set of passages across both processes and publish disagreements. That would address the source of the large process contrast more directly than another pooled total.

Methods and sources

This is a fixed October 9, 2026 research edition, based on source revision e65bc3c0b6. The score studies reconstruct published records rather than commission fresh assessments. The direct slogan counts retain the September 5 review of 187 transcripts; this edition recalculates those saved records and adds a separately labeled current score check. No new slogan reviews or debate assessments were performed.

The source manifest records file fingerprints. The analysis uses 20,000 deterministic resampling draws with baseline seed 20260904, retaining it for edition comparability. Whole debates are drawn for paired debate estimates; shared speakers for cross-role estimates. Ranking models and their dependence checks use the units explained in paper seven. These calculations do not capture every selection effect, repeated-speaker dependency, caption error or systematic judging bias.

The [data and methods page](#) links the source records, classification decisions, calculations, chart inputs and reading guides. Download the [current research package](#) and [historical direct slogan study](#). Earlier editions remain preserved in the repository's history. Debate links identify published assessments, not fresh independent evidence for this research.