

OCTOBER 2026 / EXPANDED-CATALOGUE EDITION

Do Slugfester Rankings Measure Stable Performance?

A larger fixed field, repeatable broad patterns, and overlapping exact places.

0.87

median order agreement
between random halves

0.20

typical score gap between
neighboring means

57

people with at least
three eligible appearances

The broad order repeats better than exact places.

PHIL STILWELL · SLUGFESTER

October 9, 2026 | Frozen archive: 308 assessments

Recalculated scores; direct slogan evidence retains its September scope.

Executive summary

The broad order remains repeatable, while precise places remain fragile. This edition has 57 people with at least three comparable appearances, totaling 416 appearances. The September paper had 50 people and 334 appearances. More records enlarge the field as well as strengthen some individual samples.

The median gap between neighboring averages is only 0.20 score points. The typical middle-95% rank range has a width of 14 places when drawing again from each person's recorded scores, and 21 places under a model that allows score variation estimated across the wider group. These widths were 12 and 19 in September. A larger evidence base need not narrow every range when the competitive field also changes.

Randomly splitting the records of sufficiently frequent participants into halves gives a median order agreement of 0.87. That is evidence of a broad signal within this archive, not proof that the AI is correct or that the next debate will preserve a person's rank. Rankings are best read alongside sample size, opponent context, score distribution, and the source-linked assessments.

Reader's guide

Executive summary	2
1. Define the field before interpreting its order	4
2. Why a repeatable broad order can have uncertain places	4
3. Two rank ranges with different assumptions	5
4. Read real rows rather than a single headline	6
5. Check dependence, process differences, and sample size	7
6. The complete fixed field	8
7. Use ranks as an index into the evidence	11
Methods and sources	11

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. Score results refer to the October 9 snapshot. Direct slogan counts retain their September source coverage; they have not been extended to new transcripts.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. Define the field before interpreting its order

The comparable one-on-one records contain 582 appearances by 183 people. Only 57 meet the paper's three-appearance minimum. Team scores and other non-comparable formats are not credited as if each participant individually earned the team result.

The field is fixed to the October 9 source snapshot. This paper is not a live leaderboard: new assessments may subsequently change the website. Average scores use every eligible appearance in the research set, and displayed places sort by average, then appearance count and name for deterministic ties. The uncertainty calculations use average ranks for tied means, avoiding a claim that alphabetical order carries evidence.




Figure 1. The minimum appearance rule excludes many people because evidence is sparse, not because their scores are low.

How to read this graph. Gray bars are below the three-appearance research threshold; teal bars enter the fixed field. A small sample is not a poor score.

A three-appearance threshold is a practical publication rule, not a guarantee of stability. Three similar scores might reflect a limited range of topics and opponents. A person with many appearances has a broader observed record, but those appearances are not randomized tests of equal difficulty.

2. Why a repeatable broad order can have uncertain places

Imagine, solely as a teaching example, two lists A-B-C-D-E and A-C-B-D-E. Most of the order agrees even though B and C exchange exact places. A measure called Spearman rank correlation compares the ordering of all entries: +1 means identical order, -1 reversed order, and values near zero little overall ordering agreement. The example has agreement +0.90, not “90% of people kept their places.”

For the actual archive, 31 people have at least six eligible appearances. In each of 3,000 calculations, each person's scores are randomly divided into two nearly equal groups and averaged. The resulting rank

correlations have a median of 0.87 and a middle-95% range of 0.81-0.92. Splitting by assessment-number order instead gives 0.84. Assessment order is not necessarily recording chronology.

These results describe agreement between subsets of existing scored performances. They do not validate the judges, remove shared-opponent effects, or measure forecast accuracy on unseen debates. A consistent bias can create repeatable rankings. Nevertheless, a strong broad ordering is compatible with many close neighboring means exchanging places.

3. Two rank ranges with different assumptions

Resampled ranks are obtained by drawing, with replacement, from each person's recorded scores, retaining their observed appearance count, then ranking the resulting means. Repeat this 20,000 times. A person with only a few nearly identical scores can get a deceptively narrow range because the procedure has no other performances to draw.

Model ranks use a simple normal model that estimates within-person performance variation across the ranked group and variation between people's underlying means. Small-sample averages are pulled somewhat toward the overall mean; the model then generates 20,000 possible sets of underlying means and ranks them. This permits variation not present in a tiny personal sample, but it adds assumptions about common variation and the model's shape.

Both ranges keep the middle 95% of calculated positions and hold the field fixed. The model conditions on its fitted variance estimates rather than accounting for every uncertainty in those estimates. Neither range is a probability of winning the next debate, and neither is a guarantee that the “true rank” lies inside the printed interval.




Figure 2. The leading 15 observed averages have overlapping uncertainty ranges. Calculations still include the full eligible field.

How to read this graph. Dark dots are displayed places. Teal lines contain the middle 95% of ranks obtained by drawing again from recorded scores. Dashed rust lines use a model allowing variation estimated across the group. Parentheses count appearances. Neither line is a future-win probability.

Typical widths are 14 and 21 places. Width means upper endpoint minus lower endpoint: a range of 2-9 has width seven, even though it includes eight integer places. Half-places arise from tied ranks and interpolation; they are not a hidden extra level of measurement precision.

4. Read real rows rather than a single headline

Joseph Schmid has six eligible appearances and an average of 87.33, first in this fixed research field. His resampled rank range is 1-5; the model range is 1-11. First place is a correct description of the observed mean, not a well-established separation from every nearby speaker.

Matt Dillahunty has 32 appearances and an average of 85.62 when rounded to two decimals. His resampled range is 3-11 and model range 2-11. The larger record supports a more constrained estimate than a small

sample at a similar location, but it does not make the result independent of his particular opponents and topics.

Susan Blackmore has three appearances with little observed score variation. Her resampled range is 5-11.5, whereas the model allows 1-24. The disagreement is instructive: drawing only from three closely spaced scores can understate unobserved variation; borrowing a wider group's variation can overstate it if that person is genuinely more consistent. Neither procedure resolves that question with the evidence available.

These examples explain the methods, not personal psychological qualities. Religious identities and biographies play no role in this ranking calculation. The source records, not a person's present worldview, determine which assessed performances enter.

5. Check dependence, process differences, and sample size

The independent resampling calculation redraws each person's appearances separately, breaking some dependence between opponents who shared a debate. A sensitivity check instead assigns a common positive random weight to each debate for both participants. All people remain in the fixed field. The typical rank-range width is 13 places under this shared-weight check, compared with 14 under independent drawing. It preserves more pairing but still generates no genuinely new performances.

The earlier and later assessment processes have different score levels. Subtracting each process's mean before ranking produces order agreement of 0.99 with the observed table. That suggests substantial broad agreement, not that every individual place is unaffected. Process-centering is an exploratory sensitivity exercise, not an authorized adjustment to public scores.




Figure 3. Appearance counts matter, but proximity to other averages and ordinary score variation also affect rank precision.

How to read this graph. Each dot is a person. Height is the upper rank endpoint minus the lower: 2-9 has width 7. Lower means a narrower calculated range, not a lower performance score. Rank density and score variation matter as well as sample size.

The model estimates a within-person standard deviation of about 3.09 score points and single-appearance repeatability of 0.61. Averaging three appearances gives a modeled repeatability of 0.83, and ten gives 0.94. These are conditional model projections, not observed improvement in people as they debate more often.

6. The complete fixed field

Read the columns together. “Place” is the ordering of observed means; “Appearances” is the eligible sample size; “Average / 100” is the unadjusted score mean. “Resampled ranks” and “Model ranks” are the different middle-95% ranges described above. The downloadable table additionally contains process counts, leave-one-out results, and the shared-debate-weight sensitivity.

Place	Speaker	Appearances	Average / 100	Resampled ranks	Model ranks
1	Joseph Schmid	6	87.33	1-5	1-11
2	Alex Malpass	7	87.14	1-13	1-11
3	Scott Clifton	5	86.40	1-12	1-15
4	Joe Folley	6	86.17	1-11	1-14
5	Sean Carroll	3	86.00	1-14.5	1-21

Place	Speaker	Appearances	Average / 100	Resampled ranks	Model ranks
6	Matt Dillahunty	32	85.62	3-11	2-11
7	Bart Ehrman	14	85.57	3-13	2-13
8	Alex O'Connor	28	85.46	3-12	2-12
9	Susan Blackmore	3	85.33	5-11.5	1-24
10	Justin Schieber	4	85.00	4-14.5	2-23
11	Lance Bush	5	84.80	2-19	2-22
12	Paulogia	5	84.60	3.5-20	2-23
13	Graham Oppy	12	84.08	8-18	6-20
14	Stephen Law	5	83.40	5.5-25	5-29
15	Daniel Dennett	7	83.00	9-27	8-29
16	Max Baker-Hytch	3	83.00	13-21	5-36
17	Michael Ruse	3	83.00	11-23	5-36
18	Richard Dawkins	9	82.33	12-33.5	11-32
19	Michael Huemer	7	82.00	11-37	12-35
20	Randal Rauser	7	82.00	12.5-36	11-35
21	Philip Goff	5	81.60	16-33	12-38
22	Christopher Hitchens	15	81.47	18-31	15-34
23	Sam Harris	8	81.12	18.5-36	15-38
24	Michael Licona	7	81.00	19-36	15-39
25	Stephen Woodford	6	81.00	20-36	14-40
26	Michael Shermer	8	80.88	18-38	16-39
27	James White	5	80.80	15-39	14-42
28	Bernardo Kastrup	6	80.67	18.5-41	15-41

Place	Speaker	Appearances	Average / 100	Resampled ranks	Model ranks
29	Ben Watkins	3	80.67	6-49	13-44
30	Josh Rasmussen	5	80.60	18-40	15-42
31	Richard Carrier	8	80.50	19.5-41	17-41
32	Lawrence Krauss	3	80.33	14-47	13-45
33	Shabir Ally	3	80.33	11-44.5	13-45
34	James Tour	4	80.25	24-38	15-44
35	Miles Donahue	3	80.00	27-38.5	14-46
36	Robert Price	4	79.75	28-40	17-46
37	David Wood	5	79.60	20-47	19-46
38	Jonathan McLatchie	6	79.50	22-45	20.975-45
39	William Lane Craig	30	78.83	32.5-43	33-44
40	Glen Scrivener	6	78.33	29.5-48	27-49
41	Richard Swinburne	5	78.00	33.5-48	27-50
42	Trent Horn	9	77.67	30-51	34-50
43	Tom Jump	6	77.67	37-48.5	31-51
44	Ross Douthat	3	77.67	14.5-57	24-52
45	Dan Barker	7	76.86	38-52	36-52
46	John Lennox	10	76.70	39-51	39-52
47	Aron Ra	6	76.50	41-51	37-53
48	Gregory Pine	3	76.00	39-53	33-54
49	Michael Jones	14	75.93	41-53	43-53
50	Keith Ward	3	75.33	34.5-57	36-55
51	Dinesh D'Souza	7	75.00	45-54	44-55

Place	Speaker	Appearances	Average / 100	Resampled ranks	Model ranks
52	Jordan Peterson	6	74.83	47-54	43-56
53	Frank Turek	8	74.12	49-55	47-56
54	Alister McGrath	4	74.00	42-56.5	43-56
55	Eric Hernandez	4	72.00	52.5-57	49-57
56	Peter Atkins	6	71.67	52-57	52-57
57	Matt Slick	4	71.25	53.5-57	51-57

7. Use ranks as an index into the evidence

Reason 1: different subsets reproduce much of the broad order. **Reason 2:** neighboring means are close and rank ranges overlap. **Reason 3:** unequal samples, changing fields, opponents, topics, and assessment processes limit individual comparisons.

Therefore: a leaderboard can be useful for navigation without supporting a precise hierarchy of personal reasoning ability. Readers should compare distributions and source-linked performances before treating a one-place difference as meaningful.

Therefore: the site should make sample size and uncertainty easy to see, retain a minimum-appearance rule, and avoid future-win language. A stronger future study would evaluate predictions on genuinely new debates and matched opponents, not merely rerank the same observations. More data are valuable, but their independence, comparability, and diversity matter as much as the total.

Methods and sources

This is a fixed October 9, 2026 research edition, based on source revision e65bc3c0b6. The score studies reconstruct published records rather than commission fresh assessments. The direct slogan counts retain the September 5 review of 187 transcripts; this edition recalculates those saved records and adds a separately labeled current score check. No new slogan reviews or debate assessments were performed.

The source manifest records file fingerprints. The analysis uses 20,000 deterministic resampling draws with baseline seed 20260904, retaining it for edition comparability. Whole debates are drawn for paired debate estimates; shared speakers for cross-role estimates. Ranking models and their dependence checks use the units explained in paper seven. These calculations do not capture every selection effect, repeated-speaker dependency, caption error or systematic judging bias.

The [data and methods page](#) links the source records, classification decisions, calculations, chart inputs and reading guides. Download the [current research package](#) and [historical direct slogan study](#). Earlier editions remain preserved in the repository's history. Debate links identify published assessments, not fresh independent evidence for this research.