

Do Slugfester Rankings Measure Stable Performance?

How much of the order holds up, why nearby places can change, and how to read both kinds of rank range.

0.86

typical order agreement
between random halves

0.17

typical score-point gap
between neighboring averages

50

speakers with at least
three eligible appearances

The broad order repeats better than the exact places. Typical rank ranges have a width of 12 places when recorded scores are drawn again, and 19 when a model also allows the score variation seen across the wider group.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

The leaderboard's broad order holds up better than its exact places. In the September 4 archive, 50 speakers have the three appearances needed to enter the table. Together they have 334 eligible appearances. Among the 23 speakers with at least six appearances, we repeatedly divide each person's scores into two random halves and make a ranking from each half. The two lists usually put the same people near the top and bottom. Their typical agreement score is **0.86**, where +1 means the same order, 0 means no clear tendency to keep or reverse the order, and -1 means the reverse order. Section 5 illustrates this measure, called Spearman rank correlation.

That agreement does not mean that every neighboring place marks a real difference in performance. The typical gap between neighboring averages is only **0.17 score points**. Repeatedly drawing from each person's recorded scores gives rank ranges with a typical width of **12 places**. A second method also allows for the score changes seen across the wider group; its typical width is **19 places**. In the appendix, these are called **Resampled ranks** and **Model ranks**. Both warn against treating an exact place as settled.

The number of appearances matters. Sean Carroll averages 86.00 from three appearances; Matt Dillahunty averages 85.93 from 28. A difference of 0.07 points gives us little reason to call one the stronger debater. Under the second method, Carroll's ranks run from 1 to 17, while Dillahunty's run from 2 to 9. More appearances do not guarantee a higher score. They usually give us a firmer basis for describing the performances we have recorded.

The model-based estimate of single-appearance repeatability is **0.63**. Under its assumptions, averaging three appearances raises repeatability to about **0.84**, five to 0.89, and ten to 0.94. These are not probabilities that the AI is correct. Stable scores can reflect stable judging preferences, recurring topics, or opponent selection as well as persistent reasoning skill.

The full ranking analysis covers **237 comparable debates, 474 appearances, and 158 speakers**. Seventy-six speakers appear only once and 32 twice; they are not in the default ranked field. The leaderboard therefore describes a selected repeat-appearance sample, not the world's best thinkers or all participants on the site.

The recommendation is to **retain the ranking as a guide, while reading it through appearance counts, uncertainty ranges, and source context**. Broad bands are informative. Exact places, especially in tightly packed regions or small samples, remain provisional.

Reader's guide

Executive summary	2
1. What should a ranking mean?	4
2. Who enters the leaderboard - and who does not?	4
3. Broad differences are larger than ordinary within-person variation	5
4. Why averaging helps - and what it cannot fix	6
5. Do two halves of the record put people in a similar order?	8
6. "Resampled ranks": what happens when we reuse the scores we have?	9
7. "Model ranks": allowing for variation missing from a short record	10
8. More appearances help, but the density of the field also matters	12
9. The top ten are not equally secure	13
10. Assessment process and opponent context	14
11. What an honest leaderboard should communicate	14
12. What the next review should test	15
13. Conclusion: trust the broad pattern more than the exact place	15
Appendix: the 50 speakers and how to read their rows	16
Methods and sources	20

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. All results refer to the September 4 snapshot, not later additions to the site.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. What should a ranking mean?

A useful ranking should tell us more than who happened to receive a high score once. It should summarize a pattern that is likely to recur under reasonably similar conditions. At the same time, it should not pretend that a small difference between two averages reveals a deep difference between two people.

SLUGFESTER scores performance in debate in selected debate transcripts. The public speaker average is therefore an average of those assessed performances. It is not a measure of intelligence, moral worth, academic expertise, or the truth of a person's worldview. A brilliant researcher can argue poorly in a particular exchange; a speaker with a false conclusion can still display care in the reasoning they present.

We ask two questions. Do the same people tend to score higher across different appearances? And do we have enough evidence to place one person clearly above their neighbors? The first answer can be yes while the second is no. The broad order can hold up even when many exact places can change.

This paper treats September 4, 2026 as a frozen reference point. Its 253-assessment archive includes 237 comparable locked one-on-one assessments. The underlying scores are labeled 5.6 Sol. The Astra-era work reanalyzes the rankings; it does not provide an independent new-model reassessment of each speaker.

2. Who enters the leaderboard - and who does not?

The analysis includes one appearance for each side of each eligible one-on-one debate: **474 appearances by 158 speakers**. Speaker names are resolved using the site's canonical identity records so that variants such as Joe and Joseph Schmid do not become separate competitors. Multi-speaker formats are not silently treated as equivalent individual performances.

The default three-appearance minimum leaves **50 ranked speakers and 334 appearances**. The other 108 speakers have only one or two eligible appearances. Their absence from the default field is a sample-size rule, not a conclusion that they are weaker than the ranked participants.




Figure 1. Most speakers have very limited coverage. The three-appearance threshold excludes 108 of the 158 eligible speakers from the default ranked field. A missing rank therefore often means insufficient repeated evidence, not poor performance.

How to read this graph. Each bar counts people with that many eligible appearances. Gray groups fall below the three-appearance minimum; teal groups enter the ranking. All bars together contain 158 people; the teal bars contain 50. A gray bar does not mean poor performance - it means too few appearances for the default table.

Seventy-six speakers appear once, 32 twice, 27 three to five times, 16 six to ten times, and seven at least 11 times. The uneven coverage means that the word “average” carries very different support across the field. An average based on three narrow discussions is not as broad a performance sample as an average based on 28.

The minimum of three is defensible as a simple publication rule, but it is not a magical threshold at which a stable ability suddenly becomes measurable. A three-debate sample can still be unusually favorable or unfavorable. A larger sample can still be biased if all appearances cover one specialty or face similar opponents.

Raising the minimum would increase the amount of evidence per ranked person while excluding more of the field. Lowering it would make the table more inclusive while increasing uncertainty. Neither choice removes the need to show counts and explain what the ranking summarizes.

3. Broad differences are larger than ordinary within-person variation

The repeatability calculation separates two kinds of variation: differences between speaker averages and fluctuations among the appearances of the same speaker. It uses a simple model in which a person's observed score combines a persistent speaker-associated level with ordinary appearance-to-appearance variation.

The model puts the ordinary spread of one person's appearance scores at about **3.05 points** around their usual level. The name for this kind of spread is **standard deviation**. It is not a rule that every score must stay within 3.05 points, and it is not the average gap between two opponents.

To calculate how much of the variation repeats, the model uses squared spreads, called **variances**. Squaring avoids letting differences above and below an average cancel each other out. The model assigns 15.81 to lasting differences between speaker levels and 9.33 to appearance-to-appearance changes. Readers who want to check the arithmetic can follow the ratio below; the main point is the distinction between differences that repeat and changes from one appearance to another.

The ratio $15.81 / (15.81 + 9.33)$ is approximately **0.63**. This is the estimated single-appearance repeatability. It says that a substantial part of the modeled variation is associated with which speaker is being assessed rather than which of that speaker's appearances happened to be sampled.

It does not tell us why the persistent component exists. It may include genuine differences in reasoning, but also stable topic expertise, recurring reasoning roles, opponent selection, recognizable presentation styles, or systematic preferences in the judging process. A repeatable measurement can be biased. Reliability is not the same thing as truth or fairness.

4. Why averaging helps - and what it cannot fix

If appearance-to-appearance fluctuations are approximately independent, averaging reduces their contribution. In the fitted model, a three-appearance mean has repeatability around **0.84**; a five-appearance mean around **0.89**; and a ten-appearance mean around **0.94**. The model does not predict that a speaker will improve as they appear more often. It predicts that their average will be less affected by ordinary variation.




Figure 2. The curve is a conditional projection from the fitted repeatability model, not a newly observed experiment at each sample size. It assumes comparable appearances with approximately independent ordinary variation.

How to read this graph. The teal curve is a model's expected stability as more appearances are averaged, not observed improvement over time. Dots label examples at 1, 3, 5, and 10 appearances. The 0-1 scale measures the share of modeled variation linked to lasting differences between speakers. It is not the probability that the AI is correct.

An intuitive approximation uses the square-root rule. If the ordinary spread is 3.05 points, the uncertainty in a mean shrinks roughly in proportion to 1 divided by the square root of the number of appearances. Three appearances therefore leave much more ordinary variation in the average than ten or 28. Doubling the sample does not halve uncertainty; reducing it by half typically requires about four times as many comparable observations.

That rule is helpful, but it is conditional. If all new appearances come from the same narrow topic or share a systematic judging preference, they may add less independent information than their count suggests. Repeating a biased procedure can make an estimate more stable without making it more accurate.

More data can narrow uncertainty without correcting bias. A scale that consistently reads two kilograms high becomes no more accurate because a person steps on it 20 times. The average is more repeatable, but the offset remains. Debate rankings need both repeated evidence and scrutiny of the scoring instrument.

This is why paper six's process audit matters to the leaderboard. Absolute score levels differ between the earlier and later assessment groups. A speaker's average can be affected by how many appearances come from each. Large samples help with ordinary score changes, not automatically with that comparability problem.

5. Do two halves of the record put people in a similar order?

The repeatability model is not the only evidence. For the **23 speakers with at least six appearances**, the analysis repeatedly divides each person's scores into two random halves, computes a mean from each half, and compares the resulting rankings. If the leaderboard were mostly random fluctuation, high placement in one half would reveal little about placement in the other.

Across 3,000 random splits, the middle agreement score is **0.86**. The formal name is **Spearman rank correlation**. The question it answers is simple: **if someone is high on one list, are they also high on the other?** It compares places, not the size of the score gaps.

The scale runs from -1 to $+1$. A value of $+1$ means that both lists put everyone in the same order. A value of -1 means that the second list exactly reverses the first. A value near 0 means that knowing someone's place on one list gives little help in guessing whether they will be near the top or bottom of the other. Zero does not mean that the lists have no matching places.




Figure 3. These five speakers are made up to explain the calculation. The left and right lists match perfectly in the first panel. In the middle panel, B and C exchange neighboring places, so the agreement falls slightly. The last panel reverses every place. None of these panels is an extra debate assessment.

How to read this graph. Each panel compares two lists of five made-up speakers, A through E. First place is at the top. Follow a letter's line from the first list to the second; crossing lines show changes of order. Teal is B and rust is C only to highlight the swap. The agreement number compares the whole order, not the percentage who kept their place.

Here is the small calculation behind the middle panel. B moves one place and C moves one place; the other three do not move. Squaring those changes and adding gives $1 + 1 = 2$. For five names with no ties, the formula gives $1 - (6 \times 2) / (5 \times 24) = 0.90$. Larger and more frequent changes of place reduce the result. When scores tie, the calculation uses shared average places instead; it does not arbitrarily put one tied person first.

0.86 does not mean “86% of speakers kept their place.” It is one summary of how closely two entire orders agree. Nor does it mean that 86% of the AI's judgments are correct. Adding five points to every score on one list would leave its order unchanged, even though none of the score values would match the old list.

In the actual split-record test, the middle 95% of the 3,000 results runs from about **0.79 to 0.92**. This means that 2.5% of the results are below the lower end and 2.5% above the upper end. Most splits therefore show fairly close agreement about broad order.

Those limits describe variation across the chosen splits, not a universal confidence interval for how well the leaderboard predicts future performance. The same recorded scores are reused in different divisions. The exercise tests internal repeatability within the available archive, not external validity on a new, independently selected set of debates.

A second check divides those speakers' appearances by archive sequence: earlier-numbered versus later-numbered records. Its rank correlation is **0.84**. This is a useful harder separation than a purely random split, but it is not a test of chronological improvement or decline. The archive sequence is not the filming date, and different assessment processes appear unevenly across it.

Taken together, the model and split checks support a persistent broad signal. They do not establish that the displayed rank of every individual is stable, or that the signal isolates a general ability independent of topics and judging conventions.

6. “Resampled ranks”: what happens when we reuse the scores we have?

The typical gap between neighboring speaker averages is **0.17 points**. For comparison, the estimated spread of one person's scores around their own usual level is about **3.05 points**. These are different measures, but their size helps explain the problem: a small change in the scores can switch several neighboring places. The displayed order can be correctly calculated without proving that one neighbor is reliably better than another.

The table orders Joseph Schmid first at **87.33 from six appearances**, Alex Malpass second at **87.17 from six**, and Scott Clifton third at **86.75 from four**. Those are accurate descriptive means for the frozen sample. They should not be translated into strong claims that the first person would reliably outperform the second or third in a matched new debate.

Resampled ranks means “places produced by repeatedly drawing again from the recorded scores.” Think of putting a person's scores on slips of paper in a bag. If they have six appearances, draw six slips, putting each slip back after drawing it. The same score can appear more than once, and another score may not appear at all. Average the six draws. Do this for every one of the 50 people and rank the new averages together. Then repeat the whole exercise **20,000 times**.

For a made-up record of 80, 84, and 88, one draw might be 80, 80, 88; another might be 84, 88, 88. We are not inventing new assessed performances. We are asking how much the order changes when the performances already observed receive different influence.

For each speaker, the table keeps the middle **95%** of the resulting places. A range of **1-5** means first through fifth place in that part of the repeated-draw results. It does not mean a score between 1 and 5. It does not guarantee a future place between first and fifth.

The typical range width is **12 places**, measured as the upper end minus the lower end. For example, a range from 8 to 20 has width 12, although it contains 13 whole-number places if both endpoints are counted. Throughout this paper, “typical” here means the middle value after all 50 widths are sorted from narrowest to widest, also called the median.

There is a complication: two opposing scores come from the same debate. A second check gives each whole debate a random amount of influence, shared by both sides, and ranks the speakers again. Its typical width is also 12 places. Keeping opponents linked in this way does not remove the warning about exact places.

7. “Model ranks”: allowing for variation missing from a short record

Drawing again from recorded scores can only reuse what we have seen. Suppose someone has three scores of 85. Every draw from those scores still averages 85. That does not prove the person's next assessed performance will score 85. The short record may simply have missed the ups and downs that more appearances would reveal.

Model ranks means “places from a calculation that also uses the score variation seen across the wider group.” The calculation learns how much scores usually change from one appearance to another across the 50 ranked speakers' 334 appearances. It combines that shared information with each person's own average and appearance count. It then makes 20,000 sets of possible underlying average scores and ranks the same 50 people each time.

In this calculation, short records are moved somewhat toward the group's usual level. Longer records stay closer to their recorded averages. This does not change a public score or punish someone for having few appearances. It avoids treating a very short record as complete knowledge of their usual assessed performance. The printed average remains the actual recorded average.

The middle 95% of these calculated ranks forms the **Model ranks** column. Its typical width is **19 places**, compared with 12 for Resampled ranks. This is a useful second warning, not proof that 19 is the uniquely correct answer. The model assumes roughly bell-shaped differences between people's usual score levels and a similar amount of ordinary score variation across people. It also treats its estimated amounts of variation as fixed. If those assumptions are wrong, its ranges can be misleading.

Current leaders: two ways to check their places




Figure 4. The dark dots show the current order of the 15 highest recorded averages. Solid teal lines show Resampled ranks; dashed rust lines show Model ranks. Numbers beside names count appearances. Both methods rank the same 50 people, not just the 15 shown here.

How to read this graph. Dark dots mark the current displayed place. Solid teal lines are Resampled ranks, made by drawing again from recorded scores. Dashed rust lines are Model ranks, also using score variation across the wider group. Lines keep the middle 95% of calculated ranks. Parentheses count appearances. First place is left; all calculations use the full fixed group of 50.

The comparison between Sean Carroll and Matt Dillahunty is especially instructive. Carroll's average is **86.00 from three appearances**; Dillahunty's is **85.93 from 28**. The raw gap is only 0.07 points. Their model ranges are 1-17 and 2-9, respectively. The apparent advantage in the displayed averages is much smaller than the uncertainty surrounding it.

The model range is **not always wider for every person**. Alex Malpass has Resampled ranks of 1-12 and Model ranks of 1-9. Each calculation changes all 50 people's possible averages together, and the model uses information differently. The ranges need not fit inside one another. We should compare the assumptions and the overall warning, not choose whichever range makes a favored speaker look best.

Neither column contains newly judged debates. Neither gives the probability of winning a future debate. Neither removes mistakes in the source scores, unfair selection of topics, or repeated preferences in the

judging process. They show how exact places can change **under two stated ways of using this particular record**.

8. More appearances help, but the density of the field also matters

Rank uncertainty depends on two things: how uncertain a speaker's average is and how tightly packed the surrounding means are. A one-point change near a crowded middle can move a speaker many places. The same change near an isolated end of the field may move them very little.

Sample size helps, but position in the field matters too




Figure 5. More appearances generally provide firmer evidence, but rank-range width also depends on where a speaker sits among neighboring averages. A crowded part of the field can produce a broad rank range even when score uncertainty is moderate.

How to read this graph. Each teal dot is one of the 50 ranked speakers. Horizontal position counts appearances. Vertical position is the upper Model ranks endpoint minus the lower endpoint; a range of 2-9 has width 7. Lower dots have narrower calculated ranges, not lower scores. Overlapping dots may hide one another.

This is why rank ranges should be read alongside score ranges and appearance counts. A person can have a reasonably stable average and still have an unstable exact rank if many competitors have nearly the same average. Conversely, an extreme recorded average can make a rank look relatively stable within the current field even when the underlying performance sample is small.

A leave-one-appearance-out check makes the vulnerability concrete. For Ross Douthat, whose mean uses three appearances, omitting one observed score at a time while holding the other speakers' means fixed gives ranks from 15.5 to 45. For Ben Watkins, also with three, the range is 12 to 34. Fractional positions arise from tied means in the calculation.

These are not proposals to delete inconvenient debates, and they are not confidence intervals. They ask how much influence a single recorded appearance has on the descriptive rank. The answer can be substantial

near the minimum sample threshold. That is exactly where readers should resist treating a displayed integer as a durable personal rating.

9. The top ten are not equally secure

The same model counts how often each speaker reaches the top ten across 20,000 calculated rankings. For example, 80% would mean reaching the top ten in about 16,000 of those calculations. It would **not** mean an 80% chance of being one of the world's ten best debaters, or an 80% chance of winning the next debate.




Figure 6. Some current top-ten positions are far more secure under the model than others. High inclusion is conditional on the recorded field, its selected debates, the scoring system, and the model assumptions.

How to read this graph. Each bar shows the share of 20,000 model calculations putting that speaker in the top ten of the same 50-person group. For example, 80% means about 16,000 calculations. The chart shows the current top 12 averages, not the whole field. These are not future win probabilities or a measure of the AI's accuracy.

Seven of the current ten have top-ten inclusion above 90% under this model. Sean Carroll's is about 79.7%, Lance Bush's 66.4%, and Graham Oppy's 52.6%. Stephen Law, currently just outside the top ten, has about 36.2% inclusion. The boundary is therefore a zone of uncertainty, not a sharp divide between two established classes of performer.

Such probabilities are often more useful than a claim that a speaker is “ranked tenth.” They show that a boundary can be unstable even when the broad set of leading performers is reasonably persistent. But they should never be detached from the assumptions. A different selection of debates, scoring process, or model of variation could change them.

The result also depends on keeping the field fixed. Adding a previously unranked speaker with several strong performances can move everyone else down without any of their own scores changing. Rank is relational: it describes a position among the current competitors, not a property carried by a person independently of the field.

10. Assessment process and opponent context

Paper six finds that later-process score levels are lower on average. To test whether that alone reorganizes the leaderboard, we subtract each assessment group's mean from the appearances produced by that family and recompute speaker ordering. The correlation between the original and centered ranks is **0.99**. The broad ordering survives this simple adjustment.

Individual changes still matter. Christopher Hitchens moves from displayed place 17 to centered place 22; Richard Dawkins from 19 to 14; Ben Watkins from 24 to 19; and James White from 32 to 27. These movements are results under different choices, not a corrected official ranking. Subtracting a group average cannot establish what the same debate would have scored under another process.

Opponent context is another unresolved influence. Some speakers repeatedly face highly experienced critics; others appear against a narrower or less demanding set of opponents. A score is intended to reflect the argument's performance, not simply whether the opponent was strong, but opponent quality affects which objections arise and how much work needed to support a claim is required.

The current analysis does not supply a validated opponent-adjusted ranking. With many sparse speaker records and a selected debate network, a complicated adjustment could appear precise while depending heavily on assumptions. The safer conclusion is to make context inspectable and reserve stronger claims for matched comparisons or better-connected samples.

Topic specialization matters too. A speaker repeatedly debating their strongest subject may have a higher mean than someone taking on a wider range of difficult motions. Neither average is false; they summarize different portfolios. A general ability ranking would need more deliberate balancing than this archive currently supplies.

11. What an honest leaderboard should communicate

The observed average should remain visible, along with the number of eligible appearances. Readers should be able to distinguish a three-debate result from a 28-debate result immediately. A displayed order can be convenient for browsing, but the surrounding explanation should make clear that nearby places are often not meaningfully separated.

Broad performance bands would better reflect the evidence than treating every ordinal place as a distinct achievement. Such bands should not be invented from arbitrary cutoffs without testing. One practical

starting point is to show uncertainty ranges and let readers see overlapping evidence rather than presenting a false categorical separation.

The site should also make each speaker's underlying debate list and mix of assessment groups easy to inspect. A user ought to be able to ask whether an average comes from many topics, one favored subject, one opponent, or several assessment groups. Those are important questions, not merely statistical housekeeping.

Changes over time should preserve the reason for the change. A rank can move because a speaker receives a new assessment, because competitors enter the field, because a source correction changes an old score, or because the assessment process changes. “Up five places” means very different things under those four scenarios.

Finally, a leaderboard should invite scrutiny rather than substitute for it. The productive use is to identify patterns and select debates for closer examination. The unproductive use is to cite a rank as if it settles the truth of a philosophical claim or a person's general intellectual worth.

12. What the next review should test

The next major GPT-model transition is the intended main review point for these frozen papers. A stronger ranking study would start with the scale calibration proposed in paper six: shared transcripts assessed under old and new procedures, with source fidelity and judgments reviewed independently.

It would then test prediction on genuinely new material. A ranking built from one set of debates should be evaluated against later, independently selected appearances, with the selection rules fixed before outcomes are known. Reusing the same scores in many random splits is useful for internal stability, but it is not the same as predicting a fresh sample.

Where possible, new collection should reduce the largest gaps in coverage. More appearances by scarcely observed speakers are valuable, but merely adding many similar debates by already frequent speakers will not solve uneven representation. Diversity of opponents, topics, roles, and formats matters alongside count.

The study should report whether conclusions survive reasonable alternative assumptions about ordinary variation. Some speakers may be more variable than others; some topics may be systematically harder to score; and recurring opponents can connect observations. No single model should be treated as an unquestionable machine for turning a selected archive into a universal ranking.

13. Conclusion: trust the broad pattern more than the exact place

Argument one: the broad ordering contains useful information

1. If a ranking records a recurring pattern, separate parts of the record should often put the same people near the top and bottom.

2. For the 23 speakers with at least six appearances, 3,000 random splits of their records give a typical order-agreement score of **0.86**, fairly close to perfect agreement at +1.

3. This test reuses existing scores. It shows agreement within this archive, not that the AI's judgments are correct or that the order will predict a fresh set of debates.

Therefore: the broad order is a useful guide to the assessed performances already on the site. It is not yet a proven ranking of general debating ability.

Argument two: nearby places deserve much less confidence

1. A tiny difference between averages is weak grounds for a strong claim when the recorded scores can vary enough to reverse that difference.

2. The typical neighboring gap is **0.17 points**. Estimated ordinary score variation within a person's record is about **3.05 points**. The typical rank-range width is **12 places** when recorded scores are drawn again and **19 places** when the calculation also uses the wider group's variation.

3. Carroll's three-appearance average exceeds Dillahunty's 28-appearance average by only **0.07 points**; their model ranges overlap greatly. The decimal difference alone cannot settle who would perform better in a fair new comparison.

Therefore: keep the averages, but read them with the appearance counts and both rank ranges. First rather than second, or tenth rather than eleventh, is often more precise than the evidence allows.

The practical rule is simple: use the ranking to choose arguments to examine, not to avoid examining them. A high place does not make a claim true, and a low place does not make it false. The reasons offered in the debate still have to do that work.

Appendix: the 50 speakers and how to read their rows

This is the **complete eligible ranking group**, not the complete list of people on SLUGFESTER and not a list of the world's 50 best debaters. There are 158 speakers in the comparable one-on-one records. Only 50 have at least three eligible appearances. The other 108 have one or two appearances and are excluded by the minimum-count rule, not because this paper found them less capable. The 16 other-format assessments do not enter these individual averages.

The table starts with the highest **recorded average**, using appearance count and then name to settle display ties. The first column gives the resulting displayed place. Every average and range refers to the same fixed group of 50 people. Adding new eligible speakers could change places even if no existing score changed.

What each column tells you

Place: the current order of the recorded averages. This is a convenient display order, not a guaranteed future position.

Speaker: the person whose eligible appearances are combined. Name variants are joined so that one person does not occupy two rows.

Appearances: how many assessed debate-side performances support the average. Three gives much less evidence than 28, even when the averages nearly match.

Average / 100: the ordinary average of that person's eligible overall scores, on the site's 0-100 scale. This column does not contain a model-adjusted score.

Resampled ranks: the middle 95% of places from 20,000 repeats that draw again from each person's recorded scores, with replacement. Each person keeps the same number of draws as their original appearance count. All 50 new averages are ranked together each time. The method tests the influence of the performances we have seen; it cannot reveal score variation absent from a short record.

Model ranks: the middle 95% of places from 20,000 calculations that also use the ordinary score variation estimated across all 334 appearances in this ranked group. This lets a short record leave room for ups and downs it may have missed. The result depends on the model's assumptions, explained in section 7. It is not a second AI's new judgment of the debates.

Read three real rows

Joseph Schmid: six appearances, average 87.33, Resampled ranks 1-5, Model ranks 1-9. He has the highest recorded average. Drawing again from the observed scores puts his middle 95% of ranks between first and fifth. Adding the wider group's information gives a middle range from first to ninth. Neither range guarantees where he will rank after new debates.

Sean Carroll: three appearances, average 86.00, Resampled ranks 1-12, Model ranks 1-17. The average is high, but rests on only three appearances. The second calculation leaves more room for variation that this short record may not show.

Matt Dillahunty: 28 appearances, average 85.93, Resampled ranks 2-9, Model ranks 2-9. His recorded average is almost the same as Carroll's, but has much more evidence behind it. Here the two methods happen to give the same rounded endpoints; they are still different calculations.

Read “1-5” as “first through fifth place,” not as a score. The ranges leave out the lowest 2.5% and highest 2.5% of calculated ranks. That does not establish a 95% guarantee about a person's true or future ability. A fractional place can arise from a tie: two people tied for fifth and sixth receive place 5.5 in a calculation. It does not mean half a person.

The full ranked group

Place	Speaker	Appearances	Average / 100	Resampled ranks	Model ranks
1	Joseph Schmid	6	87.33	1-5	1-9
2	Alex Malpass	6	87.17	1-12	1-9
3	Scott Clifton	4	86.75	1-10	1-12
4	Joe Folley	6	86.17	1.5-9	1-12
5	Sean Carroll	3	86.00	1-12	1-17
6	Matt Dillahunty	28	85.93	2-9	2-9
7	Alex O'Connor	28	85.46	3-10	2-10
8	Bart Ehrman	12	85.42	3-11	2-12
9	Lance Bush	5	84.80	2-16	2-18
10	Graham Oppy	12	84.08	7-15	5-16
11	Stephen Law	4	83.75	3-22.5	4-24
12	Daniel Dennett	6	83.17	7-23.5	7-24
13	Max Baker-Hytech	3	83.00	10-17.5	5-30
14	Michael Ruse	3	83.00	9-19	5-30
15	Michael Huemer	7	82.00	9-31	10-29
16	Bernardo Kastrup	5	81.80	14-24.5	10-31
17	Christopher Hitchens	11	81.73	13-25	12-28
18	Philip Goff	5	81.60	13-28	10-32
19	Richard Dawkins	7	81.43	11-32	11-31
20	Sam Harris	7	81.00	15-31	12-33
21	Michael Shermer	6	81.00	12-33	12-34
22	Stephen Woodford	6	81.00	16-30	12-34

Place	Speaker	Appearances	Average / 100	Resampled ranks	Model ranks
23	Michael Licon	6	80.83	15-31.5	13-34
24	Ben Watkins	3	80.67	6-42	11-38
25	Josh Rasmussen	5	80.60	14-34	13-36
26	Richard Carrier	7	80.43	15-36	14-35
27	Lawrence Krauss	3	80.33	11.5-39.5	11-39
28	Robert Price	3	80.33	22-30	12-38
29	James Tour	4	80.25	20-31.5	13-38
30	William Lane Craig	24	79.50	23-35	23-36
31	David Wood	4	79.00	16.5-43	17-41
32	James White	3	79.00	28-36	15-42
33	Jonathan McLatchie	4	78.75	18-42	18-42
34	Tom Jump	5	78.00	29-41	23-43
35	Richard Swinburne	4	78.00	25-42	21-44
36	Ross Douthart	3	77.67	12-50	20-45
37	Glen Scrivener	5	77.60	28.5-43	25-44
38	John Lennox	9	77.11	31-43	31-44
39	Trent Horn	6	76.83	30-46	30-46
40	Aron Ra	4	76.75	33-45	27-47
41	Gregory Pine	3	76.00	33-46	28-48
42	Dinesh D'Souza	6	75.50	37-47	35-48
43	Keith Ward	3	75.33	29-50	31-49
44	Michael Jones	12	74.83	38.5-48	40-49
45	Jordan Peterson	6	74.83	40-47	37-49

Place	Speaker	Appearances	Average / 100	Resampled ranks	Model ranks
46	Frank Turek	7	74.14	41-49	40-50
47	Alister McGrath	4	74.00	35-50	37-50
48	Peter Atkins	4	73.00	42-50	40-50
49	Matt Slick	3	72.33	46-50	40-50
50	Eric Hernandez	4	72.00	45.5-50	43-50

Methods and sources

The source revision is **76d006b37**, frozen September 4, 2026. Eligible appearances come from 237 comparable locked one-on-one assessments and use canonical site speaker identities. The default field requires at least three appearances, leaving 50 speakers and 334 observations. The entire eligible appearance set has 158 speakers and 474 observations.

The one-way repeatability model estimates between-speaker and within-speaker variance from unequal group sizes. Estimated repeatability for an average of n appearances is between variance divided by [between variance + within variance / n]. This describes the fitted observational sample under the model; it does not identify general ability or the correctness of the AI judgments.

Observed-score rank ranges use 20,000 independent within-speaker replacement resamples. A supplementary paired-weight analysis assigns one positive exponential random weight to each debate, shared by its two appearances, and re-ranks the fixed field. The model using the wider record combines a normal distribution of persistent speaker levels with a common within-speaker variance; variance estimates are treated as fixed. Its 20,000 draws describe conditional uncertainty in latent means, not calibrated probabilities of future wins. Simulated ties receive average ranks.

Random split checks use 3,000 divisions among 23 speakers with at least six appearances. The sequence split orders by archive number, not recording date. Leave-one-out checks remove one appearance from one speaker at a time while holding the rest of the field fixed. Process-centering subtracts the relevant group mean before re-ranking and is a check using different choices, not a causal correction.

The [analysis directory](#) provides the complete ranking data, formulas, simulations, source hashes, and figure contracts. The [Backend library](#) links the six companion papers, especially the scale audit on which any stronger cross-debate interpretation depends. No public assessment or official speaker score is altered by this paper.