

Do Slugfester Rankings Measure Stable Performance?

Why broad score bands are defensible while exact numbered positions remain provisional

REPORT DATE · SEPTEMBER 1, 2026 | RANKING CORPUS · 212 DEBATES / 424 APPEARANCES | CURRENT FIELD · 42 SPEAKERS WITH 3+ DEBATES

Answer. Yes, the published rankings contain a substantial repeatable speaker-performance signal. Among the 42 speakers who meet the current three-debate minimum, a one-way random-effects model assigns 60.0% of single-appearance score variance to persistent between-speaker differences. The estimated reliability rises to 81.8% for a three-debate mean, 88.2% for five, and 93.7% for ten.

The ordinal positions are much less stable. The first and tenth raw averages differ by only 3.53 points, 38 of 41 adjacent gaps are under one point, and the median model-based 95% rank interval spans 17.5 places. Removing one observed debate moves a speaker across a median four-place range; 19 of 42 span at least five places.

The correct interpretation is therefore layered. The scorecards support broad claims about consistently stronger and weaker recorded performance, especially for speakers with larger samples. They do not support treating rank 6 as demonstrably better than rank 7, or a three-debate leader as a settled estimate of context-independent debating ability.

Strongest defensible conclusion

Slugfester's rankings are **useful descriptive summaries with a real stable component, but exact rank numbers overstate precision**. The defensible unit is a sample-qualified performance band, not an ability ladder whose neighboring positions have been statistically separated.

The analysis concerns scores in the selected Slugfester corpus. It does not measure truth, expertise, general intelligence, persuasion, or performance in debates that were not assessed.

Technical summary

The public Rankings page orders speakers by the arithmetic mean of their published overall side scores in ranking-eligible one-on-one debates. This paper separates three questions: whether scores contain persistent speaker signal, whether a speaker's estimated mean is reliable, and whether the resulting exact rank position is stable.

0.60 single-appearance intraclass correlation under the one-way model	0.83 median split-half rank correlation among 19 speakers with at least six debates	17.5 median width in places of the model-based 95% rank interval
---	---	--

What is stable

- Between-speaker variation is larger than within-speaker variation: estimated SDs are 3.83 and 3.13 score points, respectively.
- Across 20,000 random split halves for 19 speakers with six or more appearances, median rank correlation is 0.83, with a central 95% range from 0.74 to 0.91.
- The two halves share an average 4.51 of the top five speakers; all five match in 54.4% of partitions.

What is provisional

- Only 42 of 157 observed speakers meet the public three-debate threshold, and 14 of those 42 sit exactly at three appearances.
- The top ten are compressed into 3.53 points. Four are based on exactly three debates, and only six have at least a 90% modeled probability of remaining in the top ten.
- Changing from raw average to within-debate mean margin preserves a strong overall association but replaces three members of the top ten.

Decision-useful reading

Use the rankings to identify broad, repeatedly observed performance bands and to open the underlying scorecards. Do not interpret a one-place difference as evidence of a real ability difference unless the samples and uncertainty support it.

A ranking can carry signal without fixing an exact order

Three distinct properties are often compressed into the word *stable*. **Score signal** asks whether the same speaker tends to receive similar scores across appearances. **Mean reliability** asks how much a speaker's observed average would move under another comparable sample. **Rank stability** asks whether the ordering against every other eligible speaker would remain the same. The third requirement is much harder because small mean changes can reorder a dense field.

Property	Question	Paper's diagnostic
Persistent signal	Do speakers differ beyond appearance-to-appearance variation?	Variance components and single-score ICC
Mean reliability	Does averaging more appearances reduce noise?	Reliability at 3, 5, and 10 debates
Ordering reproducibility	Would another half-sample preserve the ranking?	20,000 split-half partitions
Position uncertainty	How wide is each plausible rank band?	Empirical bootstrap and hierarchical rank draws
Context sensitivity	Does another reasonable performance summary reorder speakers?	Margin, generation, and shrinkage comparisons

The central distinction

A three-debate mean can be a fairly reliable estimate on a 50-100 score scale while still producing a wide rank interval. Reliability concerns the value; rank precision also depends on how tightly every speaker's value is packed around it.

What the public number means


The displayed average is the arithmetic mean of a speaker's published side scores. Team and panel scorecards are excluded. Opponent quality, topic, affirmative burden, debate year, and assessment generation remain visible context but are not coefficients in the public average. The site already labels samples of one to three as limited, four to six as developing, and larger samples as established; this paper tests how much those cautions matter.

The current ranked field is selective and sample-thin

The September 1 snapshot contains 212 ranking-eligible one-on-one assessments and 424 speaker appearances from 157 unique speakers. Eighty-three speakers appear once and 32 appear twice. The default minimum therefore excludes 115 speakers - 73.2% of everyone observed - and retains 42.

Speakers retained at each minimum-appearance threshold

157 unique speakers across 212 ranking-eligible one-on-one assessments



The public default of three debates ranks 42 speakers and excludes 115

Figure 1. Thresholds change both field size and interpretive confidence. Raising the minimum from three to five nearly halves the field, from 42 to 22; only seven speakers have ten or more appearances.

The ranking is therefore not a census of public debaters. It is a summary of speakers who happen to recur in a selected debate catalogue. A high minimum improves stability but magnifies selection toward prolific figures and heavily represented topic communities.

Coverage tradeoff

Lower thresholds broaden participation but produce more provisional estimates. Higher thresholds improve precision but answer a narrower question about the small group of speakers repeatedly represented in the corpus.

Individual speakers vary meaningfully from debate to debate

Across the 42 ranked speakers, the median within-speaker SD is 3.07 points and the pooled SD is 3.13. The median observed score range is 7 points. Twelve speakers span at least 10 points and one spans at least 15.

Observed score ranges for selected repeated speakers

Mean marker and minimum-maximum line; selection spans sample sizes, ranks, and observed variability




Figure 2. The selected speakers illustrate why one lifetime average cannot erase debate context. Larger samples can reveal wider ranges simply because they expose more topics and opponents; a narrow three-case range is not proof of universal consistency.

The variation is substantively plausible. A speaker can face a different burden, opponent, format, or topic, and can perform better or worse on the day. The score is designed to assess the argument in a transcript, not to recover a context-free personal trait.

Averaging several appearances recovers a strong speaker signal

The one-way random-effects decomposition estimates between-speaker variance of 14.67 score points squared and within-speaker variance of 9.79. That corresponds to SDs of 3.83 and 3.13. The single-appearance intraclass correlation is 0.60: a majority of variance in this selected repeated-speaker sample aligns with persistent speaker differences rather than within-speaker fluctuation.

Random-effects reliability by number of appearances

Share of score-mean variance attributed to persistent between-speaker differences under the one-way model




Figure 3. Under the model, averaging reduces the within-speaker component by the number of appearances. The estimated reliability of the mean reaches 81.8% at three debates, 88.2% at five, and 93.7% at ten.

This is evidence against the claim that rankings are arbitrary. It is not a warrant for exact positions. The model assumes appearances are exchangeable observations around a speaker-specific mean, while the actual catalogue selects topics, opponents, and generations non-randomly.

Signal is not separation

A reliable average may still be statistically indistinguishable from several nearby averages. Reliability tells us that the number carries information; separation tells us whether its precise ordinal location is defensible.

Independent-looking halves largely reproduce the broad order

For the 19 speakers with at least six debates, the analysis randomly divides each speaker's observed scores into two halves and ranks the resulting half-means. Across 20,000 partitions, the median Spearman correlation is 0.83 and the central 95% range is 0.74 to 0.91.

Random split-half reproducibility among speakers with at least six debates

20,000 within-speaker random partitions; 19 speakers; Spearman rank correlation and top-five overlap




Figure 4. The split halves agree strongly on the overall order and share an average 4.51 of the top five speakers. The exact top five match only 54.4% of the time, illustrating the difference between broad reproducibility and exact membership.

The test deliberately focuses on speakers with enough appearances to form two minimally useful halves. It cannot validate the 23 ranked speakers with only three to five debates, and it is not a temporal forecast because the partitions are random rather than earlier-versus-later career periods.

What this establishes


The strongest repeatedly observed speakers tend to remain strong when their existing records are divided. The test supports performance tiers more clearly than it supports the exact composition or order of a short leaderboard.

The top of the published list is densely packed

The current leader averages 88.33 and the tenth-ranked speaker 84.80, a difference of 3.53 points. Across the full 42-speaker list, 38 of 41 adjacent gaps are under one point; the median adjacent gap is only 0.33. Tiny changes in an average can therefore cross several neighbors.

Published mean scores for the current top ten

Current three-debate minimum; focused 84-89 scale with exact means and sample sizes



The first-to-tenth difference is 3.53 points; four of the ten have exactly three appearances

Figure 5. The focused scale makes the actual separation visible. Four of the current top ten - Joseph Schmid, Scott Clifton, Sean Carroll, and Lance Bush - have exactly three appearances, while Matt Dillahunty and Alex O'Connor have 27 and 26.

A one-decimal display is arithmetically accurate but can be read as more epistemically precise than the sample permits. The rank number magnifies that problem because it converts a 0.1-point change into a visible ordinal event.

Precision mismatch


The calculation can distinguish 86.1 from 86.0 without the evidence establishing that the underlying speakers differ. Numerical resolution is not the same as inferential resolution.

Model-based rank bands are wide even near the top

A hierarchical one-way model shrinks small samples toward the ranked-field mean and draws each speaker's latent mean from its posterior uncertainty. It is more conservative than simply resampling the three observed scores of a three-debate speaker, because it uses the pooled within-speaker variance rather than treating a narrow three-case range as complete knowledge.

Model-based rank intervals for the current top fifteen

One-way random-effects posterior; line is 95% rank interval, dot is current raw rank



Intervals quantify latent-rank uncertainty inside the selected corpus, not uncertainty about all possible debates

Figure 6. The intervals do not say every rank is equally plausible. They show that several current top-ten speakers plausibly trade places with nearby speakers. Even the strongest records support bands more naturally than exact ordinal claims.


Across all 42 speakers, the median interval spans 17.5 places and 35 intervals span at least 10. The model produces 1,340 distinct top-ten sets in 20,000 draws; the most common complete set appears only 21.7% of the time.

Larger samples narrow rank uncertainty, but do not eliminate it

Appearance count and hierarchical interval width have a Spearman association of -0.59. The median width is 21.5 places among speakers with three or four appearances, 18 among those with five to nine, and 10 among the seven speakers with at least ten.

Appearance count and model-based rank-interval width

42 ranked speakers; interval width is upper minus lower endpoint of the 95% latent-rank interval



Most current leaders remain plausible, but the boundary is porous

The model gives six of the current top ten at least a 90% probability of remaining in the top ten: Joseph Schmid, Scott Clifton, Alex Malpass, Matt Dillahunty, Bart Ehrman, and Alex O'Connor. The remaining four are credible leaders but sit closer to the boundary under uncertainty.

Probability of top-ten membership for the current top fifteen

20,000 latent-mean draws from the one-way random-effects model




Figure 8. Probability falls gradually rather than at the published tenth-place line. Graham Oppy, Lance Bush, and Sean Carroll retain substantial modeled top-ten probability, while Stephen Law and Daniel Dennett remain credible entrants from just below it.

This is exactly the situation in which tiers communicate better than a hard cutoff. A leading band can acknowledge strong evidence without pretending that the tenth and eleventh observations occupy categorically different performance classes.

One observed debate can move many speakers several places

The leave-one-debate-out check recomputes one speaker's mean after omitting each appearance in turn while leaving the other 41 means fixed. The median best-to-worst span is four places, and 19 speakers span at least five. A three-debate speaker would also fall below the public eligibility threshold if any appearance disappeared.

Largest leave-one-debate-out rank spans

Current 42-speaker field retained; line shows best to worst rank after omitting one observed appearance




Figure 9. Ross Douthat has the largest observed range and correspondingly large rank sensitivity. Several other speakers move far enough to cross common top-ten or top-half boundaries when one debate is omitted.

This check is deterministic and easy to audit. It does not imagine new debate contexts; it only asks how much influence each already-published assessment has. For small samples it is therefore a minimum warning, not a complete uncertainty estimate.

Practical display rule

When a speaker's rank moves materially under leave-one-out analysis, the site should foreground the score range and underlying debates rather than the exact place number.

A different reasonable performance question changes the order

Raw average score asks how highly the speaker's sides were scored. Mean margin asks how far those sides finished above or below their opponents in the same debates. The two rankings are strongly related, with Spearman rho 0.879, but only seven current top-ten speakers remain in the margin top ten.

Raw-score rank and within-debate-margin rank for the largest shifts

Same 42 speakers; blue is published mean-score rank, rust is mean score-minus-opponent rank




Figure 10. Sean Carroll's raw rank is 7 but his mean-margin rank is 21 because his three debates were high-scoring on both sides. Michael Shermer moves from 19 to 6, and Richard Dawkins from 18 to 8, because their mean margins are stronger than their raw averages alone suggest.

Neither metric is the uniquely correct measure of ability. Raw score rewards absolute rubric performance; margin embeds opponent and debate context. Their disagreement demonstrates that a leaderboard is partly a choice about what question deserves ordering.

Generation centering and shrinkage preserve the broad raw ranking

The earlier assessment generation averages 81.32 per appearance and the later generation 78.17. Centering each appearance on its generation mean produces $\rho = 0.985$ with the raw order and preserves all ten current top-ten members. A one-way empirical-Bayes shrinkage adjustment produces $\rho = 0.998$ and also preserves all ten.

Rank agreement under three reasonable adjustments

Spearman correlation with raw mean-score order; labels report current top-ten overlap




Figure 11. The raw ranking is robust to two measurement-focused adjustments. It is less invariant to the substantive choice between absolute score and within-debate relative performance.

Adjustment	Rank ρ	Top-ten overlap	Largest observed rank movement
Generation-centered score	0.985	10 of 10	Richard Dawkins: 18 to 12
Random-effects shrinkage	0.998	10 of 10	Two places
Mean within-debate margin	0.879	7 of 10	Sean Carroll: 7 to 21

Robust broad signal

The ranking is not an artifact of the known assessment-generation shift or of unshrunk three-case extremes. Its main vulnerability is not disappearance of the speaker signal; it is overinterpretation of exact positions and under-description of debate context.

Methods, uncertainty, limitations, and recommended practice

Analysis steps

- Reproduce the public ranking eligibility rule: one avatar-recognized speaker per side, no scorecard explicitly marked ineligible, and at least three appearances for the default field.
- Extract 424 side-level appearances from 212 one-on-one assessments, including raw score, opponent score, score margin, and locked assessment generation.
- Estimate a one-way random-effects variance decomposition across the 277 appearances belonging to the 42 currently ranked speakers.
- Run 20,000 random split-half partitions for 19 speakers with at least six appearances; compute Spearman order agreement and top-five overlap.
- Run 20,000 empirical resamples and 20,000 hierarchical latent-mean draws; report model-based rank intervals rather than treating empirical three-case variation as exhaustive.
- Recompute ranks under leave-one-out omission, generation centering, random-effects shrinkage, and within-debate mean margin.

Principal limitations

Limitation	Consequence
Selected debate catalogue	Results do not generalize automatically to all speakers or future debates
Non-random topics and opponents	Within-speaker variation is not pure measurement error
One-way normal model	Rank intervals depend on exchangeability and distribution assumptions
Random, not temporal, split halves	Reproducibility is not a career-stability or forecasting test
Known assessment-generation shift	Raw absolute averages should retain generation context
AI-assisted published scores	All estimates inherit model judgment error and later reassessment risk

Recommended practice and next test

Keep the public averages, but pair them with model-based rank bands, leave-one-out sensitivity, and a stronger visual distinction between limited, developing, and established samples. Repeat the paper after the planned full-corpus reassessment, then prospectively test whether present bands predict scores in newly added debates.

Conclusion: the rankings measure performance, but not to the precision their exact order implies

The current Slugfester rankings are neither arbitrary nor final. They detect a real, repeatable speaker-level pattern in the published scorecards. Persistent between-speaker variation exceeds within-speaker variation, modeled mean reliability rises quickly with repeated appearances, and broad rank order reproduces strongly across random halves of the better-sampled records.

That positive result makes the qualification more important, not less. Rank is a relative statistic: it depends not only on one speaker's uncertainty but on every nearby speaker's uncertainty. The top ten occupy only 3.53 score points, almost every adjacent gap is below one point, and the median modeled rank interval spans 17.5 places. A calculation can be exact to one decimal while the ordering it produces remains provisional.

The evidence therefore supports a hierarchy of claims. It strongly supports saying that some repeatedly observed speakers occupy a leading performance band and others a lower band in this corpus. It moderately supports top-ten membership for most current leaders. It weakly supports claims about neighboring positions, especially when one or both speakers have only three to five debates. It does not support a context-free ranking of debating ability, philosophical truth, expertise, or personal worth.

Alternative specifications reinforce this reading. Correcting the known generation shift or shrinking small samples preserves the broad raw order and the complete top ten. Switching to within-debate margin changes three top-ten members, showing that context and metric choice matter more than the measurement correction itself. The stable object is the broad score pattern, not one uniquely privileged ordinal list.

Strongest defensible conclusion

Slugfester's rankings measure a meaningful and repeatable component of recorded argumentative performance. They should be presented as **uncertain, sample-qualified performance bands**. Exact place numbers are useful navigation and description, but they are not demonstrated separations between neighboring speakers.

Recommended public interpretation

- Preserve raw averages and access to every underlying scorecard.
- Display rank intervals or tiers beside exact positions, especially below ten appearances.
- Make limited, developing, and established sample status visually primary rather than auxiliary.
- Offer within-debate margin and assessment-generation context as alternate views, not silent replacements.
- Rerun the analysis after the planned roughly twice-yearly corpus reassessment and test new debates prospectively.

Recommended citation: SLUGFESTER. *Do Slugfester Rankings Measure Stable Performance? Why Broad Score Bands Are Defensible While Exact Numbered Positions Remain Provisional*. Corpus-level analysis report, September 1, 2026.