

Does the CON Side Have an Inherent Advantage?

Why higher scores for CON do not, by themselves, show that the role gives a built-in advantage.

4.70

average CON advantage
in 237 comparable debates

0.80

same-person weighted gap;
its range includes zero

31

people with assessed debates
on both PRO and CON sides

The full role gap mixes together the people and claims assigned to each side. Comparing the same people leaves a much smaller average difference. A modest CON advantage remains possible and deserves a fair test.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

CON scores higher than PRO by an average of **4.70 points across 237 comparable one-on-one assessments**. It leads in 176 debates, loses in 51, and ties in ten. This is a large and consistent raw pattern in the archive. It is not, by itself, evidence that being assigned CON causes a speaker to perform better or causes the AI to award more points.

The positions assigned to the two roles are highly uneven. In the broad religious-versus-skeptical set, the religious side occupies PRO in **164 of 187 debates**. CON's mean advantage in that group is 6.77 points. In the 23 debates where the religious side occupies CON, the role contrast reverses: **CON scores 3.26 points lower**. Much of the apparent role pattern therefore travels with the positions defended and speakers, rather than remaining attached to the word CON.

The expanded data also contain evidence that should not be dismissed. Outside the religious comparison, CON has a **1.54-point advantage across 50 debates**, with a 95% range from repeated draws of debates of 0.12 to 2.92. Equalizing the influence of the two religious-role orientations leaves a 1.76-point contrast. These estimates make a smaller residual role advantage plausible, although neither comparison randomly assigns roles or fully balances claims and participants.

The closest available same-person check is smaller and less certain. Among **31 speakers observed in both roles**, the equal-speaker estimate is 0.84 points and the weighted estimate 0.80. The weighted interval runs from approximately **-0.25 to 1.83 points**. That is compatible with no average role difference and with a modest CON advantage. It is not compatible with confidently declaring an inherent 4.70-point bonus on the basis of these data alone.

The practical conclusion is selective: **do not interpret the raw CON advantage as a demonstrated property of the label, but do not conclude that role effects have been disproved**. The next decisive test should separate label bias, differences in reasoning burden, and the people or claims selected for each side.

Reader's guide

Executive summary	2
1. Three different meanings of “CON advantage”	4
2. The raw pattern is real in the recorded data	4
3. Why combined differences can mislead	5
4. The position contrast reverses the role result	6
5. The residual advantage deserves attention	7
6. What happens when the same people switch roles?	8
7. Why even the same-person check is not decisive	10
8. Does criticism sometimes have a lighter burden?	11
9. Distinguishing label bias from a real task difference	11
10. How readers should use the current evidence	12
11. Conclusion: a higher average does not establish a built-in bonus	12
Methods and sources	13

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. All results refer to the September 4 snapshot, not later additions to the site.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. Three different meanings of “CON advantage”

The phrase can describe three distinct possibilities. First, CON could have an advantage in the arithmetic of this archive: its recorded scores are higher on average. Second, arguing against a claim could sometimes require less work than establishing it. Third, the scoring process could favor a label or a recognizable critical style even when the quality of the arguments is otherwise comparable.

The first possibility is directly observed. The second and third require explanation and testing. Collapsing them produces an easy but invalid reasoning step: CON wins often, therefore CON is inherently easier, therefore the score difference tells us little about the arguments. Each step adds a claim that the raw averages do not establish.

The question also depends on what “inherent” means. An effect that holds across every claim, speaker, and format is much stronger than an average advantage in selected discussions. A formal label can remain unchanged while the practical burden changes radically. Defending “religion is always harmful” is not merely the act of withholding belief; it is a strong affirmative claim about a broad class of practices.

This paper therefore uses **PRO and CON for the stored debate roles**, not as synonyms for believer and skeptic, advocate and critic, or constructive and destructive thinking. It uses the explicit religious-versus-skeptical classification only when examining the positions actually defended. Keeping those labels separate is the first safeguard against a misleading conclusion.

2. The raw pattern is real in the recorded data

The September 4 archive contains 253 assessments. The role analysis uses the 237 with comparable locked one-on-one scores. Sixteen assessments in other formats are not mixed into this comparison. Each eligible debate contributes one difference: CON score minus PRO score. Positive values favor CON; negative values favor PRO.




Figure 1. CON's score is usually higher, but the distribution contains many PRO advantages. The average is 4.70 points and the middle value five. The dashed line marks the mean; zero marks equal overall scores.

How to read this graph. Each bar counts debates with that CON-minus-PRO difference. Teal to the right means CON higher; rust to the left means PRO higher; gray at zero means a tie. The dashed line is the 4.70-point average. The labels PRO and CON identify debate roles, not religious belief.

The average difference is **4.70 points**, with a middle value of five. CON leads in **176 of 237 debates**, or 74.3% of the full eligible set. Excluding the ten ties, it leads in 77.5% of the remaining comparisons. These are two ways of reporting the same counts, not two independent pieces of evidence.

The 95% range from repeated draws of debates is **3.85 to 5.55 points**. Repeatedly selecting whole debates from the archive therefore preserves a substantial positive average. This makes the raw result resistant to ordinary changes in the mix of observed cases. It does not show that the result would survive changing the archive's selection process, balancing the positions, or correcting a systematic judging preference.

The scores themselves are existing 5.6 Sol assessments. This Astra-era paper recalculates and reinterprets them; it does not supply a second model's independent judgment of all 237 debates. Any reasoning step about a label preference must be tested against the underlying judging process, not inferred from the fact that the calculations reproduce correctly.

3. Why combined differences can mislead

Suppose an archive contains two speakers. One consistently scores 86 and the other 76. Neither improves or declines when switching roles. Now suppose the stronger speaker appears much more often as CON, while the weaker speaker appears much more often as PRO.

Hypothetical speaker	Mean as PRO	Mean as CON	Appearances by role
Stronger speaker	86	86	2 PRO; 8 CON
Weaker speaker	76	76	8 PRO; 2 CON
Combined role average	78	84	10 appearances per role

The combined archive shows a six-point CON advantage even though **neither person has any role advantage at all**. The difference is produced by who occupies the roles. These numbers are a worked example, not SLUGFESTER observations.

The same issue can arise with topics and claims. If one role more often carries a difficult or ambitious conclusion, the role average can differ even if the scoring rules treat equivalent reasoning identically. An archive of debates is not a randomized experiment: participants often choose motions, opponents, or formats in which they believe they can make a strong case.

This problem is sometimes discussed under the heading of composition effects: the makeup of the groups changes the comparison. The everyday question is simpler and more useful: **are we comparing the same kinds of people doing the same kinds of work needed to support a claim?** If not, the role label may be standing in for several other differences.

4. The position contrast reverses the role result

In 164 of the 187 religious-versus-skeptical debates, the theist side is PRO. That is **87.7%** of the relevant comparison set. Because the skeptical side usually scores higher in these debates, the imbalance gives CON many opportunities to inherit a skeptical-side advantage.




Figure 2. The archive is divided into three non-overlapping groups. CON leads when the theist is PRO, trails when the theist is CON, and has a smaller advantage outside the religious comparison. The role label is not carrying the same meaning across these groups.

How to read this graph. Dots show average CON-minus-PRO gaps; lines hold the middle 95% from repeated draws of whole debates. Parentheses count debates. Teal dots right of zero favor CON; the rust dot left of zero favors PRO. The groups do not overlap. They change the people and claims as well as who occupies each role.

When the theist is PRO, the average CON advantage is **6.77 points**. When the theist is CON, the average CON difference is **-3.26 points**. In the latter group, PRO scores higher in 18 of 23 debates. A universal claim that the stated CON side is inherently favored is difficult to reconcile with this reversal without introducing further conditions.

The reversal does not show that role is irrelevant. The skeptical advantage is smaller when the skeptic is PRO, and different people or claims may occupy the two orientations. But it does show why identifying the side actually defended matters. A reader who automatically translates PRO as theism and CON as atheism will misread exactly the cases that help test the simple role explanation.

The [Folley-Scrivener debate, number 247](#), is an accessible example. The God-independent position is PRO and leads 85 to 76. The [O'Connor-Baker-Hytch debate, number 251](#), also places the skeptical challenge on PRO and ends 84 to 83. These examples differ sharply in margin, but both prevent the stated role from being mistaken for a worldview.

5. The residual advantage deserves attention

Outside the 187-debate religious comparison, CON averages **1.54 points higher across 50 debates**. It leads in 29, trails in 18, and ties in three. The range from repeated draws, 0.12 to 2.92, lies just above zero. This is a change from the earlier paper's smaller comparison set, whose interval included zero.

The update matters. It would be misleading to carry forward the old claim that the outside-group estimate is plainly inconclusive under the same simple resampling method. The expanded data provide somewhat

stronger evidence of a modest positive difference in that selected group. The paper should change when its evidence changes.

At the same time, “outside the religious comparison” is not a neutral experimental control. It includes secular philosophical disputes, religion-adjacent subjects, and debates between different religious or purposive accounts. Those discussions differ in their burdens and participants. The group's positive mean therefore leaves open whether the residual comes from the role itself, its associated tasks, the selected speakers, or the judging process.

A second check gives equal total influence to the two religious-role orientations. Instead of allowing 164 theist-PRO debates to outweigh 23 theist-CON debates, it averages the two group averages. The result is **1.76 points**, with a range from repeated draws of about 0.52 to 3.10.

The balancing calculation: $(6.77 + -3.26) / 2 = 1.76$ points after rounding. It asks what the contrast looks like when the two observed orientation groups have equal influence. It does not tell us what would happen if the same people were randomly assigned to defend the opposite claims.

This is a useful correction to the most obvious imbalance. It is not a complete adjustment for who appears in each group or what they argue. The two estimates - 1.54 outside the religious set and 1.76 after orientation balancing - make a modest residual advantage plausible. Neither establishes an inherent or universal CON bonus.

6. What happens when the same people switch roles?

The archive contains **31 speakers with eligible appearances on both PRO and CON**. For each person, we compare their average CON score with their average PRO score. This removes one major source of distortion: some people are generally stronger or weaker assessed performers than others.

The equal-speaker average difference is **0.84 points**, with an interval of -0.40 to 2.04. Nineteen speakers have a higher mean as CON and 12 as PRO. The pattern is mildly tilted toward CON, but the variation between speakers is large enough that the average remains uncertain under this check.

31 speakers observed in both roles



Figure 3. Each equal-sized dot represents one of the 31 speakers observed in both roles. Dots above the diagonal have higher averages as CON; dots below it have higher averages as PRO. The axes show only the part of the 0-100 scale containing these averages.

How to read this graph. One equal-sized teal dot represents one person. Left-to-right position is their average as PRO; bottom-to-top position is their average as CON. The dashed diagonal marks equal averages. Above it means higher as CON; below it means higher as PRO. Both axes enlarge the 67-91 part of the 0-100 scale.

The weighted calculation gives more influence to speakers who provide more information about the contrast. A person with ten appearances on each side supplies a more balanced comparison than a person with 19 on one side and one on the other. The weight is proportional to the product of the two role counts divided by their total. This yields an estimate of **0.80 points**, with an interval of **-0.25 to 1.83**.

This calculation is closely related to a within-person regression, often called a speaker fixed-effect comparison. The technical name matters less than what it does: it asks whether the same person's scores differ across roles after removing their overall average. It does not assume that two people with the same role are equally capable.

The result is much smaller than the raw 4.70-point gap. But subtracting 0.80 from 4.70 and announcing that a particular percentage of the role advantage has been “explained away” would imply a causal decomposition the design cannot support. The comparisons use different weighting and a restricted set of crossover speakers. They are diagnostics, not interchangeable stages of one experiment.

7. Why even the same-person check is not decisive

People do not switch roles at random. A speaker may defend claims they know especially well and criticize claims they find unusually vulnerable. The speaker may face different opponents, operate under different time constraints, or appear in debates drawn from different stages of their career. The archive may also select their most visible exchanges rather than a representative set of performances.

The assessment process adds another complication. Earlier and later assessment groups differ in absolute score levels. A speaker whose CON appearances are mostly earlier and PRO appearances mostly later could show a role contrast partly produced by that imbalance. The current analysis separates the main assessment groups in the broader results, but it does not establish full comparability across every process detail.

The same-person result therefore makes a large, universal label advantage less convincing; it does not prove that the role has no effect. The upper end of the weighted interval still allows a modest average CON advantage in this kind of comparison. The lower end allows a small PRO advantage. The honest conclusion is uncertainty about a relatively small residual, not certainty that the true effect is exactly zero.

Different comparisons answer different role questions




Figure 4. The estimates shrink as the comparison changes, but each row answers a different question. Intervals crossing zero are compatible with no average difference under the stated resampling assumptions. None of these rows is a randomized role experiment.

How to read this graph. Dots show four ways of comparing CON with PRO. Lines hold the middle 95% from repeated draws of debates or speakers, as appropriate. Right of zero favors CON; a line crossing zero allows no average advantage. Rows answer different questions and reuse records; do not average or add them.

This distinction is an important statistical habit. **“Not clearly different from zero” does not mean “shown to equal zero.”** A study can be unable to distinguish a small positive effect from no effect because its sample is limited or its observations vary. Conversely, an interval just above zero does not make an effect large, universal, or shown to have a particular cause.

8. Does criticism sometimes have a lighter burden?

There is a genuine logical difference between establishing a complex argument and showing that the argument has not established its conclusion. A positive case may depend on several starting claims. A critic may need to identify a serious weakness in only one essential starting claim or connection. This can make criticism efficient without making it unfair.

Suppose a conclusion requires A, B, and C. The defender must provide adequate support for all three and explain why they jointly lead to the conclusion. The critic need not produce a complete alternative theory of the universe to show that C is unsupported. Requiring that would change the burden from evaluating the offered argument to defending a different, much larger thesis.

However, the critic acquires a new burden when making a stronger claim. “This argument does not establish God” is weaker than “God does not exist.” “The evidence does not discriminate between these explanations” is weaker than “the natural explanation is overwhelmingly probable.” A critic who advances the stronger claim must support it. The favorable role pattern does not supply a license for unsupported denial.

Similarly, a defender need not prove more than the agreed motion requires. If the issue is compatibility, a coherent possible model can be relevant and sufficient for that limited task. If the issue is probability, merely preserving possibility is not enough. The fair assessment should track these differences rather than assigning a constant burden according to a label.

The archive's role results are consistent with some difference in reasoning tasks. They do not establish its exact contribution. To measure that contribution, the analysis would need more explicit information about what each side undertook to show and how difficult that undertaking was, not simply which side was stored as PRO.

9. Distinguishing label bias from a real task difference

A label bias would occur if identical reasoning material received a different score merely because it was called PRO instead of CON. That is different from scoring two people differently because one offered weak support and the other identified the weakness. The latter may be a correct judgment of different performances.

The cleanest label test would take the same source material, hide identities where possible, and vary only the formal labels presented to independent judging runs. The claims and burdens must remain logically equivalent. If a harmless relabeling systematically moved scores, the result would be evidence of a label sensitivity in the measuring process.

A real burden test would be harder. It would compare the same speakers under carefully matched motions, with each arguing both affirmative and negative positions across conditions. The subject matter, preparation time, speaking time, opponent quality, and scoring procedure would need to be balanced. Even then, a person's genuine beliefs and expertise could interact with the assigned task.

A third test would compare judging styles. Does the model reward cautious language or the appearance of criticism even when the actual reply is incomplete? Blinded source-based review could separate justified

restraint from empty skepticism. This matters because an assessment system can mistake a style associated with good reasoning for good reasoning itself.

These are proposals for a future major-model review, not studies already carried out for this paper. They should be specified before results are known. Running many alternative tests and reporting only the one that supports a preferred interpretation would recreate the selection problem in a new form.

10. How readers should use the current evidence

For viewers, the role label should be a prompt to identify the motion, not a shortcut to the winner. Ask what PRO must establish and what CON has actually denied. Then examine whether either side has quietly moved to a stronger claim, a weaker target, or a different burden during the exchange.

For speakers, the lesson is not simply to choose CON whenever possible. A narrow critical burden can be handled badly, and a demanding positive case can be handled well. The archive includes clear PRO advantages and theist victories. The useful strategy is to make the target claim explicit, support the starting claims it requires, and avoid treating an opponent's incompleteness as proof of one's own stronger worldview.

For SLUGFESTER, the findings argue against presenting the raw average as an inherent feature of debate. A site-level explanation should distinguish the descriptive 4.70-point contrast from the smaller adjusted comparisons. It should retain the updated outside-group result, acknowledge remaining uncertainty, and encourage source inspection where a score may reflect a mistaken assignment of burden.

The result also supports caution in using combined scores to compare intellectual communities. If roles, claims, and participants are entangled, the score difference does not isolate a single group trait. The appropriate unit remains the actual argument as presented in its context.

11. Conclusion: a higher average does not establish a built-in bonus

Argument one: the full gap cannot safely be credited to the CON role

1. To show that a role causes higher scores, we must separate the role from the people and claims usually assigned to it.
2. CON averages **4.70 points higher** across 237 debates. But theists occupy PRO in **164 of 187** religious comparisons. When theists instead occupy CON, CON averages **3.26 points lower**. The pattern does not remain attached to the CON label.
3. For the **31 speakers who appear in both roles**, the weighted average difference is only **0.80 points**. Its repeated-draw range runs from about **-0.25 to 1.83**. These are different debates by the same people, not a controlled swap of roles.

Therefore: the raw 4.70-point gap does not establish a built-in 4.70-point advantage. Who argues, and what they defend, must be part of the explanation.

Argument two: a smaller role advantage remains possible

1. Failing to prove a large effect does not prove that there is no effect.
2. Outside the religious comparison, CON averages **1.54 points higher across 50 debates**, with a repeated-draw range of **0.12 to 2.92**. Giving the two religious-role groups equal influence also leaves a smaller positive gap.
3. The same-person range allows both no average role difference and a modest CON advantage. None of these checks holds the argument, opponent, topic, and assessment process fully constant.

Therefore: a smaller advantage deserves a fair test. The current evidence neither proves that CON inevitably benefits nor establishes that the process is wholly free of role effects.

The next test should change one thing at a time. Swap harmless role labels while keeping the argument unchanged to test label preference. Use matched claims and speakers to investigate differences in the work each side must do. Until then, treat the role average as a question about the debates, not an explanation that makes their reasoning irrelevant.

Methods and sources

The analysis uses source revision **76d006b37**, frozen September 4, 2026. All 237 eligible one-on-one overall score pairs are verified against their underlying records. The 187 broad religious-versus-skeptical comparisons divide into 164 theist-PRO and 23 theist-CON cases; the remaining 50 form a mixed outside group. The underlying comparable assessments are labeled 5.6 Sol.

Raw and subgroup means average CON-minus-PRO scores once per debate. Orientation balancing averages the two observed religious-role group averages with equal weight. Same-speaker estimates compare within-person role means for 31 speakers with appearances in both roles. The weighted estimate uses $n(\text{PRO}) \times n(\text{CON}) / [n(\text{PRO}) + n(\text{CON})]$ for each speaker. Debate-based intervals resample debates; same-speaker intervals resample speakers. All use 20,000 draws with fixed documented seeds. Shared debates and repeated opponents create dependence not fully removed by these procedures.

The analyses are observational and exploratory. They do not randomize role, identify every local burden, adjust fully for opponent strength, or remove all assessment-process differences. Confidence intervals describe the stated resampling exercise, not all uncertainty about a difference caused by the role.

The [analysis directory](#) contains the classification, speaker-role comparisons, calculations, and figure data. The [Backend library](#) links the companion position, topic, slogan, fallacy, scale, and ranking papers. This edition supersedes the earlier smaller-sample role analysis; in particular, it updates the outside-group result rather than repeating its former uncertainty statement.