

Does the CON Side Have an Inherent Advantage?

The residual role association is now clearer under one weighting - but still not causal.

4.44

raw CON-minus-PRO gap
in comparable debates

1.33

weighted same-person gap
positive resampling range

37

people observed
in both roles

The weighted same-speaker CON gap is now positive across its displayed range.

PHIL STILWELL · SLUGFESTER

October 9, 2026 | Frozen archive: 308 assessments

Recalculated scores; direct slogan evidence retains its September scope.

Executive summary

CON still scores higher on average, but the full difference cannot be assigned to the role itself. Across 291 comparable debates, the average CON-minus-PRO difference is 4.44 points, down from 4.70 in September. The role labels identify positions in a particular debate, not the speakers' permanent beliefs.

Comparing 37 people observed in both roles gives an equal-speaker gap of 1.09 points and a weighted gap of 1.33. The weighted estimate's middle-95% resampling range is 0.32 to 2.30 - now wholly positive. The equal-speaker range still narrowly includes zero.

This is a substantive change from the earlier report, not a reason to announce a causal discovery. The comparisons do not hold topics, opponents, debate difficulty, preparation, or assessment process fixed. The better conclusion is that a modest residual CON association deserves closer investigation; neither an inevitable advantage nor complete absence of an effect has been established.

Reader's guide

Executive summary	2
1. What does “advantage” mean?	4
2. Reverse the positions and the contrast changes	4
3. Follow the same people across roles	5
4. Why the change from September matters	6
5. A reply can be easier, but it still has a burden	6
6. A fair test and a bounded conclusion	7
Methods and sources	7

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. Score results refer to the October 9 snapshot. Direct slogan counts retain their September source coverage; they have not been extended to new transcripts.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. What does “advantage” mean?

There are at least three questions. Do CON sides receive higher recorded scores? Do the same people tend to score higher when occupying CON? Would assigning an otherwise comparable person to CON rather than PRO cause a higher score? The first is descriptive, the second reduces one important difference between groups, and the third requires a design this archive does not provide.

To answer the first, subtract PRO from CON within every comparable debate and average the differences. There are 211 CON-higher results, 68 PRO-higher results, and 12 ties. The median difference is 4 points. The raw association is real in these records; the uncertainty concerns interpretation and extension, not whether the arithmetic can be calculated.

The selected debate sample matters. If more experienced skeptical participants often occupy CON and many religious advocates occupy PRO, the role label will carry some of those differences. Conversely, CON may sometimes face the stronger claim: arguing that suffering is evidence against God can put the skeptic in PRO and the religious response in CON.

2. Reverse the positions and the contrast changes

Among 202 religious-versus-skeptical debates with the religious side in PRO, CON averages 6.21 points higher. Where the religious side occupies CON, the CON difference is -2.75 across 32 comparisons. The other 57 comparable debates average 2.18 points in the CON direction.



Figure 1. The role gap changes markedly when the substantive positions change. These groups do not overlap.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population.

Giving the two religious orientations equal influence produces a 1.73-point CON gap. This asks what the arithmetic would look like under an equal mixture of those two groups. It does not create matched debates,

randomize roles, or hold the topic distribution constant. It should be described as a sensitivity check, not a corrected estimate of the true role effect.

The [Rauser-Malpass debate](#) is a useful reminder. The skeptical argument advances a positive comparative-evidence claim; the religious response contests its force given other background beliefs. Skepticism is not always mere refusal to accept a claim, and the defending side is not always the party taking on the largest positive burden.

3. Follow the same people across roles

The stronger comparison here uses only people with eligible appearances in both roles. For each person, calculate their average CON score minus their average PRO score. It is still a comparison of different debate performances, not a repeat judgment of the same performance.

Giving every person equal influence yields 1.09 points, with a range of -0.02 to 2.21 under repeated draws of speakers. Weighting each person by $n(\text{PRO}) \times n(\text{CON}) \div \text{total appearances}$ gives 1.33 points, with a range of 0.32 to 2.30.

The weight rewards evidence on both sides of the comparison. A person with many PRO appearances and only one CON appearance has a less balanced bridge than a person with substantial records in both roles. It does not guarantee equal precision: score variability and opponent composition also matter.




Figure 2. The weighted same-speaker interval is now positive, while the equal-speaker interval still crosses zero. These estimates answer different questions.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population. Right of zero favors CON. The weighted same-speaker range is now above zero, but no comparison randomly assigns roles, opponents or topics.

Resampling speakers preserves each person's pair of role averages. It does not preserve every shared-opponent dependency, and a systematic judgment error remains in every draw. These are exploratory stability intervals, not fully adjusted causal estimates.

4. Why the change from September matters

The earlier weighted result was approximately 0.80 with a range crossing zero. It would now be inaccurate to retain the sentence “the weighted interval includes zero.” The updated calculation is about 1.33, and its range is positive. At the same time, the equal-speaker estimate is about 1.09 and its lower endpoint is slightly negative. Weighting choices therefore matter to the strength of a directional claim.

Neither result should be rounded into “no effect.” Nor should the positive weighted range be promoted into “CON gets an automatic bonus.” The archive added six people to the cross-role comparison and more performances for returning speakers. Those changes alter the estimate without randomly separating the role from the other features of a debate.

The raw gap is much larger than either within-person estimate. That reduction is compatible with differences in speaker and position composition accounting for a substantial portion of the raw association. It is not a numerical decomposition of causation: the estimates use different records and different weighting rules.

5. A reply can be easier, but it still has a burden

Rejecting a specific argument may require less than proving a competing worldview. A speaker can successfully show that the supplied reasons do not establish a conclusion while leaving the underlying question open. A rubric should recognize that limited success rather than demand proof of the exact negation in every exchange.

However, rejecting an argument is not the same as establishing that its conclusion is false. If a respondent moves to a stronger claim - such as the impossibility of God, the certainty of a natural explanation, or a general claim about religion - the respondent acquires a corresponding burden. The word CON does not erase it.

A useful source-level audit would mark the conclusion each move actually defends, the reason needed for it, and whether the response answers that reason. It would then examine whether unsupported assertions receive different treatment when uttered by a familiar skeptic or a familiar religious advocate. The present averages identify a question for that audit, not its answer.

6. A fair test and a bounded conclusion

A stronger test would use closely matched motions and opponents, have the same participants take different sides where feasible, standardize preparation and format, and obtain independent assessments blind to the proposed result. It should distinguish success in defeating one argument from support for a wider alternative and publish disagreement cases for inspection.

Reason 1: the archive shows a positive raw CON association. **Reason 2:** controlling partially for position mixture or person reduces it substantially. **Reason 3:** the updated weighted same-person range is positive, but other design differences remain uncontrolled.

Therefore: a modest residual CON advantage is more plausible in the current recorded comparisons than a blanket statement that no role association remains. Its cause and generality are unresolved.

Therefore: readers should interpret a score as success at the particular burden undertaken, not assume that the side label determines how hard the job was. The appropriate next step is a matched, source-sensitive check, not an automatic correction to all PRO scores.

Methods and sources

This is a fixed October 9, 2026 research edition, based on source revision e65bc3c0b6. The score studies reconstruct published records rather than commission fresh assessments. The direct slogan counts retain the September 5 review of 187 transcripts; this edition recalculates those saved records and adds a separately labeled current score check. No new slogan reviews or debate assessments were performed.

The source manifest records file fingerprints. The analysis uses 20,000 deterministic resampling draws with baseline seed 20260904, retaining it for edition comparability. Whole debates are drawn for paired debate estimates; shared speakers for cross-role estimates. Ranking models and their dependence checks use the units explained in paper seven. These calculations do not capture every selection effect, repeated-speaker dependency, caption error or systematic judging bias.

The [data and methods page](#) links the source records, classification decisions, calculations, chart inputs and reading guides. Download the [current research package](#) and [historical direct slogan study](#). Earlier editions remain preserved in the repository's history. Debate links identify published assessments, not fresh independent evidence for this research.