

How Much Do Opponents Affect Recorded Performance?

Repeated matchups and within-person comparisons do not support a dependable schedule correction.

291

comparable debates

21

repeated speaker pairs

+0.06

same-procedure check: points

The small opponent-score association is not stable when the eligible comparison set changes.

PHIL STILWELL / SLUGFESTER

October 9, 2026 | Frozen source archive: 308 assessments

Source-linked editorial analysis. No debate scores changed.

Summary

Facing a demanding opponent could lower a speaker's assessed performance by exposing weaknesses, or improve it by eliciting more precise answers. A strong opponent also raises the standard needed to finish ahead. Those are different effects: a smaller winning margin need not mean a lower score for the speaker's own arguments.

This study therefore examines each speaker's own recorded score, not just wins and losses. It uses 291 comparable one-on-one debates, 582 appearances and 183 speakers. Fifty-five people faced at least three distinct opponents. Twenty-one pairs met more than once, producing 43 assessments of repeated matchups. Only two of those pairs show a strict reversal in which side scores higher. Neither reversal occurs within the 12 pair groups that have repeat meetings under the same assessment procedure.

A descriptive model compares the same person's performances against opponents with different average scores elsewhere. It excludes all meetings of the focal pair when calculating the opponent's average. After accounting for the speaker, assessment procedure and PRO/CON role, a five-point higher opponent average is associated with a 0.49-point lower own score. Requiring same-procedure opponent experience reduces the sample and yields a 0.06-point increase. On those same remaining records, the original pooled opponent measure already gives a near-zero 0.02-point decrease. The result depends strongly on which appearances qualify, not simply on changing the procedure measure.

Main conclusion: The archive does not currently support a dependable opponent penalty or a new opponent-adjusted ranking. Repeated-matchup results mostly preserve their direction, while the small score association changes substantially with how opponent experience is made comparable.

1. Why opponents deserve a separate study

A speaker's public average combines debates on different topics, against different opponents and under different assessment procedures. It is reasonable to ask whether some people face a harder schedule. But “harder” needs a definition. Here it means opponents who have higher assessed scores in other debates, not opponents deemed philosophically correct or naturally talented.

Using the opponent's score in the same debate would mix several things: a shared topic, the interaction itself, and the assessment of both performances. Instead, we construct an outside-matchup average. For an appearance by A against B, B's proxy is the average of B's scores against people other than A. Every A-B meeting is excluded, not just the current one. This prevents a repeated pair's own results from defining its supposed outside difficulty.

The proxy still is not a pure measure of ability. It reflects B's own opponents, topics, dates and judging conditions. It uses the whole frozen snapshot, including records added later than the focal assessment, so it describes the archive retrospectively; it is not a forecast available before the debate.

2. The evidence and the eligible comparison set

The frozen archive contains 308 assessments, but 17 do not meet the comparable locked one-on-one criteria used in the existing research edition. The retained 291 debates contain 179 earlier-process and 112 later-process assessments. No video URL repeats within this comparable set. Repeated speakers remain common, however, so 582 appearances are not 582 independent people or experiments.

Evidence set	Count
Comparable one-on-one debates	291
Speaker appearances / distinct people	582 / 183
People facing at least three distinct opponents	55
Distinct unordered speaker pairs	269
Repeated pairs / their debates	21 / 43
Main fitted comparison	295 appearances, 80 people, 188 debates
Same-procedure opponent-proxy check	180 appearances, 54 people, 132 debates

For the main fitted comparison, an opponent must have at least three scored appearances against other people. A focal speaker must have at least two appearances remaining after that requirement is applied, so there is some within-person comparison. Three appearances do not necessarily mean three distinct outside opponents. A stricter check requires five outside appearances. These restrictions explain why the fitted study uses fewer than all 582 appearances.

Some people can meet the two-appearance fitting rule without facing three distinct opponents. The count of 80 fitted speakers therefore does not conflict with the count of 55 people who faced three or more distinct opponents. The two figures answer different questions.

3. What happens when the same pair meets again?

Repeated pairs provide a direct descriptive check. We fix the order of the two names and subtract the second person's score from the first person's score in each meeting. A strict reversal requires both a positive and a negative difference. Reaching a tie is not counted as a reversal.

Only two of the 21 repeated pairs reverse: Alex O'Connor-Max Baker-Hytch and Michael Licona-Richard Carrier. Both comparisons cross the earlier/later process boundary. Across all repeated pairs, the median distance from the smallest to largest margin is three points. Within the 12 pair-by-procedure groups with repeated meetings, it is two points, with no strict reversal.

Most repeated pairs keep the same lead

Score margin: first named person minus second




Figure 1. Repeated pairs usually preserve the direction of their score difference.

How to read it. Each row is one pair. A mark gives the first named person's score minus the second's in one meeting. Circles use the earlier procedure; squares use the later one. The line joins the smallest and largest margins. Red rows cross zero with strictly opposite leads. Overlapping marks can hide repeated identical margins. These are observed meetings, not uncertainty intervals.

These observations support stability of relative recorded performance within this small repeated-matchup set. They do not show that a second meeting is a replication of the first: subjects, source quality, formats and the speakers' preparation can change. Nor can the two cross-procedure reversals be attributed to the procedure change. Topic, time and procedure move together in those observations.

Concrete examples

Alex O'Connor and Trent Horn meet three times in the comparable record. O'Connor's margins are eight, six and ten points in assessments 60, 122 and 215. Justin Schieber and Randal Rauser have the same seven-point margin in later-process assessments 291 and 308, although both people's scores rise by two points. John Lennox and Peter Atkins likewise retain a seven-point margin in assessments 4 and 51 even as both scores fall by five points.

Those last two examples show why margins and individual performance must be distinguished. A stable margin can coexist with substantial changes in both scores. A schedule adjustment aimed at wins is not automatically an adjustment to the quality of each speaker's arguments.

The two reversals are small in their margins. O'Connor moves from two points behind Baker-Hytch to one point ahead; Licona moves from three points behind Carrier to one point ahead. These changes matter for exact ordering but do not demonstrate a wholesale inversion of performance. All 21 pairs and their scores are listed in the accompanying data rather than only the most striking examples.

4. Comparing the same person against different opponents

The main model uses a separate starting level for each speaker. In ordinary terms, it asks whether a person scores differently when facing opponents whose outside-matchup averages differ, instead of confusing differences between people with differences between their opponents. It also includes the earlier/later assessment procedure and whether the speaker is PRO or CON.

Technically, this is an ordinary least-squares model with speaker fixed effects - separate intercepts for the speakers - and common coefficients for the other recorded factors. The outcome is the person's overall score on the 100-point scale. The reported coefficient is multiplied by five so it describes a five-point difference in the opponent's outside average. The source code and saved appearance table make the calculation reproducible.

Comparison	Appearances	Own-score change for a five-point higher opponent average
Same speaker only	295	-0.42
Plus procedure and PRO/CON role	295	-0.49
Plus broad research topic	295	-0.82
Give each focal person equal total weight	295	-0.38
Require five outside appearances for opponent	203	-0.33

Comparison	Appearances	Own-score change for a five-point higher opponent average
Earlier-process focal appearances only	151	-0.50
Later-process focal appearances only	82	-0.36
Opponent average from same procedure only	180	+0.06

The main estimate is small relative to the five-point contrast in the predictor. The underlying within-person association is also weak: after the main adjustments, the correlation is about -0.11 . This is not the fraction of variation caused by opponents, and the coefficient is not a causal effect. It is a fitted association within selected records.

The broad-topic check uses the eight research groups preserved by the earlier studies plus an Other questions group. These are not the current browsing categories and do not resolve every topic difference. Equal-person weighting gives every included person the same total influence, rather than allowing people with more eligible appearances to contribute more rows.

The small association is not dependable

Own-score change for a 5-point higher opponent average




Figure 2. The opponent-score association depends on how comparability is handled.

How to read it. Each dot is a fitted association, not a causal effect or uncertainty interval. Negative values mean a lower own score with a higher opponent average elsewhere. All meetings of the pair are excluded from that opponent average. The red final row also requires those outside scores to use the same procedure; its eligible sample is smaller. There are no error bars because sampling uncertainty is not estimated.

The final row uses only an opponent's outside meetings assessed under the same procedure as the focal meeting. A further check holds its 180 appearances fixed but restores the original pooled opponent average:

the estimate is -0.02 , compared with $+0.06$ for the same-procedure average. Both are close to zero. Thus the reduced sample already removes most of the original -0.49 association; we should not attribute its disappearance to procedure matching alone. The negative result is not dependable across these eligibility choices.

5. Why this is not a corrected ranking

An automatic correction would require more than a plausible story about tough opponents. We would need evidence that the opponent measure captures a repeatable performance constraint and that the proposed adjustment improves comparisons. The present small and definition-sensitive association does not meet that standard.

Several obstacles remain. Matchups are selected, not assigned at random. A speaker may choose a specialist opponent on an unusually difficult subject. Strong opponents can elicit better explanations as well as expose weaknesses. An opponent's outside average is itself an estimate based on an uneven schedule. Assessment procedures and source conditions differ. With many one-time speakers, there is not enough overlap to estimate separate effects for every person, topic, period and interaction reliably.

For that reason, the table reports descriptive estimates rather than significance tests or confidence intervals. The spread between specifications is not a 95% uncertainty interval. The rows reuse many of the same debates, and changing a restriction changes both the comparison and sometimes the people represented. Absence of a stable negative coefficient is not proof that opponents have no influence in any debate.

Paired-comparison models such as Bradley-Terry address a related but different question: the relative tendency to prevail in comparisons. Turner and Firth describe extensions for contest-specific factors and structured abilities. Such models can be useful, but turning rubric scores into wins discards information about absolute performance and makes a tied or one-point result qualitatively different from a wider margin. We have not fitted a Bradley-Terry ranking here, and invoking that model would not remove selection or judging bias.

6. What readers can conclude now

First, repeated meetings provide a reassuring check on the direction of recorded score differences: strict reversals are rare, and none appears in the small same-procedure repeat set. This complements the existing split-half ranking study, though both use overlapping archive data and therefore are not independent replications.

Second, the specific claim that a stronger opponent systematically suppresses a speaker's own score receives only weak, unstable support here. The main association is less than half a point for a five-point opponent contrast. The restricted sample gives essentially no association with either the original or same-procedure proxy. These results do not warrant a schedule penalty or a revised public ranking.

Third, opponent context still belongs beside averages. Readers should inspect who was faced, how often, on which questions and under which assessment procedure. That is more defensible than displaying a supposedly corrected number whose adjustment has not been validated.

7. A useful next test

A stronger future study would record its comparison rules before examining new scores, use one assessment procedure, and select speakers with enough overlapping opponents. It could then test whether an opponent-based adjustment predicts held-out recorded performances better than a person's own prior average and topic history. Entire debates, not individual sides, should be held out together. Opponent averages must be computed without the held-out data.

That would be a test of predictive usefulness, still not a randomized estimate of how much an opponent causes performance to change. A separate reliability study would be needed to investigate judging errors. The immediate practical decision is simpler: retain the current unadjusted scores, present sample sizes and opponent records, and do not convert this exploratory association into a new league table.

Methods and sources

The study uses the frozen October 9, 2026 research edition at source revision e65bc3cob6. The comparable set is identical to that used by the existing corpus studies. It includes all 291 locked one-on-one assessments, excludes the 17 other formats, and changes no scores. The main outcome is the published overall score, not a rescored move average or a win indicator. Procedure labels describe how the assessment was produced, not the date on which the original debate occurred.

The [data and methods page](#) supplies the full appearance table, every repeated pair, the eight specifications and the additional same-record comparison. The [analysis package](#) specifies exclusions, source fingerprints and definitions. Model coefficients are independently checked using a full design-matrix fit in the publication verification, separate from the residualized calculation used to produce them.

Methodological reference: Heather Turner and David Firth (2012), [Bradley-Terry Models in R: The BradleyTerry2 Package](#), *Journal of Statistical Software*, 48(9). This is background on paired-comparison models, not a method claimed to have been applied in this paper. Related Slugfester study: [Do Slugfester Rankings Measure Stable Performance?](#).