

Evidence, faith, and fair assessment

Understanding the evidence gap in SLUGFESTER's Charts

SLUGFESTER · Methods and interpretation · October 8, 2026 · Clarified edition

The Charts page shows a substantial evidence-score gap between arguments supporting religious claims and arguments challenging them. Does it reflect the arguments, the judging system, the selection of debates, or some combination? An accusation of anti-Christian AI bias deserves a substantive answer. So does the possibility that some religious arguments leave more of their evidential burden unmet.

The defensible conclusion is specific: in this selected archive, under SLUGFESTER's published criteria, supporting arguments receive lower evidence scores on average. The difference is broad and survives several descriptive checks. It is neither a finding about every Christian nor an independent demonstration that the judging model is impartial. Where the recorded reasoning leaves essential evidential checks undone, this document identifies evidence neglect directly.

Recorded measure	Supporting	Challenging
Mean evidence score	66.5 / 100	74.2 / 100
Mean share below 70	60.6%	25.8%
Debates with higher evidence mean	38	188

These figures give each debate equal weight and count each selected move once. They describe 226 included debates and 5,405 assessed moves in the October 8 revision 2 snapshot. They are evidence-dimension summaries, not official overall debate scores. [1, 2]

Why take the assessments seriously? The process connects judgments to source passages, states the burden each side undertakes, separates initial reviews, records disagreement resolution, and calculates scores by fixed rules. A skeptical label earns no automatic credit; a religious label warrants no automatic deduction. The underlying examples include strong Christian arguments and weak anti-religious arguments.

Why retain uncertainty? The selected debates are not a representative sample, repeated speakers influence the archive, and all 226 included assessment records identify the same model, 5.6 Sol. Separate reviews by one model can reproduce shared mistakes. Mathematical reproducibility establishes what was calculated; it does not establish that every judgment was correct.

This explanation defends an inspectable method and examines whether some understandings of faith reduce the pressure to supply public evidence. That proposal merits investigation, but these charts do not measure its causal effect.

Prepared with AI assistance. This document adds descriptive calculations and interpretation of existing records; it supplies neither new debate judgments nor independent human validation.

1. What the recorded gap actually is

The chart's positions are **Supporting** and **Challenging** the religious claim at issue. They are not interchangeable with the website's original PRO and CON columns, and they are not classifications of a speaker's entire identity. A Christian can challenge another religious claim; a skeptic can make a positive claim that carries its own burden. The archive also includes non-Christian religious positions. A Christianity-only conclusion would require an explicitly defined, separately reviewed subset. [1]

Across all 226 included debates, the equal-debate mean evidence score is **66.5 for supporting moves and 74.2 for challenging moves**: a 7.7-point difference on the recorded 0-100 scale. Challenging moves have the higher debate-level evidence mean in 188 debates; supporting moves have the higher mean in 38. There are no ties in these unrounded evidence means. These counts are not counts of overall debate victories. [2]

Dimension	Supporting	Challenging	Difference
Logical coherence	76.3	82.7	+6.4
Evidence & support	66.5	74.2	+7.7
Responsiveness	83.1	89.0	+5.9
Relevance & burden	82.4	84.6	+2.1
Precision & clarity	81.8	86.3	+4.5
Calibration & charity	74.1	81.5	+7.4

Table 1. Equal-debate means, all included moves. Difference = challenging minus supporting. All scores use a 0-100 scale. Source: frozen snapshot and calculation record [1, 2].

Evidence is the lowest mean dimension for both positions and has the largest mean difference of the six dimensions in this summary. This gives the concern a concrete location: support for premises and the connection between evidence and conclusion. It does not establish that the other dimensions independently validate the evidence judgments. They arise from the same assessment system and can share errors.

The direction is also present in each broad topic group. For God and theism, the evidence means are 65.8 and 74.0 across 136 debates. For specific religious claims, they are 68.6 and 75.0 across 62. For religion, meaning, and society, they are 64.9 and 73.4 across 28. These are descriptive comparisons within the selected archive, not estimates of how all believers or all skeptics reason. [2]

The visible consistency is worth explaining. It is not a license to describe Christians as uniformly careless, dishonest, or incapable of rigorous thought. Nor does a repeated difference become an unbiased result merely because it is repeated.

2. How to read the evidence heatmap

An evidence score assesses the adequacy of a move's support for its particular conclusion. It is not a count of citations, a probability that Christianity is false, or a measure of how much evidence exists outside the debate. A short, well-supported objection can outperform a long list of claims; an uncited but correctly defended conceptual distinction can be strong. [3]

At threshold 70, a cell reports the percentage of selected moves with scores **strictly below 70**, first within each debate and then averaged equally across represented debates. A score of exactly 70 does not count. The page uses the same calculation at thresholds from 50 to 100. Darker cells mean more below-threshold scores, not stronger evidence against the worldview.

For example, if one debate has one of two moves below the threshold and another has one of ten below it, the equal-debate rate is $(50\% + 10\%) / 2 = 30\%$. Pooling the moves would give $2 / 12 = 16.7\%$, a different measure. Equal weighting stops a debate with many selected moves from automatically dominating the chart.

Share of moves below the selected evidence threshold

Equal-debate average (%) · all 226 included debates




Figure 1. Both positions converge at the high end of the slider. These are descriptive shares, not probabilities that a claim is false or confidence intervals. Sources [1, 2].

For all included moves, the below-70 rates are 60.6% and 25.8%. The difference remains in the same direction at thresholds 50 through 90. At 95 and 100, both positions reach 100%, so those settings cease to distinguish their evidence scores. A high threshold will classify competent arguments as below threshold; it does not turn them into wholly unsupported arguments. [2]

The figure above summarizes unique moves across the whole included set. The website's individual family cells use their own family subsets, so this aggregate is not an average of the ten displayed family percentages. Families overlap: one move can concern both morality and suffering. Moreover, a family can appear on only one side of a debate, producing different supporting and challenging denominators. Those labels must be read with the cell percentages.

The snapshot is manually refreshed. Changing the slider recalculates a view of saved ratings; it does not call an AI model or reassess a debate. Later debate-page corrections or additions do not silently change this edition.

3. What the rubric asks of both sides

SLUGFESTER assesses the argument presented, its burden, and its response to the live exchange. It does not award points for the ultimate truth of a religion or for conformity to atheism. The six fixed dimension weights are 25% logical coherence, 20% evidence and warrant, 20% responsiveness, 15% relevance and burden, 10% precision and clarity, and 10% calibration and charity. A warrant is the reason the offered evidence supports the conclusion. [3, 4]

Evidence and warrant. Strong support must be relevant and sufficiently specific, reliable, and connected to the claim. The rubric asks whether selective examples, an unexamined authority, a speculative mechanism, an uncertain probability estimate, or an undefended bridge is doing more work than it can bear. The standard concerns adequacy, not the number of references. Missing an inline citation is not automatically a defect; its importance depends on what needs verification.

The evidence scale gives 90 or above to particularly well-defended, sourceable support; 75-89 to appropriate support with remaining gaps; 50-74 to support relying substantially on unchecked assertions or similar weaknesses; and below 50 to effectively unsupported or contradicted fact-dependent conclusions. The slider's default of 70 is a display threshold, not a verdict on a religion.

Reasoning is not restricted to laboratory experiments. Textual interpretation, historical testimony, conceptual analysis, mathematical reasoning, and philosophical argument can provide support appropriate to their claims. Scripture can establish what a passage says without first establishing divine authorship. Establishing divine authorship, a miracle, or a universal metaphysical conclusion requires additional reasoning. A rule that discounted a valid inference merely because it was theological would misapply this standard.

Logical coherence. A conclusion must follow from the premises with stable meanings. An unsupported premise belongs primarily under evidence; it is not automatically a logical fallacy. The same weakness should not generate multiple deductions unless it produces distinct, documented failures.

Responsiveness and burden. A reply must address the actual objection, while a constructive opening is assessed against its motion and known burden. A critic who argues only that a proposed proof fails need not supply a complete cosmology. A critic who asserts that no gods exist, that all religious experience is delusion, or that science explains every relevant phenomenon undertakes a stronger claim and must defend it. A Christian can likewise defeat a claim of logical incompatibility without proving that Christianity is true.

Precision, calibration, and charity. Terms and qualifications must be clear; confidence must fit the support; opposing arguments must be represented accurately. Fluency, applause, reputation, and polite manners do not substitute for these qualities. A sweeping dismissal of religion can lose credit under exactly the standard that penalizes a sweeping dismissal of unbelievers.

One public standard for both sides. Every participant must justify the claims they ask an audience to accept with evidence and reasoning that can be examined and challenged. Private certainty, claimed revelation, or an asserted self-authenticating experience cannot substitute for that support. Reports of personal experience can be assessed as testimony; their reliability and the conclusions drawn from them receive no exemption from scrutiny. Equal scrutiny applies to each side's actual claims, including a critic's positive assertions.

4. How an assessment reaches the chart

The documented workflow has several separations designed to prevent the desired outcome from driving the score. The public process guide describes the standard September 2026 workflow; retained historical records preserve their own versions and exceptions. An explanatory guide should not be read as a claim that every historical control was identical. [4-7]

1. Secure the source and define the burden. The recording and transcript establish what was said and by whom. The workflow checks coverage and material gaps, preserves timestamps and file fingerprints, and records the motion and each side's burden. A fingerprint can reveal that a file changed; it cannot prove that an automatic transcript heard every word correctly.

2. Select and lock the argumentative moves. Important claims, replies, concessions, and unresolved objections are connected to source passages. Section weights and move importance are fixed before scoring. This limits later cherry-picking to obtain a preferred result. It does not eliminate judgment from deciding which contributions are important.

3. Obtain two isolated reviews. Reviewers receive controlled evidence and rules without the other review, old scorecards, winner labels, or published explanatory prose. Separate contexts reduce copying and anchoring on a previous verdict. They do not make the models unrelated, and they do not conceal every identity or religious position visible in the passage. The chart's included records all identify 5.6 Sol.

4. Resolve disputes and preserve the record. Fixed rules identify disagreements. A fresh adjudication context examines disputed evidence and anonymous options without calculated overall totals. Required attribution uncertainties trigger audio checks. The record retains the initial reviews and resolutions. For example, the archived Batch 11 manifest records twenty review contexts for ten debates, and a later standalone adapter explicitly links its two passes and final record. [6, 7]

5. Calculate rather than improvise totals. Software derives a move score from its six dimension ratings, a side's section score from importance-weighted moves, and an overall score from prelocked section weights and the permitted burden-completion adjustment. The default adjustment is zero; the documented bound is -5 to +5 for consequential burden completion not already counted. Stability rules constrain changes from the original reviews; they are consistency safeguards, not proof of impartiality.

6. Explain and publish the locked result. Critiques must explain the recorded assessment. AI-generated strengthened arguments are separately labeled and do not affect participant scores. Fallacy and bias labels do not impose additional numerical penalties. Published records undergo source, calculation, quotation, and presentation checks.

7. Build a frozen chart snapshot. The chart builder checks selected published moves against compatible assessment records and copies the recorded dimensions, or reconstructs them through the recorded deterministic rules. It does not infer new scores from the displayed summaries. The October 8 edition includes 226 of 300 published debates at capture; its 74 exclusions have recorded decisions. Excluding incompatible records protects comparability, but can also affect sample composition. [1, 2, 8]

5. Why unequal burdens can yield unequal scores

A public debate need not distribute identical argumentative tasks to its participants. Establishing a proposition and showing that one offered argument fails are different tasks. Fairness requires judging the burden actually undertaken, not forcing each side to make equally expansive claims.

Consider an illustrative claim: the universe has an explanation; therefore the Christian God exists. The conclusion requires more than a general explanatory principle. Depending on the case presented, the speaker may need to justify a cause beyond the universe, agency, relevant divine attributes, and the link to Christianity. A critic who identifies a missing link can make a warranted objection without having solved the origin of the universe. That does not establish atheism; it limits what the proposed argument has established.

The rule works in the other direction. If a skeptic says evolution logically excludes divine guidance, a Christian can challenge that entailment by distinguishing biological randomness from the absence of intention. Establishing compatibility can defeat the alleged contradiction without independently establishing guidance. The Plantinga example in the examples section receives a high evidence score for just such a bounded contribution.

Possibility, probability, and explanation must remain distinct. A possible explanation may defeat a strict impossibility claim. It does not automatically show that the explanation is likely, superior to rivals, or what actually occurred. This matters for proposed divine purposes in suffering, harmonizations of historical accounts, speculative naturalistic origins, and multiverse proposals alike.

Similarly, quoting a text may strongly support a claim about that text while only weakly supporting the historical occurrence it reports. Explaining how a belief could arise naturally does not by itself establish that the belief is false. Showing that a source has errors does not establish that every claim it contains is false. Neither side receives permission to skip these inferential steps.

Burden asymmetry is therefore a plausible contributor to the aggregate gap. It is not a complete explanation established by the chart. In a supplementary comparison restricted to debates where both positions have constructive moves, the evidence means remain 65.0 and 72.8 across 152 debates. Among 225 debates with replies on both sides, they are 66.7 and 74.2. Matching these broad roles reduces one simple comparison problem, but it does not match the difficulty, specificity, or subject matter of every claim. [2]

The appropriate response is to inspect whether the critic earned the narrower burden and whether the advocate received full credit for what was actually established. Artificially equalizing the scores would not answer either question.

6. What the accusation of AI bias would require

An unfavorable outcome is compatible with bias, but it does not establish bias by itself. Nor can the method be defended by saying that an AI has no feelings and therefore no prejudice. A model can systematically favor familiar wording, conventional assumptions, a style of argument, or particular content without having a human motive.

Research on language models used as judges has documented sensitivity to answer position and verbosity and investigated self-preference and reasoning limitations. Those findings concern other evaluation tasks and models; they do not measure anti-Christian bias in SLUGFESTER. They establish why a judging system needs targeted validation rather than a presumption of neutrality. [9]

Three different questions should remain separate. **Procedural fairness** asks whether the same stated rules and source controls apply. **Measurement validity** asks whether the ratings track the kind of argumentative quality the site intends to assess. **Outcome disparity** asks whether one position scores lower. Stronger procedures help, but they do not logically settle the other two questions.

SLUGFESTER has grounds to defend traceability, fixed calculations, separate initial reviews, and criticism anchored in specific moves. These controls constrain arbitrary scoring and make objections answerable. They are more informative than an unexplained winner label. The contrasting examples in this document also show that the implementation is capable of rewarding Christian contributions and penalizing anti-religious overstatement.

Those examples do not establish equal error rates. The initial reviewers and adjudicator can share assumptions. A stability rule may retain a shared initial mistake. The inventory can omit a strong argument or compress an important distinction before either reviewer sees it. A transcript can misattribute speech. A critic can be credited with caution while an equally cautious believer is treated as making a stronger claim. These are concrete possibilities for an audit, not findings that any particular error occurred.

The rubric itself also deserves scrutiny, but SLUGFESTER's subject is public argument. Neither side may bypass assessment by declaring its knowledge private or immune to criticism. An assertion that a conviction is divinely guaranteed, intuitively certain, or self-authenticating supplies no additional evidence unless the speaker defends an assessable reason to accept it. A belief's possible personal significance cannot repair an unsupported public inference. Readers can challenge the rubric and its application; private assurance is not an exception to its evidential requirements.

The responsible defense is provisional: the gap is a documented result under explicit rules; it is not a demonstrated artifact of anti-Christian prejudice, and it is not yet a demonstrated absence of such prejudice. The right next step for a bias allegation is an identifiable scoring error or a controlled test. A critic need not produce an entire alternative scoring system before a specific error deserves attention.

7. Identifying evidence neglect in the record

The saved assessments identify concrete omissions that warrant the description **evidence neglect**: evidential work needed for the stated conclusion is left undone. Uncertainty about why an omission occurred does not erase the omission. A commitment to fair assessment requires identifying these failures in religious arguments as directly as in skeptical arguments.

Here, evidence neglect describes a failure to supply or engage relevant support, check a consequential evidential assumption, or address counterevidence material to the claim actually made. It need not be deliberate. It is a criticism of the presented argument, not a diagnosis of the speaker's motives or every belief they hold. The following examples identify the missing work rather than inferring neglect merely from a low score.

A probability claim without its probability justification

In Debate 01 at 94:39, Craig argues that life-permitting constants are extraordinarily unlikely by accident and favors intelligent selection. The recorded **evidence score is 61**, while the overall move score is 68. The published critique identifies no articulated probability measure, prior distribution, treatment of observation selection, or systematic comparison with the relevant alternatives. These omissions matter because the inference depends on how likely the observations are under competing explanations. [18]

This is a specific failure to justify the evidential comparison on which the stated conclusion depends. Calling a life-permitting range narrow does not, by itself, establish its probability or make intention the superior explanation. Invoking personal certainty about Christianity would supply none of that missing work. The criticism concerns the version presented, not a claim that every possible fine-tuning argument fails. The Carrier example later applies the same demand to a naturalistic probability claim.

A broad empirical claim without representative support

In the same debate at 79:24, Craig invokes Michael Murray's book while denying reflective awareness of pain across familiar mammals. The recorded **evidence score is 58**, while the overall move score is 71. The critique acknowledges the relevance of distinguishing pain from reflective self-awareness, but finds that the generalization across species exceeds the supporting reference. It calls for representative empirical support and an explanation of why reduced reflective complexity sufficiently answers the suffering objection. [18]

The point is not that Craig names no source: he does. The failure is that naming that source does not adequately establish the breadth or argumentative consequence of the empirical claim as presented. A potentially relevant distinction receives credit; the omitted evidential work remains a defect.

These are grounds for identifying evidence neglect in particular assessed contributions. They do not establish deliberate evasion, that every Christian neglects evidence, or that faith caused these omissions. Those are additional claims requiring additional support. Nor are the saved critiques beyond correction: a reader who identifies the missing justification in the relevant exchange has grounds to challenge them. That openness does not require withholding criticism until every possible alternative explanation has been ruled out.

8. Could an emphasis on faith weaken evidential habits?

One proposed explanation is that some Christian environments praise faith in ways that normalize insufficient attention to evidence. Stated carefully, this is a hypothesis about learned standards of justification, not an assertion that Christians have an inherent defect or that all uses of the word faith mean the same thing.

The proposed mechanism. If maintaining a religious conclusion despite unresolved objections is treated as a virtue, while revising it is treated as disloyalty or spiritual failure, a person may face less incentive to identify disconfirming evidence, compare rival explanations, or specify what would change the conclusion. Arguments can become tools for defending an already protected commitment rather than tests that could genuinely revise it. In public debate, that could appear as confidence outrunning support or repeated reliance on premises that the audience has not been given reason to accept.

Public claims require public support. SLUGFESTER's debates ask listeners to evaluate publicly stated claims. Both sides must supply premises, sources, testimony, and reasoning adequate to the claims they make and open to examination. Sincerity or psychological certainty does not discharge that burden. A report of an experience can be considered as testimony, but its reliability, proposed cause, and relevance to the conclusion must be defended rather than insulated from scrutiny.

Craig's own discussion is relevant evidence, not an excuse. He distinguishes self-authenticating knowledge through the Holy Spirit from public demonstration, yet treats his own witness as genuinely divine and incompatible witnesses as mistaken. He expects contrary believers' confidence to yield to evidence while defending the rational security of the genuine witness. [10] As a public defense, that presupposes the disputed reliability of his own experience instead of establishing it. This illustrates the concern: a claimed source of certainty is protected before the public comparison is made. **The knowing/showing distinction does not discharge the burden.** Craig's public claims still require assessable support; inaccessible assurance cannot repair a missing premise or answer counterevidence. Endorsing a role for evidence does not establish that adequate evidence was supplied.

Christianity also contains explicit defenses of reason. In *Fides et Ratio*, John Paul II criticizes fideism for neglecting rational knowledge and philosophical discourse. He nevertheless maintains a role for revelation. That is a counterexample to treating the Christian concept of faith as uniformly hostile to inquiry. It is a statement of theological commitments, not evidence that every Christian practice meets those commitments. [11]

The strongest version of the hypothesis therefore concerns **insulation from correction**: whether a community or speaker makes a favored conclusion resistant to adverse evidence. A profession of trust in reasons earns no exemption either: the reasons must actually be supplied and withstand scrutiny. Skeptical, political, or secular commitments can also be insulated from correction. General research on motivated reasoning offers a possible mechanism for selective evaluation across people, not a finding that Christians in these debates were more affected by it. [12]

To test the proposal, independent readers would need to code actual uses of faith, willingness to revise, treatment of counterevidence, and demands for public justification without seeing the existing scores. The comparison would need to account for topic, argumentative role, speaker repetition, and judging errors. The present snapshot contains no such faith variable. It cannot establish that faith caused the evidence gap, let alone determine what share of the gap it caused.

9. Other explanations worth investigating

The faith hypothesis should compete with alternatives rather than inherit the entire gap by default. Several explanations could operate together. Their relative importance has not been established by this analysis.

The size of the claim. A defense of a personal, morally perfect, intervening God, or of a particular revelation, can require several distinct bridges. Evidence for a first cause is not automatically evidence for all of those attributes. An opposing move may need to expose only one unsupported bridge. This predicts smaller gaps when the claims and burdens are narrowly matched.

Transfer of authority to a mixed audience. A premise supported by scriptural or ecclesial authority may be intelligible within a tradition but remain unestablished for an audience questioning that authority. A speaker accustomed to shared commitments may compress the very premise a critic needs defended. Secular advocates can make the same mistake by assuming disputed philosophical naturalism or a narrow account of knowledge. This predicts lower scores when contested background commitments are left unargued, whichever side uses them.

Historical distance and limited discrimination between explanations. Some religious arguments concern unique ancient events, unobservable intentions, or ultimate origins. Their available evidence may leave competing explanations difficult to distinguish. That can constrain the public case even when the advocate is careful and well informed. Naturalistic origin stories and sweeping accounts of religion's evolution face analogous demands. The prediction concerns the type and ambition of the claim, not the speaker's religion.

Debate incentives and compression. A public performance may reward memorable formulations, confidence, speed, or reassurance to an existing audience. Detailed source criticism and explicit uncertainty take time. The archive does not measure audience reward, speaker motives, or preparation, so these remain possible explanations. They should be tested against actual omissions and format differences, not inferred from a low score alone.

Selection and repeated speakers. The site's chosen debates may overrepresent certain apologetic traditions, successful critics, popular formats, or recurring topics. Removing each speaker's debates in turn leaves the aggregate evidence gap positive, ranging from 6.7 to 8.2 points. That shows that no single represented speaker is necessary for the direction of this result. It does not eliminate network effects, selection bias, or a skewed selection shared across many speakers. [2]

Assessment conventions and shared model error. The model may recognize compressed skeptical objections more readily than unfamiliar theological reasoning, accept a contestable background premise too easily, or mistake an argument's intended scope. Conversely, a familiar named theological argument may receive undue credit. Separate reviews and automatic arithmetic do not settle these possibilities. A matched human and multi-model audit is needed.

The current evidence supports a difference in recorded argumentative performance. The causes could include genuine differences in the support offered, task asymmetry, the selected sample, judging error, and cultural habits of argument. Assigning a precise contribution to any one of them would go beyond the data.

10. Examples that keep the standard symmetrical

These four cases illustrate high evidence ratings for Christian-side contributions and low ratings for skeptical overreach. They were selected for that explanatory purpose, not as a random sample or proof that the system is unbiased. The descriptions below paraphrase the published assessment. The displayed evidence score is one dimension, not the move's overall score. [13-16]

A Christian compatibility argument receives 88

In Debate 144, Alvin Plantinga distinguishes biological randomness from metaphysical absence of intention. The assessed contribution blocks an inference from random mutation to the impossibility of divine guidance. Its **evidence score is 88**, and its overall move score is 91. The critique also limits its reach: compatibility does not establish that guidance actually occurred. The favorable rating rewards the bounded argument without granting the larger religious conclusion. [13]

A Christian textual reply receives 91

In Debate 120, Jonathan McLatchie points to Mary's plural wording in John as compatible with accounts describing several women at the tomb. Its **evidence score is 91**, and its overall move score is 93. The critique distinguishes resolving this particular tension from proving the general reliability or independence of the resurrection accounts. A narrowly supported reply can therefore score highly while contributing only a limited amount to the overall case. [14]

An anti-religious generalization receives 40

In Debate 85, Peter Atkins argues that science and religion are intellectually incompatible and characterizes religion in terms of superstition and wishful thinking. Its **evidence score is 40**, and its overall move score is 59. The critique requires identification of particular incompatible claims and engagement with qualified religious positions, rather than treating the entire category as evidential failure. This is the same objection that should be made to an unsupported blanket condemnation of atheists. [15]

A speculative naturalistic origin receives 44

In Debate 14, Richard Carrier proposes a minimally described, lawless origin that allegedly makes outcomes equally likely. Its **evidence score is 44**, and its overall move score is 65. The critique asks for a specified range of possibilities and a justified probability distribution; lack of laws does not itself establish equal likelihood. Invoking a naturalistic mechanism is not a substitute for defending it. [16]

These cases show that religious contributions can receive high evidence scores and skeptical contributions can receive low ones. They do not show how the same arguments would fare without their religious or skeptical associations, or establish equal treatment on average. That requires matched comparisons and review of disputed cases, including cases less favorable to the site's existing conclusion.

11. How to challenge and improve an assessment

The most useful objection identifies the debate, source passage, selected move, dimension, and reason the recorded judgment is wrong. A Christian reader might show that the critique overlooked an explicit qualification, treated a compatibility defense as a probability claim, ignored a source supplied in the debate, or imposed a burden the speaker never accepted. A skeptical reader can identify precisely the same kinds of errors.

The assessment should then be checked against the complete relevant exchange, not defended merely because its score contributes to the aggregate pattern. Corrections should preserve the old record, explain the reason for the change, and enter a later chart edition through the manual refresh process. The site's correction channel is linked in the references. [17]

A stronger bias audit

A prospective audit should select cases before examining the desired outcome and publish its rules in advance. It should include strong and weak arguments from both positions, positive cases and replies, Christian and other religious claims, and several topics. It should retain disagreements rather than reporting only accepted or favorable cases.

Test irrelevant labels and presentation. Where meaning can genuinely be preserved, vary speaker identity labels, ordering, and stylistic presentation while keeping the argument and context fixed. Blind reviewers to previous ratings and the hypothesis being tested. Do not mechanically swap theological content in a way that changes the argument, its evidence, or its burden, then mistake the resulting score change for prejudice.

Use diverse, qualified judges. Include human reviewers with Christian and non-Christian perspectives, relevant philosophical or historical expertise, and more than one model family. Record their reasons and agreement on specific defects. Human consensus is not infallible ground truth, but it supplies a source of scrutiny that repeated calls to one model cannot provide.

Measure errors and their consequences. Examine missed support, invented claims, burden misassignment, citation mistakes, inconsistent thresholds, and confidence that outruns the source. Compare both the direction and size of rating changes. Group repeated speakers and related debates appropriately rather than treating thousands of moves as independent people.

Separate method revision from outcome targets. If a criterion is defective, improve it and apply the change consistently. If an individual judgment is wrong, correct it. Do not raise Christian scores or lower skeptical scores simply to manufacture parity; do not refuse a correction merely because the aggregate gap might narrow.

The fair conclusion remains open to revision: the saved assessments contain a substantial and inspectable evidence gap, with plausible argumentative and cultural explanations. Their best defense is the ability to show what earned each judgment and to correct demonstrated errors. Neither religious identity nor an AI-generated verdict should be exempt from that standard.

12. Calculation notes and sensitivity checks

This document uses the **October 8, 2026 revision 2** frozen archive, containing 226 debates, 2,686 supporting moves, and 2,719 challenging moves. The catalogue contained 300 debates at capture. Later catalogue growth does not change those denominators. The snapshot's exact file fingerprint and the per-debate inputs accompany the reproducible analysis. [1, 2]

For each debate and position, the new summaries average the evidence dimension over unique assessed moves without importance weights. They then average the debate means equally. This matches the Charts page's equal-debate descriptive approach; it is not a reconstruction of official section-weighted or importance-weighted overall scores. The reported gap is challenging minus supporting. The threshold summaries average each debate's fraction strictly below the chosen value.

Comparison	Debates	Supporting	Challenging	Difference
All included debates	226	66.5	74.2	+7.7
Both sides: constructive	152	65.0	72.8	+7.8
Both sides: replies	225	66.7	74.2	+7.5
Earlier record format	146	67.2	75.6	+8.4
Later record format	80	65.2	71.6	+6.5

Table 2. Evidence means and challenging-minus-supporting differences. The recorded format labels identify methodological subsets for this check, not visitor-facing filters. Source: calculation record [2].

The first three subsets retain only debates containing both positions in the stated comparison. The two assessment-record subsets contain different debates, so their differences cannot isolate a change in judging practice. The earlier representation derives some dimensions from structured findings; later records preserve resolved dimensions more directly. The chart builder's recorded conversion rules matter to comparability. Both subsets show the evidence-gap direction, which does not establish that the scale is perfectly invariant. [8]

In the leave-one-speaker-out check, every debate featuring the omitted speaker is removed, rather than deleting only that speaker's moves. Across the represented speakers, the retained mean evidence difference ranges from 6.7 to 8.2 points. This is a sensitivity range, not a confidence interval. It does not control for all shared opponents, topics, selection effects, or model bias.

The calculations are descriptive. No random sampling design, causal identification, Christian-only estimate, or external human rating standard is supplied. The two role labels do not measure equivalent task difficulty. Scores and mechanisms should therefore be discussed separately.

Reproduction requires the archived snapshot and matching assessment adapters. The analysis script verifies their fingerprints and checks the scope-level means and threshold rates using a second grouping calculation. Reproducing the arithmetic cannot reproduce or validate the full original judgment process by itself.

Sources and reading links

Data, scoring rules, and process

- [1] SLUGFESTER. Charts snapshot, October 8, 2026, revision 2. Frozen inputs, positions, selected moves, dimensions, classifications, exclusions, and assessment fingerprints. [Download the archived JSON](#).
- [2] SLUGFESTER. Descriptive evidence-gap analysis for this explanation, October 8, 2026. Includes scope and role comparisons, threshold sensitivity, record-format checks, and whole-debate speaker omissions. [Calculation record and reproduction instructions](#). [Analysis program](#).
- [3] SLUGFESTER. Reassessment Rubric v2.1. Dimension definitions, shared anchors, burden rules, and penalty controls. [Read the rubric](#).
- [4] SLUGFESTER. From debate to scorecard, September 26, 2026. Public explanation of the standard workflow and its limits. [Read the process guide](#). [Production workflow source](#).
- [5] SLUGFESTER. Reassessment campaign closure and calibration-promotion records. Source and calculation integrity checks are not independent validation of philosophical judgments. [Read the archive overview](#).
- [6] SLUGFESTER. Batch 11 independent-judgment and final-ledger records. Example of retained separate reviews and adjudication provenance. [Review preparation](#). [Final-ledger manifest](#).
- [7] SLUGFESTER. Debate 196 assessment adapter, linking the source lock, inventory, two reviews, disputes, adjudication, final record, and calculated results. [Inspect the saved adapter](#).
- [8] SLUGFESTER. Chart snapshot builder and method. Describes inclusion, source checks, dimension conversions, equal-debate weighting, and manual refresh. [Builder](#). [Chart method](#).

Sources and reading links

External research and Christian primary sources

[9] Zheng, L., et al. (2023). Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. NeurIPS Datasets and Benchmarks. Documents limitations in model judging; it does not test SLUGFESTER, 5.6 Sol, or anti-Christian bias. [Paper and bibliographic record](#).

[10] Craig, W. L. (2022). Evidence, Arguments, and the Witness of the Holy Spirit. Reasonable Faith podcast transcript. Primary statement of Craig's knowing/showing distinction and treatment of incompatible spiritual witnesses. The discussion examines its public evidential deficiency. [Read the transcript](#).

[11] John Paul II (1998). Fides et Ratio, especially sections 52-55. Primary Catholic statement supporting reason's role in faith and criticizing fideism. [Official Vatican text](#).

[12] Kunda, Z. (1990). The Case for Motivated Reasoning. Psychological Bulletin, 108(3), 480-498. A general account of how desired conclusions can influence evaluation; not evidence of a Christian-specific causal effect in this archive. [Journal article PDF](#).

Source-linked assessment examples and corrections

[13] Debate 144. Alvin Plantinga and Daniel Dennett, Science, Religion, and Evolution. Move at 9:57, identifier pro-biological-randomness-distinction. Evidence 88; move 91. [Published assessment](#).

[14] Debate 120. Jonathan McLatchie and Alex O'Connor, Is Christianity True? Move at 25:36, identifier pro-john-plural-women-reconciliation. Evidence 91; move 93. [Published assessment](#).

[15] Debate 85. Peter Atkins and Keith Fox, God vs Science Debate. Move at 43:19, identifier pro-science-religion-intellectual-incompatibility. Evidence 40; move 59. The original PRO label here belongs to the challenging position. [Published assessment](#).

[16] Debate 14. Michael Jones and Richard Carrier, Does God Exist? Move at 65:58, identifier con-minimal-lawless-field-origin. Evidence 44; move 65. [Published assessment](#).

[17] SLUGFESTER. Feedback and corrections. Include the debate, timestamp, move, disputed judgment, relevant source context, and proposed correction. [Open the correction channel](#).

[18] Debate 01. William Lane Craig and Alex O'Connor, Does God Exist? Fine-tuning move at 94:39, evidence 61; animal-awareness move at 79:24, evidence 58. [Published critiques](#). [Saved assessment and source excerpts](#).

External sources checked October 8, 2026. The example descriptions are paraphrases of the saved assessments, not newly verified transcript quotations. The snapshot and pinned repository records preserve this edition's values if live pages later change.