SLUGFESTER / RESEARCH

OCTOBER 2026 / ARGUMENTS AND OPPONENTS

What Makes a Reply Actually Answer an Argument?

A complete response-link audit and eight source-linked close readings.

2,444
recorded replies

58.2%
links aimed at another reply

8
worked examples

A useful reply changes the support for the claim actually being made - not merely a nearby claim.

PHIL STILWELL / SLUGFESTER

October 9, 2026 | Frozen source archive: 308 assessments

Source-linked editorial analysis. No debate scores changed.

Summary

A reply answers an argument when it changes what the opponent's stated reasons can establish. Merely discussing the same subject is not enough. A useful reply identifies the actual claim, challenges a necessary premise or inference, or explains why the objection attacks a stronger claim than the speaker defended. An honest concession can also be a successful reply: it removes a failed objection while making the remaining disagreement more exact.

This study combines a complete audit of recorded response links with eight close readings. The audit covers 112 later-process one-on-one debates containing 3,609 scored moves, including 2,444 replies. Every reply has at least one recorded target. There are 2,978 reply-to-target links, of which 1,733, or 58.2%, point to another reply. The practical lesson is immediate: to understand most recorded links, readers need the exchange leading up to a statement, not just the opening case.

The eight examples illustrate real differences between attacking a premise, questioning an inference, clarifying a conclusion, conceding a point, and answering an inaccurately reconstructed argument. They were deliberately chosen to display these differences, not sampled to measure how often each speaker or viewpoint replies well. The existing scores appear for orientation; they are not new ratings and do not independently validate this interpretation.

Main conclusion: The decisive question is not whether a reply sounds forceful, but whether its reasoning changes the support for the claim actually under discussion. A strong local reply can leave the wider debate unsettled.

1. What the complete response record shows

The unit counted here is a scored move: one recorded argument or reply. A reply can address several earlier moves, so the number of links exceeds the number of replies. Among 2,444 replies, 472 address more than one target (19.3%). Of the 2,978 links, 2,968 cross to the other side and ten remain on the same side. All targets exist, and every target's recorded excerpt starts no later than the reply's excerpt.

These checks establish that the stored response network can be followed. They do not establish that every link is philosophically apt, that every reply answers its target, or that captions are error-free. Same-side links may represent follow-up development rather than opposition; the mechanical audit does not recode them. Earlier-process assessments are excluded from this link census so differences in their recording structure are not mistaken for differences in debate behavior.

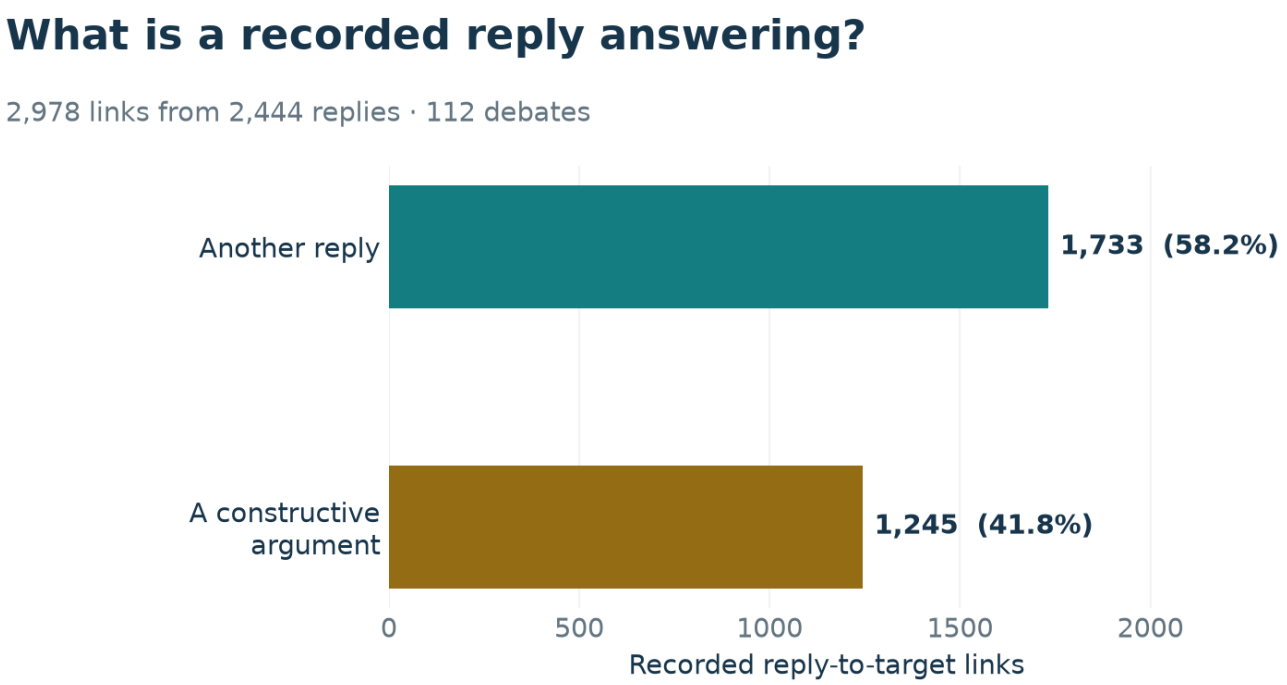


Figure 1. Replies often address another reply, not a constructive argument.

How to read it. Bar length counts recorded links. A reply can have more than one target. Teal means its target is another reply; gold means a constructive argument. Both bars start at zero. These categories describe the stored links, not whether a reply succeeds.

Recorded feature	Count and interpretation
Later-process debates	112; the complete eligible later-process set
Constructive moves / replies	1,165 / 2,444
Replies with multiple targets	472; 19.3% of replies
Links to other replies	1,733 of 2,978; 58.2% of links
Links to constructive moves	1,245; 41.8% of links
Missing targets / later-starting targets	0 / 0; a structural check, not a quality judgment

2. A practical test for an answer

First identify the target as a reasoned claim, not merely a topic. Write down what is being asserted and why. Then ask what the reply changes: the acceptability of a premise, the connection between reasons and conclusion, or the scope of the conclusion. Finally ask what remains unresolved after granting the reply its strongest reasonable reading.

These are overlapping functions, not mutually exclusive bins. A reply may clarify a term and then challenge the resulting inference. We do not turn these categories into a new numerical grading scale or infer their prevalence from eight examples.

Function	What the reader should look for
Challenge a premise	A reason to reject or doubt a claim the argument needs
Challenge an inference	An explanation of why the stated reasons do not establish the conclusion
Clarify the scope	A distinction between the claim actually defended and a stronger claim attributed to it
Concede and revise	Explicit acceptance of a successful objection, with an identified remaining question
Miss the target	A reply to a view the opponent did not defend, leaving the relevant support untouched

3. Eight worked examples

A. Correct the target before answering it

Justin Schieber, debate 308, 19:38. Existing move score: 87/100. Randal Rauser challenges a landlord analogy partly because a theist need not believe that God gives contradictory revelations. Schieber distinguishes God issuing contradictions from people sincerely attributing incompatible messages to God. His argument concerns why a relationship-seeking God leaves serious misunderstandings uncorrected.

This reply removes an unnecessary premise from the objection: Schieber need not prove that God authored both messages. He also distinguishes minor disagreements, which may be educational, from disagreements about salvation, moral authority and God's nature. The remaining burden is substantive: why should a good God correct this amount and kind of error, rather than allow it for other reasons? Clarifying the target keeps that burden visible; it does not discharge it automatically.

B. Separate the phenomenon explained from a nearby phenomenon

Stephen Meyer, debate 306, 13:37. Existing score: 85/100. Peter Atkins appeals to local order emerging through physical processes as an example of information without an intelligent agent. Meyer accepts that spontaneous order occurs, then distinguishes repetitive physical order from a sequence whose particular arrangement performs a biological function. He asks which process explains the relevant sequence, rather than order in general.

The reply engages the explanatory target directly. A demonstration that vortices form is not yet a demonstration of the origin of a particular functional sequence. But the distinction also limits Meyer's success: identifying a gap in this reply does not establish that natural processes cannot fill it, or that an intelligent designer did. This is a useful example of a theistic-side reply earning substantial local credit without proving the larger design claim.

C. Locate the disputed logical step

William Lane Craig, debate 303, 55:19. Existing score: 84/100. Michael Nugent says that if God knows future acts, either people lack freedom or God knowingly creates agents who will do evil. Craig challenges an inference from infallible knowledge to necessity: certainty that an act occurs is not, by itself, a demonstration that the act could not have been otherwise.

This addresses one branch of Nugent's challenge. It does not settle every account of freedom and foreknowledge, and it does not answer the second branch about responsibility for knowingly creating agents who do evil. The distinction matters because a reply can successfully expose a missing inference without answering every issue raised in a compound objection. Credit should track the part actually answered.

D. Show why the objection does not require the alleged assumption

Justin Schieber, debate 308, 36:08. Existing score: 85/100. Rauser argues that the disagreement objection imports exclusivism: the idea that mistaken religious beliefs lead to severe eternal consequences. Schieber identifies harms that do not depend on that premise. False beliefs can obstruct an informed relationship with God and create conflict between believers even if salvation is more inclusive.

This is a scope clarification with argumentative force. Removing eternal punishment does not remove every reason a truth-loving God might have to correct fundamental misunderstandings. Schieber must still compare these reasons with possible benefits of permitting disagreement. The reply preserves a narrower argument; it does not show that every form of religious disagreement is decisive evidence against God.

E. A concession can improve an exchange

Philip Goff, debate 245, 2:12:05. Existing score: 86/100. In a hypothetical about a meaningful book matching the initial structure of the universe, Joe Folley distinguishes two probability questions. One varies the universe's initial conditions; another holds those conditions fixed and considers the later, independently produced book and language. Goff acknowledges that he had understood the objection differently and that Folley may have a good response to this case.

The valuable feature is precise updating. Goff identifies what he misunderstood, grants the force of the revised example, and reserves judgment about other examples. He does not pretend the entire fine-tuning dispute is settled. A concession is therefore not equivalent to losing the debate; here it makes the remaining disagreement easier to investigate. The recorded score is consistent with that reading but is not evidence independent of the original assessment.

F. Relevant material can still leave the main question unanswered

Randal Rauser, debate 308, 31:05. Existing score: 77/100. Rauser answers Schieber's team-building analogy by invoking human error, tailored instruction and a shared moral framework. These are relevant possibilities. He also returns to the point that God need not reveal incompatible messages, even though Schieber has already distinguished that claim from failures to correct incompatible beliefs.

The analogy itself makes Rauser's concern understandable: Schieber imagines a manager giving different instructions. Yet the explicit target includes the clarification about uncorrected misunderstanding. A complete reply needs to explain why tailored instruction or shared morality addresses persistent

disagreement on the fundamental issues Schieber identifies. Naming possible goods is not the same as showing that they require this pattern of confusion. The reply makes contact but leaves an important connection unfinished.

G. An attribution error can disable a wider criticism

William Lane Craig, debate 303, 1:23:43. Existing score: 64/100. Craig challenges Nugent's rational-committee account of morality while describing Nugent as denying objective moral values. The earlier saved passage explicitly presents unjustified harm as objectively wrong and offers a naturalistic account of moral reasoning. Nugent's proposal may still face difficulties, but treating it as a denial of objective morality does not answer the view he stated.

The relevant criticism would need to show why the proposed standards fail to establish objective obligations. Merely combining a naturalistic explanation of moral capacities with an accusation of nihilism bypasses that task. The case illustrates why accurate attribution is not a cosmetic virtue: when the premise attributed to the opponent is wrong, the alleged inconsistency may disappear. This conclusion concerns the recorded exchange, not Craig's character or every version of his moral argument.

H. Expose an arbitrary probability partition

Dan Linford, debate 265, 109:35. Existing score: 91/100. Miles Donahue distinguishes a few broad ways God might create life, including naturally embodied life, miraculous sustenance and unembodied beings. Linford points out that the number of named categories depends on how finely possibilities are divided. Naturalistic possibilities can also be grouped into just two broad alternatives.

This reply identifies a specific missing inference: counting categories does not assign their probabilities. Three labels do not automatically yield one-third probabilities, and grouping many physical hypotheses under one heading does not make them collectively improbable. Donahue's more modest suggestion that embodied life has some substantial probability still needs independent support. Linford's objection does not show that the probability is tiny; it shows why the category count alone cannot establish it.

4. What the scores can and cannot tell us

The selected scores range from 64 to 91, but that range is not a statistical result about reply types. The cases were chosen after inspecting the arguments, and the same assessment system already rewards responsiveness, accurate representation and calibrated conclusions. It would be circular to treat agreement between these examples and those scores as independent proof that the rubric is correct.

The paper instead contributes an explicit, checkable explanation of argumentative contact. Readers can disagree with an assessment by identifying the target, showing which inference the reply changes, and explaining what remains. This is more informative than appealing only to a high or low number. The structural audit supports navigation through exchanges; the case readings support distinctions about reasoning. Neither substitutes for the other.

5. Conclusion and use

Three conclusions follow. First, exchanges are genuinely sequential in this record: 58.2% of the recorded response links lead to another reply. Second, good engagement is not confined to one worldview or to refusing concessions; Meyer, Craig, Schieber, Goff and Linford supply different useful examples. Third, local success must be separated from overall success. Refuting an analogy, narrowing an objection or acknowledging a mistake rarely decides the complete motion.

For a reader, the most useful three questions are: **What exactly is being answered? Which supporting step changes? What remains to be shown?** For a future larger study, those questions could form a coding guide applied to a randomly selected set of exchanges by independent AI agents called by the administrative AI. That would estimate frequencies and disagreement between those AI reviewers. This paper supplies the worked distinctions and source trail, not the results of that unperformed study.

Methods and sources

AI review roles. Slugfester's assessment reviewers, judges and adjudicators are independent AI agents called by the administrative AI. Independence means separate review contexts with controlled inputs; the agents may use the same model and share biases. Review arrangements vary across assessment periods. The direct slogan study used one AI review per transcript. AI-role wording clarified October 22, 2026; research snapshot and results unchanged.

Snapshot: October 9, 2026, source revision e65bc3c0b6. The full archive contains 308 assessments; the reply audit uses all 112 later-process comparable one-on-one cases. We verified the saved ledgers against their digital fingerprints and extracted recorded targets and excerpt start times. The eight close readings use saved caption excerpts for reply and target, with public assessments as context. They are AI-assisted editorial readings, not a fresh assessment by blinded independent AI agents, newly corrected transcripts, or a full-video rewatch. All passage descriptions above are paraphrases.

Read the source-linked assessments: debate 308: Schieber-Rauser, debate 306: Meyer-Atkins, debate 303: Craig-Nugent, debate 245: Goff-Folley, and debate 265: Donahue-Linford. Exact debate paths, video URLs, timestamps, move identifiers and saved excerpts for every example appear in the selected evidence file.

The study's data and methods link the complete 2,978-edge audit, calculation instructions, editorial coding and source fingerprints. No debate scores were changed.