

Where Is the Theist Disadvantage Largest?

Stable research categories, more cases, and an uncertain topic order.

7.28

largest observed topic gap:
religion, culture and meaning

3.04

smallest observed topic gap:
resurrection and history

8

stable research groups
not the browsing categories

The topic differences remain useful; exact topic rankings remain uncertain.

PHIL STILWELL · SLUGFESTER

October 9, 2026 | Frozen archive: 308 assessments

Recalculated scores; direct slogan evidence retains its September scope.

Executive summary

Religion, culture, and meaning still has the largest observed average gap, but a settled league table of topics would overstate the evidence. Its 25 comparisons average 7.28 points in the skeptical direction. Resurrection and historical evidence is closest, at 3.04 points across 24 comparisons. Both figures are smaller than the September edition's 8.27 and 3.53.

The expanded set contains 234 comparisons across eight stable research groups. They are intentionally different from the website's more detailed browsing categories. Keeping the research groups stable makes the earlier and current editions comparable; it does not mean that creator-only arguments, personal theism, and religion's social effects are interchangeable.

Support-score differences remain in both constructive cases and replies. The practical lesson is to inspect the transition a speaker needs to defend: from usefulness to truth, from cause to agency, from testimony to a miracle, or from possible divine reasons to an explanation of observed suffering. The topic average is a starting point for that reading, not a verdict on the subject.

Reader's guide

Executive summary	2
1. Preserve a useful comparison without freezing the website	4
2. The current topic map	4
3. Look at the burden, not just the category name	6
4. Is the problem confined to opening a case?	7
5. What should change in a reader's approach?	9
Methods and sources	9

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. Score results refer to the October 9 snapshot. Direct slogan counts retain their September source coverage; they have not been extended to new transcripts.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. Preserve a useful comparison without freezing the website

The site now offers more topic distinctions than the September research did. Replacing the research categories silently would mix two changes: added evidence and a new classification scheme. This edition preserves eight exploratory groups, assigns each included assessment once, and publishes the decisions. The reader can therefore compare editions without mistaking a relabeling exercise for a change in performance.

“Primary topic” is a summary of the main burden, not a denial that other subjects arise. A broad God debate can discuss morality, cosmology, and evil. It belongs to General theism and naturalism here if no single specialized burden adequately describes the assessment. That group has grown to 62 comparisons; the other groups range from 21 to 27. Unequal group sizes affect how stable their averages are.

The research classification follows the position defended. It does not put every statement by a religious person into the religious side. Cases without an explicit enough religious-versus-skeptical contrast are excluded, while broad historical defenses receive narrower-set exclusions where appropriate. The classification table records both choices.

2. The current topic map

Research topic	Cases	Mean gap	95% resampling range
Religion, culture & meaning	25	7.28	5.08-9.44
Evil, suffering & hiddenness	24	6.83	5.21-8.38
Morality & moral foundations	24	6.79	4.54-9.08
Mind, reason & logic	21	6.24	2.81-9.81
Scripture, revelation & doctrine	27	6.00	3.52-8.52
Cosmology, science & design	27	5.89	2.89-8.63
General theism & naturalism	62	4.98	3.55-6.40
Resurrection & historical evidence	24	3.04	1.50-4.58

Topic gaps: ordering remains uncertain




Figure 1. All eight observed mean gaps favor skeptical-side scores, but the displayed ranges overlap substantially.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population. Counts are in parentheses. These stable research groups are not the current browsing categories.

Each difference is skeptical minus religious-side score, with one vote per debate. The uncertainty lines show the middle 95% of repeated-draw averages within the same topic, not a guarantee about a universal population. They omit systematic judging error, and repeated speakers limit the independence of the observations.

The culture group remains highest in the observed table, yet it is highest in only about 32% of the resampling calculations. Other groups can take the lead when their recorded cases are weighted differently. This percentage is not a probability that culture is the hardest subject in the world. It describes an exercise performed on this particular archive.




Figure 2. A topic can lead the observed table without consistently leading when its recorded cases are drawn again.

How to read this graph. Each bar counts calculations in which the topic has the highest gap. This is a stability exercise, not the probability that it is truly the hardest topic in the wider world.

The resurrection group remains comparatively close. That does not establish that resurrection is true, that its evidence is uniformly strong, or that every resurrection defense outperforms other religious arguments. It says that the assessed sides are closer on average in these selected comparisons. The narrower historical sincerity case is included in the broad group but separately flagged because sincerity alone is not resurrection.

3. Look at the burden, not just the category name

Religion's value and religion's truth

Community, consolation, and a shared moral vocabulary can be relevant to a motion about religion's effects. They do not automatically establish a supernatural explanation. Conversely, showing a disputed supernatural premise does not alone establish that a social institution has no benefits. A useful critique identifies which conclusion is actually under discussion and prevents either side from trading it for an easier one.

The [Sacks-Dawkins discussion](#) makes this distinction concrete. Compatibility between science and religion, the interpretation of scripture, and the independent support for cosmic purpose are different questions. A

coherent division of roles can answer a conflict claim without proving the larger religious picture. Readers should track which burden each reply addresses.

A creator is not yet a complete theology

Origins arguments can attempt to establish an initial cause, an intelligent source, or a necessary foundation. A further step is needed to identify that source with a personal, morally good, interventionist God. The [Loke-Linford assessment](#) examines causes, physical beginnings, timeless grounding, and agency. A criticism of one step need not settle the rest.

The [Meyer-Atkins assessment](#) asks whether design is scientifically eligible as an explanation of biological origins. It does not establish every religious attribute of a designer. The research group includes it as a broad design-versus-naturalistic comparison and offers an exclusion check in paper one. Keeping this distinction visible respects the difference between creator arguments and fuller theism.

A possible explanation is not yet the better explanation

For suffering and hiddenness, it is important to distinguish showing logical consistency from explaining how likely the observed world is under competing accounts. A possible divine reason can defeat a claim of strict contradiction without defeating an evidential argument. A skeptical account must likewise do more than restate an intuition about what a deity would do.

The [Rauser-Malpass discussion](#) concerns comparative evidence, background beliefs, and suffering. The central reading question is what makes the observations more expected under one account than another. A score gap cannot replace that comparison.

Sincerity, reliability, and a miracle

Witnesses can be sincere yet mistaken, and a report can preserve a sincere belief without establishing its supernatural explanation. Conversely, identifying uncertainty in one source does not establish a complete rival history. The [Horn-Paulogia assessment](#) distinguishes evidence for named leaders from confidence about all the original witnesses. Its scope problem is not interchangeable with the truth of resurrection.

4. Is the problem confined to opening a case?

No such simple pattern appears. The average support-score gap is 7.56 points across all eligible moves, 7.84 for high-importance moves, 7.55 for constructive moves, and 7.33 for replies.




Figure 3. Both constructive arguments and replies retain support-score differences. Each row has its own eligible debate count.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population. Each side's qualifying moves are averaged before comparing sides. Rows overlap; these are not official weighted overall scores.

These comparisons average the qualifying moves within each side before subtracting sides. They do not reproduce the first paper's importance- and section-weighted official scores. Both sides must have the relevant move kind: the counts are 234, 233, 159, and 233 respectively. The rows overlap and must not be added.

A reply has its own burden. Pointing out that an opponent has not proved everything may be relevant, but a speaker who advances a rival positive claim must still support it. A constructive case must also address credible alternatives. The data therefore favor examining the inferential work in both roles, rather than attributing the whole gap to having to speak first.

5. What should change in a reader's approach?

Reason 1: the topic differences remain descriptively useful but are not sharply ordered. **Reason 2:** new cases can alter the average without establishing historical improvement or decline. **Reason 3:** the recurring interpretive issue concerns the strength of the step between premise and conclusion, not merely the subject label.

Therefore: use the topic map as a guide to questions. In cultural debates, separate value from truth. In origins debates, distinguish cause, agency, and theology. In historical debates, distinguish sincerity, reliability, and the proposed event. In evil debates, distinguish possibility from comparative expectation.

Therefore: do not use this table to choose a topic winner in advance. A fair extension would review classifications without seeing the score gaps, publish reasonable alternatives, and compare matched opponents and burdens. The current categories are exploratory, not a scheme specified before any evidence was examined.

Methods and sources

This is a fixed October 9, 2026 research edition, based on source revision e65bc3c0b6. The score studies reconstruct published records rather than commission fresh assessments. The direct slogan counts retain the September 5 review of 187 transcripts; this edition recalculates those saved records and adds a separately labeled current score check. No new slogan reviews or debate assessments were performed.

The source manifest records file fingerprints. The analysis uses 20,000 deterministic resampling draws with baseline seed 20260904, retaining it for edition comparability. Whole debates are drawn for paired debate estimates; shared speakers for cross-role estimates. Ranking models and their dependence checks use the units explained in paper seven. These calculations do not capture every selection effect, repeated-speaker dependency, caption error or systematic judging bias.

The [data and methods page](#) links the source records, classification decisions, calculations, chart inputs and reading guides. Download the [current research package](#) and [historical direct slogan study](#). Earlier editions remain preserved in the repository's history. Debate links identify published assessments, not fresh independent evidence for this research.