

Where Is the Theist Disadvantage Largest?

Topics, burdens, and the argumentative transitions where a public case most often falls short.

8.27

largest observed topic mean:
religion, culture, and meaning

3.53

mean gap in resurrection debates:
the smallest of eight topic groups

187

debates mapped into eight
explicitly defined topic groups

The largest observed gap concerns religion, culture, and meaning. The more durable lesson is not a league table of topics, but the need to justify the step from an attractive possibility to a preferred explanation.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

The theist disadvantage is not identical across subjects. In the September 4 archive, **religion, culture, and meaning** has the largest observed mean non-theist advantage: **8.27 points across 22 debates**. The non-theist side scores higher in 21 of those comparisons; the remaining one is tied. Scripture and doctrine follow at 7.55 points, then evil and hiddenness at 7.15. Resurrection debates are closer, averaging 3.53 points across 17 comparisons.

The location of the largest observed mean should not be confused with a proven ranking of subject areas. The groups are small, their uncertainty intervals overlap, and several debates could reasonably be assigned to more than one topic. When the observed debates are repeatedly reweighted within their topic groups, religion, culture, and meaning ranks highest in only **46.2%** of the resamples. The leading category is worth investigating; its first-place position is not secure enough to be treated as a law.

The more durable finding is about the **work an argument has to do**. A speaker may establish that religion provides meaning without establishing that its religious explanation is true; that a first cause is possible without identifying it as a particular God; or that suffering is compatible with divine goodness without showing that the observed suffering is likely under that hypothesis. The assessment gap often opens at these transitions.

Move-level checks reinforce that interpretation. In debates containing high-importance moves on both sides, the average evidence-and-warrant difference is **8.83 dimension points**. Among replies, the responsiveness difference is **6.76 points**. These are paired within-debate comparisons, not counts that allow a few unusually long assessments to dominate. The findings point toward incomplete justification of central claims and incomplete engagement with objections.

This paper maps **187 religious-versus-skeptical debates and 4,086 assessed moves**, selected from the full 253-assessment archive. It does not rank the truth of religions or philosophical positions. Its central recommendation is practical: **use the map to locate vulnerable argumentative transitions, not to decide in advance which side must lose a debate about a given topic.**

Reader's guide

Executive summary	2
1. A map of performance, not a map of truth	4
2. Eight topics, with their boundaries visible	4
3. The largest observed gap concerns meaning and culture	5
4. A high mean and a consistent direction are different facts	7
5. Why the topic order is not secure	8
6. Different topics, recurring missing steps	9
7. Why resurrection debates are closer	11
8. Burden and argumentative ambition	11
9. Central claims and replies deserve special attention	12
10. What survives the main sensitivity checks?	13
11. Conclusion: the most useful map shows transitions	13
Methods and sources	14

Reading the evidence

Start with the summary and figures for the main argument. The worked examples explain the statistical ideas; the final methods and sources show how the numbers can be checked. Counts and results refer to the frozen September 4 archive, not to later additions to the site.

1. A map of performance, not a map of truth

A single mean compresses different kinds of discussion. A debate about Gospel sources may turn on the independence of testimony. A debate about morality may turn on whether an explanation of moral knowledge also explains moral obligation. A debate about suffering may turn on the contrast between logical possibility and comparative expectation. Calling all these “arguments for God” can conceal the actual work left undone.

The aim here is to recover that detail. Paper one establishes the broad 6.34-point mean difference and decomposes it into scoring dimensions. This paper asks where that difference is concentrated and what a reader should look for when moving from a whole-corpus finding back to an individual exchange.

The unit is the assessed debate, and the outcome is the non-theist-side score minus the theist-side score. A positive number favors the skeptical performance; a negative number favors the religious performance. Every debate has equal influence on its topic's mean. The source scores are the existing 5.6 Sol judgments, not new Astra judgments. The new work is the expanded analysis, interpretation, presentation, and set of checks.

The archive is selected rather than randomly sampled from all debates. The topic means therefore describe the material collected and assessed by SLUGFESTER. A subject can look difficult because its claims are hard to defend, because its sampled defenders are weak, because its sampled critics are strong, because of how the discussions are framed, or because of the scoring process. The map cannot separate those explanations on its own.

2. Eight topics, with their boundaries visible

Each of the 187 eligible debates is assigned one primary topic. This makes the counts add up and prevents a multi-topic debate from quietly appearing several times in the headline calculation. It also imposes a simplification: a general God debate may contain sustained discussion of suffering, morality, and cosmology even when it appears only under general theism.

Primary topic	Debates	Mean gap	Non-theist higher
Religion, culture, and meaning	22	8.27	21
Scripture, revelation, and doctrine	22	7.55	20
Evil, suffering, and hiddenness	20	7.15	19
Morality and moral foundations	21	6.67	18
Mind, reason, and logic	20	6.50	15
Cosmology, science, and design	24	5.88	21
General theism and naturalism	41	5.46	33
Resurrection and historical evidence	17	3.53	13

“Mean gap” means score points on the 0-100 scale, not a percentage of false claims. The final column counts only strictly higher non-theist scores; the remaining debates include both ties and theist advantages. The topic classification is published debate by debate in the companion data.

The main population is deliberately broader than explicit God-existence motions. It includes religious defenses of meaning, scripture, and social arrangements facing skeptical challenges. It does not automatically include every dispute involving a religious person. Historical Jesus versus mythicism, secular moral realism versus anti-realism, and soul existence alone are not treated as theism-versus-atheism comparisons. Comparisons between two religious readings are also excluded when they lack the intended contrast.

This boundary is not beyond criticism. For that reason, paper one also reports a narrower God-related and supernatural-truth subset. The present paper retains the broad category because religious meaning and public justification are themselves central to the question under study. Readers should compare the category labels to their substantive interests, not assume the title describes one perfectly homogeneous kind of dispute.

3. The largest observed gap concerns meaning and culture

Religion, culture, and meaning has the highest observed mean, **8.27 points**, with a resampling interval of **6.32 to 10.27**. The original 20-debate group averaged 8.65. Adding the Oldfield-Marriott and Peterson-Harris comparisons reduces the mean slightly rather than increasing it. The category remains prominent, but the update is not a story of an ever-widening disadvantage.

Mean score difference by primary topic




Figure 1. The dots show topic means and the lines show 95% debate-resampling intervals. Parentheses give the number of debates. All eight observed means favor the non-theist side, but their relative ordering is much less certain than their shared direction.

Scope: 187 debates assigned once across eight themes.

One recurring challenge in this category is the move from *usefulness* to *truth*. A religious narrative may help people organize sacrifice, cooperation, personal identity, or purpose. Those are potentially important achievements. Yet their existence does not independently establish the narrative's supernatural claims or show that no secular alternative can deliver comparable goods.

In [Peterson-Harris, debate 239](#), the assessment credits Peterson for identifying how finite agents need structures that make some facts practically relevant. Harris is credited for separating practical value, literal truth, and religious necessity. The score difference - 81 to 75 - does not imply that narrative is worthless. It indicates that a case for narrative's value did not fully establish the stronger claims about religious truth or necessity.

That distinction suggests a constructive way to improve the argument. A defender could first state the modest claim: certain stories help people coordinate action. Then the defender could identify the additional evidence needed for a stronger claim: those stories are indispensable, their benefits cannot be reproduced without religious belief, or their effectiveness is best explained by a supernatural source. The critic, in turn, must supply more than the assertion that secular substitutes are available. Their actual effectiveness and limits require evidence too.

The difficulty is not caused by discussing meaning instead of laboratory science. It arises when one kind of success is asked to establish a different kind without the connecting argument. A religious defense can avoid that failure by matching its conclusion to the evidence it actually presents.

4. A high mean and a consistent direction are different facts

The cosmology-and-design group illustrates why counting outcomes adds information. Its mean gap of 5.88 is smaller than several other topic means, but the non-theist side scores higher in **21 of 24 debates**. The mind-and-logic group has a somewhat larger mean of 6.50 but only **15 non-theist advantages in 20 debates**, with one tie and four theist advantages.

One category can contain many modest differences in the same direction. Another can contain a mixture of close contests and a few large gaps. An average alone does not reveal which pattern generated it. A reader deciding what to watch next should care about both: the mean tells us the typical balance of points, while the direction counts help show how consistently that balance appears.

Direction counts, not just averages




Figure 2. Each bar shows all assessed debates in a topic. Counts are printed inside the segments. A category's mean can be raised or lowered by the size of a few outcomes even when its direction count remains similar.

Scope: 187 classified debates.

Consider the added [Anderson-Jump debate, number 250](#). Anderson's theistic position scores 79 to Jump's 74. A reader who uses “mind and logic favors skeptics on average” as a substitute for examining the debate will miss the result. Conversely, one theist advantage cannot establish that the entire topic is favorable to theism. The case and the aggregate answer different questions.

The same restraint applies to unusually large losses. A striking presuppositional debate can illustrate a weakness very clearly, but it should not become the mental prototype for every philosophical discussion of reason or God. The more variable a topic's results, the less sensible it is to summarize its participants with a single rhetorical judgment.

5. Why the topic order is not secure

The highest observed mean is a factual description of this archive. Whether the same topic would remain highest after collecting different debates is a separate question. The uncertainty intervals overlap substantially. That does not by itself prove that no topic differences exist; it does show why the gaps between neighboring means should not be treated as settled.

To make the issue visible, we repeatedly draw replacement samples within each topic and recalculate the eight means. Religion, culture, and meaning is highest in 46.2% of these resamples; scripture and doctrine in 22.2%; mind and logic in 13.1%; and evil and hiddenness in 8.5%. The remaining categories divide the smaller share.




Figure 3. This exercise asks which category is highest when the observed cases are reweighted within each topic. These are resampling frequencies, not probabilities that a topic is truly the most difficult in the wider world. They do not account for changing the taxonomy or collecting different kinds of speakers.

Scope: 20,000 independent within-topic resamples; exploratory, not population probabilities.

The lesson is not that the leading category is uninteresting. It is that **a useful lead can be more secure than an exact rank**. Several high-gap topics share a demand for careful comparison between a worldview and alternatives. A cluster of related vulnerabilities is more defensible than a declaration that one subject is conclusively the weakest part of all theistic argument.

A small lesson in selection: if you compare eight averages and then focus on the largest, you have selected an unusually high result by design. Even when all groups have the same underlying mean, one sample mean must come first. A follow-up should examine the leading category in new data, not treat its first-place finish in the discovery data as a completed test.

There are also classification choices. Should a debate emphasizing consciousness and evil belong under mind, suffering, or general theism? A genuinely strong future analysis would specify its rules before examining the outcomes, allow independent coders to disagree, and report whether the main picture survives plausible alternative assignments. The current taxonomy is transparent and useful, but it is not the only reasonable map.

6. Different topics, recurring missing steps

The score decomposition offers a way past the ranking question. For each topic, we calculate how much of the overall gap comes from each of the six dimensions after applying the official weights. The resulting map asks whether a topic's disadvantage is mainly a failure of evidence, coherence, replies, or something else.




Figure 4. Cells show contributions in overall score points, not raw dimension scores. The official weights are already applied. A darker cell therefore means a larger contribution to that topic's mean gap. Small rounding differences are not shown.

Scope: 187 debates; official weights; rounding residual omitted.

Scripture: what kind of reliability has been established?

Textual preservation, authorship, historical accuracy, theological consistency, and divine authority are distinct claims. Evidence that a passage is early or substantially recoverable does not by itself establish that every event it describes occurred. Evidence that an event occurred does not by itself establish a particular theology. A careful defense can succeed on one level while remaining incomplete on the next.

This helps explain why a scripture debate can contain detailed scholarship and still earn a lower overall score. The issue is not a shortage of technical vocabulary. It is whether the cited evidence supports the exact conclusion being used in the larger argument. The skeptical side also incurs a burden: uncertainty about some details does not automatically establish wholesale unreliability.

Evil and hiddenness: possibility versus expectation

A possible morally sufficient divine reason can answer the claim that a good God and suffering are logically incompatible. The question “Would we expect suffering of this amount and distribution?” demands more. A proposed reason must help explain the pattern, not merely leave the pattern possible.

The newer [Jones-O'Connor discussion, debate 252](#), sharpens this distinction. The assessment credits a developed account of divine permission while retaining objections about the actual forms and distribution of suffering. The 85-74 difference should lead a reader to inspect the comparison between permission, purpose, and alternatives - not to assume that merely raising suffering defeats every conception of God.

Morality: explain the dependence, not just the importance

The seriousness of moral obligation does not itself show that obligation depends on God. Nor does showing that nonbelievers can behave well settle the question of what makes moral claims true. Good debate keeps conduct, motivation, knowledge, and ultimate grounding separate long enough to identify which issue is actually contested.

In [Folley-Scrivener, debate 247](#), the God-independent position occupies PRO and leads 85 to 76. That reversal of formal orientation is useful: the religious disadvantage is not simply the advantage of saying “no.” The assessment concerns whether the proposed Christian foundation earns the dependence claim it carries.

Cosmology and mind: identify the proposed explanation

The first-cause question, the source of order, and the existence of consciousness can expose real explanatory difficulties. But a difficulty is not yet an identification. A necessary foundation need not have been shown to be a person; a mind-first picture need not have been shown to imply the Christian God; and a mathematical description of complexity need not define a valid probability comparison.

A better exchange asks which features the proposed explanation predicts, what alternatives remain, and what additional assumptions are being introduced. That task can be done in philosophical reasoning as well as empirical science. It does not require reducing every metaphysical question to a laboratory test.

7. Why resurrection debates are closer

Resurrection has the smallest observed mean gap, **3.53 points**, with a median of two. The skeptical side leads in 13 of 17 debates, the religious side in three, and one is tied. It would be a mistake to translate this into “resurrection is the most probable religious claim.” It is a result about how the arguments were presented and assessed.

One plausible interpretation is that these debates often force both sides onto a relatively explicit evidential field: sources, dates, testimony, dependence between documents, alternative explanations, and background assumptions about miracles. Those constraints make it harder for either side to rely entirely on a general worldview slogan. This interpretation fits the format, but the present analysis has not experimentally isolated its effect.

The [Licona-Ehrman debate](#), number 237, ends at 78 to 79. A one-point difference should encourage a reader to inspect where each historian's burden begins and ends. Licona must show why resurrection outperforms alternatives; Ehrman must clarify what historical methods can and cannot establish without replacing that methodological issue with an unexplained impossibility claim.

The two later-process resurrection debates average only a half-point gap. That is an interesting observation, but **two debates are not a reliable basis for a new topic-level trend**. The sample is too small and too selected. The correct response is to preserve the observation and seek more comparable cases, not promote it into a reversal of the broader pattern.

8. Burden and argumentative ambition

Across all relevant topics, the mean skeptical advantage is 6.77 points when the theist occupies PRO and 3.26 when the theist occupies CON. That contrast is consistent with affirmative defenses often carrying more work: establishing a claim can require several supported premises, while undermining one essential step may be enough to block it.

Yet PRO and CON are imperfect indicators of that burden. A skeptical PRO side may defend a strong claim that religion is harmful or that theism is unreasonable. A theistic CON side may still need to offer a positive alternative. The actual motion, the concessions made, and the local exchange determine what must be shown.

Holding the orientation constant does not erase the topic pattern. Among theist-PRO debates, religion, culture, and meaning averages 8.72 points; scripture and doctrine 8.47; evil and hiddenness 8.33; resurrection 3.53. This is useful evidence against the simplest claim that the whole topic pattern is merely a mix of formal sides. It is not a complete adjustment for the ambition or difficulty of the propositions.

The key distinction is between a **modest claim competently defended** and an **ambitious claim only partly defended**. A speaker can improve an argument by narrowing it honestly. Establishing a limited conclusion is not intellectual surrender. It can be better reasoning than announcing a sweeping conclusion supported by the same evidence.

9. Central claims and replies deserve special attention

A whole-debate score can conceal whether the weakness lies in peripheral comments or load-bearing claims. The recorded importance weights allow a narrower check. Among high-importance moves - those with the top recorded importance level - the paired evidence-and-warrant gap is **8.83 dimension points across 186 debates**.




Figure 5. Each row first compares the two sides within eligible debates. The evidence gap persists among central moves and replies. The constructive row has fewer debates because both sides do not always have recorded constructive moves; these rows are not identical populations.

Scope: Only debates containing eligible moves on both sides; equal debate weights.

This analysis differs slightly from the officially weighted decomposition. Within each eligible move type, moves are averaged within each side, and then the side difference is averaged across debates. It answers a question about the typical paired subset, not a reconstruction of the published total. The small difference between the 8.41-point all-move gap here and the 8.48-point officially weighted gap in paper one is therefore expected.

Reply responsiveness shows a separate **6.76-point** skeptical advantage across 186 eligible debates. Among constructive moves, the corresponding responsiveness gap is 2.64 points across 116 debates. It is tempting to call that a precise “reply penalty,” but that would overstate the comparison. The debates and the purposes of the moves differ. The defensible point is that the religious-side weakness remains when attention is restricted to exchanges in which an objection must be answered.

For a viewer, this suggests a productive reading habit. Mark the exact challenge before listening to the response. Then ask which part of it was answered. A forceful return to the speaker's central theme may be relevant without being responsive. The test is whether the answer protects the threatened inference, not whether it sounds recognizably like the speaker's worldview.

10. What survives the main sensitivity checks?

Removing Dillahunty, O'Connor, Hitchens, and Oppy leaves religion, culture, and meaning at 7.91 points across 11 debates. Morality is 6.83 and evil/hiddenness 6.50. The leading cultural category therefore does not disappear when several prolific skeptical speakers are removed. Its sample, however, is then only half its original size, and the remaining differences should be read with wider caution.

The earlier-process cultural group averages 7.80 points across 15 debates; the later-process group averages 9.29 across seven. The direction is consistent, but the difference between those two means is not evidence of a temporal worsening of religious argument. The people, propositions, recordings, and assessment process differ. Debate number is an archive sequence, not a measure of when a speaker became better or worse.

These checks preserve the broad interpretation while limiting its precision. They support sustained attention to meaning, scripture, and suffering; they do not establish that every apparent difference between those topics is real, independent, or caused by the topic itself. The next review should keep that distinction central.

11. Conclusion: the most useful map shows transitions

The largest observed mean disadvantage appears in religion, culture, and meaning. Resurrection debates are closer. Several other topic groups sit between them with overlapping uncertainty. That is a useful map of this archive - but the map becomes more illuminating when we ask what speakers are trying to move from and to.

Again and again, the difficult step is a transition: from benefit to truth, from possibility to probability, from an unexplained fact to a preferred explanation, from a necessary foundation to a personal God, or from some permissible suffering to the suffering actually observed. The numerical differences are not proof that these transitions always fail. They show that, in the recorded assessments, religious defenders often leave more of that work unfinished.

The strongest lesson for the reader is therefore not “expect the theist to lose on topic X.” It is **“locate the bridge on which the conclusion depends.”** Ask what would justify crossing it, whether the speaker supplied that justification, and whether the critic answered the actual claim rather than a stronger or easier substitute.

For a defender, that discipline can turn a broad affirmation into a credible argument. For a critic, it prevents a statistical tendency from becoming an excuse for careless dismissal. For SLUGFESTER, it makes the corpus useful as a guide to better reasoning rather than merely a scoreboard of opposing identities.

Methods and sources

This paper uses the September 4, 2026 snapshot at source revision **76d006b37**. The full archive contains 253 assessments; 237 have comparable locked one-on-one scoring records; 187 meet the broad religious-versus-skeptical classification. The 169 earlier topic assignments are preserved, and 18 new eligible debates are assigned explicitly. Each debate enters one primary topic. The complete table includes every inclusion, exclusion, side orientation, and topic decision.

All mean gaps use one observation per debate. Weighted contribution figures follow the public rubric's move-importance and section weights. Move-subset figures use equal-debate averages of paired within-debate move means. Topic intervals use 20,000 replacement resamples. The highest-topic exercise uses independent resamples within each of the eight groups. It is exploratory, is not a multiple-comparison-adjusted significance test, and does not model uncertainty in the topic assignments themselves.

The [analysis directory](#) supplies calculations, results, classification data, source-file hashes, and chart source contracts. The [Backend research library](#) links all seven companion papers. The discussion of argument quality is grounded in the linked scorecards; it does not independently verify the historical, scientific, or metaphysical claims made in the source debates.

The papers are a stable Astra-era analysis of existing judgments, not a new model's reassessment of every recording. A future major-model review should retest the taxonomy, scoring comparability, and source interpretations rather than assuming that larger samples or newer models automatically remove bias.