

Why Do the Theist Sides Score Lower?

A strong test of the substantiation hypothesis - and a careful account of what the scores cannot establish.

6.34

mean non-theist advantage
in score points out of 100

187

relevant debate comparisons
from the 253-assessment archive

4,086

assessed argumentative moves
in the position comparison

The gap is persistent. Its largest raw dimension difference is in evidence and warrant, while logic, warrant, and replies together account for about three-quarters of the overall score gap.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

The theist sides still score lower, and the enlarged archive makes the location of the weakness clearer. Among **187 relevant comparisons drawn from 253 published assessments**, the mean non-theist advantage is **6.34 points on a 0-100 scale**. The non-theist side scores higher in 160 debates, the theist side in 20, and seven are tied. The median difference is six points. This is a broad pattern in the assessed archive, not an average manufactured by a few extraordinary defeats.

The largest difference before applying the scoring weights is in **evidence and warrant**: the reasons offered for accepting a claim and the connection between those reasons and the conclusion. The mean difference is 8.48 dimension points. Logical coherence differs by 7.11 points and responsiveness by 6.48. After applying the rubric's weights, these three dimensions account for about three-quarters of the overall gap. The recurring problem is not simply that theists state a religious conclusion. It is that a conclusion often travels farther than the publicly supplied reasons can carry it.

This finding supports the observable part of Phil Stilwell's **epistemic bleed-through hypothesis**. In everyday language, the hypothesis is that permissive standards for sustaining a religious commitment may carry into public debate, where an unconvinced listener needs independent reasons, a fair comparison with alternatives, and an answer to the actual objection. The scores contain the predicted pattern: weaker substantiation accompanied by overextended inferences and incomplete replies. They do not directly measure anyone's private standards of belief, upbringing, emotional needs, or motives. The proposed cause remains an explanation to test, not an established psychological diagnosis.

The result survives important changes in selection. A narrower set of 146 God-related or supernatural-truth comparisons gives a 6.27-point gap. Removing four frequent skeptical speakers leaves a 5.21-point gap across 115 debates. The earlier and later assessment processes produce similar within-debate gaps, despite differing in absolute score levels. The 18 newly eligible comparisons average 6.28 points, almost matching the original 169-debate result.

The practical conclusion is demanding but even-handed: belief, coherence, comfort, possibility, and public justification are different achievements. A speaker earns a stronger assessment by showing the missing steps between them. The same demand applies to a secular claim that science explains everything, that religious people are irrational, or that an unspecified natural alternative is overwhelmingly probable.

Reader's guide

Executive summary	2
1. What this edition measures	4
2. Who counts as theist or non-theist here?	4
3. How large and how consistent is the difference?	5
4. Where the missing points actually come from	6
5. Five recurring failures of substantiation	8
6. Does the evidence problem depend on an arbitrary cutoff?	10
7. How the hypothesis fares against rival explanations	11
8. Counterexamples are part of the result	12
9. What would make the hypothesis more convincing - or weaker?	12
10. Conclusion: the burden is in the bridge	13
Methods and sources	13

Reading the evidence

Start with the summary and figures for the main argument. The worked examples explain the statistical ideas; the final methods and sources show how the numbers can be checked. Counts and results refer to the frozen September 4 archive, not to later additions to the site.

1. What this edition measures

This is the September 4, 2026 research edition of the SLUGFESTER corpus papers. It takes the 253-assessment archive as a fixed reference point. Of those assessments, 237 have comparable one-on-one scoring records, covering 5,282 scored argumentative moves. Sixteen other-format assessments remain in the archive but are not mixed into the detailed position comparison. The relevant religious-versus-skeptical subset contains 187 debates and 4,086 moves.

An *argumentative move* is a recorded contribution to a debate: a claim, supporting reason, objection, or reply. It is not a sentence count, a speaking-time measure, or a complete inventory of every thought expressed in the recording. The analysis therefore concerns the arguments selected and assessed in the site's records. The source videos and scorecards remain necessary for judging whether the selection and interpretation were fair.

The research was rewritten at the introduction of Astra, but the underlying 237 comparable assessments are labeled **5.6 Sol**. Astra has not freshly judged all 253 debates for this edition. Recalculating statistics is different from asking a new model to reassess the source material. Conflating those operations would exaggerate both the independence and the novelty of the evidence.

Slugfester intends to rerun debate assessments roughly twice a year, when better models, source access, and quality controls permit. These papers have a separate publication cycle: this edition is intended to remain a stable baseline until the next major GPT model. That review should test whether the new model is more accurate, consistent, and fair; a newer name alone establishes none of those properties. Genuine corrections should not wait for either cycle.

2. Who counts as theist or non-theist here?

The labels follow the **position defended in the assessment**, not a permanent classification of the person. A critic of a particular argument need not claim that every conception of God is false. A defender of biblical reliability need not be arguing the same proposition as a defender of cosmic design. Treating all of these as interchangeable would obscure what the comparison actually tests.

The main set preserves the earlier papers' broad category: a religious or theistic defense facing a skeptical or non-theistic challenge. It includes disputes about religious meaning and scripture as well as God's existence. The title uses the familiar shorthand "theist sides," but the population is more accurately described as **religious-versus-skeptical debate positions**. The broader label is especially important in cultural and historical debates.

Every addition from debate 227 through 253 received an explicit inclusion or exclusion decision. Historical Jesus versus mythicism is excluded: an ordinary historical Jesus would not establish a divine Jesus. Soul debates are excluded when neither position establishes a God contrast. Secular moral realism is not treated as theism. Comparisons between two religious interpretations or two purposive cosmologies are not automatically counted as theism versus non-theism.

The narrower sensitivity set removes religion-and-culture debates and selected historical, social, or internal-doctrine disputes. It is not a perfectly neutral boundary - no classification of this material could be -

but it answers a useful objection: is the headline being driven by debates about religion's social role rather than its truth claims? The remaining 146 comparisons still show a 6.27-point mean disadvantage. The full decision list is published with the analysis, so readers can inspect contested boundaries instead of accepting them on trust.

3. How large and how consistent is the difference?

The mean theist-side score is **77.30**; the mean non-theist-side score is **83.65**. Subtracting the first from the second gives 6.34 points after rounding. Each debate receives equal weight in this average. A long debate with many scored moves does not count as several debates merely because its assessment is more detailed.




Figure 1. Most differences lie on the non-theist side of zero, but the distribution also contains theist advantages. The dashed line marks the mean; it is not a threshold for winning an individual debate.

Scope: 187 broad religious-versus-skeptical debates.

The non-theist side scores higher in **85.6% of all 187 comparisons**, or **88.9% of the 180 non-tied comparisons**. Both statements are correct, but their denominators differ. Reporting only the larger percentage without explaining the removal of ties would make the result look slightly more decisive than the full count warrants.

The spread matters as much as the average. A typical difference is six points, but individual results range from clear theist advantages to large skeptical advantages. This means the average should guide expectations about this archive, not replace attention to the case in front of the reader. A population pattern cannot tell us that a particular Lennox, Craig, Oppy, or Dillahunty contribution is sound or unsound.

What does “6.34 points” mean? It is the average difference between two assessed performances on a 0-100 scoring scale. It is not a 6.34% probability that theism is false, a percentage of arguments that are wrong, or a percentile ranking of intelligence. A performance score and a proposition's truth are different quantities.

The 95% debate-resampling interval is approximately **5.46 to 7.22 points**. To obtain it, the calculation repeatedly makes new lists by drawing whole debates from the observed set, allowing some to appear more than once. It then recalculates the mean 20,000 times. The middle 95% of those means shows how sensitive the estimate is to the particular mix of observed debates.

This is not uncertainty about the arithmetic mean of the frozen archive: that mean is known. Nor does the interval include every source of uncertainty. It does not cover transcript errors, mistaken judgments, selective video collection, disputed classifications, or a model's systematic preference. Recurring speakers also make the debates less independent than the simplest resampling exercise assumes. The interval is a useful stability diagnostic, not a certificate of objectivity.

4. Where the missing points actually come from

SLUGFESTER combines six dimensions. Coherence receives 25% of a move's weight; evidence and warrant and responsiveness each receive 20%; relevance and burden receive 15%; precision and clarity and confidence/fairness each receive 10%. “Confidence and fairness” is the reader-friendly label here for the rubric's combined calibration-and-charity dimension: confidence proportionate to support, together with fair representation of the other side.

Move scores are combined using the recorded importance of the moves and the weight of each debate section. Reconstructing these operations matters. Simply averaging every dimension score in the archive would let debates with more recorded moves dominate the answer and would not reproduce the public totals.

How the 6.34-point gap adds up



Figure 2. The officially weighted dimensions reconstruct the mean score difference. Coherence, evidence, and responsiveness contribute 4.77 points together. A small negative rounding residual reconciles the dimension calculation with the published whole-number scores; the mean burden-adjustment difference is zero.

Scope: 187 debates; official section and move weights.

The strongest raw dimension gap is evidence and warrant, at **8.48 points**. Coherence contributes slightly more to the final score gap - 1.78 points versus 1.70 - because it has the larger weight. Those are not conflicting findings. “Largest underlying difference” and “largest contribution after weighting” answer different questions.

Responsiveness adds 1.30 overall points. Confidence and fairness add 0.80, precision and clarity 0.51, and relevance and burden 0.29. The small contribution from relevance suggests that the typical theist contribution is not simply off-topic. It often addresses an important issue while failing to justify the stronger conclusion drawn from it or to answer an important qualification.




Figure 3. Evidence and warrant show the largest unweighted dimension difference. These are dimension points, not contributions to the final score. Each interval resamples complete debates.
Scope: 187 paired, officially weighted debate comparisons.

This decomposition is **an accounting identity, not a causal experiment**. The dimensions help determine the final score, so it is unsurprising that weaker dimension ratings accompany lower totals. Furthermore, one defect can affect several dimensions. An unsupported leap can count as weak warrant, questionable coherence, and excessive confidence. Agreement among those ratings is therefore not three independent discoveries of the same fact.

The informative question is more specific: what kinds of shortcomings did the assessments record, and do the source passages make that diagnosis reasonable? The following patterns connect the numbers to arguments a reader can inspect.

5. Five recurring failures of substantiation

A possibility is promoted into a preferred explanation

A proposal can be logically possible without being likely, and likely without being the best-supported option. “A mind could explain this” establishes less than “a divine mind explains this better than the alternatives.” The stronger conclusion requires a comparison: what would we expect if the proposed explanation were true, and what would we expect otherwise?

In debate 244, Michael Jones argues toward a necessary divine mind, while Joseph Schmid challenges the identification of a necessary foundation with God. Even if some foundation is granted, the further description - mental, personal, good, unique, or religiously recognizable - still requires support. The weakness is an unfinished bridge between claims, not necessarily a failure to notice the original philosophical problem. See the [Jones-Schmid scorecard](#).

A gap in a rival account becomes evidence for one's own account

An unresolved problem for naturalism may justify less confidence in a particular natural explanation. It does not automatically establish intelligent agency. Both explanations can remain incomplete. The crucial issue is whether the proposed agency account explains the evidence more specifically, or merely places a familiar label over the unknown.

The earlier assessment of debate 3 records a move from difficulty explaining cellular origins toward agency. The corresponding burden is not satisfied by making the difficulty vivid. A stronger case would identify an observable feature more expected under a specified agency hypothesis than under relevant non-agency alternatives. The [Hitchens-D'Souza debate](#) provides a concrete setting in which to examine that distinction.

The benefits of a belief stand in for the truth of its explanation

A belief can comfort people, organize a community, or support admirable behavior without its account of reality being true. Conversely, a true claim need not be comforting. The inference from “this gives people purpose” to “its supernatural explanation is correct” needs a premise connecting those two achievements. Repeating the first achievement more eloquently does not supply the missing premise.

This is especially relevant to affirmations of God-given human worth. A theist may offer a coherent account of worth *within the worldview*. The public argumentative task is different: explain why a listener should accept that worldview, why creation confers the claimed moral status, and why alternatives fail. It is possible to respect the moral aspiration while finding that the presented inference has not been completed.

Compatibility is treated as sufficient evidence

Showing that suffering and God are not logically contradictory can answer an argument that alleges a contradiction. It does not, by itself, answer the claim that the amount or distribution of suffering makes a particular God less likely. A possible divine reason may defeat an impossibility claim while doing little against an evidential comparison.

The distinction protects good theistic reasoning too. In debate 144, Alvin Plantinga explicitly limits a move to the compatibility of theism and Darwinism. That contribution receives a score of 92, with evidence/warrant at 86. The assessment recognizes that a modest conclusion can be well defended. The problem is not arguing for possibility; it is silently upgrading possibility into probability or actuality. See [Plantinga-Dennett, 81:23](#).

The reply answers an easier claim

If a critic asks why *this distribution* of suffering is expected, replying that *some suffering* permits growth does not finish the exchange. If the question is whether testimony is reliable, replying that testimony is widespread leaves reliability open. If an argument concerns whether a premise follows, an attack on the critic's worldview may miss the local objection entirely.

These failures explain why responsiveness belongs beside warrant in the diagnosis. A case is not strengthened merely by restating its central theme after a challenge. It must identify what the challenge threatens and show why the supporting route remains intact. The reader should ask not only “Was an answer given?” but “What proposition did the answer actually address?”

6. Does the evidence problem depend on an arbitrary cutoff?

The mean evidence difference uses the whole scoring scale, so it does not depend on choosing a boundary between adequate and inadequate support. Thresholds nevertheless make the pattern easier to picture. For each debate, we calculate the share of each side's moves below a chosen evidence score, then average those shares across debates.




Figure 4. Changing the evidence threshold changes the size of the percentage gap but not its direction. The percentages describe an average debate's move shares, not the percentage of all theist or non-theist people who reason poorly.

Scope: 187 paired debates; each debate has equal weight.

Below 70, the average share is **60.5% for the theist side and 23.1% for the non-theist side**, a 37.4-percentage-point difference. Below 60, the corresponding shares are 24.1% and 5.5%. Below 80, they are 92.3% and 68.5%. The substantial skeptical share at the broader boundary is worth noticing: the archive does not depict the non-theist side as a storehouse of fully demonstrated claims.

A cutoff is an analytical convention, not a discovery that a score of 69 is fundamentally different from 70. The point of checking several cutoffs is to see whether the result depends on one convenient choice. Here it does not. The pattern remains, although its numerical description changes with the question being asked.

An evidence score also assesses what was supplied *in the recorded contribution*. It should not be translated into “no evidence exists anywhere.” A speaker may know an extensive literature and fail to present the relevant part. That is a performance limitation, not proof that the literature cannot support the claim.

7. How the hypothesis fares against rival explanations

The bleed-through hypothesis has explanatory force because it predicts a linked pattern: confidence sustained by commitments that are not independently established for the audience; resistance to inconvenient alternatives; and replies that preserve the worldview more successfully than they resolve the objection. The assessments contain many moves with precisely that shape.

However, several other processes could produce some of the same observations. A more ambitious conclusion requires more supporting steps. An affirmative speaker may accept a heavier burden. Experienced skeptics may be overrepresented. The selected recordings may favor particular formats or speakers. The model may consistently reward a style of cautious criticism. And a scoring process can reproduce a systematic preference even when its numerical calculations are flawless.




Figure 5. The score difference survives a narrower truth-claim comparison, removal of four frequent skeptical speakers, and separation by assessment process. These are robustness checks, not randomized tests of faith or debating role.

Scope: Each row states its own denominator.

Equalizing the influence of speakers provides another perspective. Giving each of the 84 religious-side speakers equal weight yields an average gap of 6.50 points; giving each of the 48 skeptical-side speakers equal weight yields 4.92. Neither is the uniquely correct estimate. They show that the headline depends somewhat on whose appearances are numerous, but the broad direction does not disappear when prolific speakers lose their extra influence.

The role check is similarly informative. When the theist occupies PRO, the mean skeptical advantage is 6.77 points. When the theist occupies CON, it is 3.26. Formal position matters to the observed contrast, but it does not account for its entire direction. Paper four examines this issue in detail and finds that same-speaker comparisons leave a smaller and less certain role difference.

Most importantly, none of these checks measures the hypothesized transfer of private belief standards into public reasoning. They reduce particular alternative explanations; they do not turn a plausible interpretation into a measured cause. A strong defense of the hypothesis should make that boundary visible rather than conceal it.

8. Counterexamples are part of the result

The theist side scores higher in 20 comparisons. Some advantages are substantial. Keith Fox leads Peter Atkins by 15 points in debate 85; William Lane Craig leads Atkins by ten in debate 166; John Lennox leads Atkins by seven in debate 51. In the new material, James Anderson leads Tom Jump by five in debate 250. These are not exceptions to be hidden after the headline. They show that the rubric can reward theistic defenses and penalize skeptical overstatement.

The lesson is not that those results cancel the mean gap. Twenty theist advantages do not erase 160 skeptical advantages. The lesson is that the proposed mechanism must explain variation as well as direction. Some religious arguments supply bounded claims, clear alternatives, and appropriate qualifications; some non-religious arguments make sweeping assertions that their evidence does not support.

A more useful claim than “theists are sloppy” is therefore: **certain ways of protecting a worldview make public substantiation harder, and those ways are disproportionately recorded on the religious side of this archive.** That formulation identifies conduct that can change. It does not turn a group label into an explanation of a person's abilities or motives.

9. What would make the hypothesis more convincing - or weaker?

A stronger test would separate the identification of the proposed mechanism from the assignment of scores. Independent readers could examine randomly selected passages without seeing the public score, speaker name, or debate winner. They would classify whether the speaker appeals to revelation, private experience, inherited authority, practical benefit, or an independently checkable comparison. Other readers would judge the quality of the support. Agreement and disagreement between those teams would be reported.

The strongest prediction is not simply that theist sides score lower again. It is that moves relying on the hypothesized permissive standards should receive weaker independent support ratings, even after matching topic, burden, and argumentative ambition. Within theist speakers, more disciplined treatment of those standards should predict better performance. Within non-theist speakers, similarly protected secular assumptions should incur comparable penalties.

The hypothesis would weaken if independent, blinded readers failed to find the proposed pattern; if the apparent difference largely vanished when the conclusion's ambition was matched; or if model changes strongly altered which side was judged to have supplied adequate support. It would also weaken if the source passages consistently supplied qualifications that the assessment summaries omitted.

These are genuine risks to the interpretation. A hypothesis that accommodates every possible outcome would imitate the very defect it criticizes. The next major-model review should publish the comparison rules before seeing the new results and retain both the original and revised judgments for inspection.

10. Conclusion: the burden is in the bridge

The enlarged archive supports a firm conclusion: **the lower theist-side scores are substantially associated with incomplete public justification, not merely with the presence of religious language or the absence of rhetorical polish.** The gap is broad, persists in the added material, and remains visible when the comparison is narrowed or frequent speakers are down-weighted. Evidence, inference, and direct engagement with objections are its central locations.

The bleed-through hypothesis is a serious explanation of that pattern. A standard that is sufficient to sustain personal faith may be insufficient to persuade someone who does not share the faith's starting points. The failure becomes visible when a comforting conclusion is treated as a demonstrated one, a possible explanation as a preferred one, or an unresolved objection as evidence of the critic's blindness. Those are identifiable argumentative transitions, not insults.

But the archive has not measured the private causes of those transitions. It establishes a pattern of assessed performance and supplies source-based reasons to take the proposed mechanism seriously. It does not establish that religious people are generally incapable of rigorous thought, that skeptical arguments are automatically sound, or that a debate score settles God's existence.

The best response to the result is not a more forceful affirmation from either camp. It is a completed bridge: from the evidence to the claim, from the claim to its proper strength, and from the objection to an answer that actually meets it. **A worldview becomes more credible in public when its defenders make clear what supports it, what limits it, and what could require them to think again.**

Methods and sources

The frozen source revision is **76d006b37**, containing 253 unique debate identifiers and unique video URLs. The analysis verifies 5,282 move scores against both their six dimensions and the public cards, 2,558 side-section scores, all 474 comparable overall side scores, and 58 later-process final-record file hashes. These checks establish arithmetic and record consistency, not the correctness of the judgments themselves.

The main average gives one vote to each debate. Dimension averages preserve the official within-debate importance and section weights. Move-type and threshold comparisons first summarize each side within a debate and then average paired differences. Intervals use 20,000 resamples and the fixed random seed 20260904. Repeated speakers, selective source inclusion, and systematic judging errors are not fully represented by those intervals. The exploratory topic and sensitivity choices were not preregistered.

The broader religious category, the narrower sensitivity flag, exclusions, normalized move data, source hashes, and executable calculations are available in the [September 4 analysis directory](#). The [Backend page](#) explains the scoring framework and links the companion papers. Paper two examines topics; paper three separates measurable slogan risk from literal non-falsifiability; paper four tests role effects; papers six and seven examine scale and ranking limitations.

The qualitative examples are interpretations of the named scorecards and locked source passages, not a new independent coding of every transcript. All numerical claims in this paper concern the recorded performances in this snapshot. They should not be generalized without qualification to all theists, all non-theists, or all public debates.