

Why Do the Theist Sides Score Lower?

A smaller gap in a larger archive - and the limits of explaining it.

5.74

average skeptical-side advantage
in score points / 100

234

classified comparisons
from 308 assessments

47

new eligible comparisons
since September

The score gap persists, but newer additions are closer.

PHIL STILWELL · SLUGFESTER

October 9, 2026 | Frozen archive: 308 assessments

Recalculated scores; direct slogan evidence retains its September scope.

Executive summary

The score difference persists, but it is smaller in the expanded archive. Across 234 classified religious-versus-skeptical comparisons, the non-theist side averages 5.74 points higher on the 0-100 scale. It leads in 193 assessments, the religious side leads in 33, and 8 are tied. The September edition's 6.34-point gap has fallen by 0.60 points.

The change is not a correction to earlier scores. All earlier overall scores and assessment-process classifications reproduce the September records. The 47 newly eligible comparisons average a smaller 3.34-point gap. They include 13 religious-side advantages and one tie. The archive has grown by 55 assessments, but only the additions meeting this study's substantive and scoring criteria enter this comparison.

Support for claims remains the largest raw scoring-area difference. Its gap is 7.60 points. Consistent reasoning, support, and replies together contribute 4.28 points, or 74.6% of the overall gap after the official weights are applied. This locates the recorded weakness; it does not independently prove that the judges interpreted the arguments fairly.

The practical finding is narrower than a verdict on religion. A possibility is not yet a probability; a useful belief is not thereby true; a difficulty for a rival is not automatically evidence for one's preferred explanation. These distinctions apply just as much to skeptical arguments. The faith-to-debate hypothesis remains an explanation to test, not a demonstrated cause or a diagnosis of speakers' private motives.

Reader’s guide

Executive summary	2
1. What changed, and what did not	4
2. Which positions belong in this comparison?	4
3. What the average does and does not say	5
4. Follow the missing points	6
5. Read the transitions, including the counterexamples	7
6. What follows for the hypothesis?	8
Methods and sources	8

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. Score results refer to the October 9 snapshot. Direct slogan counts retain their September source coverage; they have not been extended to new transcripts.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. What changed, and what did not

The October 9, 2026 edition covers 308 published assessments containing 7,264 assessed moves. Of these, 291 have comparable one-on-one locked scoring records, with 7,032 moves. A move is an assessed claim, reason, objection, or reply, not a sentence or a measure of speaking time. This refresh reconstructs those records; it does not ask a new model to judge the debates again.

There are 307 distinct video links. Assessments 13 and 125 share one source video but assess it under different formats. The full-archive count is therefore a count of assessments, not independent recordings. Only assessment 13 enters this paper's comparable set; no source URL repeats inside that set. Multiple appearances by the same speakers remain a separate source of dependence.

The change from 187 to 234 comparisons is analytically important. New additions are not random draws, and recently added videos are not necessarily recently recorded debates. A smaller average gap in the additions cannot be called an improvement over time in religious argumentation. It describes a different selected mixture of speakers, subjects, and assessed performances.

2. Which positions belong in this comparison?

Classify the argument, not the identity of the speaker. The broad set includes defenses of religion, scripture, purpose, supernatural claims, and creator or design arguments facing skeptical challenges. The title retains the familiar phrase “theist sides,” but this is not a pure sample of arguments for a personal, interventionist God.

The distinction between a creator and a fuller religious conception matters. Intelligence as an origin explanation does not by itself establish omniscience, goodness, revelation, or a particular religion. The new Meyer-Atkins and Loke-Linford cases are assigned to origins and design, with their narrower burdens recorded. Removing both leaves 232 comparisons and a 5.74-point gap, essentially the same headline.

By contrast, secular moral realism, free will, a historical human Jesus, and consciousness-first metaphysics do not by themselves establish the required religious-versus-skeptical contrast. The new Copan-Rauser conquest debate and Wood-Ally resurrection debate are excluded from this position study because they compare religious interpretations rather than a clean religious-versus-nonreligious contrast. Their assessments remain in the overall archive and other applicable studies.

The Horn-Paulogia sincerity debate enters the broad historical set, but not the narrower truth-claim check: sincerity is distinct from resurrection. The narrower selection also removes cultural-benefit and specified historical or doctrinal comparisons. Its 186 cases retain a 5.77-point gap. Every decision is available in the classification table; reasonable disagreement about a boundary is a reason for a sensitivity check, not for concealing it.

3. What the average does and does not say

Religious-side scores average 77.40; skeptical-side scores average 83.14. Subtract within each debate first, then average with one vote per debate. This prevents a long or densely annotated debate from outweighing a shorter one solely because it contains more assessed moves. The middle difference is 6 points.

The middle 95% of averages from 20,000 repeated draws of whole debates runs from 4.96 to 6.52 points. Resampling means drawing again from the existing records, with replacement. It checks sensitivity to the observed mixture. It cannot repair caption errors, mistaken scores, a selective catalogue, or systematic judging bias. Recurring speakers mean these are not fully independent observations.

A 5.74-point performance difference is not a probability that a worldview is false, a percentage of arguments that are invalid, or a measure of anyone's intelligence. It is a difference between recorded assessments under a particular rubric.

The difference is smaller in the additions




Figure 1. The original comparison and the additions are now materially different. Other checks retain a positive mean difference.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population. Parentheses count comparisons. Rows overlap and are not independent replications.

Removing Dillahunty, O'Connor, Hitchens, and Oppy from the skeptical side leaves 154 comparisons averaging 4.65 points. The result is not solely a by-product of those frequent participants. This is still a selected subset, not a controlled test of equal speaker ability or equal debate difficulty.

4. Follow the missing points

The six areas are Logic (25%), Support (20%), Replies (20%), Task (15%), Clarity (10%), and Care (10%). Care combines fitting confidence to the evidence and representing the other side fairly. These short names do not alter the official rubric. Move importance and section weights are preserved before comparing the sides.




Figure 2. The support difference is largest before weighting. The intervals concern each area's own points, not contributions to the total.

How to read this graph. Dots show averages; lines contain the middle 95% from 20,000 repeated draws of the stated records. Zero means no average difference. These ranges describe sensitivity to the saved observations, not all sources of judging error or a representative population. Right of zero favors the skeptical side. These are area scores before weighting, not additive contributions.

Area key: Logic = consistent reasoning; Support = evidence and reasons; Replies = answering the other side; Task = staying on the question and doing what the claim requires; Clarity = clear, exact wording; Care = fitting confidence to support and treating the other side fairly.

Logic contributes 1.58 overall points, Support 1.52, and Replies 1.18. Logic's higher weight makes its contribution slightly larger despite a smaller raw area gap. The remaining areas contribute 0.31, 0.44, and 0.73 points respectively. A rounding residual of -0.021 and a mean burden-adjustment difference of 0.00 reconcile the contributions with the published totals.




Figure 3. The weighted components explain the arithmetic of the overall gap, not six separate causes.

How to read this graph. Bar lengths show additive weighted contributions, not independent causes. The small rounding residual and any burden adjustment reconcile their sum with the overall gap; both are reported in the text.

Area key: Logic = consistent reasoning; Support = evidence and reasons; Replies = answering the other side; Task = staying on the question and doing what the claim requires; Clarity = clear, exact wording; Care = fitting confidence to support and treating the other side fairly.

An unsupported leap can lower support, coherence, and confidence ratings together. Their agreement is therefore not three independent confirmations of the same claim. Source-linked critiques are necessary to determine whether the scored deficiency is actually present.

5. Read the transitions, including the counterexamples

In [Jones-Schmid on contingency](#), the central issue includes the move from a necessary foundation to a divine mind. Granting a foundation does not automatically grant mentality, uniqueness, goodness, or religious identity. A stronger case supplies reasons for each extra property and compares alternatives at that same level of specificity.

The new [Meyer-Atkins discussion](#) is an important counterexample to a one-directional narrative. The published assessment gives the design side a seven-point advantage. Its critiques distinguish spontaneous order from functionally specified sequence, origins from subsequent evolution, and scientific eligibility from proof of a designer. That does not establish intelligent design as true. It shows why a skeptical response can score poorly when it fails to engage the particular inference being defended.

In [Rauser-Aron Ra on revelation](#), readers should likewise inspect the actual conditional claims rather than assume that a religious defense must be the weaker side. Examples are illustrations, not additional statistical

evidence. The linked classification and casebook files make the wider record inspectable, including the cases that run against the average.

6. What follows for the hypothesis?

Reason 1: the gap remains across the broad set, narrower truth claims, both assessment processes, and several exclusions. **Reason 2:** support and inferential work account for much of the recorded difference. **Reason 3:** new comparisons reduce the gap and add substantial counterexamples, requiring a less sweeping narrative than the earlier edition.

Therefore: the assessments support a bounded claim about publicly supplied reasons in this archive. They do not establish that personal faith causes the weakness. Topic selection, opponent difficulty, rubric design, judging preferences, and repeated speakers remain alternative explanations or contributing influences.

A fair next test would judge matched arguments independently, conceal identities where feasible, apply the same burdens to skeptical and religious conclusions, and inspect disagreements against full source context. Evidence that those checks remove the gap would weaken the proposed explanation. Merely adding more similarly selected records would not settle it.

Therefore: readers should use the pattern to ask better questions of an argument, not to pre-judge a speaker. What exactly is claimed? What makes it more likely than its rivals? Which objection was answered? Which would change the conclusion? Those questions remain useful on either side.

Methods and sources

This is a fixed October 9, 2026 research edition, based on source revision e65bc3c0b6. The score studies reconstruct published records rather than commission fresh assessments. The direct slogan counts retain the September 5 review of 187 transcripts; this edition recalculates those saved records and adds a separately labeled current score check. No new slogan reviews or debate assessments were performed.

The source manifest records file fingerprints. The analysis uses 20,000 deterministic resampling draws with baseline seed 20260904, retaining it for edition comparability. Whole debates are drawn for paired debate estimates; shared speakers for cross-role estimates. Ranking models and their dependence checks use the units explained in paper seven. These calculations do not capture every selection effect, repeated-speaker dependency, caption error or systematic judging bias.

The [data and methods page](#) links the source records, classification decisions, calculations, chart inputs and reading guides. Download the [current research package](#) and [historical direct slogan study](#). Earlier editions remain preserved in the repository's history. Debate links identify published assessments, not fresh independent evidence for this research.