ASTRA-ERA EDITION / PLAIN-LANGUAGE REVISION

01

Why Do the Theist Sides Score Lower?

Where the arguments lack support, and what that tells us about the proposed link between faith and public debate.

6.34
average non-theist advantage in score points out of 100

187
relevant debate comparisons from 253 assessments

4,086
assessed claims and replies in those comparisons

The difference persists across several checks. Its largest gap is in support for claims. Support, consistent reasoning, and answers to objections together make up about three-quarters of the overall score gap.

PHIL STILWELL · SLUGFESTER

September 4, 2026 | Frozen archive: 253 assessments

New analysis and writing; not an Astra re-judging of the source debates.

Executive summary

The theist sides still score lower, and the enlarged archive makes the location of the weakness clearer. Among **187 relevant comparisons drawn from 253 published assessments**, the average non-theist advantage is **6.34 points on a 0-100 scale**. The non-theist side scores higher in 160 debates, the theist side in 20, and seven are tied. The middle value difference is six points. This is a broad pattern in the assessed archive, not an average manufactured by a few extraordinary defeats.

The largest difference before applying the scoring weights is in **evidence and supporting reasons**: the reasons offered for accepting a claim and the connection between those reasons and the conclusion. The average difference is 8.48 scoring area points. Consistent reasoning differs by 7.11 points and answering the other side by 6.48. After applying the rubric's weights, these three scoring areas account for about three-quarters of the overall gap. The recurring problem is not simply that theists state a religious conclusion. It is that a conclusion often travels farther than the publicly supplied reasons can carry it.

This finding supports the observable part of Phil Stilwell's **faith-to-debate hypothesis**. In everyday language, the hypothesis is that permissive standards for sustaining a religious commitment may carry into public debate, where an unconvinced listener needs independent reasons, a fair comparison with alternatives, and an answer to the actual objection. The scores contain the predicted pattern: weaker support for the claim accompanied by overextended reasoning steps and incomplete replies. They do not directly measure anyone's private standards of belief, upbringing, emotional needs, or motives. The proposed cause remains an explanation to test, not an established psychological diagnosis.

The result survives important changes in selection. A narrower set of 146 God-related or supernatural-truth comparisons gives a 6.27-point gap. Removing four frequent skeptical speakers leaves a 5.21-point gap across 115 debates. The earlier and later assessment processes produce similar within-debate gaps, despite differing in absolute score levels. The 18 newly eligible comparisons average 6.28 points, almost matching the original 169-debate result.

The practical conclusion is demanding but even-handed: belief, consistent reasoning, comfort, possibility, and public support for a claim are different achievements. A speaker earns a stronger assessment by showing the missing steps between them. The same demand applies to a secular claim that science explains everything, that religious people are irrational, or that an unspecified natural alternative is overwhelmingly probable.

Reader’s guide

Reading the evidence

Start with the summary and graphs. Each graph has a reading key for its colors, marks, units, and limits. The conclusion gives numbered reasons leading to a stated finding. All results refer to the September 4 snapshot, not later additions to the site.

An **average**, also called a mean, adds the values and divides by their number. A **middle value**, or median, is halfway through the sorted list. A **range from repeated draws** shows how the result changes when we draw again from existing records; it does not cover every possible judging mistake. The examples in the paper explain these ideas where they matter.

1. What this edition measures

This is the September 4, 2026 research edition of the SLUGFESTER archive papers. It takes the 253-assessment archive as a fixed reference point. Of those assessments, 237 have comparable one-on-one scoring records, covering 5,282 scored reasoning moves. Sixteen other-format assessments remain in the archive but are not mixed into the detailed position comparison. The relevant religious-versus-skeptical subset contains 187 debates and 4,086 moves.

A *reasoning move* is a recorded contribution to a debate: a claim, supporting reason, objection, or reply. It is not a sentence count, a speaking-time measure, or a complete inventory of every thought expressed in the recording. The analysis therefore concerns the arguments selected and assessed in the site's records. The source videos and scorecards remain necessary for judging whether the selection and interpretation were fair.

The research was rewritten at the introduction of Astra, but the underlying 237 comparable assessments are labeled **5.6 Sol**. Astra has not freshly judged all 253 debates for this edition. Recalculating statistics is different from asking a new model to reassess the source material. Conflating those operations would exaggerate both the independence and the novelty of the evidence.

Slugfester intends to rerun debate assessments roughly twice a year, when better models, source access, and quality controls permit. These papers have a separate publication cycle: this edition is intended to remain a stable baseline until the next major GPT model. That review should test whether the new model is more accurate, consistent, and fair; a newer name alone establishes none of those properties. Genuine corrections should not wait for either cycle.

2. Who counts as theist or non-theist here?

The labels follow the **position defended in the assessment**, not a permanent classification of the person. A critic of a particular argument need not claim that every conception of God is false. A defender of biblical reliability need not be arguing the same claim as a defender of cosmic design. Treating all of these as interchangeable would obscure what the comparison actually tests.

The main set preserves the earlier papers' broad category: a religious or theistic defense facing a skeptical or non-theistic challenge. It includes disputes about religious meaning and scripture as well as God's existence. The title uses the familiar shorthand “theist sides,” but the population is more accurately described as **religious-versus-skeptical debate positions**. The broader label is especially important in cultural and historical debates.

Every addition from debate 227 through 253 received an explicit inclusion or exclusion decision. Historical Jesus versus mythicism is excluded: an ordinary historical Jesus would not establish a divine Jesus. Soul debates are excluded when neither position establishes a God contrast. Secular moral realism is not treated as theism. Comparisons between two religious interpretations or two purposive cosmologies are not automatically counted as theism versus non-theism.

The narrower comparison removes religion-and-culture debates and selected historical, social, or internal-doctrine disputes. It is not a perfectly neutral boundary - no classification of this material could be -

but it answers a useful objection: is the headline being driven by debates about religion's social role rather than its truth claims? The remaining 146 comparisons still show a 6.27-point average disadvantage. The full decision list is published with the analysis, so readers can inspect contested boundaries instead of accepting them on trust.

3. How large and how consistent is the difference?

The average theist-side score is **77.30**; the average non-theist-side score is **83.65**. Using the full, unrounded numbers gives a difference of **6.34 points** after rounding. Subtracting the two rounded averages printed here instead gives 6.35; that one-hundredth difference comes only from when rounding occurs. Each debate receives equal weight. A long debate with many scored moves does not count as several debates merely because its assessment is more detailed.

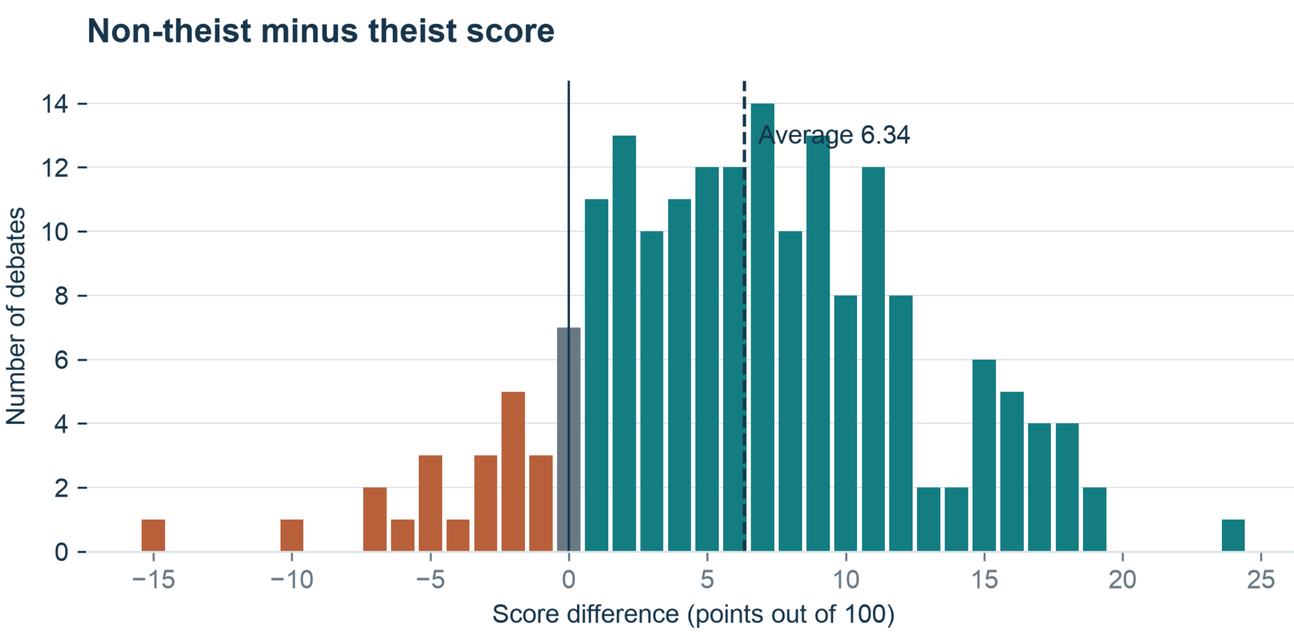

Figure 1. Most differences lie on the non-theist side of zero, but the distribution also contains theist advantages. The dashed line marks the mean; it is not a threshold for winning an individual debate.

How to read this graph. Each bar counts debates with that score difference. Teal, right of zero: non-theist higher. Rust, left of zero: theist higher. Gray at zero: ties. The dashed line is the average, not a winning threshold. One horizontal step is one score point; height is a debate count.

The non-theist side scores higher in **85.6% of all 187 comparisons**, or **88.9% of the 180 non-tied comparisons**. Both statements are correct, but their totals being counted differ. Reporting only the larger percentage without explaining the removal of ties would make the result look slightly more decisive than the full count warrants.

The spread matters as much as the average. A typical difference is six points, but individual results range from clear theist advantages to large skeptical advantages. This means the average should guide expectations about this archive, not replace attention to the case in front of the reader. A group pattern cannot tell us that a particular Lennox, Craig, Oppy, or Dillahunty contribution is sound or unsound.

> **What does "6.34 points" mean?** It is the average difference between two assessed performances on a 0-100 scoring scale. It is not a 6.34% probability that theism is false, a percentage of arguments that are wrong, or a percentile ranking of intelligence. A performance score and a claim's truth are different quantities.

The 95% range from repeated draws of debates is approximately **5.46 to 7.22 points**. To obtain it, the calculation repeatedly makes new lists by drawing whole debates from the observed set, allowing some to appear more than once. It then recalculates the average 20,000 times. The middle 95% of those means shows how sensitive the estimate is to the particular mix of observed debates.

This is not uncertainty about the calculated average of the frozen archive: that average is known. Nor does the interval include every source of uncertainty. It does not cover transcript errors, mistaken judgments, selective video collection, disputed classifications, or a model's systematic preference. Recurring speakers also make the debates less independent than the simplest resampling exercise assumes. The interval is a useful check of how much the result changes, not a certificate of objectivity.

4. Where the missing points actually come from

SLUGFESTER combines six scoring areas. The charts use short names, explained here. **Logic** asks whether the reasoning holds together (25% of a move's score). **Support** asks whether the evidence and reasons back up the claim (20%). **Replies** asks whether the speaker answers the other side (20%). **Task** asks whether the speaker stays on the question and does the work their claim requires (15%). **Clarity** asks whether the wording is clear and exact (10%). **Care** asks whether confidence fits the support and the opponent's view is represented fairly (10%). These names shorten the official labels; they do not change the scoring rules.

Move scores are combined using the recorded importance of the moves and the weight of each debate section. Reconstructing these operations matters. Simply averaging every scoring area score in the archive would let debates with more recorded moves dominate the answer and would not reproduce the public totals.

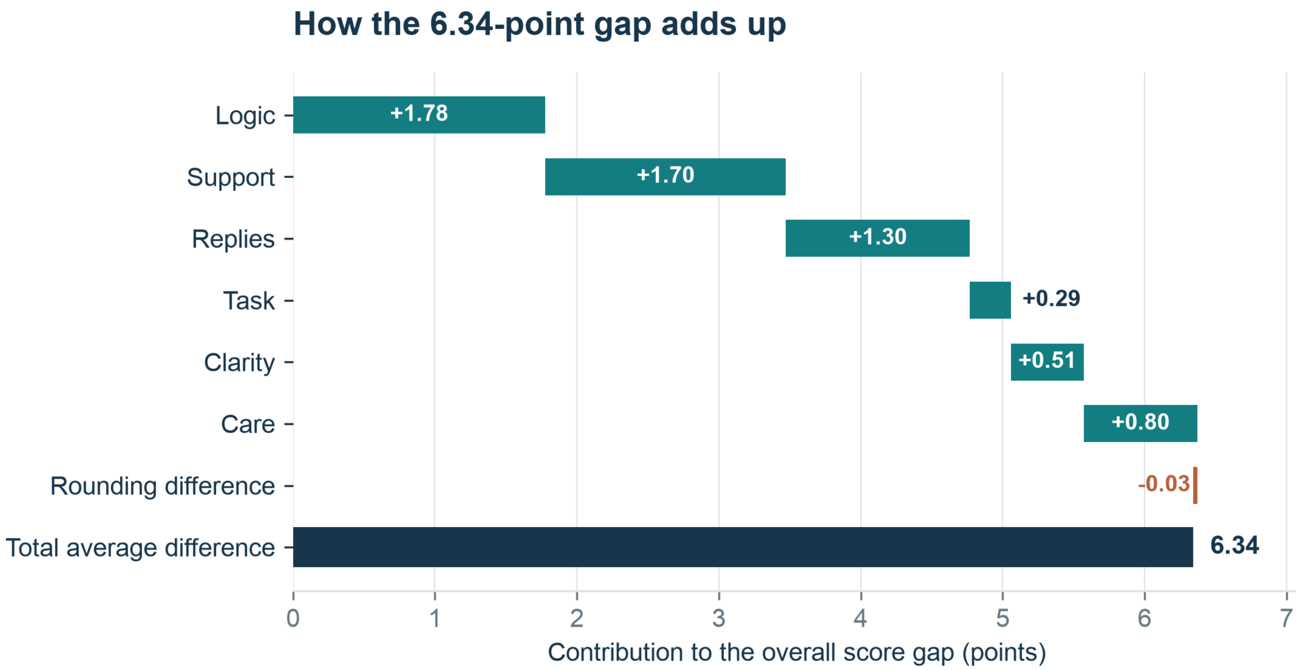

Figure 2. The officially weighted scoring areas reconstruct the average score difference. Consistent reasoning, evidence, and answering the other side contribute 4.77 points together. A small negative rounding residual reconciles the scoring area calculation with the published whole-number scores; the average burden-adjustment difference is zero.

How to read this graph. Each teal segment adds one scoring area's weighted contribution to the gap. The tiny rust segment subtracts the rounding difference. The dark final bar is the total, not another contribution. A segment's length, not where it starts, shows its contribution. These parts add to the score; they do not identify separate causes.
Area key: Logic = consistent reasoning; Support = evidence and reasons; Replies = answering the other side; Task = staying on the question and doing what the claim requires; Clarity = clear, exact wording; Care = fitting confidence to support and treating the other side fairly.

The strongest raw scoring area gap is evidence and supporting reasons, at **8.48 points**. Consistent reasoning contributes slightly more to the final score gap - 1.78 points versus 1.70 - because it has the larger weight. Those are not conflicting findings. “Largest underlying difference” and “largest contribution after weighting” answer different questions.

Answering the other side adds 1.30 overall points. Confidence and fairness add 0.80, precision and clarity 0.51, and relevance and burden 0.29. The small contribution from relevance suggests that the typical theist contribution is not simply off-topic. It often addresses an important issue while failing to justify the stronger conclusion drawn from it or to answer an important limit or condition.

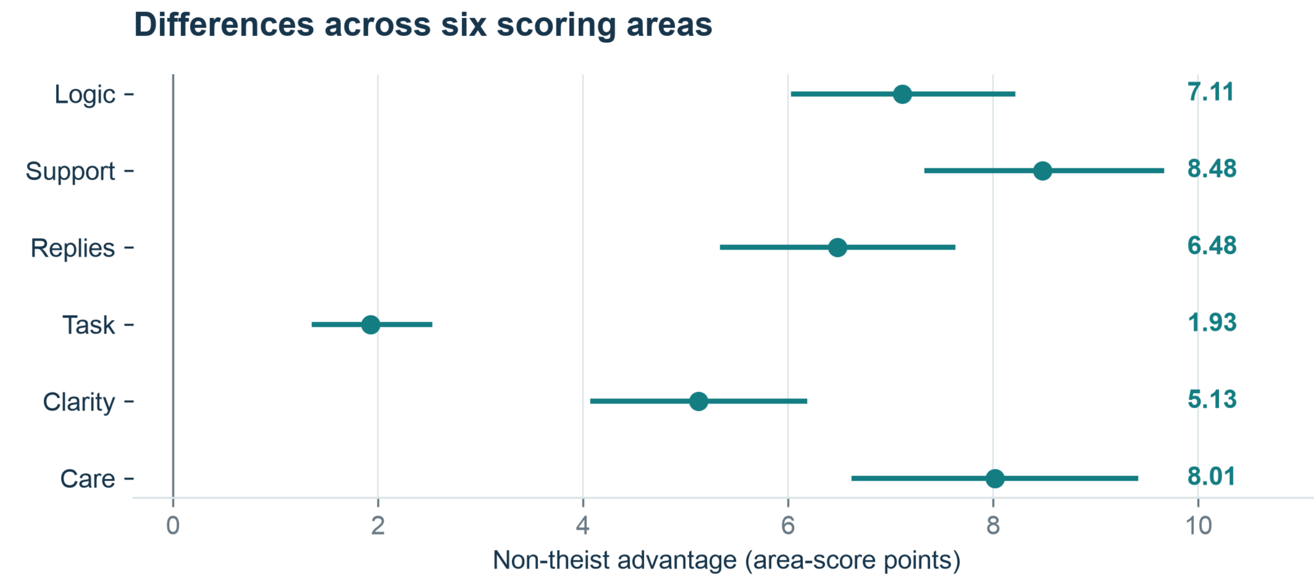

Figure 3. Support shows the largest difference before applying the six areas' weights. These are points on each area's own 0-100 scale, not contributions to the final score. Each range comes from repeatedly drawing complete debates.

How to read this graph. Each dot is an average non-theist-minus-theist difference. The line holds the middle 95% of averages from 20,000 repeated draws of whole debates, with replacement. Zero means equal averages. Right of zero favors the non-theist side. The units are each area's own 0-100 points, before weighting the areas together.
Area key: Logic = consistent reasoning; Support = evidence and reasons; Replies = answering the other side; Task = staying on the question and doing what the claim requires; Clarity = clear, exact wording; Care = fitting confidence to support and treating the other side fairly.

This decomposition is **a way of adding the scores, not a test of what caused them**. The scoring areas help determine the final score, so it is unsurprising that weaker scoring area ratings accompany lower totals. Furthermore, one defect can affect several scoring areas. An unsupported leap can count as weak support, questionable consistent reasoning, and excessive confidence. Agreement among those ratings is therefore not three independent discoveries of the same fact.

The informative question is more specific: what kinds of shortcomings did the assessments record, and do the source passages make that diagnosis reasonable? The following patterns connect the numbers to arguments a reader can inspect.

5. Five recurring failures of support for the claim

A possibility is promoted into a preferred explanation

A proposal can be logically possible without being likely, and likely without being the best-supported option. “A mind could explain this” establishes less than “a divine mind explains this better than the alternatives.” The stronger conclusion requires a comparison: what would we expect if the proposed explanation were true, and what would we expect otherwise?

In debate 244, Michael Jones argues toward a necessary divine mind, while Joseph Schmid challenges the identification of a necessary foundation with God. Even if some foundation is granted, the further description - mental, personal, good, unique, or religiously recognizable - still requires support. The weakness is an unfinished bridge between claims, not necessarily a failure to notice the original philosophical problem. See the Jones-Schmid scorecard.

A gap in a rival account becomes evidence for one's own account

An unresolved problem for naturalism may justify less confidence in a particular natural explanation. It does not automatically establish intelligent agency. Both explanations can remain incomplete. The crucial issue is whether the proposed agency account explains the evidence more specifically, or merely places a familiar label over the unknown.

The earlier assessment of debate 3 records a move from difficulty explaining cellular origins toward agency. The corresponding burden is not satisfied by making the difficulty vivid. A stronger case would identify an observable feature more expected under a specified agency hypothesis than under relevant non-agency alternatives. The Hitchens-D'Souza debate provides a concrete setting in which to examine that distinction.

The benefits of a belief stand in for the truth of its explanation

A belief can comfort people, organize a community, or support admirable behavior without its account of reality being true. Conversely, a true claim need not be comforting. The reasoning step from “this gives people purpose” to “its supernatural explanation is correct” needs a starting claim connecting those two achievements. Repeating the first achievement more eloquently does not supply the missing starting claim.

This is especially relevant to affirmations of God-given human worth. A theist may offer a coherent account of worth *within the worldview*. The public reasoning task is different: explain why a listener should accept that worldview, why creation confers the claimed moral status, and why alternatives fail. It is possible to respect the moral aspiration while finding that the presented reasoning step has not been completed.

Compatibility is treated as sufficient evidence

Showing that suffering and God are not logically contradictory can answer an argument that alleges a contradiction. It does not, by itself, answer the claim that the amount or distribution of suffering makes a particular God less likely. A possible divine reason may defeat an impossibility claim while doing little against an evidence-based comparison.

The distinction protects good theistic reasoning too. In debate 144, Alvin Plantinga explicitly limits a move to the compatibility of theism and Darwinism. That contribution receives a score of 92, with support at 86. The assessment recognizes that a modest conclusion can be well defended. The problem is not arguing for possibility; it is silently upgrading possibility into probability or actuality. See Plantinga-Dennett, 81:23.

The reply answers an easier claim

If a critic asks why *this distribution* of suffering is expected, replying that *some suffering* permits growth does not finish the exchange. If the question is whether testimony is reliable, replying that testimony is widespread leaves reliability open. If an argument concerns whether a starting claim follows, an attack on the critic's worldview may miss the local objection entirely.

These failures explain why answering the other side belongs beside support in the diagnosis. A case is not strengthened merely by restating its central theme after a challenge. It must identify what the challenge threatens and show why the supporting route remains intact. The reader should ask not only “Was an answer given?” but “What proposition did the answer actually address?”

6. Does the evidence problem depend on an arbitrary cutoff?

The average support difference uses the whole scoring scale, so it does not depend on choosing a boundary between adequate and inadequate support. Thresholds nevertheless make the pattern easier to picture. For each debate, we calculate the share of each side's moves below a chosen evidence score, then average those shares across debates.

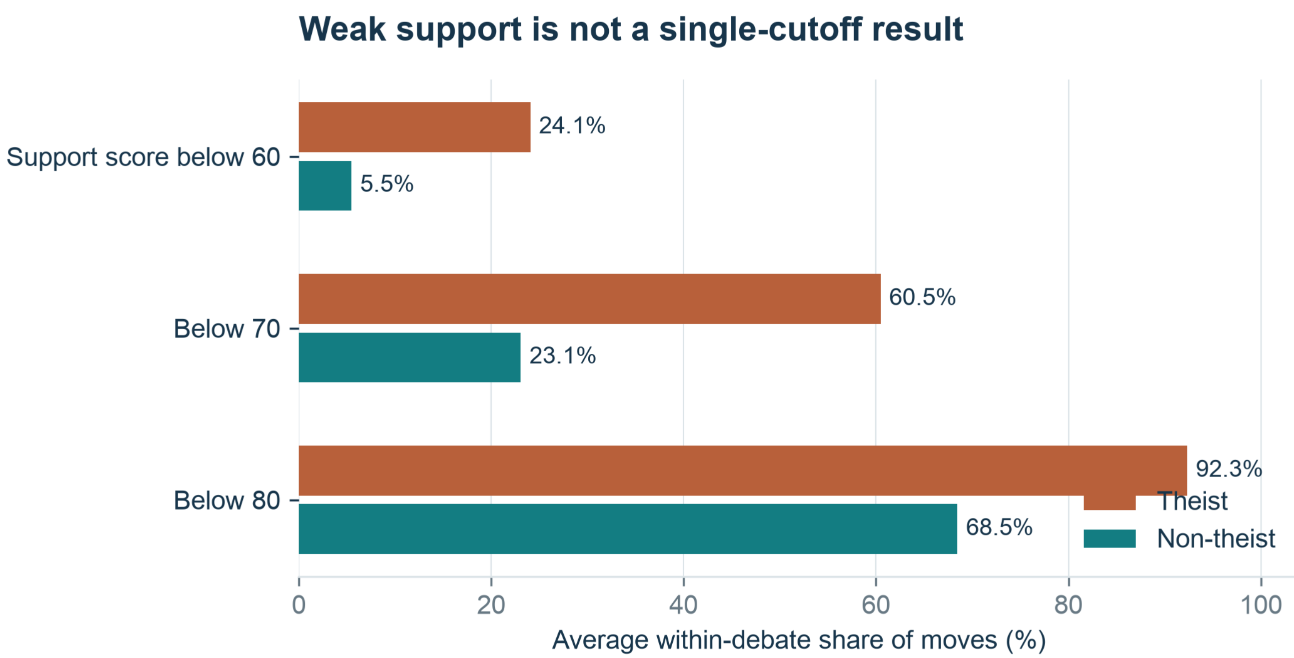

Figure 4. Changing the evidence threshold changes the size of the percentage gap but not its direction. The percentages describe an average debate's move shares, not the percentage of all theist or non-theist people who reason poorly.

How to read this graph. Rust bars are theist; teal bars are non-theist. Each percentage is the average share of a side's moves below the stated support score, giving every debate equal influence. A move below 60 is also below 70 and 80, so do not add the rows. The horizontal scale is percent, not score points.

Below 70, the average share is **60.5% for the theist side and 23.1% for the non-theist side**, a 37.4-percentage-point difference. Below 60, the corresponding shares are 24.1% and 5.5%. Below 80, they are 92.3% and 68.5%. The substantial skeptical share at the broader boundary is worth noticing: the archive does not depict the non-theist side as a storehouse of fully demonstrated claims.

A cutoff is an analytical convention, not a discovery that a score of 69 is fundamentally different from 70. The point of checking several cutoffs is to see whether the result depends on one convenient choice. Here it does not. The pattern remains, although its numerical description changes with the question being asked.

An evidence score also assesses what was supplied *in the recorded contribution*. It should not be translated into "no evidence exists anywhere." A speaker may know an extensive literature and fail to present the relevant part. That is a performance limitation, not proof that the literature cannot support the claim.

7. How the hypothesis fares against rival explanations

The faith-to-debate hypothesis has explanatory force because it predicts a linked pattern: confidence sustained by commitments that are not independently established for the audience; resistance to inconvenient alternatives; and replies that preserve the worldview more successfully than they resolve the objection. The assessments contain many moves with precisely that shape.

However, several other processes could produce some of the same observations. A more ambitious conclusion requires more supporting steps. An affirmative speaker may accept a heavier burden. Experienced skeptics may be overrepresented. The selected recordings may favor particular formats or speakers. The model may consistently reward a style of cautious criticism. And a scoring process can reproduce a systematic preference even when its numerical calculations are flawless.

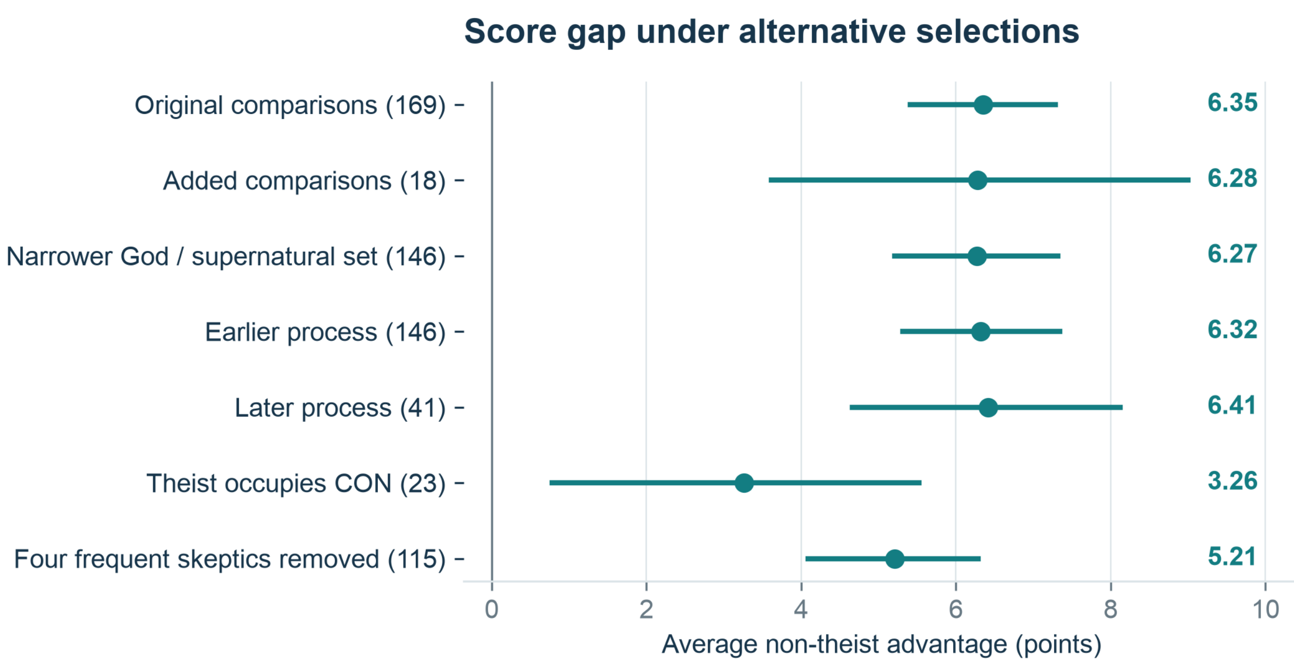

Figure 5. The score difference survives a narrower truth-claim comparison, removal of four frequent skeptical speakers, and separation by assessment process. These are checks using different selections, not randomized tests of faith or debating role.

How to read this graph. Each dot is an average gap for a different selection of debates. The horizontal line holds the middle 95% of averages from repeated draws of that selection. Parentheses give debate counts. Right of zero favors the non-theist side. Overlapping rows reuse many of the same debates; they are not independent tests.

Equalizing the influence of speakers provides another perspective. Giving each of the 84 religious-side speakers equal weight yields an average gap of 6.50 points; giving each of the 48 skeptical-side speakers equal weight yields 4.92. Neither is the uniquely correct estimate. They show that the headline depends somewhat on whose appearances are numerous, but the broad direction does not disappear when frequent speakers lose their extra influence.

The role check is similarly informative. When the theist occupies PRO, the average skeptical advantage is 6.77 points. When the theist occupies CON, it is 3.26. Formal position matters to the observed contrast, but it does not account for its entire direction. Paper four examines this issue in detail and finds that same-speaker comparisons leave a smaller and less certain role difference.

Most importantly, none of these checks measures the hypothesized transfer of private belief standards into public reasoning. They reduce particular alternative explanations; they do not turn a plausible interpretation into a measured cause. A strong defense of the hypothesis should make that boundary visible rather than conceal it.

8. Counterexamples are part of the result

The theist side scores higher in 20 comparisons. Some advantages are substantial. Keith Fox leads Peter Atkins by 15 points in debate 85; William Lane Craig leads Atkins by ten in debate 166; John Lennox leads Atkins by seven in debate 51. In the new material, James Anderson leads Tom Jump by five in debate 250. These are not exceptions to be hidden after the headline. They show that the rubric can reward theistic defenses and penalize skeptical overstatement.

The lesson is not that those results cancel the average gap. Twenty theist advantages do not erase 160 skeptical advantages. The lesson is that the proposed mechanism must explain variation as well as direction. Some religious arguments supply bounded claims, clear alternatives, and appropriate limits or conditions; some non-religious arguments make sweeping assertions that their evidence does not support.

A more useful claim than “theists are sloppy” is therefore: **certain ways of protecting a worldview make public support for the claim harder, and those ways are disproportionately recorded on the religious side of this archive**. That formulation identifies conduct that can change. It does not turn a group label into an explanation of a person's abilities or motives.

9. What would make the hypothesis more convincing - or weaker?

A stronger test would separate the identification of the proposed mechanism from the assignment of scores. Independent readers could examine randomly selected passages without seeing the public score, speaker name, or debate winner. They would classify whether the speaker appeals to revelation, private experience, inherited authority, practical benefit, or an independently checkable comparison. Other readers would judge the quality of the support. Agreement and disagreement between those teams would be reported.

The strongest prediction is not simply that theist sides score lower again. It is that moves relying on the hypothesized permissive standards should receive weaker independent support ratings, even after matching topic, burden, and how much a claim tries to establish. Within theist speakers, more disciplined treatment of those standards should predict better performance. Within non-theist speakers, similarly protected secular assumptions should incur comparable penalties.

The hypothesis would weaken if independent, blinded readers failed to find the proposed pattern; if the apparent difference largely vanished when the conclusion's ambition was matched; or if model changes strongly altered which side was judged to have supplied adequate support. It would also weaken if the source passages consistently supplied limits or conditions that the assessment summaries omitted.

These are genuine risks to the interpretation. A hypothesis that accommodates every possible outcome would imitate the very defect it criticizes. The next major-model review should publish the comparison rules before seeing the new results and retain both the original and revised judgments for inspection.

10. Conclusion: stronger claims need stronger support

Argument one: the score gap chiefly concerns the reasoning

1. A public argument must give listeners enough reason to accept its conclusion. Stating a belief firmly or clearly does not supply the missing evidence.

2. Across 187 relevant debates, the non-theist side averages **6.34 points higher** and leads in **160 comparisons**. The largest difference among the six scoring areas is in support for claims: **8.48 points** before applying the scoring weights.

3. Support, consistent reasoning, and answers to objections together account for about three-quarters of the overall gap under the site's scoring rules. The gap remains when we narrow the comparison or remove four frequent skeptical speakers.

Therefore: the records locate much of the theist-side disadvantage in claims that lack enough support, steps that do not establish the conclusion, and replies that leave objections unanswered. This is a finding about the assessed reasoning, not merely about presentation.

Argument two: the proposed cause fits the pattern, but is not proved

1. If standards that sustain faith carry into public debate, we would expect some speakers to treat shared belief, comfort, or mere possibility as enough support for a stronger claim.

2. The scores and examples contain that predicted pattern. They give the hypothesis evidence to explain.

3. The assessments do not measure the speakers' private standards of belief or establish why they argued as they did. Selection of speakers, topics, and judging preferences remain possible contributors. The theist side also scores higher in **20 debates**.

Therefore: the hypothesis deserves further testing; the claim that faith caused the lower scores goes beyond the present evidence. Neither a general judgment about religious people nor a verdict on God's existence follows from these scores.

The useful response is to complete the argument. Show why the evidence supports this conclusion rather than another; state how strong the conclusion really is; answer the objection that was raised. Those demands apply to both sides. A repeated affirmation becomes no better supported merely because its audience already believes it.

Methods and sources

The frozen source revision is **76d006b37**, containing 253 unique debate identifiers and unique video URLs. The analysis verifies 5,282 move scores against both their six scoring areas and the public cards, 2,558 side-section scores, all 474 comparable overall side scores, and 58 later-process final-record file hashes. These checks establish arithmetic and record consistency, not the correctness of the judgments themselves.

The main average gives one vote to each debate. Scoring area averages preserve the official within-debate importance and section weights. Move-type and threshold comparisons first summarize each side within a debate and then average paired differences. Intervals use 20,000 resamples and the fixed random seed 20260904. Repeated speakers, selective source inclusion, and systematic judging errors are not fully represented by those intervals. The exploratory topic and sensitivity choices were not set out before seeing the results.

The broader religious category, the marker for the narrower comparison, exclusions, normalized move data, source hashes, and executable calculations are available in the September 4 analysis directory. The Backend page explains the scoring framework and links the companion papers. Paper two examines topics; paper three separates measurable slogan risk from literal non-falsifiability; paper four tests role effects; papers six and seven examine scale and ranking limitations.

The qualitative examples are interpretations of the named scorecards and locked source passages, not a new independent coding of every transcript. All numerical claims in this paper concern the recorded performances in this snapshot. They should not be generalized without further evidence to all theists, all non-theists, or all public debates.